

MPI-2: Extensions to the Message-Passing Interface

Message Passing Interface Forum

July 18, 1997

This work was supported in part by NSF and DARPA under NSF contract CDA-9115428 and Esprit under project HPC Standards (21111).

Abstract

This document describes the **MPI-1.2** and **MPI-2** standards. They are both extensions to the **MPI-1.1** standard. The **MPI-1.2** part of the document contains clarifications and corrections to the **MPI-1.1** standard and defines **MPI-1.2**. The **MPI-2** part of the document describes additions to the **MPI-1** standard and defines **MPI-2**. These include miscellaneous topics, process creation and management, one-sided communications, extended collective operations, external interfaces, I/O, and additional language bindings.

©1995, 1996, 1997 University of Tennessee, Knoxville, Tennessee. Permission to copy without fee all or part of this material is granted, provided the University of Tennessee copyright notice and the title of this document appear, and notice is given that copying is by permission of the University of Tennessee.

Contents

Acknowledgments	ix
1 Introduction to MPI-2	1
1.1 Background	1
1.2 Organization of this Document	2
2 MPI-2 Terms and Conventions	5
2.1 Document Notation	5
2.2 Naming Conventions	5
2.3 Procedure Specification	6
2.4 Semantic Terms	7
2.5 Data Types	8
2.5.1 Opaque Objects	8
2.5.2 Array Arguments	10
2.5.3 State	10
2.5.4 Named Constants	10
2.5.5 Choice	11
2.5.6 Addresses	11
2.5.7 File Offsets	11
2.6 Language Binding	11
2.6.1 Deprecated Names and Functions	12
2.6.2 Fortran Binding Issues	13
2.6.3 C Binding Issues	14
2.6.4 C++ Binding Issues	15
2.7 Processes	18
2.8 Error Handling	18
2.9 Implementation Issues	19
2.9.1 Independence of Basic Runtime Routines	20
2.9.2 Interaction with Signals	20
2.10 Examples	20
3 Version 1.2 of MPI	21
3.1 Version Number	21
3.2 MPI-1.0 and MPI-1.1 Clarifications	22
3.2.1 Clarification of MPI_INITIALIZED	22
3.2.2 Clarification of MPI_FINALIZE	22
3.2.3 Clarification of status after MPI_WAIT and MPI_TEST	25
3.2.4 Clarification of MPI_INTERCOMM_CREATE	25

3.2.5	Clarification of <code>MPI_INTERCOMM_MERGE</code>	26
3.2.6	Clarification of Binding of <code>MPI_TYPE_SIZE</code>	26
3.2.7	Clarification of <code>MPI_REDUCE</code>	26
3.2.8	Clarification of Error Behavior of Attribute Callback Functions . . .	26
3.2.9	Clarification of <code>MPI_PROBE</code> and <code>MPI_IREQ</code>	26
3.2.10	Minor Corrections	27
4	Miscellany	37
4.1	Portable MPI Process Startup	37
4.2	Passing <code>NULL</code> to <code>MPI_Init</code>	39
4.3	Version Number	39
4.4	Datatype Constructor <code>MPI_TYPE_CREATE_INDEXED_BLOCK</code>	40
4.5	Treatment of <code>MPI_Status</code>	40
4.5.1	Passing <code>MPI_STATUS_IGNORE</code> for <code>Status</code>	40
4.5.2	Non-destructive Test of <code>status</code>	41
4.6	Error Class for Invalid Keyval	42
4.7	Committing a Committed Datatype	42
4.8	Allowing User Functions at Process Termination	42
4.9	Determining Whether MPI Has Finished	43
4.10	The <code>Info</code> Object	43
4.11	Memory Allocation	47
4.12	Language Interoperability	49
4.12.1	Introduction	49
4.12.2	Assumptions	50
4.12.3	Initialization	50
4.12.4	Transfer of Handles	51
4.12.5	Status	54
4.12.6	MPI Opaque Objects	54
4.12.7	Attributes	57
4.12.8	Extra State	59
4.12.9	Constants	59
4.12.10	Interlanguage Communication	60
4.13	Error Handlers	61
4.13.1	Error Handlers for Communicators	62
4.13.2	Error Handlers for Windows	63
4.13.3	Error Handlers for Files	64
4.14	New Datatype Manipulation Functions	65
4.14.1	Type Constructors with Explicit Addresses	66
4.14.2	Extent and Bounds of Datatypes	68
4.14.3	True Extent of Datatypes	69
4.14.4	Subarray Datatype Constructor	70
4.14.5	Distributed Array Datatype Constructor	72
4.15	New Predefined Datatypes	77
4.15.1	Wide Characters	77
4.15.2	Signed Characters and Reductions	77
4.15.3	Unsigned long long Type	78
4.16	Canonical <code>MPI_PACK</code> and <code>MPI_UNPACK</code>	78
4.17	Functions and Macros	80

4.18	Profiling Interface	80
5	Process Creation and Management	81
5.1	Introduction	81
5.2	The MPI-2 Process Model	82
5.2.1	Starting Processes	82
5.2.2	The Runtime Environment	82
5.3	Process Manager Interface	84
5.3.1	Processes in MPI	84
5.3.2	Starting Processes and Establishing Communication	84
5.3.3	Starting Multiple Executables and Establishing Communication	89
5.3.4	Reserved Keys	91
5.3.5	Spawn Example	92
5.4	Establishing Communication	94
5.4.1	Names, Addresses, Ports, and All That	94
5.4.2	Server Routines	95
5.4.3	Client Routines	97
5.4.4	Name Publishing	99
5.4.5	Reserved Key Values	101
5.4.6	Client/Server Examples	101
5.5	Other Functionality	104
5.5.1	Universe Size	104
5.5.2	Singleton MPI_INIT	104
5.5.3	MPI_APPNUM	105
5.5.4	Releasing Connections	106
5.5.5	Another Way to Establish MPI Communication	107
6	One-Sided Communications	109
6.1	Introduction	109
6.2	Initialization	110
6.2.1	Window Creation	110
6.2.2	Window Attributes	112
6.3	Communication Calls	113
6.3.1	Put	114
6.3.2	Get	116
6.3.3	Examples	116
6.3.4	Accumulate Functions	119
6.4	Synchronization Calls	121
6.4.1	Fence	126
6.4.2	General Active Target Synchronization	127
6.4.3	Lock	130
6.4.4	Assertions	132
6.4.5	Miscellaneous Clarifications	134
6.5	Examples	134
6.6	Error Handling	136
6.6.1	Error Handlers	136
6.6.2	Error Classes	137
6.7	Semantics and Correctness	137

6.7.1	Atomicity	140
6.7.2	Progress	140
6.7.3	Registers and Compiler Optimizations	142
7	Extended Collective Operations	145
7.1	Introduction	145
7.2	Intercommunicator Constructors	145
7.3	Extended Collective Operations	149
7.3.1	Intercommunicator Collective Operations	149
7.3.2	Operations that Move Data	151
7.3.3	Reductions	162
7.3.4	Other Operations	163
7.3.5	Generalized All-to-all Function	164
7.3.6	Exclusive Scan	166
8	External Interfaces	169
8.1	Introduction	169
8.2	Generalized Requests	169
8.2.1	Examples	173
8.3	Associating Information with Status	175
8.4	Naming Objects	177
8.5	Error Classes, Error Codes, and Error Handlers	181
8.6	Decoding a Datatype	184
8.7	MPI and Threads	193
8.7.1	General	193
8.7.2	Clarifications	194
8.7.3	Initialization	195
8.8	New Attribute Caching Functions	198
8.8.1	Communicators	199
8.8.2	Windows	202
8.8.3	Datatypes	204
8.9	Duplicating a Datatype	207
9	I/O	209
9.1	Introduction	209
9.1.1	Definitions	209
9.2	File Manipulation	211
9.2.1	Opening a File	211
9.2.2	Closing a File	213
9.2.3	Deleting a File	214
9.2.4	Resizing a File	215
9.2.5	Preallocating Space for a File	215
9.2.6	Querying the Size of a File	216
9.2.7	Querying File Parameters	216
9.2.8	File Info	218
9.3	File Views	221
9.4	Data Access	223
9.4.1	Data Access Routines	223

9.4.2	Data Access with Explicit Offsets	226
9.4.3	Data Access with Individual File Pointers	230
9.4.4	Data Access with Shared File Pointers	235
9.4.5	Split Collective Data Access Routines	240
9.5	File Interoperability	246
9.5.1	Datatypes for File Interoperability	248
9.5.2	External Data Representation: “external32”	250
9.5.3	User-Defined Data Representations	251
9.5.4	Matching Data Representations	255
9.6	Consistency and Semantics	255
9.6.1	File Consistency	255
9.6.2	Random Access vs. Sequential Files	258
9.6.3	Progress	259
9.6.4	Collective File Operations	259
9.6.5	Type Matching	259
9.6.6	Miscellaneous Clarifications	259
9.6.7	<code>MPI_Offset</code> Type	260
9.6.8	Logical vs. Physical File Layout	260
9.6.9	File Size	260
9.6.10	Examples	261
9.7	I/O Error Handling	265
9.8	I/O Error Classes	265
9.9	Examples	266
9.9.1	Double Buffering with Split Collective I/O	266
9.9.2	Subarray Filetype Constructor	268
10	Language Bindings	271
10.1	C++	271
10.1.1	Overview	271
10.1.2	Design	271
10.1.3	C++ Classes for <code>MPI</code>	272
10.1.4	Class Member Functions for <code>MPI</code>	273
10.1.5	Semantics	273
10.1.6	C++ Datatypes	275
10.1.7	Communicators	278
10.1.8	Exceptions	280
10.1.9	Mixed-Language Operability	281
10.1.10	Profiling	281
10.2	Fortran Support	284
10.2.1	Overview	284
10.2.2	Problems With Fortran Bindings for <code>MPI</code>	284
10.2.3	Basic Fortran Support	291
10.2.4	Extended Fortran Support	291
10.2.5	Additional Support for Fortran Numeric Intrinsic Types	292
	Bibliography	301

A	Language Binding	303
A.1	Introduction	303
A.2	Defined Values and Handles	303
A.2.1	Defined Constants	303
A.2.2	Info Keys	307
A.2.3	Info Values	308
A.3	MPI-1.2 C Bindings	308
A.4	MPI-1.2 Fortran Bindings	308
A.5	MPI-1.2 C++ Bindings	309
A.6	MPI-2 C Bindings	309
A.6.1	Miscellany	309
A.6.2	Process Creation and Management	311
A.6.3	One-Sided Communications	312
A.6.4	Extended Collective Operations	312
A.6.5	External Interfaces	312
A.6.6	I/O	314
A.6.7	Language Bindings	316
A.6.8	User Defined Functions	317
A.7	MPI-2 Fortran Bindings	317
A.7.1	Miscellany	317
A.7.2	Process Creation and Management	320
A.7.3	One-Sided Communications	321
A.7.4	Extended Collective Operations	322
A.7.5	External Interfaces	323
A.7.6	I/O	325
A.7.7	Language Bindings	329
A.7.8	User Defined Subroutines	330
A.8	MPI-2 C++ Bindings	331
A.8.1	Miscellany	331
A.8.2	Process Creation and Management	333
A.8.3	One-Sided Communications	334
A.8.4	Extended Collective Operations	334
A.8.5	External Interfaces	336
A.8.6	I/O	337
A.8.7	Language Bindings	341
A.8.8	User Defined Functions	341
B	MPI-1 C++ Language Binding	342
B.1	C++ Classes	342
B.2	Defined Constants	342
B.3	Typedefs	346
B.4	C++ Bindings for Point-to-Point Communication	347
B.5	C++ Bindings for Collective Communication	349
B.6	C++ Bindings for Groups, Contexts, and Communicators	351
B.7	C++ Bindings for Process Topologies	352
B.8	C++ Bindings for Environmental Inquiry	353
B.9	C++ Bindings for Profiling	353
B.10	C++ Bindings for Status Access	353

B.11 C++ Bindings for New 1.2 Functions	354
B.12 C++ Bindings for Exceptions	354
B.13 C++ Bindings on all MPI Classes	354
B.13.1 Construction / Destruction	355
B.13.2 Copy / Assignment	355
B.13.3 Comparison	355
B.13.4 Inter-language Operability	355
B.13.5 Function Name Cross Reference	356
MPI Function Index	360

Acknowledgments

This document represents the work of many people who have served on the MPI Forum. The meetings have been attended by dozens of people from many parts of the world. It is the hard and dedicated work of this group that has led to the MPI standard.

The technical development was carried out by subgroups, whose work was reviewed by the full committee. During the period of development of the Message Passing Interface (MPI-2), many people helped with this effort. Those who served as the primary coordinators are:

- Ewing Lusk, Convener and Meeting Chair
- Steve Huss-Lederman, Editor
- Ewing Lusk, Miscellany
- Bill Saphir, Process Creation and Management
- Marc Snir, One-Sided Communications
- Bill Gropp and Anthony Skjellum, Extended Collective Operations
- Steve Huss-Lederman, External Interfaces
- Bill Nitzberg, I/O
- Andrew Lumsdaine, Bill Saphir, and Jeff Squyres, Language Bindings
- Anthony Skjellum and Arkady Kanevsky, Real-Time

The following list includes some of the active participants who attended MPI-2 Forum meetings and are not mentioned above.

Greg Astfalk	Robert Babb	Ed Benson	Rajesh Bordawekar
Pete Bradley	Peter Brennan	Ron Brightwell	Maciej Brodowicz
Eric Brunner	Greg Burns	Margaret Cahir	Pang Chen
Ying Chen	Albert Cheng	Yong Cho	Joel Clark
Lyndon Clarke	Laurie Costello	Dennis Cottel	Jim Cownie
Zhenqian Cui	Suresh Damodaran-Kamal	Raja Daoud	Judith Devaney
David DiNucci	Doug Doeffler	Jack Dongarra	Terry Dontje
Nathan Doss	Anne Elster	Mark Fallon	Karl Feind
Sam Fineberg	Craig Fischberg	Stephen Fleischman	Ian Foster
Hubertus Franke	Richard Frost	Al Geist	Robert George
David Greenberg	John Hagedorn	Kei Harada	Leslie Hart
Shane Hebert	Rolf Hempel	Tom Henderson	Alex Ho
Hans-Christian Hoppe	Joefon Jann	Terry Jones	Karl Kesselman
Koichi Konishi	Susan Kraus	Steve Kubica	Steve Landherr
Mario Lauria	Mark Law	Juan Leon	Lloyd Lewins
Ziyang Lu	Bob Madahar	Peter Madams	John May
Oliver McBryan	Brian McCandless	Tyce McLarty	Thom McMahon
Harish Nag	Nick Nevin	Jarek Nieplocha	Ron Oldfield
Peter Ossadnik	Steve Otto	Peter Pacheco	Yoonho Park
Perry Partow	Pratap Pattnaik	Elsie Pierce	Paul Pierce
Heidi Poxon	Jean-Pierre Prost	Boris Protopopov	James Pruyve
Rolf Rabenseifner	Joe Rieken	Peter Rigsbee	Tom Robey
Anna Rounbehler	Nobutoshi Sagawa	Arindam Saha	Eric Salo
Darren Sanders	Eric Sharakan	Andrew Sherman	Fred Shirley
Lance Shuler	A. Gordon Smith	Ian Stockdale	David Taylor
Stephen Taylor	Greg Tensa	Rajeev Thakur	Marydell Tholburn
Dick Treumann	Simon Tsang	Manuel Ujaldon	David Walker
Jerrell Watts	Klaus Wolf	Parkson Wong	Dave Wright

The MPI Forum also acknowledges and appreciates the valuable input from people via e-mail and in person.

The following institutions supported the MPI-2 effort through time and travel support for the people listed above.

Argonne National Laboratory
 Bolt, Beranek, and Newman
 California Institute of Technology
 Center for Computing Sciences
 Convex Computer Corporation
 Cray Research
 Digital Equipment Corporation
 Dolphin Interconnect Solutions, Inc.
 Edinburgh Parallel Computing Centre
 General Electric Company
 German National Research Center for Information Technology

Hewlett-Packard
Hitachi
Hughes Aircraft Company
Intel Corporation
International Business Machines
Khoral Research
Lawrence Livermore National Laboratory
Los Alamos National Laboratory
MPI Software Technology, Inc.
Mississippi State University
NEC Corporation
National Aeronautics and Space Administration
National Energy Research Scientific Computing Center
National Institute of Standards and Technology
National Oceanic and Atmospheric Administration
Oak Ridge National Laboratory
Ohio State University
PALLAS GmbH
Pacific Northwest National Laboratory
Pratt & Whitney
San Diego Supercomputer Center
Sanders, A Lockheed-Martin Company
Sandia National Laboratories
Schlumberger
Scientific Computing Associates, Inc.
Silicon Graphics Incorporated
Sky Computers
Sun Microsystems Computer Corporation
Syracuse University
The MITRE Corporation
Thinking Machines Corporation
United States Navy
University of Colorado
University of Denver
University of Houston
University of Illinois
University of Maryland
University of Notre Dame
University of San Francisco
University of Stuttgart Computing Center
University of Wisconsin

MPI-2 operated on a very tight budget (in reality, it had no budget when the first meeting was announced). Many institutions helped the MPI-2 effort by supporting the efforts and travel of the members of the MPI Forum. Direct support was given by NSF and DARPA under NSF contract CDA-9115428 for travel by U.S. academic participants and Esprit under project HPC Standards (21111) for European participants.

Chapter 1

Introduction to MPI-2

1.1 Background

Beginning in March 1995, the MPI Forum began meeting to consider corrections and extensions to the original MPI Standard document [5]. The first product of these deliberations was Version 1.1 of the MPI specification, released in June of 1995 (see <http://www.mpi-forum.org> for official MPI document releases). Since that time, effort has been focused in five types of areas.

1. Further corrections and clarifications for the MPI-1.1 document.
2. Additions to MPI-1.1 that do not significantly change its types of functionality (new datatype constructors, language interoperability, etc.).
3. Completely new types of functionality (dynamic processes, one-sided communication, parallel I/O, etc.) that are what everyone thinks of as “MPI-2 functionality.”
4. Bindings for Fortran 90 and C++. This document specifies C++ bindings for both MPI-1 and MPI-2 functions, and extensions to the Fortran 77 binding of MPI-1 and MPI-2 to handle Fortran 90 issues.
5. Discussions of areas in which the MPI process and framework seem likely to be useful, but where more discussion and experience are needed before standardization (e.g. 0-copy semantics on shared-memory machines, real-time specifications).

Corrections and clarifications (items of type 1 in the above list) have been collected in Chapter 3 of this document, “Version 1.2 of MPI.” This chapter also contains the function for identifying the version number. Additions to MPI-1.1 (items of types 2, 3, and 4 in the above list) are in the remaining chapters, and constitute the specification for MPI-2. This document specifies Version 2.0 of MPI. Items of type 5 in the above list have been moved to a separate document, the “MPI Journal of Development” (JOD), and are not part of the MPI-2 Standard.

This structure makes it easy for users and implementors to understand what level of MPI compliance a given implementation has:

- MPI-1 compliance will mean compliance with MPI-1.2. This is a useful level of compliance. It means that the implementation conforms to the clarifications of MPI-1.1 function behavior given in Chapter 3. Some implementations may require changes to be MPI-1 compliant.

- MPI-2 compliance will mean compliance with all of MPI-2.
- The MPI Journal of Development is not part of the MPI Standard.

It is to be emphasized that forward compatibility is preserved. That is, a valid MPI-1.1 program is both a valid MPI-1.2 program and a valid MPI-2 program, and a valid MPI-1.2 program is a valid MPI-2 program.

1.2 Organization of this Document

This document is organized as follows:

- Chapter 2, **MPI-2 Terms and Conventions**, explains notational terms and conventions used throughout the MPI-2 document.
- Chapter 3, **Version 1.2 of MPI**, contains the specification of MPI-1.2, which has one new function and consists primarily of clarifications to MPI-1.1. It is expected that some implementations will need modification in order to become MPI-1 compliant, as the result of these clarifications.

The rest of this document contains the MPI-2 Standard Specification. It adds substantial new types of functionality to MPI, in most cases specifying functions for an extended computational model (e.g., dynamic process creation and one-sided communication) or for a significant new capability (e.g., parallel I/O).

The following is a list of the chapters in MPI-2, along with a brief description of each.

- Chapter 4, **Miscellany**, discusses items that don't fit elsewhere, in particular language interoperability.
- Chapter 5, **Process Creation and Management**, discusses the extension of MPI to remove the static process model in MPI. It defines routines that allow for creation of processes.
- Chapter 6, **One-Sided Communications**, defines communication routines that can be completed by a single process. These include shared-memory operations (put/get) and remote accumulate operations.
- Chapter 7, **Extended Collective Operations**, extends the semantics of MPI-1 collective operations to include intercommunicators. It also adds more convenient methods of constructing intercommunicators and two new collective operations.
- Chapter 8, **External Interfaces**, defines routines designed to allow developers to layer on top of MPI. This includes generalized requests, routines that decode MPI opaque objects, and threads.
- Chapter 9, **I/O**, defines MPI-2 support for parallel I/O.
- Chapter 10, **Language Bindings**, describes the C++ binding and discusses Fortran-90 issues.

The Appendices are:

- Annex A, **Language Bindings**, gives bindings for MPI-2 functions, and lists constants, error codes, etc.
- Annex B, **MPI-1 C++ Language Binding**, gives C++ bindings for MPI-1.

The **MPI Function Index** is a simple index showing the location of the precise definition of each MPI-2 function, together with C, C++, and Fortran bindings.

MPI-2 provides various interfaces to facilitate interoperability of distinct MPI implementations. Among these are the canonical data representation for MPI I/O and for `MPI_PACK_EXTERNAL` and `MPI_UNPACK_EXTERNAL`. The definition of an actual binding of these interfaces that will enable interoperability is outside the scope of this document.

A separate document consists of ideas that were discussed in the MPI Forum and deemed to have value, but are not included in the MPI Standard. They are part of the “Journal of Development” (JOD), lest good ideas be lost and in order to provide a starting point for further work. The chapters in the JOD are

- Chapter 2, **Spawning Independent Processes**, includes some elements of dynamic process management, in particular management of processes with which the spawning processes do not intend to communicate, that the Forum discussed at length but ultimately decided not to include in the MPI Standard.
- Chapter 3, **Threads and MPI**, describes some of the expected interaction between an MPI implementation and a thread library in a multi-threaded environment.
- Chapter 4, **Communicator ID**, describes an approach to providing identifiers for communicators.
- Chapter 5, **Miscellany**, discusses Miscellaneous topics in the MPI JOD, in particular single-copy routines for use in shared-memory environments and new datatype constructors.
- Chapter 6, **Toward a Full Fortran 90 Interface**, describes an approach to providing a more elaborate Fortran 90 interface.
- Chapter 7, **Split Collective Communication**, describes a specification for certain non-blocking collective operations.
- Chapter 8, **Real-Time MPI**, discusses MPI support for real time processing.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 2

MPI-2 Terms and Conventions

This chapter explains notational terms and conventions used throughout the MPI-2 document, some of the choices that have been made, and the rationale behind those choices. It is similar to the MPI-1 Terms and Conventions chapter but differs in some major and minor ways. Some of the major areas of difference are the naming conventions, some semantic definitions, file objects, Fortran 90 *vs* Fortran 77, C++, processes, and interaction with signals.

2.1 Document Notation

Rationale. Throughout this document, the rationale for the design choices made in the interface specification is set off in this format. Some readers may wish to skip these sections, while readers interested in interface design may want to read them carefully. (*End of rationale.*)

Advice to users. Throughout this document, material aimed at users and that illustrates usage is set off in this format. Some readers may wish to skip these sections, while readers interested in programming in MPI may want to read them carefully. (*End of advice to users.*)

Advice to implementors. Throughout this document, material that is primarily commentary to implementors is set off in this format. Some readers may wish to skip these sections, while readers interested in MPI implementations may want to read them carefully. (*End of advice to implementors.*)

2.2 Naming Conventions

MPI-1 used informal naming conventions. In many cases, MPI-1 names for C functions are of the form `Class_action_subset` and in Fortran of the form `CLASS_ACTION_SUBSET`, but this rule is not uniformly applied. In MPI-2, an attempt has been made to standardize names of new functions according to the following rules. In addition, the C++ bindings for MPI-1 functions also follow these rules (see Section 2.6.4). C and Fortran function names for MPI-1 have not been changed.

1. In C, all routines associated with a particular type of MPI object should be of the form `Class_action_subset` or, if no subset exists, of the form `Class_action`. In Fortran,

all routines associated with a particular type of MPI object should be of the form `CLASS_ACTION_SUBSET` or, if no subset exists, of the form `CLASS_ACTION`. For C and Fortran we use the C++ terminology to define the `Class`. In C++, the routine is a method on `Class` and is named `MPI::Class::Action_subset`. If the routine is associated with a certain class, but does not make sense as an object method, it is a static member function of the class.

2. If the routine is not associated with a class, the name should be of the form `Action_subset` in C and `ACTION_SUBSET` in Fortran, and in C++ should be scoped in the `MPI` namespace, `MPI::Action_subset`.
3. The names of certain actions have been standardized. In particular, `Create` creates a new object, `Get` retrieves information about an object, `Set` sets this information, `Delete` deletes information, `Is` asks whether or not an object has a certain property.

C and Fortran names for MPI-1 functions violate these rules in several cases. The most common exceptions are the omission of the `Class` name from the routine and the omission of the `Action` where one can be inferred.

MPI identifiers are limited to 30 characters (31 with the profiling interface). This is done to avoid exceeding the limit on some compilation systems.

2.3 Procedure Specification

MPI procedures are specified using a language-independent notation. The arguments of procedure calls are marked as IN, OUT or INOUT. The meanings of these are:

- the call may use the input value but does not update an argument is marked IN,
- the call may update an argument but does not use its input value is marked OUT,
- the call may both use and update an argument is marked INOUT.

There is one special case — if an argument is a handle to an opaque object (these terms are defined in Section 2.5.1), and the object is updated by the procedure call, then the argument is marked OUT. It is marked this way even though the handle itself is not modified — we use the OUT attribute to denote that what the handle *references* is updated. Thus, in C++, IN arguments are either references or pointers to `const` objects.

Rationale. The definition of MPI tries to avoid, to the largest possible extent, the use of INOUT arguments, because such use is error-prone, especially for scalar arguments. (*End of rationale.*)

MPI's use of IN, OUT and INOUT is intended to indicate to the user how an argument is to be used, but does not provide a rigorous classification that can be translated directly into all language bindings (e.g., `INTENT` in Fortran 90 bindings or `const` in C bindings). For instance, the “constant” `MPI_BOTTOM` can usually be passed to OUT buffer arguments. Similarly, `MPI_STATUS_IGNORE` can be passed as the OUT status argument.

A common occurrence for MPI functions is an argument that is used as IN by some processes and OUT by other processes. Such an argument is, syntactically, an INOUT argument and is marked as such, although, semantically, it is not used in one call both for input and for output on a single process.

Another frequent situation arises when an argument value is needed only by a subset of the processes. When an argument is not significant at a process then an arbitrary value can be passed as an argument.

Unless specified otherwise, an argument of type OUT or type INOUT cannot be aliased with any other argument passed to an MPI procedure. An example of argument aliasing in C appears below. If we define a C procedure like this,

```
void copyIntBuffer( int *pin, int *pout, int len )
{
    int i;
    for (i=0; i<len; ++i) *pout++ = *pin++;
}
```

then a call to it in the following code fragment has aliased arguments.

```
int a[10];
copyIntBuffer( a, a+3, 7);
```

Although the C language allows this, such usage of MPI procedures is forbidden unless otherwise specified. Note that Fortran prohibits aliasing of arguments.

All MPI functions are first specified in the language-independent notation. Immediately below this, the ANSI C version of the function is shown followed by a version of the same function in Fortran and then the C++ binding. Fortran in this document refers to Fortran 90; see Section 2.6.

2.4 Semantic Terms

When discussing MPI procedures the following semantic terms are used.

nonblocking A procedure is nonblocking if the procedure may return before the operation completes, and before the user is allowed to reuse resources (such as buffers) specified in the call. A nonblocking request is **started** by the call that initiates it, e.g., `MPI_SEND`. The word complete is used with respect to operations, requests, and communications. An **operation completes** when the user is allowed to reuse resources, and any output buffers have been updated; i.e. a call to `MPI_TEST` will return `flag = true`. A **request is completed** by a call to wait, which returns, or a test or get status call which returns `flag = true`. This completing call has two effects: the status is extracted from the request; in the case of test and wait, if the request was nonpersistent, it is **freed**. A **communication completes** when all participating operations complete.

blocking A procedure is blocking if return from the procedure indicates the user is allowed to reuse resources specified in the call.

local A procedure is local if completion of the procedure depends only on the local executing process.

non-local A procedure is non-local if completion of the operation may require the execution of some MPI procedure on another process. Such an operation may require communication occurring with another user process.

collective A procedure is collective if all processes in a process group need to invoke the procedure. A collective call may or may not be synchronizing. Collective calls over the same communicator must be executed in the same order by all members of the process group.

predefined A predefined datatype is a datatype with a predefined (constant) name (such as `MPI_INT`, `MPI_FLOAT_INT`, or `MPI_UB`) or a datatype constructed with `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL`, or `MPI_TYPE_CREATE_F90_COMPLEX`. The former are **named** whereas the latter are **unnamed**.

derived A derived datatype is any datatype that is not predefined.

portable A datatype is portable, if it is a predefined datatype, or it is derived from a portable datatype using only the type constructors `MPI_TYPE_CONTIGUOUS`, `MPI_TYPE_VECTOR`, `MPI_TYPE_INDEXED`, `MPI_TYPE_INDEXED_BLOCK`, `MPI_TYPE_CREATE_SUBARRAY`, `MPI_TYPE_DUP`, and `MPI_TYPE_CREATE_DARRAY`. Such a datatype is portable because all displacements in the datatype are in terms of extents of one predefined datatype. Therefore, if such a datatype fits a data layout in one memory, it will fit the corresponding data layout in another memory, if the same declarations were used, even if the two systems have different architectures. On the other hand, if a datatype was constructed using `MPI_TYPE_CREATE_HINDEXED`, `MPI_TYPE_CREATE_HVECTOR` or `MPI_TYPE_CREATE_STRUCT`, then the datatype contains explicit byte displacements (e.g., providing padding to meet alignment restrictions). These displacements are unlikely to be chosen correctly if they fit data layout on one memory, but are used for data layouts on another process, running on a processor with a different architecture.

equivalent Two datatypes are equivalent if they appear to have been created with the same sequence of calls (and arguments) and thus have the same typemap. Two equivalent datatypes do not necessarily have the same cached attributes or the same names.

2.5 Data Types

2.5.1 Opaque Objects

MPI manages **system memory** that is used for buffering messages and for storing internal representations of various MPI objects such as groups, communicators, datatypes, etc. This memory is not directly accessible to the user, and objects stored there are **opaque**: their size and shape is not visible to the user. Opaque objects are accessed via **handles**, which exist in user space. MPI procedures that operate on opaque objects are passed handle arguments to access these objects. In addition to their use by MPI calls for object access, handles can participate in assignments and comparisons.

In Fortran, all handles have type `INTEGER`. In C and C++, a different handle type is defined for each category of objects. In addition, handles themselves are distinct objects in C++. The C and C++ types must support the use of the assignment and equality operators.

Advice to implementors. In Fortran, the handle can be an index into a table of opaque objects in a system table; in C it can be such an index or a pointer to the

object. C++ handles can simply “wrap up” a table index or pointer.

(End of advice to implementors.)

Opaque objects are allocated and deallocated by calls that are specific to each object type. These are listed in the sections where the objects are described. The calls accept a handle argument of matching type. In an allocate call this is an OUT argument that returns a valid reference to the object. In a call to deallocate this is an INOUT argument which returns with an “invalid handle” value. MPI provides an “invalid handle” constant for each object type. Comparisons to this constant are used to test for validity of the handle.

A call to a deallocate routine invalidates the handle and marks the object for deallocation. The object is not accessible to the user after the call. However, MPI need not deallocate the object immediately. Any operation pending (at the time of the deallocate) that involves this object will complete normally; the object will be deallocated afterwards.

An opaque object and its handle are significant only at the process where the object was created and cannot be transferred to another process.

MPI provides certain predefined opaque objects and predefined, static handles to these objects. The user must not free such objects. In C++, this is enforced by declaring the handles to these predefined objects to be `static const`.

Rationale. This design hides the internal representation used for MPI data structures, thus allowing similar calls in C, C++, and Fortran. It also avoids conflicts with the typing rules in these languages, and easily allows future extensions of functionality. The mechanism for opaque objects used here loosely follows the POSIX Fortran binding standard.

The explicit separation of handles in user space and objects in system space allows space-reclaiming and deallocation calls to be made at appropriate points in the user program. If the opaque objects were in user space, one would have to be very careful not to go out of scope before any pending operation requiring that object completed. The specified design allows an object to be marked for deallocation, the user program can then go out of scope, and the object itself still persists until any pending operations are complete.

The requirement that handles support assignment/comparison is made since such operations are common. This restricts the domain of possible implementations. The alternative would have been to allow handles to have been an arbitrary, opaque type. This would force the introduction of routines to do assignment and comparison, adding complexity, and was therefore ruled out. *(End of rationale.)*

Advice to users. A user may accidentally create a dangling reference by assigning to a handle the value of another handle, and then deallocating the object associated with these handles. Conversely, if a handle variable is deallocated before the associated object is freed, then the object becomes inaccessible (this may occur, for example, if the handle is a local variable within a subroutine, and the subroutine is exited before the associated object is deallocated). It is the user’s responsibility to avoid adding or deleting references to opaque objects, except as a result of MPI calls that allocate or deallocate such objects. *(End of advice to users.)*

Advice to implementors. The intended semantics of opaque objects is that opaque objects are separate from one another; each call to allocate such an object copies

all the information required for the object. Implementations may avoid excessive copying by substituting referencing for copying. For example, a derived datatype may contain references to its components, rather than copies of its components; a call to `MPI_COMM_GROUP` may return a reference to the group associated with the communicator, rather than a copy of this group. In such cases, the implementation must maintain reference counts, and allocate and deallocate objects in such a way that the visible effect is as if the objects were copied. (*End of advice to implementors.*)

2.5.2 Array Arguments

An MPI call may need an argument that is an array of opaque objects, or an array of handles. The array-of-handles is a regular array with entries that are handles to objects of the same type in consecutive locations in the array. Whenever such an array is used, an additional `len` argument is required to indicate the number of valid entries (unless this number can be derived otherwise). The valid entries are at the beginning of the array; `len` indicates how many of them there are, and need not be the size of the entire array. The same approach is followed for other array arguments. In some cases `NULL` handles are considered valid entries. When a `NULL` argument is desired for an array of statuses, one uses `MPI_STATUSES_IGNORE`.

2.5.3 State

MPI procedures use at various places arguments with *state* types. The values of such a data type are all identified by names, and no operation is defined on them. For example, the `MPI_TYPE_CREATE_SUBARRAY` routine has a state argument `order` with values `MPI_ORDER_C` and `MPI_ORDER_FORTRAN`.

2.5.4 Named Constants

MPI procedures sometimes assign a special meaning to a special value of a basic type argument; e.g., `tag` is an integer-valued argument of point-to-point communication operations, with a special wild-card value, `MPI_ANY_TAG`. Such arguments will have a range of regular values, which is a proper subrange of the range of values of the corresponding basic type; special values (such as `MPI_ANY_TAG`) will be outside the regular range. The range of regular values, such as `tag`, can be queried using environmental inquiry functions (Chapter 7 of the MPI-1 document). The range of other values, such as `source`, depends on values given by other MPI routines (in the case of `source` it is the communicator size).

MPI also provides predefined named constant handles, such as `MPI_COMM_WORLD`.

All named constants, with the exceptions noted below for Fortran, can be used in initialization expressions or assignments. These constants do not change values during execution. Opaque objects accessed by constant handles are defined and do not change value between MPI initialization (`MPI_INIT`) and MPI completion (`MPI_FINALIZE`).

The constants that cannot be used in initialization expressions or assignments in Fortran are:

```
MPI_BOTTOM
MPI_STATUS_IGNORE
MPI_STATUSES_IGNORE
MPI_ERRCODES_IGNORE
```



```

MPI_IN_PLACE
MPI_ARGV_NULL
MPI_ARGVS_NULL

```

Advice to implementors. In Fortran the implementation of these special constants may require the use of language constructs that are outside the Fortran standard. Using special values for the constants (e.g., by defining them through `parameter` statements) is not possible because an implementation cannot distinguish these values from legal data. Typically, these constants are implemented as predefined static variables (e.g., a variable in an MPI-declared `COMMON` block), relying on the fact that the target compiler passes data by address. Inside the subroutine, this address can be extracted by some mechanism outside the Fortran standard (e.g., by Fortran extensions or by implementing the function in C). (*End of advice to implementors.*)

2.5.5 Choice

MPI functions sometimes use arguments with a *choice* (or union) data type. Distinct calls to the same routine may pass by reference actual arguments of different types. The mechanism for providing such arguments will differ from language to language. For Fortran, the document uses `<type>` to represent a choice variable; for C and C++, we use `void *`.

2.5.6 Addresses

Some MPI procedures use *address* arguments that represent an absolute address in the calling program. The datatype of such an argument is `MPI_Aint` in C, `MPI::Aint` in C++ and `INTEGER (KIND=MPI_ADDRESS_KIND)` in Fortran. There is the MPI constant `MPI_BOTTOM` to indicate the start of the address range.

2.5.7 File Offsets

For I/O there is a need to give the size, displacement, and offset into a file. These quantities can easily be larger than 32 bits which can be the default size of a Fortran integer. To overcome this, these quantities are declared to be `INTEGER (KIND=MPI_OFFSET_KIND)` in Fortran. In C one uses `MPI_Offset` whereas in C++ one uses `MPI::Offset`.

2.6 Language Binding

This section defines the rules for MPI language binding in general and for Fortran, ANSI C, and C++, in particular. (Note that ANSI C has been replaced by ISO C. References in MPI to ANSI C now mean ISO C.) Defined here are various object representations, as well as the naming conventions used for expressing this standard. The actual calling sequences are defined elsewhere.

MPI bindings are for Fortran 90, though they are designed to be usable in Fortran 77 environments.

Since the word `PARAMETER` is a keyword in the Fortran language, we use the word “argument” to denote the arguments to a subroutine. These are normally referred to as parameters in C and C++, however, we expect that C and C++ programmers will understand the word “argument” (which has no specific meaning in C/C++), thus allowing us to avoid unnecessary confusion for Fortran programmers.

1 Since Fortran is case insensitive, linkers may use either lower case or upper case when
2 resolving Fortran names. Users of case sensitive languages should avoid the “mpi_” and
3 “pmpi_” prefixes.

4 5 2.6.1 Deprecated Names and Functions

6 A number of chapters refer to deprecated or replaced **MPI-1** constructs. These are constructs
7 that continue to be part of the **MPI** standard, but that users are recommended not to
8 continue using, since **MPI-2** provides better solutions. For example, the Fortran binding for
9 **MPI-1** functions that have address arguments uses **INTEGER**. This is not consistent with the
10 C binding, and causes problems on machines with 32 bit **INTEGERs** and 64 bit addresses.
11 In **MPI-2**, these functions have new names, and new bindings for the address arguments.
12 The use of the old functions is deprecated. For consistency, here and a few other cases,
13 new C functions are also provided, even though the new functions are equivalent to the
14 old functions. The old names are deprecated. Another example is provided by the **MPI-1**
15 predefined datatypes **MPI_UB** and **MPI_LB**. They are deprecated, since their use is awkward
16 and error-prone, while the **MPI-2** function **MPI_TYPE_CREATE_RESIZED** provides a more
17 convenient mechanism to achieve the same effect.

18 The following is a list of all of the deprecated constructs. Note that the constants
19 **MPI_LB** and **MPI_UB** are replaced by the function **MPI_TYPE_CREATE_RESIZED**; this is
20 because their principle use was as input datatypes to **MPI_TYPE_STRUCT** to create resized
21 datatypes. Also note that some C typedefs and Fortran subroutine names are included in
22 this list; they are the types of callback functions.
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Deprecated	MPI-2 Replacement
MPI_ADDRESS	MPI_GET_ADDRESS
MPI_TYPE_HINDEXED	MPI_TYPE_CREATE_HINDEXED
MPI_TYPE_HVECTOR	MPI_TYPE_CREATE_HVECTOR
MPI_TYPE_STRUCT	MPI_TYPE_CREATE_STRUCT
MPI_TYPE_EXTENT	MPI_TYPE_GET_EXTENT
MPI_TYPE_UB	MPI_TYPE_GET_EXTENT
MPI_TYPE_LB	MPI_TYPE_GET_EXTENT
MPI_LB	MPI_TYPE_CREATE_RESIZED
MPI_UB	MPI_TYPE_CREATE_RESIZED
MPI_ERRHANDLER_CREATE	MPI_COMM_CREATE_ERRHANDLER
MPI_ERRHANDLER_GET	MPI_COMM_GET_ERRHANDLER
MPI_ERRHANDLER_SET	MPI_COMM_SET_ERRHANDLER
MPI_Handler_function	MPI_Comm_errhandler_fn
MPI_KEYVAL_CREATE	MPI_COMM_CREATE_KEYVAL
MPI_KEYVAL_FREE	MPI_COMM_FREE_KEYVAL
MPI_DUP_FN	MPI_COMM_DUP_FN
MPI_NULL_COPY_FN	MPI_COMM_NULL_COPY_FN
MPI_NULL_DELETE_FN	MPI_COMM_NULL_DELETE_FN
MPI_Copy_function	MPI_Comm_copy_attr_function
COPY_FUNCTION	COMM_COPY_ATTR_FN
MPI_Delete_function	MPI_Comm_delete_attr_function
DELETE_FUNCTION	COMM_DELETE_ATTR_FN
MPI_ATTR_DELETE	MPI_COMM_DELETE_ATTR
MPI_ATTR_GET	MPI_COMM_GET_ATTR
MPI_ATTR_PUT	MPI_COMM_SET_ATTR

2.6.2 Fortran Binding Issues

MPI-1.1 provided bindings for Fortran 77. MPI-2 retains these bindings but they are now interpreted in the context of the Fortran 90 standard. MPI can still be used with most Fortran 77 compilers, as noted below. When the term Fortran is used it means Fortran 90.

All MPI names have an `MPI_` prefix, and all characters are capitals. Programs must not declare variables, parameters, or functions with names beginning with the prefix `MPI_`. To avoid conflicting with the profiling interface, programs should also avoid functions with the prefix `PMPI_`. This is mandated to avoid possible name collisions.

All MPI Fortran subroutines have a return code in the last argument. A few MPI operations which are functions do not have the return code argument. The return code value for successful completion is `MPI_SUCCESS`. Other error codes are implementation dependent; see the error codes in Chapter 7 of the MPI-1 document and Annex A in the MPI-2 document.

Constants representing the maximum length of a string are one smaller in Fortran than in C and C++ as discussed in Section 4.12.9.

Handles are represented in Fortran as `INTEGER`s. Binary-valued variables are of type `LOGICAL`.

Array arguments are indexed from one.

The MPI Fortran binding is inconsistent with the Fortran 90 standard in several respects. These inconsistencies, such as register optimization problems, have implications for user codes that are discussed in detail in Section 10.2.2. They are also inconsistent with

Fortran 77.

- An MPI subroutine with a choice argument may be called with different argument types.
- An MPI subroutine with an assumed-size dummy argument may be passed an actual scalar argument.
- Many MPI routines assume that actual arguments are passed by address and that arguments are not copied on entrance to or exit from the subroutine.
- An MPI implementation may read or modify user data (e.g., communication buffers used by nonblocking communications) concurrently with a user program executing outside MPI calls.
- Several named “constants,” such as `MPI_BOTTOM`, `MPI_STATUS_IGNORE`, and `MPI_ERRCODES_IGNORE`, are not ordinary Fortran constants and require a special implementation. See Section 2.5.4 on page 10 for more information.

Additionally, MPI is inconsistent with Fortran 77 in a number of ways, as noted below.

- MPI identifiers exceed 6 characters.
- MPI identifiers may contain underscores after the first character.
- MPI requires an include file, `mpif.h`. On systems that do not support include files, the implementation should specify the values of named constants.
- Many routines in MPI-2 have KIND-parameterized integers (e.g., `MPI_ADDRESS_KIND` and `MPI_OFFSET_KIND`) that hold address information. On systems that do not support Fortran 90-style parameterized types, `INTEGER*8` or `INTEGER` should be used instead.
- The memory allocation routine `MPI_ALLOC_MEM` can’t be usefully used in Fortran without a language extension that allows the allocated memory to be associated with a Fortran variable.

2.6.3 C Binding Issues

We use the ANSI C declaration format. All MPI names have an `MPI_` prefix, defined constants are in all capital letters, and defined types and functions have one capital letter after the prefix. Programs must not declare variables or functions with names beginning with the prefix `MPI_`. To support the profiling interface, programs should not declare functions with names beginning with the prefix `PMPI_`.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

Almost all C functions return an error code. The successful return code will be `MPI_SUCCESS`, but failure return codes are implementation dependent.

Type declarations are provided for handles to each category of opaque objects.

Array arguments are indexed from zero.

Logical flags are integers with value 0 meaning “false” and a non-zero value meaning “true.”

Choice arguments are pointers of type `void *`.

Address arguments are of MPI defined type `MPI_Aint`. File displacements are of type `MPI_Offset`. `MPI_Aint` is defined to be an integer of the size needed to hold any valid address on the target architecture. `MPI_Offset` is defined to be an integer of the size needed to hold any valid file size on the target architecture.

2.6.4 C++ Binding Issues

There are places in the standard that give rules for C and not for C++. In these cases, the C rule should be applied to the C++ case, as appropriate. In particular, the values of constants given in the text are the ones for C and Fortran. A cross index of these with the C++ names is given in Annex A.

We use the ANSI C++ declaration format. All MPI names are declared within the scope of a namespace called `MPI` and therefore are referenced with an `MPI::` prefix. Defined constants are in all capital letters, and class names, defined types, and functions have only their first letter capitalized. Programs must not declare variables or functions in the `MPI` namespace. This is mandated to avoid possible name collisions.

The definition of named constants, function prototypes, and type definitions must be supplied in an include file `mpi.h`.

Advice to implementors. The file `mpi.h` may contain both the C and C++ definitions. Usually one can simply use the defined value (generally `_cplusplus`, but not required) to see if one is using C++ to protect the C++ definitions. It is possible that a C compiler will require that the source protected this way be legal C code. In this case, all the C++ definitions can be placed in a different include file and the “`#include`” directive can be used to include the necessary C++ definitions in the `mpi.h` file. (*End of advice to implementors.*)

C++ functions that create objects or return information usually place the object or information in the return value. Since the language neutral prototypes of MPI functions include the C++ return value as an OUT parameter, semantic descriptions of MPI functions refer to the C++ return value by that parameter name (see Section B.13.5 on page 356). The remaining C++ functions return `void`.

In some circumstances, MPI permits users to indicate that they do not want a return value. For example, the user may indicate that the status is not filled in. Unlike C and Fortran where this is achieved through a special input value, in C++ this is done by having two bindings where one has the optional argument and one does not.

C++ functions do not return error codes. If the default error handler has been set to `MPI::ERRORS_THROW_EXCEPTIONS`, the C++ exception mechanism is used to signal an error by throwing an `MPI::Exception` object.

It should be noted that the default error handler (i.e., `MPI::ERRORS_ARE_FATAL`) on a given type has not changed. User error handlers are also permitted. `MPI::ERRORS_RETURN` simply returns control to the calling function; there is no provision for the user to retrieve the error code.

User callback functions that return integer error codes should not throw exceptions; the returned error will be handled by the MPI implementation by invoking the appropriate error handler.

- 1
- 2
- 3
- 4

6
7

8

9

10

11
12
1314
15
16

17
18
19
20
21
22
23
24
25
26

27
28

29

30

32
33
34
35
36
37
38
39
40
41

43
44
45

47

The C++ bindings presented in Annex B and throughout this document were generated by applying a simple set of name generation rules to the MPI function specifications. While these guidelines may be sufficient in most cases, they may not be suitable for all situations. In cases of ambiguity or where a specific semantic statement is desired, these guidelines may be superseded as the situation dictates.

1. All functions, types, and constants are declared within the scope of a `namespace` called `MPI`.
2. Arrays of MPI handles are always left in the argument list (whether they are IN or OUT arguments).
3. If the argument list of an MPI function contains a scalar IN handle, and it makes sense to define the function as a method of the object corresponding to that handle, the function is made a member function of the corresponding MPI class. The member functions are named according to the corresponding MPI function name, but without the “MPI_” prefix and without the object name prefix (if applicable). In addition:
 - (a) The scalar IN handle is dropped from the argument list, and `this` corresponds to the dropped argument.
 - (b) The function is declared `const`.
4. MPI functions are made into class functions (static) when they belong on a class but do not have a unique scalar IN or INOUT parameter of that class.
5. If the argument list contains a single OUT argument that is not of type `MPI_STATUS` (or an array), that argument is dropped from the list and the function returns that value.

Example 2.1 The C++ binding for `MPI_COMM_SIZE` is `int MPI::Comm::Get_size(void) const`.

6. If there are multiple OUT arguments in the argument list, one is chosen as the return value and is removed from the list.
7. If the argument list does not contain any OUT arguments, the function returns `void`.

Example 2.2 The C++ binding for `MPI_REQUEST_FREE` is `void MPI::Request::Free(void)`

8. MPI functions to which the above rules do not apply are not members of any class, but are defined in the `MPI` namespace.

Example 2.3 The C++ binding for `MPI_BUFFER_ATTACH` is `void MPI::Attach_buffer(void* buffer, int size)`.

9. All class names, defined types, and function names have only their first letter capitalized. Defined constants are in all capital letters.

10. Any IN pointer, reference, or array argument must be declared `const`.
11. Handles are passed by reference.
12. Array arguments are denoted with square brackets (`[]`), not pointers, as this is more semantically precise.

2.7 Processes

An MPI program consists of autonomous processes, executing their own code, in a MIMD style. The codes executed by each process need not be identical. The processes communicate via calls to MPI communication primitives. Typically, each process executes in its own address space, although shared-memory implementations of MPI are possible.

This document specifies the behavior of a parallel program assuming that only MPI calls are used. The interaction of an MPI program with other possible means of communication, I/O, and process management is not specified. Unless otherwise stated in the specification of the standard, MPI places no requirements on the result of its interaction with external mechanisms that provide similar or equivalent functionality. This includes, but is not limited to, interactions with external mechanisms for process control, shared and remote memory access, file system access and control, interprocess communication, process signaling, and terminal I/O. High quality implementations should strive to make the results of such interactions intuitive to users, and attempt to document restrictions where deemed necessary.

Advice to implementors. Implementations that support such additional mechanisms for functionality supported within MPI are expected to document how these interact with MPI. (*End of advice to implementors.*)

The interaction of MPI and threads is defined in Section 8.7.

2.8 Error Handling

MPI provides the user with reliable message transmission. A message sent is always received correctly, and the user does not need to check for transmission errors, time-outs, or other error conditions. In other words, MPI does not provide mechanisms for dealing with failures in the communication system. If the MPI implementation is built on an unreliable underlying mechanism, then it is the job of the implementor of the MPI subsystem to insulate the user from this unreliability, or to reflect unrecoverable errors as failures. Whenever possible, such failures will be reflected as errors in the relevant communication call. Similarly, MPI itself provides no mechanisms for handling processor failures.

Of course, MPI programs may still be erroneous. A **program error** can occur when an MPI call is made with an incorrect argument (non-existing destination in a send operation, buffer too small in a receive operation, etc.). This type of error would occur in any implementation. In addition, a **resource error** may occur when a program exceeds the amount of available system resources (number of pending messages, system buffers, etc.). The occurrence of this type of error depends on the amount of available resources in the system and the resource allocation mechanism used; this may differ from system to system. A high-quality implementation will provide generous limits on the important resources so as to alleviate the portability problem this represents.

In C and Fortran, almost all MPI calls return a code that indicates successful completion of the operation. Whenever possible, MPI calls return an error code if an error occurred during the call. By default, an error detected during the execution of the MPI library causes the parallel computation to abort, except for file operations. However, MPI provides mechanisms for users to change this default and to handle recoverable errors. The user may specify that no error is fatal, and handle error codes returned by MPI calls by himself or herself. Also, the user may provide his or her own error-handling routines, which will be invoked whenever an MPI call returns abnormally. The MPI error handling facilities are described in Chapter 7 of the MPI-1 document and in Section 4.13 of this document. The return values of C++ functions are not error codes. If the default error handler has been set to `MPI::ERRORS_THROW_EXCEPTIONS`, the C++ exception mechanism is used to signal an error by throwing an `MPI::Exception` object.

Several factors limit the ability of MPI calls to return with meaningful error codes when an error occurs. MPI may not be able to detect some errors; other errors may be too expensive to detect in normal execution mode; finally some errors may be “catastrophic” and may prevent MPI from returning control to the caller in a consistent state.

Another subtle issue arises because of the nature of asynchronous communications: MPI calls may initiate operations that continue asynchronously after the call returned. Thus, the operation may return with a code indicating successful completion, yet later cause an error exception to be raised. If there is a subsequent call that relates to the same operation (e.g., a call that verifies that an asynchronous operation has completed) then the error argument associated with this call will be used to indicate the nature of the error. In a few cases, the error may occur after all calls that relate to the operation have completed, so that no error value can be used to indicate the nature of the error (e.g., an error on the receiver in a send with the ready mode). Such an error must be treated as fatal, since information cannot be returned for the user to recover from it.

This document does not specify the state of a computation after an erroneous MPI call has occurred. The desired behavior is that a relevant error code be returned, and the effect of the error be localized to the greatest possible extent. E.g., it is highly desirable that an erroneous receive call will not cause any part of the receiver’s memory to be overwritten, beyond the area specified for receiving the message.

Implementations may go beyond this document in supporting in a meaningful manner MPI calls that are defined here to be erroneous. For example, MPI specifies strict type matching rules between matching send and receive operations: it is erroneous to send a floating point variable and receive an integer. Implementations may go beyond these type matching rules, and provide automatic type conversion in such situations. It will be helpful to generate warnings for such non-conforming behavior.

MPI-2 defines a way for users to create new error codes as defined in Section 8.5.

2.9 Implementation Issues

There are a number of areas where an MPI implementation may interact with the operating environment and system. While MPI does not mandate that any services (such as signal handling) be provided, it does strongly suggest the behavior to be provided if those services are available. This is an important point in achieving portability across platforms that provide the same set of services.

2.9.1 Independence of Basic Runtime Routines

MPI programs require that library routines that are part of the basic language environment (such as `write` in Fortran and `printf` and `malloc` in ANSI C) and are executed after `MPI_INIT` and before `MPI_FINALIZE` operate independently and that their *completion* is independent of the action of other processes in an MPI program.

Note that this in no way prevents the creation of library routines that provide parallel services whose operation is collective. However, the following program is expected to complete in an ANSI C environment regardless of the size of `MPI_COMM_WORLD` (assuming that `printf` is available at the executing nodes).

```
int rank;
MPI_Init((void *)0, (void *)0);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
if (rank == 0) printf("Starting program\n");
MPI_Finalize();
```

The corresponding Fortran and C++ programs are also expected to complete.

An example of what is *not* required is any particular ordering of the action of these routines when called by several tasks. For example, MPI makes neither requirements nor recommendations for the output from the following program (again assuming that I/O is available at the executing nodes).

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
printf("Output from task rank %d\n", rank);
```

In addition, calls that fail because of resource exhaustion or other error are not considered a violation of the requirements here (however, they are required to complete, just not to complete successfully).

2.9.2 Interaction with Signals

MPI does not specify the interaction of processes with signals and does not require that MPI be signal safe. The implementation may reserve some signals for its own use. It is required that the implementation document which signals it uses, and it is strongly recommended that it not use `SIGALRM`, `SIGFPE`, or `SIGIO`. Implementations may also prohibit the use of MPI calls from within signal handlers.

In multithreaded environments, users can avoid conflicts between signals and the MPI library by catching signals only on threads that do not execute MPI calls. High quality single-threaded implementations will be signal safe: an MPI call suspended by a signal will resume and complete normally after the signal is handled.

2.10 Examples

The examples in this document are for illustration purposes only. They are not intended to specify the standard. Furthermore, the examples have not been carefully checked or verified.

Chapter 3

Version 1.2 of MPI

This section contains clarifications and minor corrections to Version 1.1 of the MPI Standard. The only new function in MPI-1.2 is one for identifying which version of the MPI Standard the implementation being used conforms to. There are small differences between MPI-1 and MPI-1.1. There are very few differences (only those discussed in this chapter) between MPI-1.1 and MPI-1.2, but large differences (the rest of this document) between MPI-1.2 and MPI-2.

3.1 Version Number

In order to cope with changes to the MPI Standard, there are both compile-time and run-time ways to determine which version of the standard is in use in the environment one is using.

The “version” will be represented by two separate integers, for the version and subversion:

In C and C++,

```
#define MPI_VERSION    1
#define MPI_SUBVERSION 2
```

in Fortran,

```
INTEGER MPI_VERSION, MPI_SUBVERSION
PARAMETER (MPI_VERSION    = 1)
PARAMETER (MPI_SUBVERSION = 2)
```

For runtime determination,

MPI_GET_VERSION(version, subversion)

OUT	version	version number (integer)
OUT	subversion	subversion number (integer)

```
int MPI_Get_version(int *version, int *subversion)
```

```
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)
    INTEGER VERSION, SUBVERSION, IERROR
```

`MPI_GET_VERSION` is one of the few functions that can be called before `MPI_INIT` and after `MPI_FINALIZE`. Its C++ binding can be found in the Annex, Section B.11.

3.2 MPI-1.0 and MPI-1.1 Clarifications

As experience has been gained since the releases of the 1.0 and 1.1 versions of the MPI Standard, it has become apparent that some specifications were insufficiently clear. In this section we attempt to make clear the intentions of the MPI Forum with regard to the behavior of several MPI-1 functions. An MPI-1-compliant implementation should behave in accordance with the clarifications in this section.

3.2.1 Clarification of `MPI_INITIALIZED`

`MPI_INITIALIZED` returns `true` if the calling process has called `MPI_INIT`. Whether `MPI_FINALIZE` has been called does not affect the behavior of `MPI_INITIALIZED`.

3.2.2 Clarification of `MPI_FINALIZE`

This routine cleans up all MPI state. Each process must call `MPI_FINALIZE` before it exits. Unless there has been a call to `MPI_ABORT`, each process must ensure that all pending non-blocking communications are (locally) complete before calling `MPI_FINALIZE`. Further, at the instant at which the last process calls `MPI_FINALIZE`, all pending sends must be matched by a receive, and all pending receives must be matched by a send.

For example, the following program is correct:

Process 0	Process 1
-----	-----
<code>MPI_Init();</code>	<code>MPI_Init();</code>
<code>MPI_Send(dest=1);</code>	<code>MPI_Recv(src=0);</code>
<code>MPI_Finalize();</code>	<code>MPI_Finalize();</code>

Without the matching receive, the program is erroneous:

Process 0	Process 1
-----	-----
<code>MPI_Init();</code>	<code>MPI_Init();</code>
<code>MPI_Send (dest=1);</code>	
<code>MPI_Finalize();</code>	<code>MPI_Finalize();</code>

A successful return from a blocking communication operation or from `MPI_WAIT` or `MPI_TEST` tells the user that the buffer can be reused and means that the communication is completed by the user, but does not guarantee that the local process has no more work to do. A successful return from `MPI_REQUEST_FREE` with a request handle generated by an `MPI_ISEND` nullifies the handle but provides no assurance of operation completion. The `MPI_ISEND` is complete only when it is known by some means that a matching receive has completed. `MPI_FINALIZE` guarantees that all local actions required by communications the user has completed will, in fact, occur before it returns.

`MPI_FINALIZE` guarantees nothing about pending communications that have not been completed (completion is assured only by `MPI_WAIT`, `MPI_TEST`, or `MPI_REQUEST_FREE` combined with some other verification of completion).

Example 3.1 This program is correct:

```

rank 0                                rank 1
=====
...                                  ...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Barrier();
MPI_Barrier();                       MPI_Finalize();
MPI_Finalize();                      exit();
exit();

```

Example 3.2 This program is erroneous and its behavior is undefined:

```

rank 0                                rank 1
=====
...                                  ...
MPI_Isend();                          MPI_Recv();
MPI_Request_free();                   MPI_Finalize();
MPI_Finalize();                      exit();
exit();

```

If no `MPI_BUFFER_DETACH` occurs between an `MPI_BSEND` (or other buffered send) and `MPI_FINALIZE`, the `MPI_FINALIZE` implicitly supplies the `MPI_BUFFER_DETACH`.

Example 3.3 This program is correct, and after the `MPI_Finalize`, it is as if the buffer had been detached.

```

rank 0                                rank 1
=====
...                                  ...
buffer = malloc(1000000);             MPI_Recv();
MPI_Buffer_attach();                  MPI_Finalize();
MPI_Bsend();                          exit();
MPI_Finalize();
free(buffer);
exit();

```

Example 3.4 In this example, `MPI_Probe()` must return a `FALSE` flag. `MPI_Test_cancelled()` must return a `TRUE` flag, independent of the relative order of execution of `MPI_Cancel()` in process 0 and `MPI_Finalize()` in process 1.

The `MPI_Probe()` call is there to make sure the implementation knows that the “tag1” message exists at the destination, without being able to claim that the user knows about it.

```

rank 0                                rank 1
=====
MPI_Init();                          MPI_Init();
MPI_Isend(tag1);                     MPI_Barrier();
MPI_Barrier();

```

```

1          MPI_Iprobe(tag2);
2  MPI_Barrier();          MPI_Barrier();
3                          MPI_Finalize();
4                          exit();
5  MPI_Cancel();
6  MPI_Wait();
7  MPI_Test_cancelled();
8  MPI_Finalize();
9  exit();

```

Advice to implementors. An implementation may need to delay the return from MPI_FINALIZE until all potential future message cancellations have been processed. One possible solution is to place a barrier inside MPI_FINALIZE (*End of advice to implementors.*)

Once MPI_FINALIZE returns, no MPI routine (not even MPI_INIT) may be called, except for MPI_GET_VERSION, MPI_INITIALIZED, and the MPI-2 function MPI_FINALIZED. Each process must complete any pending communication it initiated before it calls MPI_FINALIZE. If the call returns, each process may continue local computations, or exit, without participating in further MPI communication with other processes. MPI_FINALIZE is collective on MPI_COMM_WORLD.

Advice to implementors. Even though a process has completed all the communication it initiated, such communication may not yet be completed from the viewpoint of the underlying MPI system. E.g., a blocking send may have completed, even though the data is still buffered at the sender. The MPI implementation must ensure that a process has completed any involvement in MPI communication before MPI_FINALIZE returns. Thus, if a process exits after the call to MPI_FINALIZE, this will not cause an ongoing communication to fail. (*End of advice to implementors.*)

Although it is not required that all processes return from MPI_FINALIZE, it is required that at least process 0 in MPI_COMM_WORLD return, so that users can know that the MPI portion of the computation is over. In addition, in a POSIX environment, they may desire to supply an exit code for each process that returns from MPI_FINALIZE.

Example 3.5 The following illustrates the use of requiring that at least one process return and that it be known that process 0 is one of the processes that return. One wants code like the following to work no matter how many processes return.

```

39  ...
40  MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
41  ...
42  MPI_Finalize();
43  if (myrank == 0) {
44      resultfile = fopen("outfile","w");
45      dump_results(resultfile);
46      fclose(resultfile);
47  }
48  exit(0);

```

3.2.3 Clarification of status after MPI_WAIT and MPI_TEST

The fields in a **status** object returned by a call to `MPI_WAIT`, `MPI_TEST`, or any of the other derived functions (`MPI_{TEST, WAIT}{ALL, SOME, ANY}`), where the **request** corresponds to a send call, are undefined, with two exceptions: The error status field will contain valid information if the wait or test call returned with `MPI_ERR_IN_STATUS`; and the returned status can be queried by the call `MPI_TEST_CANCELLED`.

Error codes belonging to the error class `MPI_ERR_IN_STATUS` should be returned only by the MPI completion functions that take arrays of `MPI_STATUS`. For the functions (`MPI_TEST`, `MPI_TESTANY`, `MPI_WAIT`, `MPI_WAITANY`) that return a single `MPI_STATUS` value, the normal MPI error return process should be used (not the `MPI_ERROR` field in the `MPI_STATUS` argument).

3.2.4 Clarification of MPI_INTERCOMM_CREATE

The Problem: The MPI-1.1 standard says, in the discussion of `MPI_INTERCOMM_CREATE`, both that

The groups must be disjoint

and that

The leaders may be the same process.

To further muddy the waters, the reason given for “The groups must be disjoint” is based on concerns about the implementation of `MPI_INTERCOMM_CREATE` that are not applicable for the case where the leaders are the same process.

The Fix: Delete the text:

(the two leaders could be the same process)

from the discussion of `MPI_INTERCOMM_CREATE`.

Replace the text:

All inter-communicator constructors are blocking and require that the local and remote groups be disjoint in order to avoid deadlock.

with

All inter-communicator constructors are blocking and require that the local and remote groups be disjoint.

Advice to users. The groups must be disjoint for several reasons. Primarily, this is the intent of the intercommunicators — to provide a communicator for communication between disjoint groups. This is reflected in the definition of `MPI_INTERCOMM_MERGE`, which allows the user to control the ranking of the processes in the created intracommunicator; this ranking makes little sense if the groups are not disjoint. In addition, the natural extension of collective operations to intercommunicators makes the most sense when the groups are disjoint. (*End of advice to users.*)

3.2.5 Clarification of MPI_INTERCOMM_MERGE

The error handler on the new intercommunicator in each process is inherited from the communicator that contributes the local group. Note that this can result in different processes in the same communicator having different error handlers.

3.2.6 Clarification of Binding of MPI_TYPE_SIZE

This clarification is needed in the MPI-1 description of MPI_TYPE_SIZE, since the issue repeatedly arises. It is a clarification of the binding.

Advice to users. The MPI-1 Standard specifies that the output argument of MPI_TYPE_SIZE in C is of type `int`. The MPI Forum considered proposals to change this and decided to reiterate the original decision. (*End of advice to users.*)

3.2.7 Clarification of MPI_REDUCE

The current text on p. 115, lines 25–28, from MPI-1.1 (June 12, 1995) says:

The **datatype** argument of MPI_REDUCE must be compatible with **op**. Predefined operators work only with the MPI types listed in Section 4.9.2 and Section 4.9.3. User-defined operators may operate on general, derived datatypes.

This text is changed to:

The **datatype** argument of MPI_REDUCE must be compatible with **op**. Predefined operators work only with the MPI types listed in Section 4.9.2 and Section 4.9.3. Furthermore, the **datatype** and **op** given for predefined operators must be the same on all processes.

Note that it is possible for users to supply different user-defined operations to MPI_REDUCE in each process. MPI does not define which operations are used on which operands in this case.

Advice to users. Users should make no assumptions about how MPI_REDUCE is implemented. Safest is to ensure that the same function is passed to MPI_REDUCE by each process. (*End of advice to users.*)

Overlapping datatypes are permitted in “send” buffers. Overlapping datatypes in “receive” buffers are erroneous and may give unpredictable results.

3.2.8 Clarification of Error Behavior of Attribute Callback Functions

If an attribute copy function or attribute delete function returns other than MPI_SUCCESS, then the call that caused it to be invoked (for example, MPI_COMM_FREE), is erroneous.

3.2.9 Clarification of MPI_PROBE and MPI_IREQ

Page 52, lines 1 thru 3 (of MPI-1.1, the June 12, 1995 version without changebars) become:

“A subsequent receive executed with the same communicator, and the source and tag returned in status by MPI_IREQ will receive the message that was matched by the probe, if no other intervening receive occurs after the probe, and the send is not successfully cancelled before the receive.”

Rationale.

The following program shows that the MPI-1 definitions of cancel and probe are in conflict:

Process 0	Process 1
-----	-----
MPI_Init();	MPI_Init();
MPI_Isend(dest=1);	
	MPI_Probe();
MPI_Barrier();	MPI_Barrier();
MPI_Cancel();	
MPI_Wait();	
MPI_Test_cancelled();	
MPI_Barrier();	MPI_Barrier();
	MPI_Recv();

Since the send has been cancelled by process 0, the wait must be local (page 54, line 13) and must return before the matching receive. For the wait to be local, the send must be successfully cancelled, and therefore must not match the receive in process 1 (page 54 line 29).

However, it is clear that the probe on process 1 must eventually detect an incoming message. Page 52 line 1 makes it clear that the subsequent receive by process 1 must return the probed message.

The above are clearly contradictory, and therefore the text "...and the send is not successfully cancelled before the receive" must be added to line 3 of page 54.

An alternative solution (rejected) would be to change the semantics of cancel so that the call is not local if the message has been probed. This adds complexity to implementations, and adds a new concept of "state" to a message (probed or not). It would, however, preserve the feature that a blocking receive after a probe is local.

(End of rationale.)

3.2.10 Minor Corrections

The following corrections to MPI-1.1 are (all page and line numbers are for the June 12, 1995 version without changebars):

- Page 11, line 36 reads
MPI_ADDRESS
but should read
MPI_ADDRESS_TYPE
- Page 19, lines 1–2 reads
for (64 bit) C integers declared to be of type **longlong int**
but should read
for C integers declared to be of type **long long**

- Page 40, line 48 should have the following text added:

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 10.2.2 of the MPI-2 Standard, pages 286 and 289. (*End of advice to users.*)

- Page 41, lines 16–18 reads

A **empty** status is a status which is set to return `tag = MPI_ANY_TAG`, `source = MPI_ANY_SOURCE`, and is also internally configured so that calls to `MPI_GET_COUNT` and `MPI_GET_ELEMENTS` return `count = 0`.

but should read

A **empty** status is a status which is set to return `tag = MPI_ANY_TAG`, `source = MPI_ANY_SOURCE`, `error = MPI_SUCCESS`, and is also internally configured so that calls to `MPI_GET_COUNT` and `MPI_GET_ELEMENTS` return `count = 0` and `MPI_TEST_CANCELLED` returns false.

- Page 52, lines 46–48 read

```
100      CALL MPI_RECV(i, 1, MPI_INTEGER, 0, 0, status, ierr)
      ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, 1, 0, status, ierr)
```

but should read

```
100      CALL MPI_RECV(i, 1, MPI_INTEGER, 0, 0, comm, status, ierr)
      ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, 1, 0, comm, status, ierr)
```

- Page 53, lines 18–23 read

```
100      CALL MPI_RECV(i, 1, MPI_INTEGER, MPI_ANY_SOURCE,
                      0, status, ierr)
      ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, MPI_ANY_SOURCE,
                      0, status, ierr)
```

but should read

```
100      CALL MPI_RECV(i, 1, MPI_INTEGER, MPI_ANY_SOURCE,
                      0, comm, status, ierr)
      ELSE
200      CALL MPI_RECV(x, 1, MPI_REAL, MPI_ANY_SOURCE,
                      0, comm, status, ierr)
```

- Page 59, line 3 should have the following text added:

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 10.2.2 of the MPI-2 Standard, pages 286 and 289. (*End of advice to users.*)

- Page 59, lines 42–45 read

```
int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
                 int dest, int sendtag, void *recvbuf, int recvcount,
                 MPI_Datatype recvtype, int source, MPI_Datatype recvtag,
                 MPI_Comm comm, MPI_Status *status)
```

but should read

```
int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
                 int dest, int sendtag, void *recvbuf, int recvcount,
                 MPI_Datatype recvtype, int source, int recvtag,
                 MPI_Comm comm, MPI_Status *status)
```

- Page 60, line 3 reads

```
SOURCE, RECV TAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
```

but should read

```
SOURCE, RECVTAG, COMM, STATUS(MPI_STATUS_SIZE), IERROR
```

- Page 70, line 16 should have the following text added:

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 10.2.2 of the MPI-2 Standard, pages 286 and 289. (*End of advice to users.*)

- Page 71, line 10 reads

and do not affect the the content of a message

but should read

and do not affect the content of a message

- Page 74, lines 39–45 read

A datatype may specify overlapping entries. The use of such a datatype in a receive operation is erroneous. (This is erroneous even if the actual message received is short enough not to write any entry more than once.)

A datatype may specify overlapping entries. If such a datatype is used in a receive operation, that is, if some part of the receive buffer is written more than once by the receive operation, then the call is erroneous.

The first part was an MPI-1.1 addition. The second part overlaps with it. The old text will be removed so it now reads

A datatype may specify overlapping entries. The use of such a datatype in a receive operation is erroneous. (This is erroneous even if the actual message received is short enough not to write any entry more than once.)

- Page 75, line 24 should have the following text added:
The **datatype** argument should match the argument provided by the receive call that set the **status** variable.

- Page 85, line 36 reads
“specified by **outbuf** and **outcount**”
but should read
“specified by **outbuf** and **outsize**.”

- Page 90, line 3 reads
MPI_Pack_size(count, MPI_CHAR, &k2);
but should read
MPI_Pack_size(count, MPI_CHAR, comm, &k2);

- Page 90, line 10 reads
MPI_Pack(chr, count, MPI_CHAR, &lbuf, k, &position, comm);
but should read
MPI_Pack(chr, count, MPI_CHAR, lbuf, k, &position, comm);

- Page 97, line 41 reads

MPI_Recv(recvbuf + disp[i] · extent(recvtype), recvcunts[i], recvtype, i, ...).

but should read

MPI_Recv(recvbuf + displs[i] · extent(recvtype), recvcunts[i], recvtype, i, ...).

- Page 109, lines 26–27 and page 110, lines 28–29 reads
The *j*th block of data sent from each process is received by every process and placed in the *j*th block of the buffer **recvbuf**.
but should read
The block of data sent from the *j*th process is received by every process and placed in the *j*th block of the buffer **recvbuf**.

- Page 117, lines 22–23 reads
MPI provides seven such predefined datatypes.
but should read
MPI provides nine such predefined datatypes.

- Page 121, line 1 reads

```
FUNCTION USER_FUNCTION( INVEC(*), INOUTVEC(*), LEN, TYPE)
```

but should read

```
SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, TYPE)
```

- Page 122, lines 35–36 read

```
MPI_OP_FREE( op)
```

IN	op	operation (handle)
----	----	--------------------

but should read

```
MPI_OP_FREE( op)
```

INOUT	op	operation (handle)
-------	----	--------------------

- Page 125, line 1 reads

```
CALL MPI_ALLREDUCE(sum, c, n, MPI_REAL, MPI_SUM, 0, comm, ierr)
```

but should read

```
CALL MPI_ALLREDUCE(sum, c, n, MPI_REAL, MPI_SUM, comm, ierr)
```

- Page 141, lines 27–27 read

IN	ranges	an array of integer triplets, of the form (first rank, last rank, stride) indicating ranks in group of processes to be included in newgroup
----	--------	---

but should read

IN	ranges	a one-dimensional array of integer triplets, of the form (first rank, last rank, stride) indicating ranks in group of processes to be included in newgroup
----	--------	--

- Page 142, line 10 reads

IN n number of elements in array ranks (integer)

but should read

IN n number of triplets in array *ranges* (integer)

- Page 194, lines 30–31 reads
to the greatest possible, extent,
but should read
to the greatest possible extent,

- Page 194, line 48 reads
MPI_ERRHANDLER_CREATE(FUNCTION, HANDLER, IERROR)
but should read
MPI_ERRHANDLER_CREATE(FUNCTION, ERRHANDLER, IERROR)

- Page 195, line 15 should have the following text added:
In the Fortran language, the user routine should be of the form:

```
SUBROUTINE HANDLER_FUNCTION(COMM, ERROR_CODE, .....)  
INTEGER COMM, ERROR_CODE
```

Advice to users. Users are discouraged from using a Fortran
HANDLER_FUNCTION since the routine expects a variable number of arguments.
Some Fortran systems may allow this but some may fail to give the correct result
or compile/link this code. Thus, it will not, in general, be possible to create
portable code with a Fortran HANDLER_FUNCTION. (*End of advice to users.*)

- Page 196, lines 1–2 reads
MPI_ERRHANDLER_FREE(errhandler)

IN errhandler MPI error handler (handle)

but should read

MPI_ERRHANDLER_FREE(errhandler)

INOUT	errhandler	MPI error handler (handle)	1
			2
			3
•	Page 197, line 25 should have added:		4
			5
	An MPI error class is a valid MPI error code. Specifically, the values defined for MPI error classes are valid MPI error codes.		6
			7
			8
•	Page 201, line 28 reads		9
	...of different language bindings is is done		10
	but should read		11
	...of different language bindings is done		12
			13
•	Page 203, line 1 reads		14
	MPI_PCONTROL(level)		15
			16
	but should read		17
	MPI_PCONTROL(LEVEL)		18
			19
•	Page 210, line 44 reads		20
	MPI_PENDING		21
	but should read		22
	MPI_ERR_PENDING		23
			24
•	Page 211, line 44 reads		25
	MPI_DOUBLE_COMPLEX		26
	but should be moved to Page 212, line 22 since it is an optional Fortran datatype.		27
			28
•	Page 212, add new lines of text at line 22 and line 25 to read:		29
	etc.		30
	Thus, the text will now read:		31
			32
			33
	/* optional datatypes (Fortran) */		34
	MPI_INTEGER1		35
	MPI_INTEGER2		36
	MPI_INTEGER4		37
	MPI_REAL2		38
	MPI_REAL4		39
	MPI_REAL8		40
	etc.		41
			42
	/* optional datatypes (C) */		43
	MPI_LONG_LONG_INT		44
	etc.		45
			46
•	Page 213, line 28. The following text should be added:		47
			48

```

1      /* Predefined functions in C and Fortran */
2      MPI_NULL_COPY_FN
3      MPI_NULL_DELETE_FN
4      MPI_DUP_FN

```

- Page 213, line 41. Add the line

```

8      MPI_Errhandler

```

- Page 214, line 9 reads

```

12     FUNCTION USER_FUNCTION( INVEC(*), INOUTVEC(*), LEN, TYPE)

```

but should read

```

18     SUBROUTINE USER_FUNCTION( INVEC, INOUTVEC, LEN, TYPE)

```

- Page 214, lines 14 and 15 read

```

23     PROCEDURE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE,
24                             ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERR)

```

but should read

```

29     SUBROUTINE COPY_FUNCTION(OLDCOMM, KEYVAL, EXTRA_STATE,
30                             ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERR)

```

- Page 214, line 21 reads

```

35     PROCEDURE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)

```

but should read

```

40     SUBROUTINE DELETE_FUNCTION(COMM, KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE, IERR)

```

- Page 214, line 23 should have the following text added:
The handler-function for error handlers should be declared like this:

```

45     SUBROUTINE HANDLER_FUNCTION(COMM, ERROR_CODE, ....)
46     INTEGER COMM, ERROR_CODE

```


- Page 216, lines 4–7 read


```
int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
                  int dest, int sendtag, void *recvbuf, int recvcount,
                  MPI_Datatype recvtype, int source, MPI_Datatype recvtag,
                  MPI_Comm comm, MPI_Status *status)
```

 but should read


```
int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,
                  int dest, int sendtag, void *recvbuf, int recvcount,
                  MPI_Datatype recvtype, int source, int recvtag,
                  MPI_Comm comm, MPI_Status *status)
```
- Page 220, lines 19–20 reads


```
int double MPI_Wtime(void)
int double MPI_Wtick(void)
```

 but should read


```
double MPI_Wtime(void)
double MPI_Wtick(void)
```
- Page 222, line 34 reads


```
INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```

 but should read


```
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```
- Page 222, line 38 reads


```
INTEGER REQUEST, COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```

 but should read


```
INTEGER COUNT, DATATYPE, DEST, TAG, COMM, REQUEST, IERROR
```
- Page 227, lines 19–20 reads


```
MPI_INTERCOMM_MERGE(INTERCOMM, HIGH, INTRACOMM, IERROR)
INTEGER INTERCOMM, INTRACOMM, IERROR
```

 but should read


```
MPI_INTERCOMM_MERGE(INTERCOMM, HIGH, NEWINTRACOMM, IERROR)
INTEGER INTERCOMM, NEWINTRACOMM, IERROR
```
- Page 228, line 46 reads


```
MPI_ERRHANDLER_CREATE(FUNCTION, HANDLER, IERROR)
```

 but should read


```
MPI_ERRHANDLER_CREATE(FUNCTION, ERRHANDLER, IERROR)
```
- Page 229, line 33 reads


```
MPI_PCONTROL(level)
```

```
but should read
MPI_PCONTROL(LEVEL)
```

```
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
```

Chapter 4

Miscellany

This chapter contains topics that do not fit conveniently into other chapters.

4.1 Portable MPI Process Startup

A number of implementations of MPI-1 provide a startup command for MPI programs that is of the form

```
mpirun <mpirun arguments> <program> <program arguments>
```

Separating the command to start the program from the program itself provides flexibility, particularly for network and heterogeneous implementations. For example, the startup script need not run on one of the machines that will be executing the MPI program itself.

Having a standard startup mechanism also extends the portability of MPI programs one step further, to the command lines and scripts that manage them. For example, a validation suite script that runs hundreds of programs can be a portable script if it is written using such a standard startup mechanism. In order that the “standard” command not be confused with existing practice, which is not standard and not portable among implementations, instead of `mpirun` MPI specifies `mpiexec`.

While a standardized startup mechanism improves the usability of MPI, the range of environments is so diverse (e.g., there may not even be a command line interface) that MPI cannot mandate such a mechanism. Instead, MPI specifies an `mpiexec` startup command and recommends but does not require it, as advice to implementors. However, if an implementation does provide a command called `mpiexec`, it must be of the form described below.

It is suggested that

```
mpiexec -n <numprocs> <program>
```

be at least one way to start `<program>` with an initial `MPI_COMM_WORLD` whose group contains `<numprocs>` processes. Other arguments to `mpiexec` may be implementation-dependent.

This is advice to implementors, rather than a required part of MPI-2. It is not suggested that this be the only way to start MPI programs. If an implementation does provide a command called `mpiexec`, however, it must be of the form described here.

Advice to implementors. Implementors, if they do provide a special startup command for MPI programs, are advised to give it the following form. The syntax is chosen in order that `mpiexec` be able to be viewed as a command-line version of `MPI_COMM_SPAWN` (See Section 5.3.4).

Analogous to `MPI_COMM_SPAWN`, we have

```

mpiexec -n    <maxprocs>
          -soft <      >
          -host <      >
          -arch <      >
          -wdir <      >
          -path <      >
          -file <      >
          ...
          <command line>

```

for the case where a single command line for the application program and its arguments will suffice. See Section 5.3.4 for the meanings of these arguments. For the case corresponding to `MPI_COMM_SPAWN_MULTIPLE` there are two possible formats:

Form A:

```

mpiexec { <above arguments> } : { ... } : { ... } : ... : { ... }

```

As with `MPI_COMM_SPAWN`, all the arguments are optional. (Even the `-n x` argument is optional; the default is implementation dependent. It might be 1, it might be taken from an environment variable, or it might be specified at compile time.) The names and meanings of the arguments are taken from the keys in the `info` argument to `MPI_COMM_SPAWN`. There may be other, implementation-dependent arguments as well.

Note that Form A, though convenient to type, prevents colons from being program arguments. Therefore an alternate, file-based form is allowed:

Form B:

```

mpiexec -configfile <filename>

```

where the lines of `<filename>` are of the form separated by the colons in Form A. Lines beginning with `#` are comments, and lines may be continued by terminating the partial line with `\`.

Example 4.1 Start 16 instances of `myprog` on the current or default machine:

```

mpiexec -n 16 myprog

```

Example 4.2 Start 10 processes on the machine called `ferrari`:

```

mpiexec -n 10 -host ferrari myprog

```

Example 4.3 Start three copies of the same program with different command-line arguments:

```
mpiexec myprog infile1 : myprog infile2 : myprog infile3
```

Example 4.4 Start the `ocean` program on five Suns and the `atmos` program on 10 RS/6000's:

```
mpiexec -n 5 -arch sun ocean : -n 10 -arch rs6000 atmos
```

It is assumed that the implementation in this case has a method for choosing hosts of the appropriate type. Their ranks are in the order specified.

Example 4.5 Start the `ocean` program on five Suns and the `atmos` program on 10 RS/6000's (Form B):

```
mpiexec -configfile myfile
```

where `myfile` contains

```
-n 5 -arch sun    ocean
-n 10 -arch rs6000 atmos
```

(End of advice to implementors.)

4.2 Passing NULL to MPI_Init

In MPI-1.1, it is explicitly stated that an implementation is allowed to require that the arguments `argc` and `argv` passed by an application to `MPI_INIT` in C be the same arguments passed into the application as the arguments to `main`. In MPI-2 implementations are not allowed to impose this requirement. Conforming implementations of MPI are required to allow applications to pass `NULL` for both the `argc` and `argv` arguments of `main`. In C++, there is an alternative binding for `MPI::Init` that does not have these arguments at all.

Rationale. In some applications, libraries may be making the call to `MPI_Init`, and may not have access to `argc` and `argv` from `main`. It is anticipated that applications requiring special information about the environment or information supplied by `mpiexec` can get that information from environment variables. *(End of rationale.)*

4.3 Version Number

The values for the `MPI_VERSION` and `MPI_SUBVERSION` for an MPI-2 implementation are 2 and 0 respectively. This applies both to the values of the above constants and to the values returned by `MPI_GET_VERSION`.

4.4 Datatype Constructor MPI_TYPE_CREATE_INDEXED_BLOCK

This function is the same as `MPI_TYPE_INDEXED` except that the `blocklength` is the same for all blocks. There are many codes using indirect addressing arising from unstructured grids where the `blocksize` is always 1 (gather/scatter). The following convenience function allows for constant `blocksize` and arbitrary displacements.

`MPI_TYPE_CREATE_INDEXED_BLOCK(count, blocklength, array_of_displacements, oldtype, newtype)`

IN	<code>count</code>	length of array of displacements (integer)
IN	<code>blocklength</code>	size of block (integer)
IN	<code>array_of_displacements</code>	array of displacements (array of integer)
IN	<code>oldtype</code>	old datatype (handle)
OUT	<code>newtype</code>	new datatype (handle)

```
int MPI_Type_create_indexed_block(int count, int blocklength,
                                int array_of_displacements[], MPI_Datatype oldtype,
                                MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_INDEXED_BLOCK(COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS,
                              OLDTYPE, NEWTYPE, IERROR)
    INTEGER COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS(*), OLDTYPE,
    NEWTYPE, IERROR
```

```
MPI::Datatype MPI::Datatype::Create_indexed_block( int count,
                                                    int blocklength, const int array_of_displacements[]) const
```

4.5 Treatment of MPI_Status

The following features add to, but do not change, the functionality associated with `MPI_STATUS`.

4.5.1 Passing MPI_STATUS_IGNORE for Status

Every call to `MPI_RECV` includes a `status` argument, wherein the system can return details about the message received. There are also a number of other MPI calls, particularly in MPI-2, where `status` is returned. An object of type `MPI_STATUS` is not an MPI opaque object; its structure is declared in `mpi.h` and `mpif.h`, and it exists in the user's program. In many cases, application programs are constructed so that it is unnecessary for them to examine the `status` fields. In these cases, it is a waste for the user to allocate a status object, and it is particularly wasteful for the MPI implementation to fill in fields in this object.

To cope with this problem, there are two predefined constants, `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE`, which when passed to a receive, wait, or test function, inform the implementation that the status fields are not to be filled in. Note that `MPI_STATUS_IGNORE`

is not a special type of `MPI_STATUS` object; rather, it is a special value for the argument. In C one would expect it to be `NULL`, not the address of a special `MPI_STATUS`.

`MPI_STATUS_IGNORE`, and the array version `MPI_STATUSES_IGNORE`, can be used everywhere a status argument is passed to a receive, wait, or test function. `MPI_STATUS_IGNORE` cannot be used when status is an IN argument. Note that in Fortran `MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are objects like `MPI_BOTTOM` (not usable for initialization or assignment). See Section 2.5.4.

In general, this optimization can apply to all functions for which `status` or an array of `statuses` is an OUT argument. Note that this converts `status` into an INOUT argument. The functions that can be passed `MPI_STATUS_IGNORE` are all the various forms of `MPI_RECV`, `MPI_TEST`, and `MPI_WAIT`, as well as `MPI_REQUEST_GET_STATUS`. When an array is passed, as in the `ANY` and `ALL` functions, a separate constant, `MPI_STATUSES_IGNORE`, is passed for the array argument. It is possible for an MPI function to return `MPI_ERR_IN_STATUS` even when `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE` has been passed to that function.

`MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE` are not required to have the same values in C and Fortran.

It is not allowed to have some of the statuses in an array of statuses for `ANY` and `ALL` functions set to `MPI_STATUS_IGNORE`; one either specifies ignoring *all* of the statuses in such a call with `MPI_STATUSES_IGNORE`, or *none* of them by passing normal statuses in all positions in the array of statuses.

There are no C++ bindings for `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`. To allow an OUT or INOUT `MPI::Status` argument to be ignored, all MPI C++ bindings that have OUT or INOUT `MPI::Status` parameters are overloaded with a second version that omits the OUT or INOUT `MPI::Status` parameter.

Example 4.6 The C++ bindings for `MPI_PROBE` are:

```
void MPI::Comm::Probe(int source, int tag, MPI::Status& status) const
void MPI::Comm::Probe(int source, int tag) const
```

4.5.2 Non-destructive Test of status

This call is useful for accessing the information associated with a request, without freeing the request (in case the user is expected to access it later). It allows one to layer libraries more conveniently, since multiple layers of software may access the same completed request and extract from it the status information.

`MPI_REQUEST_GET_STATUS(request, flag, status)`

IN	request	request (handle)
OUT	flag	boolean flag, same as from <code>MPI_TEST</code> (logical)
OUT	status	<code>MPI_STATUS</code> object if flag is true (Status)

```
int MPI_Request_get_status(MPI_Request request, int *flag,
                           MPI_Status *status)
```

```
MPI_REQUEST_GET_STATUS( REQUEST, FLAG, STATUS, IERROR)
    INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
```

LOGICAL FLAG

```

1
2 bool MPI::Request::Get_status(MPI::Status& status) const
3
4 bool MPI::Request::Get_status() const

```

Sets **flag=true** if the operation is complete, and, if so, returns in status the request status. However, unlike test or wait, it does not deallocate or inactivate the request; a subsequent call to test, wait or free should be executed with that request. It sets **flag=false** if the operation is not complete.

4.6 Error Class for Invalid Keyval

Key values for attributes are system-allocated, by `MPI_{TYPE,COMM,WIN}_CREATE_KEYVAL`. Only such values can be passed to the functions that use key values as input arguments. In order to signal that an erroneous key value has been passed to one of these functions, there is a new MPI error class: `MPI_ERR_KEYVAL`. It can be returned by `MPI_ATTR_PUT`, `MPI_ATTR_GET`, `MPI_ATTR_DELETE`, `MPI_KEYVAL_FREE`, `MPI_{TYPE,COMM,WIN}_DELETE_ATTR`, `MPI_{TYPE,COMM,WIN}_SET_ATTR`, `MPI_{TYPE,COMM,WIN}_GET_ATTR`, `MPI_{TYPE,COMM,WIN}_FREE_KEYVAL`, `MPI_COMM_DUP`, `MPI_COMM_DISCONNECT`, and `MPI_COMM_FREE`. The last three are included because **keyval** is an argument to the copy and delete functions for attributes.

4.7 Committing a Committed Datatype

In MPI-1.2, the effect of calling `MPI_TYPE_COMMIT` with a datatype that is already committed is not specified. For MPI-2, it is specified that `MPI_TYPE_COMMIT` will accept a committed datatype; in this case, it is equivalent to a no-op.

4.8 Allowing User Functions at Process Termination

There are times in which it would be convenient to have actions happen when an MPI process finishes. For example, a routine may do initializations that are useful until the MPI job (or that part of the job that being terminated in the case of dynamically created processes) is finished. This can be accomplished in MPI-2 by attaching an attribute to `MPI_COMM_SELF` with a callback function. When `MPI_FINALIZE` is called, it will first execute the equivalent of an `MPI_COMM_FREE` on `MPI_COMM_SELF`. This will cause the delete callback function to be executed on all keys associated with `MPI_COMM_SELF`, in an arbitrary order. If no key has been attached to `MPI_COMM_SELF`, then no callback is invoked. The “freeing” of `MPI_COMM_SELF` occurs before any other parts of MPI are affected. Thus, for example, calling `MPI_FINALIZED` will return **false** in any of these callback functions. Once done with `MPI_COMM_SELF`, the order and rest of the actions taken by `MPI_FINALIZE` is not specified.

Advice to implementors. Since attributes can be added from any supported language, the MPI implementation needs to remember the creating language so the correct callback is made. (*End of advice to implementors.*)

4.9 Determining Whether MPI Has Finished

One of the goals of MPI was to allow for layered libraries. In order for a library to do this cleanly, it needs to know if MPI is active. In MPI-1 the function `MPI_INITIALIZED` was provided to tell if MPI had been initialized. The problem arises in knowing if MPI has been finalized. Once MPI has been finalized it is no longer active and cannot be restarted. A library needs to be able to determine this to act accordingly. To achieve this the following function is needed:

MPI_FINALIZED(flag)

OUT	flag	true if MPI was finalized (logical)
-----	------	-------------------------------------

```
int MPI_Finalized(int *flag)
```

MPI_FINALIZED(FLAG, IERROR)

LOGICAL FLAG

INTEGER IERROR

```
bool MPI::Is_finalized()
```

This routine returns **true** if `MPI_FINALIZE` has completed. It is legal to call `MPI_FINALIZED` before `MPI_INIT` and after `MPI_FINALIZE`.

Advice to users. MPI is “active” and it is thus safe to call MPI functions if `MPI_INIT` has completed and `MPI_FINALIZE` has not completed. If a library has no other way of knowing whether MPI is active or not, then it can use `MPI_INITIALIZED` and `MPI_FINALIZED` to determine this. For example, MPI is “active” in callback functions that are invoked during `MPI_FINALIZE`. (*End of advice to users.*)

4.10 The Info Object

Many of the routines in MPI-2 take an argument `info`. `info` is an opaque object with a handle of type `MPI_Info` in C, `MPI::Info` in C++, and `INTEGER` in Fortran. It consists of (`key,value`) pairs (both `key` and `value` are strings). A key may have only one value. MPI reserves several keys and requires that if an implementation uses a reserved key, it must provide the specified functionality. An implementation is not required to support these keys and may support any others not reserved by MPI.

If a function does not recognize a key, it will ignore it, unless otherwise specified. If an implementation recognizes a key but does not recognize the format of the corresponding value, the result is undefined.

Keys have an implementation-defined maximum length of `MPI_MAX_INFO_KEY`, which is at least 32 and at most 255. Values have an implementation-defined maximum length of `MPI_MAX_INFO_VAL`. In Fortran, leading and trailing spaces are stripped from both. Returned values will never be larger than these maximum lengths. Both **key** and **value** are case sensitive.

Rationale. Keys have a maximum length because the set of known keys will always be finite and known to the implementation and because there is no reason for keys

to be complex. The small maximum size allows applications to declare keys of size `MPI_MAX_INFO_KEY`. The limitation on value sizes is so that an implementation is not forced to deal with arbitrarily long strings. (*End of rationale.*)

Advice to users. `MPI_MAX_INFO_VAL` might be very large, so it might not be wise to declare a string of that size. (*End of advice to users.*)

When it is an argument to a non-blocking routine, `info` is parsed before that routine returns, so that it may be modified or freed immediately after return.

When the descriptions refer to a key or value as being a boolean, an integer, or a list, they mean the string representation of these types. An implementation may define its own rules for how `info` value strings are converted to other types, but to ensure portability, every implementation must support the following representations. Legal values for a boolean must include the strings “true” and “false” (all lowercase). For integers, legal values must include string representations of decimal values of integers that are within the range of a standard integer type in the program. (However it is possible that not every legal integer is a legal value for a given key.) On positive numbers, + signs are optional. No space may appear between a + or – sign and the leading digit of a number. For comma separated lists, the string must contain legal elements separated by commas. Leading and trailing spaces are stripped automatically from the types of `info` values described above and for each element of a comma separated list. These rules apply to all `info` values of these types. Implementations are free to specify a different interpretation for values of other `info` keys.

`MPI_INFO_CREATE(info)`

OUT	<code>info</code>	<code>info</code> object created (handle)
-----	-------------------	---

```
int MPI_Info_create(MPI_Info *info)
```

```
MPI_INFO_CREATE(INFO, IERROR)
```

```
INTEGER INFO, IERROR
```

```
static MPI::Info MPI::Info::Create()
```

`MPI_INFO_CREATE` creates a new `info` object. The newly created object contains no key/value pairs.

`MPI_INFO_SET(info, key, value)`

INOUT	<code>info</code>	<code>info</code> object (handle)
-------	-------------------	-----------------------------------

IN	<code>key</code>	<code>key</code> (string)
----	------------------	---------------------------

IN	<code>value</code>	<code>value</code> (string)
----	--------------------	-----------------------------

```
int MPI_Info_set(MPI_Info info, char *key, char *value)
```

```
MPI_INFO_SET(INFO, KEY, VALUE, IERROR)
```

```
INTEGER INFO, IERROR
```

```
CHARACTER*(*) KEY, VALUE
```

```
void MPI::Info::Set(const char* key, const char* value)
```

`MPI_INFO_SET` adds the (key,value) pair to `info`, and overrides the value if a value for the same key was previously set. `key` and `value` are null-terminated strings in C. In Fortran, leading and trailing spaces in `key` and `value` are stripped. If either `key` or `value` are larger than the allowed maximums, the errors `MPI_ERR_INFO_KEY` or `MPI_ERR_INFO_VALUE` are raised, respectively.

`MPI_INFO_DELETE(info, key)`

INOUT	<code>info</code>	info object (handle)
IN	<code>key</code>	key (string)

`int MPI_Info_delete(MPI_Info info, char *key)`

`MPI_INFO_DELETE(INFO, KEY, IERROR)`

INTEGER	<code>INFO</code> , <code>IERROR</code>
CHARACTER*(*)	<code>KEY</code>

`void MPI::Info::Delete(const char* key)`

`MPI_INFO_DELETE` deletes a (key,value) pair from `info`. If `key` is not defined in `info`, the call raises an error of class `MPI_ERR_INFO_NOKEY`.

`MPI_INFO_GET(info, key, valuelen, value, flag)`

IN	<code>info</code>	info object (handle)
IN	<code>key</code>	key (string)
IN	<code>valuelen</code>	length of value arg (integer)
OUT	<code>value</code>	value (string)
OUT	<code>flag</code>	true if key defined, false if not (boolean)

`int MPI_Info_get(MPI_Info info, char *key, int valuelen, char *value, int *flag)`

`MPI_INFO_GET(INFO, KEY, VALUELEN, VALUE, FLAG, IERROR)`

INTEGER	<code>INFO</code> , <code>VALUELEN</code> , <code>IERROR</code>
CHARACTER*(*)	<code>KEY</code> , <code>VALUE</code>
LOGICAL	<code>FLAG</code>

`bool MPI::Info::Get(const char* key, int valuelen, char* value) const`

This function retrieves the value associated with `key` in a previous call to `MPI_INFO_SET`. If such a key exists, it sets `flag` to true and returns the value in `value`, otherwise it sets `flag` to false and leaves `value` unchanged. `valuelen` is the number of characters available in `value`. If it is less than the actual size of the value, the value is truncated. In C, `valuelen` should be one less than the amount of allocated space to allow for the null terminator.

If `key` is larger than `MPI_MAX_INFO_KEY`, the call is erroneous.

1 MPI_INFO_GET_VALUELEN(info, key, valuelen, flag)

2	IN	info	info object (handle)
3	IN	key	key (string)
4	OUT	valuelen	length of value arg (integer)
5	OUT	flag	true if key defined, false if not (boolean)

6
7
8 int MPI_Info_get_valuelen(MPI_Info info, char *key, int *valuelen,
9 int *flag)

10 MPI_INFO_GET_VALUELEN(INFO, KEY, VALUELEN, FLAG, IERROR)

11 INTEGER INFO, VALUELEN, IERROR

12 LOGICAL FLAG

13 CHARACTER*(*) KEY

14 bool MPI::Info::Get_valuelen(const char* key, int& valuelen) const

15
16
17 Retrieves the length of the **value** associated with **key**. If **key** is defined, **valuelen** is set
18 to the length of its associated value and **flag** is set to true. If **key** is not defined, **valuelen** is
19 not touched and **flag** is set to false. The length returned in C or C++ does not include the
20 end-of-string character.

21 If **key** is larger than MPI_MAX_INFO_KEY, the call is erroneous.

22
23 MPI_INFO_GET_NKEYS(info, nkeys)

24	IN	info	info object (handle)
25	OUT	nkeys	number of defined keys (integer)

26
27
28 int MPI_Info_get_nkeys(MPI_Info info, int *nkeys)

29 MPI_INFO_GET_NKEYS(INFO, NKEYS, IERROR)

30 INTEGER INFO, NKEYS, IERROR

31 int MPI::Info::Get_nkeys() const

32 MPI_INFO_GET_NKEYS returns the number of currently defined keys in **info**.

33
34 MPI_INFO_GET_NTHKEY(info, n, key)

35	IN	info	info object (handle)
36	IN	n	key number (integer)
37	OUT	key	key (string)

38
39
40 int MPI_Info_get_nthkey(MPI_Info info, int n, char *key)

41 MPI_INFO_GET_NTHKEY(INFO, N, KEY, IERROR)

42 INTEGER INFO, N, IERROR

43 CHARACTER*(*) KEY

44 void MPI::Info::Get_nthkey(int n, char* key) const

This function returns the `nth` defined key in `info`. Keys are numbered $0 \dots N - 1$ where N is the value returned by `MPI_INFO_GET_NKEYS`. All keys between 0 and $N - 1$ are guaranteed to be defined. The number of a given key does not change as long as `info` is not modified with `MPI_INFO_SET` or `MPI_INFO_DELETE`.

`MPI_INFO_DUP(info, newinfo)`

IN	<code>info</code>	info object (handle)
OUT	<code>newinfo</code>	info object (handle)

`int MPI_Info_dup(MPI_Info info, MPI_Info *newinfo)`

`MPI_INFO_DUP(INFO, NEWINFO, IERROR)`

`INTEGER INFO, NEWINFO, IERROR`

`MPI::Info MPI::Info::Dup() const`

`MPI_INFO_DUP` duplicates an existing info object, creating a new object, with the same (key,value) pairs and the same ordering of keys.

`MPI_INFO_FREE(info)`

INOUT	<code>info</code>	info object (handle)
-------	-------------------	----------------------

`int MPI_Info_free(MPI_Info *info)`

`MPI_INFO_FREE(INFO, IERROR)`

`INTEGER INFO, IERROR`

`void MPI::Info::Free()`

This function frees `info` and sets it to `MPI_INFO_NULL`. The value of an info argument is interpreted each time the info is passed to a routine. Changes to an info after return from a routine do not affect that interpretation.

4.11 Memory Allocation

In some systems, message-passing and remote-memory-access (RMA) operations run faster when accessing specially allocated memory (e.g., memory that is shared by the other processes in the communicating group on an SMP). MPI provides a mechanism for allocating and freeing such special memory. The use of such memory for message passing or RMA is not mandatory, and this memory can be used without restrictions as any other dynamically allocated memory. However, implementations may restrict the use of the `MPI_WIN_LOCK` and `MPI_WIN_UNLOCK` functions to windows allocated in such memory (see Section 6.4.3.)

1 **MPI_ALLOC_MEM**(size, info, baseptr)

2	IN	size	size of memory segment in bytes (nonnegative integer)
3	IN	info	info argument (handle)
4			
5	OUT	baseptr	pointer to beginning of memory segment allocated

6
7 **int** MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)

8 **MPI_ALLOC_MEM**(SIZE, INFO, BASEPTR, IERROR)

9 INTEGER INFO, IERROR

10 INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR

11
12 **void*** MPI::Alloc_mem(MPI::Aint size, const MPI::Info& info)

13
14 The **info** argument can be used to provide directives that control the desired location
15 of the allocated memory. Such a directive does not affect the semantics of the call. Valid
16 **info** values are implementation-dependent; a null directive value of **info** = **MPI_INFO_NULL**
17 is always valid.

18 The function **MPI_ALLOC_MEM** may return an error code of class **MPI_ERR_NO_MEM**
19 to indicate it failed because memory is exhausted.

20
21 **MPI_FREE_MEM**(base)

22	IN	base	initial address of memory segment allocated by
23			MPI_ALLOC_MEM (choice)

24
25
26 **int** MPI_Free_mem(void *base)

27 **MPI_FREE_MEM**(BASE, IERROR)

28 <type> BASE(*)

29 INTEGER IERROR

30
31 **void** MPI::Free_mem(void *base)

32
33 The function **MPI_FREE_MEM** may return an error code of class **MPI_ERR_BASE** to
34 indicate an invalid base argument.

35 *Rationale.* The C and C++ bindings of **MPI_ALLOC_MEM** and **MPI_FREE_MEM** are
36 similar to the bindings for the **malloc** and **free** C library calls: a call to
37 **MPI_Alloc_mem**(..., &base) should be paired with a call to **MPI_Free_mem**(base) (one
38 less level of indirection). Both arguments are declared to be of same type **void*** so
39 as to facilitate type casting. The Fortran binding is consistent with the C and C++
40 bindings: the Fortran **MPI_ALLOC_MEM** call returns in **baseptr** the (integer valued)
41 address of the allocated memory. The **base** argument of **MPI_FREE_MEM** is a choice
42 argument, which passes (a reference to) the variable stored at that location. (*End of*
43 *rationale.*)

44
45 *Advice to implementors.* If **MPI_ALLOC_MEM** allocates special memory, then a design
46 similar to the design of C **malloc** and **free** functions has to be used, in order to find
47 out the size of a memory segment, when the segment is freed. If no special memory is
48 used, **MPI_ALLOC_MEM** simply invokes **malloc**, and **MPI_FREE_MEM** invokes **free**.

A call to `MPI_ALLOC_MEM` can be used in shared memory systems to allocate memory in a shared memory segment. (*End of advice to implementors.*)

Example 4.7 Example of use of `MPI_ALLOC_MEM`, in Fortran with pointer support. We assume 4-byte REALs, and assume that pointers are address-sized.

```
REAL A
POINTER (P, A(100,100)) ! no memory is allocated
CALL MPI_ALLOC_MEM(4*100*100, MPI_INFO_NULL, P, IERR)
! memory is allocated
...
A(3,5) = 2.71;
...
CALL MPI_FREE_MEM(A, IERR) ! memory is freed
```

Since standard Fortran does not support (C-like) pointers, this code is not Fortran 77 or Fortran 90 code. Some compilers (in particular, at the time of writing, g77 and Fortran compilers for Intel) do not support this code.

Example 4.8 Same example, in C

```
float (* f)[100][100] ;
MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
...
(*f)[5][3] = 2.71;
...
MPI_Free_mem(f);
```

4.12 Language Interoperability

4.12.1 Introduction

It is not uncommon for library developers to use one language to develop an applications library that may be called by an application program written in a different language. MPI currently supports ISO (previously ANSI) C, C++, and Fortran bindings. It should be possible for applications in any of the supported languages to call MPI-related functions in another language.

Moreover, MPI allows the development of client-server code, with MPI communication used between a parallel client and a parallel server. It should be possible to code the server in one language and the clients in another language. To do so, communications should be possible between applications written in different languages.

There are several issues that need to be addressed in order to achieve interoperability.

Initialization We need to specify how the MPI environment is initialized for all languages.

Interlanguage passing of MPI opaque objects We need to specify how MPI object handles are passed between languages. We also need to specify what happens when an MPI object is accessed in one language, to retrieve information (e.g., attributes) set in another language.

Interlanguage communication We need to specify how messages sent in one language can be received in another language.

It is highly desirable that the solution for interlanguage interoperability be extendable to new languages, should MPI bindings be defined for such languages.

4.12.2 Assumptions

We assume that conventions exist for programs written in one language to call functions in written in another language. These conventions specify how to link routines in different languages into one program, how to call functions in a different language, how to pass arguments between languages, and the correspondence between basic data types in different languages. In general, these conventions will be implementation dependent. Furthermore, not every basic datatype may have a matching type in other languages. For example, C/C++ character strings may not be compatible with Fortran `CHARACTER` variables. However, we assume that a Fortran `INTEGER`, as well as a (sequence associated) Fortran array of `INTEGER`s, can be passed to a C or C++ program. We also assume that Fortran, C, and C++ have address-sized integers. This does not mean that the default-size integers are the same size as default-sized pointers, but only that there is some way to hold (and pass) a C address in a Fortran integer. It is also assumed that `INTEGER(KIND=MPI_OFFSET_KIND)` can be passed from Fortran to C as `MPI_Offset`.

4.12.3 Initialization

A call to `MPI_INIT` or `MPI_THREAD_INIT`, from any language, initializes MPI for execution in all languages.

Advice to users. Certain implementations use the (inout) `argc`, `argv` arguments of the C/C++ version of `MPI_INIT` in order to propagate values for `argc` and `argv` to all executing processes. Use of the Fortran version of `MPI_INIT` to initialize MPI may result in a loss of this ability. (*End of advice to users.*)

The function `MPI_INITIALIZED` returns the same answer in all languages.

The function `MPI_FINALIZE` finalizes the MPI environments for all languages.

The function `MPI_FINALIZED` returns the same answer in all languages.

The function `MPI_ABORT` kills processes, irrespective of the language used by the caller or by the processes killed.

The MPI environment is initialized in the same manner for all languages by `MPI_INIT`. E.g., `MPI_COMM_WORLD` carries the same information regardless of language: same processes, same environmental attributes, same error handlers.

Information can be added to info objects in one language and retrieved in another.

Advice to users. The use of several languages in one MPI program may require the use of special options at compile and/or link time. (*End of advice to users.*)

Advice to implementors. Implementations may selectively link language specific MPI libraries only to codes that need them, so as not to increase the size of binaries for codes that use only one language. The MPI initialization code need perform initialization for a language only if that language library is loaded. (*End of advice to implementors.*)

4.12.4 Transfer of Handles

Handles are passed between Fortran and C or C++ by using an explicit C wrapper to convert Fortran handles to C handles. There is no direct access to C or C++ handles in Fortran. Handles are passed between C and C++ using overloaded C++ operators called from C++ code. There is no direct access to C++ objects from C.

The type definition `MPI_Fint` is provided in C/C++ for an integer of the size that matches a Fortran `INTEGER`; usually, `MPI_Fint` will be equivalent to `int`.

The following functions are provided in C to convert from a Fortran communicator handle (which is an integer) to a C communicator handle, and vice versa.

```
MPI_Comm MPI_Comm_f2c(MPI_Fint comm)
```

If `comm` is a valid Fortran handle to a communicator, then `MPI_Comm_f2c` returns a valid C handle to that same communicator; if `comm = MPI_COMM_NULL` (Fortran value), then `MPI_Comm_f2c` returns a null C handle; if `comm` is an invalid Fortran handle, then `MPI_Comm_f2c` returns an invalid C handle.

```
MPI_Fint MPI_Comm_c2f(MPI_Comm comm)
```

The function `MPI_Comm_c2f` translates a C communicator handle into a Fortran handle to the same communicator; it maps a null handle into a null handle and an invalid handle into an invalid handle.

Similar functions are provided for the other types of opaque objects.

```
MPI_Datatype MPI_Type_f2c(MPI_Fint datatype)
```

```
MPI_Fint MPI_Type_c2f(MPI_Datatype datatype)
```

```
MPI_Group MPI_Group_f2c(MPI_Fint group)
```

```
MPI_Fint MPI_Group_c2f(MPI_Group group)
```

```
MPI_Request MPI_Request_f2c(MPI_Fint request)
```

```
MPI_Fint MPI_Request_c2f(MPI_Request request)
```

```
MPI_File MPI_File_f2c(MPI_Fint file)
```

```
MPI_Fint MPI_File_c2f(MPI_File file)
```

```
MPI_Win MPI_Win_f2c(MPI_Fint win)
```

```
MPI_Fint MPI_Win_c2f(MPI_Win win)
```

```
MPI_Op MPI_Op_f2c(MPI_Fint op)
```

```
MPI_Fint MPI_Op_c2f(MPI_Op op)
```

```
MPI_Info MPI_Info_f2c(MPI_Fint info)
```

```
MPI_Fint MPI_Info_c2f(MPI_Info info)
```

Example 4.9 The example below illustrates how the Fortran `MPI` function `MPI_TYPE_COMMIT` can be implemented by wrapping the C `MPI` function `MPI_Type_commit` with a C wrapper to do handle conversions. In this example a Fortran-C

interface is assumed where a Fortran function is all upper case when referred to from C and arguments are passed by addresses.

```

1  ! FORTRAN PROCEDURE
2  SUBROUTINE MPI_TYPE_COMMIT( DATATYPE, IERR)
3  INTEGER DATATYPE, IERR
4  CALL MPI_X_TYPE_COMMIT(DATATYPE, IERR)
5  RETURN
6  END
7
8  /* C wrapper */
9
10 void MPI_X_TYPE_COMMIT( MPI_Fint *f_handle, MPI_Fint *ierr)
11 {
12 MPI_Datatype datatype;
13
14 datatype = MPI_Type_f2c( *f_handle);
15 *ierr = (MPI_Fint)MPI_Type_commit( &datatype);
16 *f_handle = MPI_Type_c2f(datatype);
17 return;
18 }
19
20
21

```

The same approach can be used for all other MPI functions. The call to `MPI_XXX_f2c` (resp. `MPI_XXX_c2f`) can be omitted when the handle is an OUT (resp. IN) argument, rather than INOUT.

Rationale. The design here provides a convenient solution for the prevalent case, where a C wrapper is used to allow Fortran code to call a C library, or C code to call a Fortran library. The use of C wrappers is much more likely than the use of Fortran wrappers, because it is much more likely that a variable of type `INTEGER` can be passed to C, than a C handle can be passed to Fortran.

Returning the converted value as a function value rather than through the argument list allows the generation of efficient inlined code when these functions are simple (e.g., the identity). The conversion function in the wrapper does not catch an invalid handle argument. Instead, an invalid handle is passed below to the library function, which, presumably, checks its input arguments. (*End of rationale.*)

C and C++ The C++ language interface provides the functions listed below for mixed-language interoperability. The token `<CLASS>` is used below to indicate any valid MPI opaque handle name (e.g., `Group`), except where noted. For the case where the C++ class corresponding to `<CLASS>` has derived classes, functions are also provided for converting between the derived classes and the C MPI-`<CLASS>`.

The following function allows assignment from a C MPI handle to a C++ MPI handle.

```
MPI::<CLASS>& MPI::<CLASS>::operator=(const MPI-<CLASS>& data)
```

The constructor below creates a C++ MPI object from a C MPI handle. This allows the automatic promotion of a C MPI handle to a C++ MPI handle.

```
MPI::<CLASS>::<CLASS>(const MPI-<CLASS>& data)
```

Example 4.10 In order for a C program to use a C++ library, the C++ library must export a C interface that provides appropriate conversions before invoking the underlying C++ library call. This example shows a C interface function that invokes a C++ library call with a C communicator; the communicator is automatically promoted to a C++ handle when the underlying C++ function is invoked.

```
// C++ library function prototype
void cpp_lib_call(MPI::Comm& cpp_comm);

// Exported C function prototype
extern "C" {
void c_interface(MPI_Comm c_comm);
}

void c_interface(MPI_Comm c_comm)
{
// the MPI_Comm (c_comm) is automatically promoted to MPI::Comm
cpp_lib_call(c_comm);
}
```

The following function allows conversion from C++ objects to C MPI handles. In this case, the casting operator is overloaded to provide the functionality.

```
MPI::<CLASS>::operator MPI_<CLASS>() const
```

Example 4.11 A C library routine is called from a C++ program. The C library routine is prototyped to take an MPI_Comm as an argument.

```
// C function prototype
extern "C" {
void c_lib_call(MPI_Comm c_comm);
}

void cpp_function()
{
// Create a C++ communicator, and initialize it with a dup of
// MPI::COMM_WORLD
MPI::Intracomm cpp_comm(MPI::COMM_WORLD.Dup());
c_lib_call(cpp_comm);
}
```

Rationale. Providing conversion from C to C++ via constructors and from C++ to C via casting allows the compiler to make automatic conversions. Calling C from C++ becomes trivial, as does the provision of a C or Fortran interface to a C++ library. (*End of rationale.*)

Advice to users. Note that the casting and promotion operators return new handles by value. Using these new handles as INOUT parameters will affect the internal MPI object, but will *not* affect the original handle from which it was cast. (*End of advice to users.*)

It is important to note that all C++ objects and their corresponding C handles can be used interchangeably by an application. For example, an application can cache an attribute on `MPI_COMM_WORLD` and later retrieve it from `MPI::COMM_WORLD`.

4.12.5 Status

The following two procedures are provided in C to convert from a Fortran status (which is an array of integers) to a C status (which is a structure), and vice versa. The conversion occurs on all the information in status, including that which is hidden. That is, no status information is lost in the conversion.

```
int MPI_Status_f2c(MPI_Fint *f_status, MPI_Status *c_status)
```

If `f_status` is a valid Fortran status, but not the Fortran value of `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`, then `MPI_Status_f2c` returns in `c_status` a valid C status with the same content. If `f_status` is the Fortran value of `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`, or if `f_status` is not a valid Fortran status, then the call is erroneous.

The C status has the same source, tag and error code values as the Fortran status, and returns the same answers when queried for count, elements, and cancellation. The conversion function may be called with a Fortran status argument that has an undefined error field, in which case the value of the error field in the C status argument is undefined.

Two global variables of type `MPI_Fint*`, `MPI_F_STATUS_IGNORE` and `MPI_F_STATUSES_IGNORE` are declared in `mpi.h`. They can be used to test, in C, whether `f_status` is the Fortran value of `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`, respectively. These are global variables, not C constant expressions and cannot be used in places where C requires constant expressions. Their value is defined only between the calls to `MPI_INIT` and `MPI_FINALIZE` and should not be changed by user code.

To do the conversion in the other direction, we have the following:

```
int MPI_Status_c2f(MPI_Status *c_status, MPI_Fint *f_status)
```

This call converts a C status into a Fortran status, and has a behavior similar to `MPI_Status_f2c`. That is, the value of `c_status` must not be either `MPI_STATUS_IGNORE` or `MPI_STATUSES_IGNORE`.

Advice to users. There is not a separate conversion function for arrays of statuses, since one can simply loop through the array, converting each status. (*End of advice to users.*)

Rationale. The handling of `MPI_STATUS_IGNORE` is required in order to layer libraries with only a C wrapper: if the Fortran call has passed `MPI_STATUS_IGNORE`, then the C wrapper must handle this correctly. Note that this constant need not have the same value in Fortran and C. If `MPI_Status_f2c` were to handle `MPI_STATUS_IGNORE`, then the type of its result would have to be `MPI_Status**`, which was considered an inferior solution. (*End of rationale.*)

4.12.6 MPI Opaque Objects

Unless said otherwise, opaque objects are “the same” in all languages: they carry the same information, and have the same meaning in both languages. The mechanism described in the previous section can be used to pass references to MPI objects from language to

language. An object created in one language can be accessed, modified or freed in another language.

We examine below in more detail, issues that arise for each type of MPI object.

Datatypes

Datatypes encode the same information in all languages. E.g., a datatype accessor like `MPI_TYPE_GET_EXTENT` will return the same information in all languages. If a datatype defined in one language is used for a communication call in another language, then the message sent will be identical to the message that would be sent from the first language: the same communication buffer is accessed, and the same representation conversion is performed, if needed. All predefined datatypes can be used in datatype constructors in any language. If a datatype is committed, it can be used for communication in any language.

The function `MPI_GET_ADDRESS` returns the same value in all languages. Note that we do not require that the constant `MPI_BOTTOM` have the same value in all languages (see 4.12.9, page 59).

Example 4.12

```

! FORTRAN CODE
REAL R(5)
INTEGER TYPE, IERR
INTEGER (KIND=MPI_ADDRESS_KIND) ADDR

! create an absolute datatype for array R
CALL MPI_GET_ADDRESS( R, ADDR, IERR)
CALL MPI_TYPE_CREATE_STRUCT(1, 5, ADDR, MPI_REAL, TYPE, IERR)
CALL C_ROUTINE(TYPE)

/* C code */

void C_ROUTINE(MPI_Fint *ftype)
{
    int count = 5;
    int lens[2] = {1,1};
    MPI_Aint displs[2];
    MPI_Datatype types[2], newtype;

    /* create an absolute datatype for buffer that consists
    /*   of count, followed by R(5)

    MPI_Get_address(&count, &displs[0]);
    displs[1] = 0;
    types[0] = MPI_INT;
    types[1] = MPI_Type_f2c(*ftype);
    MPI_Type_create_struct(2, lens, displs, types, &newtype);
    MPI_Type_commit(&newtype);

    MPI_Send(MPI_BOTTOM, 1, newtype, 1, 0, MPI_COMM_WORLD);

```

```

1  /* the message sent contains an int count of 5, followed */
2  /* by the 5 REAL entries of the Fortran array R.          */
3  }

```

Advice to implementors. The following implementation can be used: MPI addresses, as returned by `MPI_GET_ADDRESS`, will have the same value in all languages. One obvious choice is that MPI addresses be identical to regular addresses. The address is stored in the datatype, when datatypes with absolute addresses are constructed. When a send or receive operation is performed, then addresses stored in a datatype are interpreted as displacements that are all augmented by a base address. This base address is (the address of) `buf`, or zero, if `buf = MPI_BOTTOM`. Thus, if `MPI_BOTTOM` is zero then a send or receive call with `buf = MPI_BOTTOM` is implemented exactly as a call with a regular buffer argument: in both cases the base address is `buf`. On the other hand, if `MPI_BOTTOM` is not zero, then the implementation has to be slightly different. A test is performed to check whether `buf = MPI_BOTTOM`. If true, then the base address is zero, otherwise it is `buf`. In particular, if `MPI_BOTTOM` does not have the same value in Fortran and C/C++, then an additional test for `buf = MPI_BOTTOM` is needed in at least one of the languages.

It may be desirable to use a value other than zero for `MPI_BOTTOM` even in C/C++, so as to distinguish it from a NULL pointer. If `MPI_BOTTOM = c` then one can still avoid the test `buf = MPI_BOTTOM`, by using the displacement from `MPI_BOTTOM`, i.e., the regular address - `c`, as the MPI address returned by `MPI_GET_ADDRESS` and stored in absolute datatypes. (*End of advice to implementors.*)

Callback Functions

MPI calls may associate callback functions with MPI objects: error handlers are associated with communicators and files, attribute copy and delete functions are associated with attribute keys, reduce operations are associated with operation objects, etc. In a multilanguage environment, a function passed in an MPI call in one language may be invoked by an MPI call in another language. MPI implementations must make sure that such invocation will use the calling convention of the language the function is bound to.

Advice to implementors. Callback functions need to have a language tag. This tag is set when the callback function is passed in by the library function (which is presumably different for each language), and is used to generate the right calling sequence when the callback function is invoked. (*End of advice to implementors.*)

Error Handlers

Advice to implementors. Error handlers, have, in C and C++, a “`stdargs`” argument list. It might be useful to provide to the handler information on the language environment where the error occurred. (*End of advice to implementors.*)

Reduce Operations

Advice to users. Reduce operations receive as one of their arguments the datatype of the operands. Thus, one can define “polymorphic” reduce operations that work for C, C++, and Fortran datatypes. (*End of advice to users.*)

Addresses

Some of the datatype accessors and constructors have arguments of type `MPI_Aint` (in C) or `MPI::Aint` in C++, to hold addresses. The corresponding arguments, in Fortran, have type `INTEGER`. This causes Fortran and C/C++ to be incompatible, in an environment where addresses have 64 bits, but Fortran `INTEGER`s have 32 bits.

This is a problem, irrespective of interlanguage issues. Suppose that a Fortran process has an address space of ≥ 4 GB. What should be the value returned in Fortran by `MPI_ADDRESS`, for a variable with an address above 2^{32} ? The design described here addresses this issue, while maintaining compatibility with current Fortran codes.

The constant `MPI_ADDRESS_KIND` is defined so that, in Fortran 90, `INTEGER(KIND=MPI_ADDRESS_KIND)` is an address sized integer type (typically, but not necessarily, the size of an `INTEGER(KIND=MPI_ADDRESS_KIND)` is 4 on 32 bit address machines and 8 on 64 bit address machines). Similarly, the constant `MPI_INTEGER_KIND` is defined so that `INTEGER(KIND=MPI_INTEGER_KIND)` is a default size `INTEGER`.

There are seven functions that have address arguments: `MPI_TYPE_HVECTOR`, `MPI_TYPE_HINDEXED`, `MPI_TYPE_STRUCT`, `MPI_ADDRESS`, `MPI_TYPE_EXTENT`, `MPI_TYPE_LB` and `MPI_TYPE_UB`.

Four new functions are provided to supplement the first four functions in this list. These functions are described in Section 4.14, page 65. The remaining three functions are supplemented by the new function `MPI_TYPE_GET_EXTENT`, described in that same section. The new functions have the same functionality as the old functions in C/C++, or on Fortran systems where default `INTEGER`s are address sized. In Fortran, they accept arguments of type `INTEGER(KIND=MPI_ADDRESS_KIND)`, wherever arguments of type `MPI_Aint` are used in C. On Fortran 77 systems that do not support the Fortran 90 `KIND` notation, and where addresses are 64 bits whereas default `INTEGER`s are 32 bits, these arguments will be of an appropriate integer type. The old functions will continue to be provided, for backward compatibility. However, users are encouraged to switch to the new functions, in Fortran, so as to avoid problems on systems with an address range $> 2^{32}$, and to provide compatibility across languages.

4.12.7 Attributes

Attribute keys can be allocated in one language and freed in another. Similarly, attribute values can be set in one language and accessed in another. To achieve this, attribute keys will be allocated in an integer range that is valid all languages. The same holds true for system-defined attribute values (such as `MPI_TAG_UB`, `MPI_WTIME_IS_GLOBAL`, etc.)

Attribute keys declared in one language are associated with copy and delete functions in that language (the functions provided by the `MPI_{TYPE,COMM,WIN}_KEYVAL_CREATE` call). When a communicator is duplicated, for each attribute, the corresponding copy function is called, using the right calling convention for the language of that function; and similarly, for the delete callback function.

Advice to implementors. This requires that attributes be tagged either as “C,” “C++” or “Fortran,” and that the language tag be checked in order to use the right calling convention for the callback function. (*End of advice to implementors.*)

The attribute manipulation functions described in Section 5.7 of the MPI-1 standard define attributes arguments to be of type `void*` in C, and of type `INTEGER`, in Fortran. On

some systems, INTEGERS will have 32 bits, while C/C++ pointers will have 64 bits. This is a problem if communicator attributes are used to move information from a Fortran caller to a C/C++ callee, or vice-versa.

MPI will store, internally, address sized attributes. If Fortran INTEGERS are smaller, then the Fortran function `MPI_ATTR_GET` will return the least significant part of the attribute word; the Fortran function `MPI_ATTR_PUT` will set the least significant part of the attribute word, which will be sign extended to the entire word. (These two functions may be invoked explicitly by user code, or implicitly, by attribute copying callback functions.)

As for addresses, new functions are provided that manipulate Fortran address sized attributes, and have the same functionality as the old functions in C/C++. These functions are described in Section 8.8, page 198. Users are encouraged to use these new functions.

MPI supports two types of attributes: address-valued (pointer) attributes, and integer valued attributes. C and C++ attribute functions put and get address valued attributes. Fortran attribute functions put and get integer valued attributes. When an integer valued attribute is accessed from C or C++, then `MPI_XXX_get_attr` will return the address of (a pointer to) the integer valued attribute. When an address valued attribute is accessed from Fortran, then `MPI_XXX_GET_ATTR` will convert the address into an integer and return the result of this conversion. This conversion is lossless if new style (MPI-2) attribute functions are used, and an integer of kind `MPI_ADDRESS_KIND` is returned. The conversion may cause truncation if old style (MPI-1) attribute functions are used.

Example 4.13 A. C to Fortran

C code

```
static int i = 5;
void *p;
p = &i;
MPI_Comm_put_attr(..., p);
....
```

Fortran code

```
INTEGER(kind = MPI_ADDRESS_KIND) val
CALL MPI_COMM_GET_ATTR(...,val,...)
IF(val.NE.5) THEN CALL ERROR
```

B. Fortran to C

Fortran code

```
INTEGER(kind=MPI_ADDRESS_KIND) val
val = 55555
CALL MPI_COMM_PUT_ATTR(...,val,ierr)
```

C code


```

int *p;
MPI_Comm_get_attr(...,&p, ...);
if (*p != 55555) error();

```

The predefined MPI attributes can be integer valued or address valued. Predefined integer valued attributes, such as `MPI_TAG_UB`, behave as if they were put by a Fortran call. I.e., in Fortran, `MPI_COMM_GET_ATTR(MPI_COMM_WORLD, MPI_TAG_UB, val, flag, ierr)` will return in `val` the upper bound for tag value; in C, `MPI_Comm_get_attr(MPI_COMM_WORLD, MPI_TAG_UB, &p, &flag)` will return in `p` a pointer to an int containing the upper bound for tag value.

Address valued predefined attributes, such as `MPI_WIN_BASE` behave as if they were put by a C call. I.e., in Fortran, `MPI_WIN_GET_ATTR(win, MPI_WIN_BASE, val, flag, ierror)` will return in `val` the base address of the window, converted to an integer. In C, `MPI_Win_get_attr(win, MPI_WIN_BASE, &p, &flag)` will return in `p` a pointer to the window base, cast to `(void *)`.

Rationale. The design is consistent with the behavior specified in MPI-1 for predefined attributes, and ensures that no information is lost when attributes are passed from language to language. (*End of rationale.*)

Advice to implementors. Implementations should tag attributes either as address attributes or as integer attributes, according to whether they were set in C or in Fortran. Thus, the right choice can be made when the attribute is retrieved. (*End of advice to implementors.*)

4.12.8 Extra State

Extra-state should not be modified by the copy or delete callback functions. (This is obvious from the C binding, but not obvious from the Fortran binding). However, these functions may update state that is indirectly accessed via extra-state. E.g., in C, extra-state can be a pointer to a data structure that is modified by the copy or callback functions; in Fortran, extra-state can be an index into an entry in a `COMMON` array that is modified by the copy or callback functions. In a multithreaded environment, users should be aware that distinct threads may invoke the same callback function concurrently: if this function modifies state associated with extra-state, then mutual exclusion code must be used to protect updates and accesses to the shared state.

4.12.9 Constants

MPI constants have the same value in all languages, unless specified otherwise. This does not apply to constant handles (`MPI_INT`, `MPI_COMM_WORLD`, `MPI_ERRORS_RETURN`, `MPI_SUM`, etc.) These handles need to be converted, as explained in Section 4.12.4. Constants that specify maximum lengths of strings (see Section A.2.1 for a listing) have a value one less in Fortran than C/C++ since in C/C++ the length includes the null terminating character. Thus, these constants represent the amount of space which must be allocated to hold the largest possible such string, rather than the maximum number of printable characters the string could contain.

Advice to users. This definition means that it is safe in C/C++ to allocate a buffer to receive a string using a declaration like

```
1      char name [MPI_MAX_NAME_STRING];
```

```
2
3      (End of advice to users.)
```

4
5 Also constant “addresses,” i.e., special values for reference arguments that are not
6 handles, such as `MPI_BOTTOM` or `MPI_STATUS_IGNORE` may have different values in different
7 languages.

8
9 *Rationale.* The current MPI standard specifies that `MPI_BOTTOM` can be used in
10 initialization expressions in C, but not in Fortran. Since Fortran does not normally
11 support call by value, then `MPI_BOTTOM` must be in Fortran the name of a predefined
12 static variable, e.g., a variable in an MPI declared `COMMON` block. On the other
13 hand, in C, it is natural to take `MPI_BOTTOM = 0` (Caveat: Defining `MPI_BOTTOM`
14 `= 0` implies that `NULL` pointer cannot be distinguished from `MPI_BOTTOM`; it may be
15 that `MPI_BOTTOM = 1` is better ...) Requiring that the Fortran and C values be the
16 same will complicate the initialization process. (*End of rationale.*)

17 4.12.10 Interlanguage Communication

18
19 The type matching rules for communications in MPI are not changed: the datatype specifi-
20 cation for each item sent should match, in type signature, the datatype specification used to
21 receive this item (unless one of the types is `MPI_PACKED`). Also, the type of a message item
22 should match the type declaration for the corresponding communication buffer location,
23 unless the type is `MPI_BYTE` or `MPI_PACKED`. Interlanguage communication is allowed if it
24 complies with these rules.

25
26 **Example 4.14** In the example below, a Fortran array is sent from Fortran and received in
27 C.

```
28
29      ! FORTRAN CODE
30      REAL R(5)
31      INTEGER TYPE, IERR, MYRANK
32      INTEGER(KIND=MPI_ADDRESS_KIND) ADDR
33
34      ! create an absolute datatype for array R
35      CALL MPI_GET_ADDRESS( R, ADDR, IERR)
36      CALL MPI_TYPE_CREATE_STRUCT(1, 5, ADDR, MPI_REAL, TYPE, IERR)
37      CALL MPI_TYPE_COMMIT(TYPE, IERR)
38
39      CALL MPI_COMM_RANK( MPI_COMM_WORLD, MYRANK, IERR)
40      IF (MYRANK.EQ.0) THEN
41          CALL MPI_SEND( MPI_BOTTOM, 1, TYPE, 1, 0, MPI_COMM_WORLD, IERR)
42      ELSE
43          CALL C_ROUTINE(TYPE)
44      END IF
45
46
47      /* C code */
48
```

```

void C_ROUTINE(MPI_Fint *fhandle)
{
MPI_Datatype type;
MPI_Status status;

type = MPI_Type_f2c(*fhandle);

MPI_Recv( MPI_BOTTOM, 1, type, 0, 0, MPI_COMM_WORLD, &status);
}

```

MPI implementors may weaken these type matching rules, and allow messages to be sent with Fortran types and received with C types, and vice versa, when those types match. I.e., if the Fortran type INTEGER is identical to the C type int, then an MPI implementation may allow data to be sent with datatype MPI_INTEGER and be received with datatype MPI_INT. However, such code is not portable.

4.13 Error Handlers

MPI-1 attached error handlers only to communicators. MPI-2 attaches them to three types of objects: communicators, windows, and files. The extension was done while maintaining only one type of error handler opaque object. On the other hand, there are, in C and C++, distinct typedefs for user defined error handling callback functions that accept, respectively, communicator, file, and window arguments. In Fortran there are three user routines.

An error handler object is created by a call to `MPI_XXX_CREATE_ERRHANDLER(function, errhandler)`, where XXX is, respectively, COMM, WIN, or FILE.

An error handler is attached to a communicator, window, or file by a call to `MPI_XXX_SET_ERRHANDLER`. The error handler must be either a predefined error handler, or an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`, with matching XXX. The predefined error handlers `MPI_ERRORS_RETURN` and `MPI_ERRORS_ARE_FATAL` can be attached to communicators, windows, and files. In C++, the predefined error handler `MPI::ERRORS_THROW_EXCEPTIONS` can also be attached to communicators, windows, and files.

The error handler currently associated with a communicator, window, or file can be retrieved by a call to `MPI_XXX_GET_ERRHANDLER`.

The MPI-1 function `MPI_ERRHANDLER_FREE` can be used to free an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER`.

Advice to implementors. High quality implementation should raise an error when an error handler that was created by a call to `MPI_XXX_CREATE_ERRHANDLER` is attached to an object of the wrong type with a call to `MPI_YYY_SET_ERRHANDLER`. To do so, it is necessary to maintain, with each error handler, information on the typedef of the associated user function. (*End of advice to implementors.*)

The syntax for these calls is given below.

4.13.1 Error Handlers for Communicators

MPI_COMM_CREATE_ERRHANDLER(function, errhandler)

IN	function	user defined error handling procedure (function)
OUT	errhandler	MPI error handler (handle)

```
int MPI_Comm_create_errhandler(MPI_Comm_errhandler_fn *function,
                               MPI_Errhandler *errhandler)
```

MPI_COMM_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

EXTERNAL FUNCTION

INTEGER ERRHANDLER, IERROR

static MPI::Errhandler

MPI::Comm::Create_errhandler(MPI::Comm::Errhandler_fn*
function)

Creates an error handler that can be attached to communicators. This function is identical to **MPI_ERRHANDLER_CREATE**, whose use is deprecated.

The user routine should be, in C, a function of type **MPI_Comm_errhandler_fn**, which is defined as

```
typedef void MPI_Comm_errhandler_fn(MPI_Comm *, int *, ...);
```

The first argument is the communicator in use, the second is the error code to be returned. This typedef replaces **MPI_Handler_function**, whose use is deprecated.

In Fortran, the user routine should be of the form:

```
SUBROUTINE COMM_ERRHANDLER_FN(COMM, ERROR_CODE, ...)
  INTEGER COMM, ERROR_CODE
```

In C++, the user routine should be of the form:

```
typedef void MPI::Comm::Errhandler_fn(MPI::Comm &, int *, ...);
```

MPI_COMM_SET_ERRHANDLER(comm, errhandler)

INOUT	comm	communicator (handle)
IN	errhandler	new error handler for communicator (handle)

```
int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
```

MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)

INTEGER COMM, ERRHANDLER, IERROR

```
void MPI::Comm::Set_errhandler(const MPI::Errhandler& errhandler)
```

Attaches a new error handler to a communicator. The error handler must be either a predefined error handler, or an error handler created by a call to **MPI_COMM_CREATE_ERRHANDLER**. This call is identical to **MPI_ERRHANDLER_SET**, whose use is deprecated.

`MPI_COMM_GET_ERRHANDLER(comm, errhandler)`

IN	<code>comm</code>	communicator (handle)
OUT	<code>errhandler</code>	error handler currently associated with communicator (handle)

`int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)`

`MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)`

`INTEGER COMM, ERRHANDLER, IERROR`

`MPI::Errhandler MPI::Comm::Get_errhandler() const`

Retrieves the error handler currently associated with a communicator. This call is identical to `MPI_ERRHANDLER_GET`, whose use is deprecated.

4.13.2 Error Handlers for Windows

`MPI_WIN_CREATE_ERRHANDLER(function, errhandler)`

IN	<code>function</code>	user defined error handling procedure (function)
OUT	<code>errhandler</code>	MPI error handler (handle)

`int MPI_Win_create_errhandler(MPI_Win_errhandler_fn *function, MPI_Errhandler *errhandler)`

`MPI_WIN_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)`

`EXTERNAL FUNCTION`

`INTEGER ERRHANDLER, IERROR`

`static MPI::Errhandler MPI::Win::Create_errhandler(MPI::Win::Errhandler_fn* function)`

The user routine should be, in C, a function of type `MPI_Win_errhandler_fn`, which is defined as

`typedef void MPI_Win_errhandler_fn(MPI_Win *, int *, ...);`

The first argument is the window in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

`SUBROUTINE WIN_ERRHANDLER_FN(WIN, ERROR_CODE, ...)`

`INTEGER WIN, ERROR_CODE`

In C++, the user routine should be of the form:

`typedef void MPI::Win::Errhandler_fn(MPI::Win &, int *, ...);`

1 MPI_WIN_SET_ERRHANDLER(win, errhandler)

2 INOUT win window (handle)

3 IN errhandler new error handler for window (handle)

5
6 int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)

7 MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)

8 INTEGER WIN, ERRHANDLER, IERROR

9
10 void MPI::Win::Set_errhandler(const MPI::Errhandler& errhandler)

11 Attaches a new error handler to a window. The error handler must be either a pre-
12 defined error handler, or an error handler created by a call to
13 MPI_WIN_CREATE_ERRHANDLER.

15
16 MPI_WIN_GET_ERRHANDLER(win, errhandler)

17 IN win window (handle)

18 OUT errhandler error handler currently associated with window (han-
19 dle)

21
22 int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)

23 MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)

24 INTEGER WIN, ERRHANDLER, IERROR

25
26 MPI::Errhandler MPI::Win::Get_errhandler() const

27 Retrieves the error handler currently associated with a window.

28 29 4.13.3 Error Handlers for Files

31
32
33 MPI_FILE_CREATE_ERRHANDLER(function, errhandler)

34 IN function user defined error handling procedure (function)

35 OUT errhandler MPI error handler (handle)

36
37
38 int MPI_File_create_errhandler(MPI_File_errhandler_fn *function,
39 MPI_Errhandler *errhandler)

40 MPI_FILE_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)

41 EXTERNAL FUNCTION

42 INTEGER ERRHANDLER, IERROR

43
44 static MPI::Errhandler

45 MPI::File::Create_errhandler(MPI::File::Errhandler_fn*
46 function)

47
48

The user routine should be, in C, a function of type `MPI_File_errhandler_fn`, which is defined as

```
typedef void MPI_File_errhandler_fn(MPI_File *, int *, ...);
```

The first argument is the file in use, the second is the error code to be returned.

In Fortran, the user routine should be of the form:

```
SUBROUTINE FILE_ERRHANDLER_FN(FILE, ERROR_CODE, ...)
  INTEGER FILE, ERROR_CODE
```

In C++, the user routine should be of the form:

```
typedef void MPI::File::Errhandler_fn(MPI::File &, int *, ...);
```

MPI_FILE_SET_ERRHANDLER(file, errhandler)

INOUT	file	file (handle)
IN	errhandler	new error handler for file (handle)

```
int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
```

```
MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
  INTEGER FILE, ERRHANDLER, IERROR
```

```
void MPI::File::Set_errhandler(const MPI::Errhandler& errhandler)
```

Attaches a new error handler to a file. The error handler must be either a predefined error handler, or an error handler created by a call to `MPI_FILE_CREATE_ERRHANDLER`.

MPI_FILE_GET_ERRHANDLER(file, errhandler)

IN	file	file (handle)
OUT	errhandler	error handler currently associated with file (handle)

```
int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
```

```
MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
  INTEGER FILE, ERRHANDLER, IERROR
```

```
MPI::Errhandler MPI::File::Get_errhandler() const
```

Retrieves the error handler currently associated with a file.

4.14 New Datatype Manipulation Functions

New functions are provided to supplement the type manipulation functions that have address sized integer arguments. The new functions will use, in their Fortran binding, address-sized `INTEGER`s, thus solving problems currently encountered when the application address range is $> 2^{32}$. Also, a new, more convenient type constructor is provided to modify the lower bound and extent of a datatype. The deprecated functions replaced by the new functions here are listed in Section 2.6.1.

4.14.1 Type Constructors with Explicit Addresses

The four functions below supplement the four corresponding type constructor functions from MPI-1. The new functions are synonymous with the old functions in C/C++, or on Fortran systems where default INTEGERS are address sized. (The old names are not available in C++.) In Fortran, these functions accept arguments of type `INTEGER(KIND=MPI_ADDRESS_KIND)`, wherever arguments of type `MPI_Aint` are used in C. On Fortran 77 systems that do not support the Fortran 90 `KIND` notation, and where addresses are 64 bits whereas default INTEGERS are 32 bits, these arguments will be of type `INTEGER*8`. The old functions will continue to be provided for backward compatibility. However, users are encouraged to switch to the new functions, in both Fortran and C.

The new functions are listed below. The use of the old functions is deprecated.

`MPI_TYPE_CREATE_HVECTOR(count, blocklength, stride, oldtype, newtype)`

IN	<code>count</code>	number of blocks (nonnegative integer)
IN	<code>blocklength</code>	number of elements in each block (nonnegative integer)
IN	<code>stride</code>	number of bytes between start of each block (integer)
IN	<code>oldtype</code>	old datatype (handle)
OUT	<code>newtype</code>	new datatype (handle)

```
int MPI_Type_create_hvector(int count, int blocklength, MPI_Aint stride,
                           MPI_Datatype oldtype, MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_HVECTOR(COUNT, BLOCKLENGTH, STIDE, OLDTYPE, NEWTYPE, IERROR)
  INTEGER COUNT, BLOCKLENGTH, OLDTYPE, NEWTYPE, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) STRIDE
```

```
MPI::Datatype MPI::Datatype::Create_hvector(int count, int blocklength,
      MPI::Aint stride) const
```

`MPI_TYPE_CREATE_HINDEXED(count, array_of_blocklengths, array_of_displacements, oldtype, newtype)`

IN	<code>count</code>	number of blocks — also number of entries in <code>array_of_displacements</code> and <code>array_of_blocklengths</code> (integer)
IN	<code>array_of_blocklengths</code>	number of elements in each block (array of nonnegative integers)
IN	<code>array_of_displacements</code>	byte displacement of each block (array of integer)
IN	<code>oldtype</code>	old datatype (handle)
OUT	<code>newtype</code>	new datatype (handle)

```
int MPI_Type_create_hindexed(int count, int array_of_blocklengths[],
                             MPI_Aint array_of_displacements[], MPI_Datatype oldtype,
```



```

        MPI_Datatype *newtype)
1
MPI_TYPE_CREATE_HINDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS,
2
        ARRAY_OF_DISPLACEMENTS, OLDTYPE, NEWTYPE, IERROR)
3
        INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), OLDTYPE, NEWTYPE, IERROR
4
        INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
5
6
MPI::Datatype MPI::Datatype::Create_hindexed(int count,
7
        const int array_of_blocklengths[],
8
        const MPI::Aint array_of_displacements[]) const
9
10
11
MPI_TYPE_CREATE_STRUCT(count, array_of_blocklengths, array_of_displacements,
12
array_of_types, newtype)
13
14
    IN          count          number of blocks (integer) — also number of entries
15
                                in arrays array_of_types, array_of_displacements and
16
                                array_of_blocklengths
17
    IN          array_of_blocklength  number of elements in each block (array of integer)
18
    IN          array_of_displacements  byte displacement of each block (array of integer)
19
    IN          array_of_types         type of elements in each block (array of handles to
20
                                datatype objects)
21
    OUT         newtype             new datatype (handle)
22
23
24
int MPI_Type_create_struct(int count, int array_of_blocklengths[],
25
        MPI_Aint array_of_displacements[],
26
        MPI_Datatype array_of_types[], MPI_Datatype *newtype)
27
28
MPI_TYPE_CREATE_STRUCT(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
29
        ARRAY_OF_TYPES, NEWTYPE, IERROR)
30
        INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_TYPES(*), NEWTYPE,
31
        IERROR
32
        INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
33
static MPI::Datatype MPI::Datatype::Create_struct(int count,
34
        const int array_of_blocklengths[], const MPI::Aint
35
        array_of_displacements[], const MPI::Datatype array_of_types[])
36
37
38
MPI_GET_ADDRESS(location, address)
39
40
    IN          location          location in caller memory (choice)
41
    OUT         address           address of location (integer)
42
43
int MPI_Get_address(void *location, MPI_Aint *address)
44
45
MPI_GET_ADDRESS(LOCATION, ADDRESS, IERROR)
46
        <type> LOCATION(*)
47
        INTEGER IERROR
48

```

```

1      INTEGER(KIND=MPI_ADDRESS_KIND) ADDRESS
2
3      MPI::Aint MPI::Get_address(void* location)
4

```

Advice to users. Current Fortran MPI codes will run unmodified, and will port to any system. However, they may fail if addresses larger than $2^{32} - 1$ are used in the program. New codes should be written so that they use the new functions. This provides compatibility with C/C++ and avoids errors on 64 bit architectures. However, such newly written codes may need to be (slightly) rewritten to port to old Fortran 77 environments that do not support KIND declarations. (*End of advice to users.*)

4.14.2 Extent and Bounds of Datatypes

The following function replaces the three functions `MPI_TYPE_UB`, `MPI_TYPE_LB` and `MPI_TYPE_EXTENT`. It also returns address sized integers, in the Fortran binding. The use of `MPI_TYPE_UB`, `MPI_TYPE_LB` and `MPI_TYPE_EXTENT` is deprecated.

```

19      MPI_TYPE_GET_EXTENT(datatype, lb, extent)
20

```

IN	datatype	datatype to get information on (handle)
OUT	lb	lower bound of datatype (integer)
OUT	extent	extent of datatype (integer)

```

25      int MPI_Type_get_extent(MPI_Datatype datatype, MPI_Aint *lb,
26                             MPI_Aint *extent)
27

```

```

28      MPI_TYPE_GET_EXTENT(DATATYPE, LB, EXTENT, IERROR)
29

```

```

30      INTEGER DATATYPE, IERROR
31

```

```

32      INTEGER(KIND = MPI_ADDRESS_KIND) LB, EXTENT
33

```

```

34      void MPI::Datatype::Get_extent(MPI::Aint& lb, MPI::Aint& extent) const
35

```

Returns the lower bound and the extent of **datatype** (as defined by the MPI-1 standard, Section 3.12.2).

MPI allows one to change the extent of a datatype, using lower bound and upper bound markers (`MPI_LB` and `MPI_UB`). This is useful, as it allows to control the stride of successive datatypes that are replicated by datatype constructors, or are replicated by the **count** argument in a send or receive call. However, the current mechanism for achieving it is painful; also it is restrictive. `MPI_LB` and `MPI_UB` are “sticky”: once present in a datatype, they cannot be overridden (e.g., the upper bound can be moved up, by adding a new `MPI_UB` marker, but cannot be moved down below an existing `MPI_UB` marker). A new type constructor is provided to facilitate these changes. The use of `MPI_LB` and `MPI_UB` is deprecated.

MPI_TYPE_CREATE_RESIZED(oldtype, lb, extent, newtype)

IN	oldtype	input datatype (handle)
IN	lb	new lower bound of datatype (integer)
IN	extent	new extent of datatype (integer)
OUT	newtype	output datatype (handle)

```
int MPI_Type_create_resized(MPI_Datatype oldtype, MPI_Aint lb, MPI_Aint
                           extent, MPI_Datatype *newtype)
```

MPI_TYPE_CREATE_RESIZED(OLDTYPE, LB, EXTENT, NEWTYPE, IERROR)

```
INTEGER OLDTYPE, NEWTYPE, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) LB, EXTENT
```

```
MPI::Datatype MPI::Datatype::Resized(const MPI::Aint lb,
                                       const MPI::Aint extent) const
```

Returns in **newtype** a handle to a new datatype that is identical to **oldtype**, except that the lower bound of this new datatype is set to be **lb**, and its upper bound is set to be **lb** + **extent**. Any previous **lb** and **ub** markers are erased, and a new pair of lower bound and upper bound markers are put in the positions indicated by the **lb** and **extent** arguments. This affects the behavior of the datatype when used in communication operations, with **count** > 1, and when used in the construction of new derived datatypes.

Advice to users. It is strongly recommended that users use these two new functions, rather than the old MPI-1 functions to set and access lower bound, upper bound and extent of datatypes. (*End of advice to users.*)

4.14.3 True Extent of Datatypes

Suppose we implement gather as a spanning tree implemented on top of point-to-point routines. Since the receive buffer is only valid on the root process, one will need to allocate some temporary space for receiving data on intermediate nodes. However, the datatype extent cannot be used as an estimate of the amount of space that needs to be allocated, if the user has modified the extent using the **MPI_UB** and **MPI_LB** values. A new function is provided which returns the true extent of the datatype.

MPI_TYPE_GET_TRUE_EXTENT(datatype, true_lb, true_extent)

IN	datatype	datatype to get information on (handle)
OUT	true_lb	true lower bound of datatype (integer)
OUT	true_extent	true size of datatype (integer)

```
int MPI_Type_get_true_extent(MPI_Datatype datatype, MPI_Aint *true_lb,
                             MPI_Aint *true_extent)
```

MPI_TYPE_GET_TRUE_EXTENT(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)

```
INTEGER DATATYPE, IERROR
INTEGER(KIND = MPI_ADDRESS_KIND) TRUE_LB, TRUE_EXTENT
```

```

1 void MPI::Datatype::Get_true_extent(MPI::Aint& true_lb,
2                                     MPI::Aint& true_extent) const

```

`true_lb` returns the offset of the lowest unit of store which is addressed by the datatype, i.e., the lower bound of the corresponding typemap, ignoring `MPI_LB` markers. `true_extent` returns the true size of the datatype, i.e., the extent of the corresponding typemap, ignoring `MPI_LB` and `MPI_UB` markers, and performing no rounding for alignment. If the typemap associated with `datatype` is

$$Typemap = \{(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})\}$$

Then

$$true_lb(Typemap) = \min_j \{disp_j : type_j \neq lb, ub\},$$

$$true_ub(Typemap) = \max_j \{disp_j + sizeof(type_j) : type_j \neq lb, ub\},$$

and

$$true_extent(Typemap) = true_ub(Typemap) - true_lb(Typemap).$$

(Readers should compare this with the definitions in Section 3.12.3 of the MPI-1 standard, which describes the function `MPI_TYPE_EXTENT`.)

The `true_extent` is the minimum number of bytes of memory necessary to hold a datatype, uncompressed.

4.14.4 Subarray Datatype Constructor

```

28 MPI_TYPE_CREATE_SUBARRAY(ndims, array_of_sizes, array_of_subsizes, array_of_starts, or-
29 der, oldtype, newtype)

```

30	IN	<code>ndims</code>	number of array dimensions (positive integer)
31	IN	<code>array_of_sizes</code>	number of elements of type <code>oldtype</code> in each dimension of the full array (array of positive integers)
32			
33	IN	<code>array_of_subsizes</code>	number of elements of type <code>oldtype</code> in each dimension of the subarray (array of positive integers)
34			
35	IN	<code>array_of_starts</code>	starting coordinates of the subarray in each dimension (array of nonnegative integers)
36			
37	IN	<code>order</code>	array storage order flag (state)
38			
39	IN	<code>oldtype</code>	array element datatype (handle)
40			
41	OUT	<code>newtype</code>	new datatype (handle)
42			

```

43 int MPI_Type_create_subarray(int ndims, int array_of_sizes[],
44                             int array_of_subsizes[], int array_of_starts[], int order,
45                             MPI_Datatype oldtype, MPI_Datatype *newtype)
46

```

```

47 MPI_TYPE_CREATE_SUBARRAY(NDIMS, ARRAY_OF_SIZES, ARRAY_OF_SUBSIZES,
48                          ARRAY_OF_STARTS, ORDER, OLDTYPE, NEWTYPE, IERROR)

```

```

INTEGER NDIMS, ARRAY_OF_SIZES(*), ARRAY_OF_SUBSIZES(*),
ARRAY_OF_STARTS(*), ORDER, OLDTYPE, NEWTYPE, IERROR
MPI::Datatype MPI::Datatype::Create_subarray(int ndims,
      const int array_of_sizes[], const int array_of_subsizes[],
      const int array_of_starts[], int order) const

```

The subarray type constructor creates an MPI datatype describing an n-dimensional subarray of an n-dimensional array. The subarray may be situated anywhere within the full array, and may be of any nonzero size up to the size of the larger array as long as it is confined within this array. This type constructor facilitates creating filetypes to access arrays distributed in blocks among processes to a single file that contains the global array.

This type constructor can handle arrays with an arbitrary number of dimensions and works for both C and Fortran ordered matrices (i.e., row-major or column-major). Note that a C program may use Fortran order and a Fortran program may use C order.

The `ndims` parameter specifies the number of dimensions in the full data array and gives the number of elements in `array_of_sizes`, `array_of_subsizes`, and `array_of_starts`.

The number of elements of type `oldtype` in each dimension of the n-dimensional array and the requested subarray are specified by `array_of_sizes` and `array_of_subsizes`, respectively. For any dimension *i*, it is erroneous to specify `array_of_subsizes[i] < 1` or `array_of_subsizes[i] > array_of_sizes[i]`.

The `array_of_starts` contains the starting coordinates of each dimension of the subarray. Arrays are assumed to be indexed starting from zero. For any dimension *i*, it is erroneous to specify `array_of_starts[i] < 0` or `array_of_starts[i] > (array_of_sizes[i] - array_of_subsizes[i])`.

Advice to users. In a Fortran program with arrays indexed starting from 1, if the starting coordinate of a particular dimension of the subarray is *n*, then the entry in `array_of_starts` for that dimension is *n*-1. (*End of advice to users.*)

The `order` argument specifies the storage order for the subarray as well as the full array. It must be set to one of the following:

`MPI_ORDER_C` The ordering used by C arrays, (i.e., row-major order)

`MPI_ORDER_FORTRAN` The ordering used by Fortran arrays, (i.e., column-major order)

A *ndims*-dimensional subarray (`newtype`) with no extra padding can be defined by the function `Subarray()` as follows:

```

newtype = Subarray(ndims, {size0, size1, ..., sizendims-1},
      {subsize0, subsize1, ..., subsizendims-1},
      {start0, start1, ..., startndims-1}, oldtype)

```

Let the typemap of `oldtype` have the form:

$$\{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}$$

where *type_i* is a predefined MPI datatype, and let *ex* be the extent of `oldtype`. Then we define the `Subarray()` function recursively using the following three equations. Equation 4.1 defines the base step. Equation 4.2 defines the recursion step when `order = MPI_ORDER_FORTRAN`, and Equation 4.3 defines the recursion step when `order = MPI_ORDER_C`.

$$\begin{aligned}
& \text{Subarray}(1, \{size_0\}, \{subsize_0\}, \{start_0\}, \\
& \quad \{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}) \\
& = \{(\text{MPI_LB}, 0), \\
& \quad (type_0, disp_0 + start_0 \times ex), \dots, (type_{n-1}, disp_{n-1} + start_0 \times ex), \\
& \quad (type_0, disp_0 + (start_0 + 1) \times ex), \dots, (type_{n-1}, \\
& \quad \quad disp_{n-1} + (start_0 + 1) \times ex), \dots \\
& \quad (type_0, disp_0 + (start_0 + subsize_0 - 1) \times ex), \dots, \\
& \quad \quad (type_{n-1}, disp_{n-1} + (start_0 + subsize_0 - 1) \times ex), \\
& \quad (\text{MPI_UB}, size_0 \times ex)\}
\end{aligned} \tag{4.1}$$

$$\begin{aligned}
& \text{Subarray}(ndims, \{size_0, size_1, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-1}\}, \text{oldtype}) \\
& = \text{Subarray}(ndims - 1, \{size_1, size_2, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_1, subsize_2, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_1, start_2, \dots, start_{ndims-1}\}, \\
& \quad \text{Subarray}(1, \{size_0\}, \{subsize_0\}, \{start_0\}, \text{oldtype}))
\end{aligned} \tag{4.2}$$

$$\begin{aligned}
& \text{Subarray}(ndims, \{size_0, size_1, \dots, size_{ndims-1}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-1}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-1}\}, \text{oldtype}) \\
& = \text{Subarray}(ndims - 1, \{size_0, size_1, \dots, size_{ndims-2}\}, \\
& \quad \{subsize_0, subsize_1, \dots, subsize_{ndims-2}\}, \\
& \quad \{start_0, start_1, \dots, start_{ndims-2}\}, \\
& \quad \text{Subarray}(1, \{size_{ndims-1}\}, \{subsize_{ndims-1}\}, \{start_{ndims-1}\}, \text{oldtype}))
\end{aligned} \tag{4.3}$$

For an example use of `MPI_TYPE_CREATE_SUBARRAY` in the context of I/O see Section 9.9.2.

4.14.5 Distributed Array Datatype Constructor

The distributed array type constructor supports HPF-like [12] data distributions. However, unlike in HPF, the storage order may be specified for C arrays as well as for Fortran arrays.

Advice to users. One can create an HPF-like file view using this type constructor as follows. Complementary filetypes are created by having every process of a group call this constructor with identical arguments (with the exception of `rank` which should be set appropriately). These filetypes (along with identical `disp` and `etype`) are then used to define the view (via `MPI_FILE_SET_VIEW`). Using this view, a collective data access operation (with identical offsets) will yield an HPF-like distribution pattern. (*End of advice to users.*)

`MPI_TYPE_CREATE_DARRAY(size, rank, ndims, array_of_gsizes, array_of_distrib,`
`array_of_dargs, array_of_psize, order, oldtype, newtype)`

IN	<code>size</code>	size of process group (positive integer)	1
IN	<code>rank</code>	rank in process group (nonnegative integer)	2
IN	<code>ndims</code>	number of array dimensions as well as process grid dimensions (positive integer)	3
IN	<code>array_of_gsizes</code>	number of elements of type <code>oldtype</code> in each dimension of global array (array of positive integers)	4
IN	<code>array_of_distrib</code>	distribution of array in each dimension (array of state)	5
IN	<code>array_of_dargs</code>	distribution argument in each dimension (array of positive integers)	6
IN	<code>array_of_psize</code>	size of process grid in each dimension (array of positive integers)	7
IN	<code>order</code>	array storage order flag (state)	8
IN	<code>oldtype</code>	old datatype (handle)	9
OUT	<code>newtype</code>	new datatype (handle)	10

```

int MPI_Type_create_darray(int size, int rank, int ndims,
    int array_of_gsizes[], int array_of_distrib[], int
    array_of_dargs[], int array_of_psize[], int order,
    MPI_Datatype oldtype, MPI_Datatype *newtype)
MPI_TYPE_CREATE_DARRAY(SIZE, RANK, NDIMS, ARRAY_OF_GSIZES, ARRAY_OF_DISTRIBS,
    ARRAY_OF_DARGS, ARRAY_OF_PSIZE, ORDER, OLDTYPE, NEWTYPE,
    IERROR)
INTEGER SIZE, RANK, NDIMS, ARRAY_OF_GSIZES(*), ARRAY_OF_DISTRIBS(*),
    ARRAY_OF_DARGS(*), ARRAY_OF_PSIZE(*), ORDER, OLDTYPE, NEWTYPE, IERROR
MPI::Datatype MPI::Datatype::Create_darray(int size, int rank, int ndims,
    const int array_of_gsizes[], const int array_of_distrib[],
    const int array_of_dargs[], const int array_of_psize[],
    int order) const

```

`MPI_TYPE_CREATE_DARRAY` can be used to generate the datatypes corresponding to the distribution of an `ndims`-dimensional array of `oldtype` elements onto an `ndims`-dimensional grid of logical processes. Unused dimensions of `array_of_psize` should be set to 1. (See Example 4.15, page 76.) For a call to `MPI_TYPE_CREATE_DARRAY` to be correct, the equation $\prod_{i=0}^{ndims-1} array_of_psize[i] = size$ must be satisfied. The ordering of processes in the process grid is assumed to be row-major, as in the case of virtual Cartesian process topologies in `MPI-1`.

Advice to users. For both Fortran and C arrays, the ordering of processes in the process grid is assumed to be row-major. This is consistent with the ordering used in virtual Cartesian process topologies in `MPI-1`. To create such virtual process topologies, or to find the coordinates of a process in the process grid, etc., users may use the corresponding functions provided in `MPI-1`. (*End of advice to users.*)

Each dimension of the array can be distributed in one of three ways:

- `MPI_DISTRIBUTE_BLOCK` - Block distribution
- `MPI_DISTRIBUTE_CYCLIC` - Cyclic distribution
- `MPI_DISTRIBUTE_NONE` - Dimension not distributed.

The constant `MPI_DISTRIBUTE_DFLT_DARG` specifies a default distribution argument. The distribution argument for a dimension that is not distributed is ignored. For any dimension i in which the distribution is `MPI_DISTRIBUTE_BLOCK`, it erroneous to specify `array_of_dargs[i] * array_of_psize[i] < array_of_gsize[i]`.

For example, the HPF layout `ARRAY(CYCLIC(15))` corresponds to `MPI_DISTRIBUTE_CYCLIC` with a distribution argument of 15, and the HPF layout `ARRAY(BLOCK)` corresponds to `MPI_DISTRIBUTE_BLOCK` with a distribution argument of `MPI_DISTRIBUTE_DFLT_DARG`.

The `order` argument is used as in `MPI_TYPE_CREATE_SUBARRAY` to specify the storage order. Therefore, arrays described by this type constructor may be stored in Fortran (column-major) or C (row-major) order. Valid values for `order` are `MPI_ORDER_FORTRAN` and `MPI_ORDER_C`.

This routine creates a new MPI datatype with a typemap defined in terms of a function called “cyclic()” (see below).

Without loss of generality, it suffices to define the typemap for the `MPI_DISTRIBUTE_CYCLIC` case where `MPI_DISTRIBUTE_DFLT_DARG` is not used.

`MPI_DISTRIBUTE_BLOCK` and `MPI_DISTRIBUTE_NONE` can be reduced to the `MPI_DISTRIBUTE_CYCLIC` case for dimension i as follows.

`MPI_DISTRIBUTE_BLOCK` with `array_of_dargs[i]` equal to `MPI_DISTRIBUTE_DFLT_DARG` is equivalent to `MPI_DISTRIBUTE_CYCLIC` with `array_of_dargs[i]` set to

$$(\text{array_of_gsize}[i] + \text{array_of_psize}[i] - 1) / \text{array_of_psize}[i].$$

If `array_of_dargs[i]` is not `MPI_DISTRIBUTE_DFLT_DARG`, then `MPI_DISTRIBUTE_BLOCK` and `MPI_DISTRIBUTE_CYCLIC` are equivalent.

`MPI_DISTRIBUTE_NONE` is equivalent to `MPI_DISTRIBUTE_CYCLIC` with `array_of_dargs[i]` set to `array_of_gsize[i]`.

Finally, `MPI_DISTRIBUTE_CYCLIC` with `array_of_dargs[i]` equal to `MPI_DISTRIBUTE_DFLT_DARG` is equivalent to `MPI_DISTRIBUTE_CYCLIC` with `array_of_dargs[i]` set to 1.

For `MPI_ORDER_FORTRAN`, an `ndims`-dimensional distributed array (`newtype`) is defined by the following code fragment:

```
oldtype[0] = oldtype;
for ( i = 0; i < ndims; i++ ) {
    oldtype[i+1] = cyclic(array_of_dargs[i],
                        array_of_gsize[i],
                        r[i],
                        array_of_psize[i],
                        oldtype[i]);
}
newtype = oldtype[ndims];
```


For `MPI_ORDER_C`, the code is:

```

oldtype[0] = oldtype;
for ( i = 0; i < ndims; i++ ) {
    oldtype[i + 1] = cyclic(array_of_dargs[ndims - i - 1],
                           array_of_gsizes[ndims - i - 1],
                           r[ndims - i - 1],
                           array_of_psize[ndims - i - 1],
                           oldtype[i]);
}
newtype = oldtype[ndims];

```

where $r[i]$ is the position of the process (with rank `rank`) in the process grid at dimension i . The values of $r[i]$ are given by the following code fragment:

```

t_rank = rank;
t_size = 1;
for (i = 0; i < ndims; i++)
    t_size *= array_of_psize[i];
for (i = 0; i < ndims; i++) {
    t_size = t_size / array_of_psize[i];
    r[i] = t_rank / t_size;
    t_rank = t_rank % t_size;
}

```

Let the typemap of `oldtype` have the form:

$$\{(type_0, disp_0), (type_1, disp_1), \dots, (type_{n-1}, disp_{n-1})\}$$

where $type_i$ is a predefined MPI datatype, and let ex be the extent of `oldtype`.

Given the above, the function `cyclic()` is defined as follows:

```

cyclic(darg, gsize, r, psize, oldtype)
= { (MPI_LB, 0),
    (type0, disp0 + r × darg × ex), ...,
    (typen-1, dispn-1 + r × darg × ex),
    (type0, disp0 + (r × darg + 1) × ex), ...,
    (typen-1, dispn-1 + (r × darg + 1) × ex),
    ...
    (type0, disp0 + ((r + 1) × darg - 1) × ex), ...,
    (typen-1, dispn-1 + ((r + 1) × darg - 1) × ex),
    (type0, disp0 + r × darg × ex + psize × darg × ex), ...,
    (typen-1, dispn-1 + r × darg × ex + psize × darg × ex),
    (type0, disp0 + (r × darg + 1) × ex + psize × darg × ex), ...,
    (typen-1, dispn-1 + (r × darg + 1) × ex + psize × darg × ex),

```

```

1      ...
2      (type0, disp0 + ((r + 1) × darg - 1) × ex + psize × darg × ex), ...,
3      (typen-1, dispn-1 + ((r + 1) × darg - 1) × ex + psize × darg × ex),
4
5      ⋮
6      (type0, disp0 + r × darg × ex + psize × darg × ex × (count - 1)), ...,
7      (typen-1, dispn-1 + r × darg × ex + psize × darg × ex × (count - 1)),
8      (type0, disp0 + (r × darg + 1) × ex + psize × darg × ex × (count - 1)), ...,
9      (typen-1, dispn-1 + (r × darg + 1) × ex
10     + psize × darg × ex × (count - 1)),
11
12     ...
13     (type0, disp0 + (r × darg + darglast - 1) × ex
14     + psize × darg × ex × (count - 1)), ...,
15     (typen-1, dispn-1 + (r × darg + darglast - 1) × ex
16     + psize × darg × ex × (count - 1)),
17
18     (MPI_UB, gsize × ex)}
19

```

where *count* is defined by this code fragment:

```

21     nblocks = (gsizes + (darg - 1)) / darg;
22     count = nblocks / psize;
23     left_over = nblocks - count * psize;
24     if (r < left_over)
25         count = count + 1;
26

```

Here, *nblocks* is the number of blocks that must be distributed among the processors.

Finally, *darg_{last}* is defined by this code fragment:

```

27
28
29
30     if ((num_in_last_cyclic = gsize % (psize * darg)) == 0)
31         darg_last = darg;
32     else
33         darg_last = num_in_last_cyclic - darg * r;
34         if (darg_last > darg)
35             darg_last = darg;
36         if (darg_last <= 0)
37             darg_last = darg;
38

```

Example 4.15 Consider generating the filetypes corresponding to the HPF distribution:

```

39
40     <oldtype> FILEARRAY(100, 200, 300)
41     !HPF$ PROCESSORS PROCESSES(2, 3)
42     !HPF$ DISTRIBUTE FILEARRAY(CYCLIC(10), *, BLOCK) ONTO PROCESSES
43

```

This can be achieved by the following Fortran code, assuming there will be six processes attached to the run:

```

44
45
46
47     ndims = 3
48     array_of_gsizes(1) = 100

```

```

array_of_distribs(1) = MPI_DISTRIBUTE_CYCLIC
array_of_dargs(1) = 10
array_of_gsizes(2) = 200
array_of_distribs(2) = MPI_DISTRIBUTE_NONE
array_of_dargs(2) = 0
array_of_gsizes(3) = 300
array_of_distribs(3) = MPI_DISTRIBUTE_BLOCK
array_of_dargs(3) = MPI_DISTRIBUTE_DFLT_ARG
array_of_psizes(1) = 2
array_of_psizes(2) = 1
array_of_psizes(3) = 3
call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
call MPI_TYPE_CREATE_DARRAY(size, rank, ndims, array_of_gsizes, &
    array_of_distribs, array_of_dargs, array_of_psizes, &
    MPI_ORDER_FORTRAN, oldtype, newtype, ierr)

```

4.15 New Predefined Datatypes

4.15.1 Wide Characters

A new datatype, `MPI_WCHAR`, is added, for the purpose of dealing with international character sets such as Unicode.

`MPI_WCHAR` is a C type that corresponds to the type `wchar_t` defined in `<stddef.h>`. There are no predefined reduction operations for `MPI_WCHAR`.

Rationale. The fact that `MPI_CHAR` is associated with the C datatype `char`, which in turn is often used as a substitute for the “missing” `byte` datatype in C makes it most natural to define this as a new datatype specifically for multi-byte characters. (*End of rationale.*)

4.15.2 Signed Characters and Reductions

MPI-1 doesn’t allow reductions on signed or unsigned `chars`. Since this restriction (formally) prevents a C programmer from performing reduction operations on such types (which could be useful, particularly in an image processing application where pixel values are often represented as “unsigned char”), we now specify a way for such reductions to be carried out.

MPI-1.2 already has the C types `MPI_CHAR` and `MPI_UNSIGNED_CHAR`. However there is a problem here in that `MPI_CHAR` is intended to represent a character, not a small integer, and therefore will be translated between machines with different character representations.

To overcome this, a new MPI predefined datatype, `MPI_SIGNED_CHAR`, is added to the predefined datatypes of MPI-2, which corresponds to the ANSI C and ANSI C++ datatype `signed char`.

Advice to users.

The types `MPI_CHAR` and `MPI_CHARACTER` are intended for characters, and so will be translated to preserve the printable representation, rather than the bit value, if sent between machines with different character codes. The types `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` should be used in C if the integer value should be preserved.

(End of advice to users.)

The types `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` can be used in reduction operations. `MPI_CHAR` (which represents printable characters) cannot be used in reduction operations. This is an extension to MPI-1.2, since MPI-1.2 does not allow the use of `MPI_UNSIGNED_CHAR` in reduction operations (and does not have the `MPI_SIGNED_CHAR` type).

In a heterogeneous environment, `MPI_CHAR` and `MPI_WCHAR` will be translated so as to preserve the printable character, whereas `MPI_SIGNED_CHAR` and `MPI_UNSIGNED_CHAR` will be translated so as to preserve the integer value.

4.15.3 Unsigned long long Type

A new type, `MPI_UNSIGNED_LONG_LONG` in C and `MPI::UNSIGNED_LONG_LONG` in C++ is added as an optional datatype.

Rationale. The ISO C9X committee has voted to include `long long` and `unsigned long long` as standard C types. *(End of rationale.)*

4.16 Canonical MPI_PACK and MPI_UNPACK

These functions read/write data to/from the buffer in the “external32” data format specified in Section 9.5.2, and calculate the size needed for packing. Their first arguments specify the data format, for future extensibility, but for MPI-2 the only valid value of the `datarep` argument is “external32.”

Advice to users. These functions could be used, for example, to send typed data in a portable format from one MPI implementation to another. *(End of advice to users.)*

The buffer will contain exactly the packed data, without headers.

`MPI_PACK_EXTERNAL(datarep, inbuf, incount, datatype, outbuf, outsize, position)`

IN	<code>datarep</code>	data representation (string)
IN	<code>inbuf</code>	input buffer start (choice)
IN	<code>incount</code>	number of input data items (integer)
IN	<code>datatype</code>	datatype of each input data item (handle)
OUT	<code>outbuf</code>	output buffer start (choice)
IN	<code>outsize</code>	output buffer size, in bytes (integer)
INOUT	<code>position</code>	current position in buffer, in bytes (integer)

```
int MPI_Pack_external(char *datarep, void *inbuf, int incount,
                     MPI_Datatype datatype, void *outbuf, MPI_Aint outsize,
                     MPI_Aint *position)
MPI_PACK_EXTERNAL(DATAREP, INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE,
                  POSITION, IERROR)
```

```

    INTEGER INCOUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) OUTSIZE, POSITION
    CHARACTER*(*) DATAREP
    <type> INBUF(*), OUTBUF(*)

void MPI::Datatype::Pack_external(const char* datarep, const void* inbuf,
    int incount, void* outbuf, MPI::Aint outsize,
    MPI::Aint& position) const

MPI_UNPACK_EXTERNAL(datarep, inbuf, incount, datatype, outbuf, outsize, position )

IN      datarep          data representation (string)
IN      inbuf            input buffer start (choice)
IN      insize           input buffer size, in bytes (integer)
INOUT   position         current position in buffer, in bytes (integer)
OUT     outbuf           output buffer start (choice)
IN      outcount         number of output data items (integer)
IN      datatype         datatype of output data item (handle)

int MPI_Unpack_external(char *datarep, void *inbuf, MPI_Aint insize,
    MPI_Aint *position, void *outbuf, int outcount,
    MPI_Datatype datatype)

MPI_UNPACK_EXTERNAL(DATAREP, INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT,
    DATATYPE, IERROR)
    INTEGER OUTCOUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) INSIZE, POSITION
    CHARACTER*(*) DATAREP
    <type> INBUF(*), OUTBUF(*)

void MPI::Datatype::Unpack_external(const char* datarep, const void* inbuf,
    MPI::Aint insize, MPI::Aint& position, void* outbuf,
    int outcount) const

MPI_PACK_EXTERNAL_SIZE( datarep, incount, datatype, size )

IN      datarep          data representation (string)
IN      incount          number of input data items (integer)
IN      datatype         datatype of each input data item (handle)
OUT     size             output buffer size, in bytes (integer)

int MPI_Pack_external_size(char *datarep, int incount,
    MPI_Datatype datatype, MPI_Aint *size)

MPI_PACK_EXTERNAL_SIZE(DATAREP, INCOUNT, DATATYPE, SIZE, IERROR)

```

```

1      INTEGER INCOUNT, DATATYPE, IERROR
2      INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
3      CHARACTER*(*) DATAREP
4
5      MPI::Aint MPI::Datatype::Pack_external_size(const char* datarep,
6          int incount) const
7
8

```

4.17 Functions and Macros

An implementation is allowed to implement `MPI_WTIME`, `MPI_WTICK`, `PMPI_WTIME`, `PMPI_WTICK`, and the handle-conversion functions (`MPI_Group_f2c`, etc.) in Section 4.12.4, and no others, as macros in C.

Advice to implementors. Implementors should document which routines are implemented as macros. (*End of advice to implementors.*)

Advice to users. If these routines are implemented as macros, they will not work with the MPI profiling interface. (*End of advice to users.*)

4.18 Profiling Interface

The profiling interface, as described in Chapter 8 of MPI-1.1, must be supported for all MPI-2 functions, except those allowed as macros (See Section 4.17). This requires, in C and Fortran, an alternate entry point name, with the prefix `PMPI_` for each MPI function. The profiling interface in C++ is described in Section 10.1.10.

For routines implemented as macros, it is still required that the `PMPI_` version be supplied and work as expected, but it is not possible to replace at link time the `MPI_` version with a user-defined version. This is a change from MPI-1.2.

Chapter 5

Process Creation and Management

5.1 Introduction

MPI-1 provides an interface that allows processes in a parallel program to communicate with one another. MPI-1 specifies neither how the processes are created, nor how they establish communication. Moreover, an MPI-1 application is static; that is, no processes can be added to or deleted from an application after it has been started.

MPI users have asked that the MPI-1 model be extended to allow process creation and management after an MPI application has been started. A major impetus comes from the PVM [7] research effort, which has provided a wealth of experience with process management and resource control that illustrates their benefits and potential pitfalls.

The MPI Forum decided not to address resource control in MPI-2 because it was not able to design a portable interface that would be appropriate for the broad spectrum of existing and potential resource and process controllers. Resource control can encompass a wide range of abilities, including adding and deleting nodes from a virtual parallel machine, reserving and scheduling resources, managing compute partitions of an MPP, and returning information about available resources. MPI-2 assumes that resource control is provided externally — probably by computer vendors, in the case of tightly coupled systems, or by a third party software package when the environment is a cluster of workstations.

The reasons for adding process management to MPI are both technical and practical. Important classes of message passing applications require process control. These include task farms, serial applications with parallel modules, and problems that require a run-time assessment of the number and type of processes that should be started. On the practical side, users of workstation clusters who are migrating from PVM to MPI may be accustomed to using PVM's capabilities for process and resource management. The lack of these features is a practical stumbling block to migration.

While process management is essential, adding it to MPI should not compromise the portability or performance of MPI applications. In particular:

- The MPI-2 process model must apply to the vast majority of current parallel environments. These include everything from tightly integrated MPPs to heterogeneous networks of workstations.
- MPI must not take over operating system responsibilities. It should instead provide a

clean interface between an application and system software.

- MPI must continue to guarantee communication determinism, i.e., process management must not introduce unavoidable race conditions.
- MPI must not contain features that compromise performance.
- MPI-1 programs must work under MPI-2, i.e., the MPI-1 static process model must be a special case of the MPI-2 dynamic model.

The MPI-2 process management model addresses these issues in two ways. First, MPI remains primarily a communication library. It does not manage the parallel environment in which a parallel program executes, though it provides a minimal interface between an application and external resource and process managers.

Second, MPI-2 does not change the concept of communicator. Once a communicator is built, it behaves as specified in MPI-1. A communicator is never changed once created, and it is always created using deterministic collective operations.

5.2 The MPI-2 Process Model

The MPI-2 process model allows for the creation and cooperative termination of processes after an MPI application has started. It provides a mechanism to establish communication between the newly created processes and the existing MPI application. It also provides a mechanism to establish communication between two existing MPI applications, even when one did not “start” the other.

5.2.1 Starting Processes

MPI applications may start new processes through an interface to an external process manager, which can range from a parallel operating system (CMOST) to layered software (POE) to an `rsh` command (p4).

`MPI_COMM_SPAWN` starts MPI processes and establishes communication with them, returning an intercommunicator. `MPI_COMM_SPAWN_MULTIPLE` starts several different binaries (or the same binary with different arguments), placing them in the same `MPI_COMM_WORLD` and returning an intercommunicator.

MPI uses the existing group abstraction to represent processes. A process is identified by a (group, rank) pair.

5.2.2 The Runtime Environment

The `MPI_COMM_SPAWN` and `MPI_COMM_SPAWN_MULTIPLE` routines provide an interface between MPI and the *runtime environment* of an MPI application. The difficulty is that there is an enormous range of runtime environments and application requirements, and MPI must not be tailored to any particular one. Examples of such environments are:

- **MPP managed by a batch queueing system.** Batch queueing systems generally allocate resources before an application begins, enforce limits on resource use (CPU time, memory use, etc.), and do not allow a change in resource allocation after a job begins. Moreover, many MPPs have special limitations or extensions, such as a limit on the number of processes that may run on one processor, or the ability to gang-schedule processes of a parallel application.

- **Network of workstations with PVM.** PVM (Parallel Virtual Machine) allows a user to create a “virtual machine” out of a network of workstations. An application may extend the virtual machine or manage processes (create, kill, redirect output, etc.) through the PVM library. Requests to manage the machine or processes may be intercepted and handled by an external resource manager.
- **Network of workstations managed by a load balancing system.** A load balancing system may choose the location of spawned processes based on dynamic quantities, such as load average. It may transparently migrate processes from one machine to another when a resource becomes unavailable.
- **Large SMP with Unix.** Applications are run directly by the user. They are scheduled at a low level by the operating system. Processes may have special scheduling characteristics (gang-scheduling, processor affinity, deadline scheduling, processor locking, etc.) and be subject to OS resource limits (number of processes, amount of memory, etc.).

MPI assumes, implicitly, the existence of an environment in which an application runs. It does not provide “operating system” services, such as a general ability to query what processes are running, to kill arbitrary processes, to find out properties of the runtime environment (how many processors, how much memory, etc.).

Complex interaction of an MPI application with its runtime environment should be done through an environment-specific API. An example of such an API would be the PVM task and machine management routines — `pvm_addhosts`, `pvm_config`, `pvm_tasks`, etc., possibly modified to return an `MPI` (group,rank) when possible. A Condor or PBS API would be another possibility.

At some low level, obviously, MPI must be able to interact with the runtime system, but the interaction is not visible at the application level and the details of the interaction are not specified by the MPI standard.

In many cases, it is impossible to keep environment-specific information out of the MPI interface without seriously compromising MPI functionality. To permit applications to take advantage of environment-specific functionality, many MPI routines take an `info` argument that allows an application to specify environment-specific information. There is a tradeoff between functionality and portability: applications that make use of `info` are not portable.

MPI does not require the existence of an underlying “virtual machine” model, in which there is a consistent global view of an MPI application and an implicit “operating system” managing resources and processes. For instance, processes spawned by one task may not be visible to another; additional hosts added to the runtime environment by one process may not be visible in another process; tasks spawned by different processes may not be automatically distributed over available resources.

Interaction between MPI and the runtime environment is limited to the following areas:

- A process may start new processes with `MPI_COMM_SPAWN` and `MPI_COMM_SPAWN_MULTIPLE`.
- When a process spawns a child process, it may optionally use an `info` argument to tell the runtime environment where or how to start the process. This extra information may be opaque to MPI.

- An attribute `MPI_UNIVERSE_SIZE` on `MPI_COMM_WORLD` tells a program how “large” the initial runtime environment is, namely how many processes can usefully be started in all. One can subtract the size of `MPI_COMM_WORLD` from this value to find out how many processes might usefully be started in addition to those already running.

5.3 Process Manager Interface

5.3.1 Processes in MPI

A process is represented in MPI by a (group, rank) pair. A (group, rank) pair specifies a unique process but a process does not determine a unique (group, rank) pair, since a process may belong to several groups.

5.3.2 Starting Processes and Establishing Communication

The following routine starts a number of MPI processes and establishes communication with them, returning an intercommunicator.

Advice to users. It is possible in MPI to start a static SPMD or MPMD application by starting first one process and having that process start its siblings with `MPI_COMM_SPAWN`. This practice is discouraged primarily for reasons of performance. If possible, it is preferable to start all processes at once, as a single MPI-1 application. (*End of advice to users.*)

`MPI_COMM_SPAWN(command, argv, maxprocs, info, root, comm, intercomm, array_of_errcodes)`

IN	command	name of program to be spawned (string, significant only at root)
IN	argv	arguments to command (array of strings, significant only at root)
IN	maxprocs	maximum number of processes to start (integer, significant only at root)
IN	info	a set of key-value pairs telling the runtime system where and how to start the processes (handle, significant only at root)
IN	root	rank of process in which previous arguments are examined (integer)
IN	comm	intracommunicator containing group of spawning processes (handle)
OUT	intercomm	intercommunicator between original group and the newly spawned group (handle)
OUT	array_of_errcodes	one code per process (array of integer)

```
int MPI_Comm_spawn(char *command, char *argv[], int maxprocs, MPI_Info info,
                  int root, MPI_Comm comm, MPI_Comm *intercomm,
```

```

        int array_of_errcodes[])
MPI_COMM_SPAWN(COMMAND, ARGV, MAXPROCS, INFO, ROOT, COMM, INTERCOMM,
        ARRAY_OF_ERRCODES, IERROR)
        CHARACTER*(*) COMMAND, ARGV(*)
        INTEGER INFO, MAXPROCS, ROOT, COMM, INTERCOMM, ARRAY_OF_ERRCODES(*),
        IERROR
MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
        const char* argv[], int maxprocs, const MPI::Info& info,
        int root, int array_of_errcodes[]) const
MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
        const char* argv[], int maxprocs, const MPI::Info& info,
        int root) const

```

`MPI_COMM_SPAWN` tries to start `maxprocs` identical copies of the MPI program specified by `command`, establishing communication with them and returning an intercommunicator. The spawned processes are referred to as children. The children have their own `MPI_COMM_WORLD`, which is separate from that of the parents. `MPI_COMM_SPAWN` is collective over `comm`, and also may not return until `MPI_INIT` has been called in the children. Similarly, `MPI_INIT` in the children may not return until all parents have called `MPI_COMM_SPAWN`. In this sense, `MPI_COMM_SPAWN` in the parents and `MPI_INIT` in the children form a collective operation over the union of parent and child processes. The intercommunicator returned by `MPI_COMM_SPAWN` contains the parent processes in the local group and the child processes in the remote group. The ordering of processes in the local and remote groups is the same as the ordering of the group of the `comm` in the parents and of `MPI_COMM_WORLD` of the children, respectively. This intercommunicator can be obtained in the children through the function `MPI_COMM_GET_PARENT`.

Advice to users. An implementation may automatically establish communication before `MPI_INIT` is called by the children. Thus, completion of `MPI_COMM_SPAWN` in the parent does not necessarily mean that `MPI_INIT` has been called in the children (although the returned intercommunicator can be used immediately). (*End of advice to users.*)

The command argument The `command` argument is a string containing the name of a program to be spawned. The string is null-terminated in C. In Fortran, leading and trailing spaces are stripped. MPI does not specify how to find the executable or how the working directory is determined. These rules are implementation-dependent and should be appropriate for the runtime environment.

Advice to implementors. The implementation should use a natural rule for finding executables and determining working directories. For instance, a homogeneous system with a global file system might look first in the working directory of the spawning process, or might search the directories in a `PATH` environment variable as do Unix shells. An implementation on top of PVM would use PVM's rules for finding executables (usually in `$HOME/pvm3/bin/$PVM_ARCH`). An MPI implementation running under POE on an IBM SP would use POE's method of finding executables. An implementation should document its rules for finding executables and determining working

directories, and a high-quality implementation should give the user some control over these rules. (*End of advice to implementors.*)

If the program named in **command** does not call **MPI_INIT**, but instead forks a process that calls **MPI_INIT**, the results are undefined. Implementations may allow this case to work but are not required to.

Advice to users. MPI does not say what happens if the program you start is a shell script and that shell script starts a program that calls **MPI_INIT**. Though some implementations may allow you to do this, they may also have restrictions, such as requiring that arguments supplied to the shell script be supplied to the program, or requiring that certain parts of the environment not be changed. (*End of advice to users.*)

The **argv** argument **argv** is an array of strings containing arguments that are passed to the program. The first element of **argv** is the first argument passed to **command**, not, as is conventional in some contexts, the command itself. The argument list is terminated by **NULL** in C and C++ and an empty string in Fortran. In Fortran, leading and trailing spaces are always stripped, so that a string consisting of all spaces is considered an empty string. The constant **MPI_ARGV_NULL** may be used in C, C++ and Fortran to indicate an empty argument list. In C and C++, this constant is the same as **NULL**.

Example 5.1 Examples of **argv** in C and Fortran

To run the program “ocean” with arguments “-gridfile” and “ocean1.grd” in C:

```
char command[] = "ocean";
char *argv[] = {"-gridfile", "ocean1.grd", NULL};
MPI_Comm_spawn(command, argv, ...);
```

or, if not everything is known at compile time:

```
char *command;
char **argv;
command = "ocean";
argv=(char **)malloc(3 * sizeof(char *));
argv[0] = "-gridfile";
argv[1] = "ocean1.grd";
argv[2] = NULL;
MPI_Comm_spawn(command, argv, ...);
```

In Fortran:

```
CHARACTER*25 command, argv(3)
command = ' ocean '
argv(1) = ' -gridfile '
argv(2) = ' ocean1.grd'
argv(3) = ' '
call MPI_COMM_SPAWN(command, argv, ...)
```

Arguments are supplied to the program if this is allowed by the operating system. In C, the `MPI_COMM_SPAWN` argument `argv` differs from the `argv` argument of `main` in two respects. First, it is shifted by one element. Specifically, `argv[0]` of `main` is provided by the implementation and conventionally contains the name of the program (given by `command`). `argv[1]` of `main` corresponds to `argv[0]` in `MPI_COMM_SPAWN`, `argv[2]` of `main` to `argv[1]` of `MPI_COMM_SPAWN`, etc. Second, `argv` of `MPI_COMM_SPAWN` must be null-terminated, so that its length can be determined. Passing an `argv` of `MPI_ARGV_NULL` to `MPI_COMM_SPAWN` results in `main` receiving `argc` of 1 and an `argv` whose element 0 is (conventionally) the name of the program.

If a Fortran implementation supplies routines that allow a program to obtain its arguments, the arguments may be available through that mechanism. In C, if the operating system does not support arguments appearing in `argv` of `main()`, the MPI implementation may add the arguments to the `argv` that is passed to `MPI_INIT`.

The `maxprocs` argument MPI tries to spawn `maxprocs` processes. If it is unable to spawn `maxprocs` processes, it raises an error of class `MPI_ERR_SPAWN`.

An implementation may allow the `info` argument to change the default behavior, such that if the implementation is unable to spawn all `maxprocs` processes, it may spawn a smaller number of processes instead of raising an error. In principle, the `info` argument may specify an arbitrary set $\{m_i : 0 \leq m_i \leq \text{maxprocs}\}$ of allowed values for the number of processes spawned. The set $\{m_i\}$ does not necessarily include the value `maxprocs`. If an implementation is able to spawn one of these allowed numbers of processes, `MPI_COMM_SPAWN` returns successfully and the number of spawned processes, m , is given by the size of the remote group of `intercomm`. If m is less than `maxproc`, reasons why the other processes were not spawned are given in `array_of_errcodes` as described below. If it is not possible to spawn one of the allowed numbers of processes, `MPI_COMM_SPAWN` raises an error of class `MPI_ERR_SPAWN`.

A spawn call with the default behavior is called *hard*. A spawn call for which fewer than `maxprocs` processes may be returned is called *soft*. See Section 5.3.4 on page 91 for more information on the `soft` key for `info`.

Advice to users. By default, requests are hard and MPI errors are fatal. This means that by default there will be a fatal error if MPI cannot spawn all the requested processes. If you want the behavior “spawn as many processes as possible, up to N ,” you should do a soft spawn, where the set of allowed values $\{m_i\}$ is $\{0 \dots N\}$. However, this is not completely portable, as implementations are not required to support soft spawning. (*End of advice to users.*)

The `info` argument The `info` argument to all of the routines in this chapter is an opaque handle of type `MPI_Info` in C, `MPI::Info` in C++ and `INTEGER` in Fortran. It is a container for a number of user-specified (`key,value`) pairs. `key` and `value` are strings (null-terminated `char*` in C, `character*(*)` in Fortran). Routines to create and manipulate the `info` argument are described in Section 4.10 on page 43.

For the `SPAWN` calls, `info` provides additional (and possibly implementation-dependent) instructions to MPI and the runtime system on how to start processes. An application may pass `MPI_INFO_NULL` in C or Fortran. Portable programs not requiring detailed control over process locations should use `MPI_INFO_NULL`.

MPI does not specify the content of the `info` argument, except to reserve a number of special `key` values (see Section 5.3.4 on page 91). The `info` argument is quite flexible and could even be used, for example, to specify the executable and its command-line arguments. In this case the `command` argument to `MPI_COMM_SPAWN` could be empty. The ability to do this follows from the fact that MPI does not specify how an executable is found, and the `info` argument can tell the runtime system where to “find” the executable “” (empty string). Of course a program that does this will not be portable across MPI implementations.

The root argument All arguments before the `root` argument are examined only on the process whose rank in `comm` is equal to `root`. The value of these arguments on other processes is ignored.

The array_of_errcodes argument The `array_of_errcodes` is an array of length `maxprocs` in which MPI reports the status of each process that MPI was requested to start. If all `maxprocs` processes were spawned, `array_of_errcodes` is filled in with the value `MPI_SUCCESS`. If only m ($0 \leq m < \text{maxprocs}$) processes are spawned, m of the entries will contain `MPI_SUCCESS` and the rest will contain an implementation-specific error code indicating the reason MPI could not start the process. MPI does not specify which entries correspond to failed processes. An implementation may, for instance, fill in error codes in one-to-one correspondence with a detailed specification in the `info` argument. These error codes all belong to the error class `MPI_ERR_SPAWN` if there was no error in the argument list. In C or Fortran, an application may pass `MPI_ERRCODES_IGNORE` if it is not interested in the error codes. In C++ this constant does not exist, and the `array_of_errcodes` argument may be omitted from the argument list.

Advice to implementors. `MPI_ERRCODES_IGNORE` in Fortran is a special type of constant, like `MPI_BOTTOM`. See the discussion in Section 2.5.4 on page 10. (*End of advice to implementors.*)

`MPI_COMM_GET_PARENT(parent)`

OUT `parent` the parent communicator (handle)

`int MPI_Comm_get_parent(MPI_Comm *parent)`

`MPI_COMM_GET_PARENT(PARENT, IERROR)`

INTEGER `PARENT`, `IERROR`

`static MPI::Intercomm MPI::Comm::Get_parent()`

If a process was started with `MPI_COMM_SPAWN` or `MPI_COMM_SPAWN_MULTIPLE`, `MPI_COMM_GET_PARENT` returns the “parent” intercommunicator of the current process. This parent intercommunicator is created implicitly inside of `MPI_INIT` and is the same intercommunicator returned by `SPAWN` in the parents.

If the process was not spawned, `MPI_COMM_GET_PARENT` returns `MPI_COMM_NULL`.

After the parent communicator is freed or disconnected, `MPI_COMM_GET_PARENT` returns `MPI_COMM_NULL`.

Advice to users. `MPI_COMM_GET_PARENT` returns a handle to a single intercommunicator. Calling `MPI_COMM_GET_PARENT` a second time returns a handle to the same intercommunicator. Freeing the handle with `MPI_COMM_DISCONNECT` or `MPI_COMM_FREE` will cause other references to the intercommunicator to become invalid (dangling). Note that calling `MPI_COMM_FREE` on the parent communicator is not useful. (*End of advice to users.*)

Rationale. The desire of the Forum was to create a constant `MPI_COMM_PARENT` similar to `MPI_COMM_WORLD`. Unfortunately such a constant cannot be used (syntactically) as an argument to `MPI_COMM_DISCONNECT`, which is explicitly allowed. (*End of rationale.*)

5.3.3 Starting Multiple Executables and Establishing Communication

While `MPI_COMM_SPAWN` is sufficient for most cases, it does not allow the spawning of multiple binaries, or of the same binary with multiple sets of arguments. The following routine spawns multiple binaries or the same binary with multiple sets of arguments, establishing communication with them and placing them in the same `MPI_COMM_WORLD`.

`MPI_COMM_SPAWN_MULTIPLE(count, array_of_commands, array_of_argv, array_of_maxprocs, array_of_info, root, comm, intercomm, array_of_errcodes)`

IN	<code>count</code>	number of commands (positive integer, significant to MPI only at root — see advice to users)
IN	<code>array_of_commands</code>	programs to be executed (array of strings, significant only at root)
IN	<code>array_of_argv</code>	arguments for commands (array of array of strings, significant only at root)
IN	<code>array_of_maxprocs</code>	maximum number of processes to start for each command (array of integer, significant only at root)
IN	<code>array_of_info</code>	info objects telling the runtime system where and how to start processes (array of handles, significant only at root)
IN	<code>root</code>	rank of process in which previous arguments are examined (integer)
IN	<code>comm</code>	intracommunicator containing group of spawning processes (handle)
OUT	<code>intercomm</code>	intercommunicator between original group and newly spawned group (handle)
OUT	<code>array_of_errcodes</code>	one error code per process (array of integer)

```
int MPI_Comm_spawn_multiple(int count, char *array_of_commands[],
                           char **array_of_argv[], int array_of_maxprocs[],
                           MPI_Info array_of_info[], int root, MPI_Comm comm,
                           MPI_Comm *intercomm, int array_of_errcodes[])
```

```

1 MPI_COMM_SPAWN_MULTIPLE(COUNT, ARRAY_OF_COMMANDS, ARRAY_OF_ARGV,
2     ARRAY_OF_MAXPROCS, ARRAY_OF_INFO, ROOT, COMM, INTERCOMM,
3     ARRAY_OF_ERRCODES, IERROR)
4     INTEGER COUNT, ARRAY_OF_INFO(*), ARRAY_OF_MAXPROCS(*), ROOT, COMM,
5     INTERCOMM, ARRAY_OF_ERRCODES(*), IERROR
6     CHARACTER*(*) ARRAY_OF_COMMANDS(*), ARRAY_OF_ARGV(COUNT, *)
7
8 MPI::Intercomm MPI::Intracomm::Spawn_multiple(int count,
9     const char* array_of_commands[], const char** array_of_argv[],
10    const int array_of_maxprocs[], const MPI::Info array_of_info[],
11    int root, int array_of_errcodes[])
12
13 MPI::Intercomm MPI::Intracomm::Spawn_multiple(int count,
14    const char* array_of_commands[], const char** array_of_argv[],
15    const int array_of_maxprocs[], const MPI::Info array_of_info[],
16    int root)

```

`MPI_COMM_SPAWN_MULTIPLE` is identical to `MPI_COMM_SPAWN` except that there are multiple executable specifications. The first argument, `count`, gives the number of specifications. Each of the next four arguments are simply arrays of the corresponding arguments in `MPI_COMM_SPAWN`. For the Fortran version of `array_of_argv`, the element `array_of_argv(i,j)` is the j th argument to command number i .

Rationale. This may seem backwards to Fortran programmers who are familiar with Fortran's column-major ordering. However, it is necessary to do it this way to allow `MPI_COMM_SPAWN` to sort out arguments. Note that the leading dimension of `array_of_argv` must be the same as `count`. (*End of rationale.*)

Advice to users. The argument `count` is interpreted by MPI only at the root, as is `array_of_argv`. Since the leading dimension of `array_of_argv` is `count`, a non-positive value of `count` at a non-root node could theoretically cause a runtime bounds check error, even though `array_of_argv` should be ignored by the subroutine. If this happens, you should explicitly supply a reasonable value of `count` on the non-root nodes. (*End of advice to users.*)

In any language, an application may use the constant `MPI_ARGVS_NULL` (which is likely to be `(char ***)0` in C) to specify that no arguments should be passed to any commands. The effect of setting individual elements of `array_of_argv` to `MPI_ARGV_NULL` is not defined. To specify arguments for some commands but not others, the commands without arguments should have a corresponding `argv` whose first element is null (`(char *)0` in C and empty string in Fortran).

All of the spawned processes have the same `MPI_COMM_WORLD`. Their ranks in `MPI_COMM_WORLD` correspond directly to the order in which the commands are specified in `MPI_COMM_SPAWN_MULTIPLE`. Assume that m_1 processes are generated by the first command, m_2 by the second, etc. The processes corresponding to the first command have ranks $0, 1, \dots, m_1 - 1$. The processes in the second command have ranks $m_1, m_1 + 1, \dots, m_1 + m_2 - 1$. The processes in the third have ranks $m_1 + m_2, m_1 + m_2 + 1, \dots, m_1 + m_2 + m_3 - 1$, etc.

Advice to users. Calling `MPI_COMM_SPAWN` multiple times would create many sets of children with different `MPI_COMM_WORLD`s whereas

`MPI_COMM_SPAWN_MULTIPLE` creates children with a single `MPI_COMM_WORLD`, so the two methods are not completely equivalent. There are also two performance-related reasons why, if you need to spawn multiple executables, you may want to use `MPI_COMM_SPAWN_MULTIPLE` instead of calling `MPI_COMM_SPAWN` several times. First, spawning several things at once may be faster than spawning them sequentially. Second, in some implementations, communication between processes spawned at the same time may be faster than communication between processes spawned separately. (*End of advice to users.*)

The `array_of_errcodes` argument is 1-dimensional array of size $\sum_{i=1}^{count} n_i$, where n_i is the i th element of `array_of_maxprocs`. Command number i corresponds to the n_i contiguous slots in this array from element $\sum_{j=1}^{i-1} n_j$ to $\left[\sum_{j=1}^i n_j\right] - 1$. Error codes are treated as for `MPI_COMM_SPAWN`.

Example 5.2 Examples of `array_of_argv` in C and Fortran

To run the program “ocean” with arguments “-gridfile” and “ocean1.grd” and the program “atmos” with argument “atmos.grd” in C:

```
char *array_of_commands[2] = {"ocean", "atmos"};
char **array_of_argv[2];
char *argv0[] = {"-gridfile", "ocean1.grd", (char *)0};
char *argv1[] = {"atmos.grd", (char *)0};
array_of_argv[0] = argv0;
array_of_argv[1] = argv1;
MPI_Comm_spawn_multiple(2, array_of_commands, array_of_argv, ...);
```

Here’s how you do it in Fortran:

```
CHARACTER*25 commands(2), array_of_argv(2, 3)
commands(1) = ' ocean '
array_of_argv(1, 1) = ' -gridfile '
array_of_argv(1, 2) = ' ocean1.grd'
array_of_argv(1, 3) = ' '

commands(2) = ' atmos '
array_of_argv(2, 1) = ' atmos.grd '
array_of_argv(2, 2) = ' '

call MPI_COMM_SPAWN_MULTIPLE(2, commands, array_of_argv, ...)
```

5.3.4 Reserved Keys

The following keys are reserved. An implementation is not required to interpret these keys, but if it does interpret the key, it must provide the functionality described.

host Value is a hostname. The format of the hostname is determined by the implementation.

arch Value is an architecture name. Valid architecture names and what they mean are determined by the implementation.

wdir Value is the name of a directory on a machine on which the spawned process(es) execute(s). This directory is made the working directory of the executing process(es). The format of the directory name is determined by the implementation.

path Value is a directory or set of directories where the implementation should look for the executable. The format of path is determined by the implementation.

file Value is the name of a file in which additional information is specified. The format of the filename and internal format of the file are determined by the implementation.

soft Value specifies a set of numbers which are allowed values for the number of processes that `MPI_COMM_SPAWN` (et al.) may create. The format of the value is a comma-separated list of Fortran-90 triplets each of which specifies a set of integers and which together specify the set formed by the union of these sets. Negative values in this set and values greater than `maxprocs` are ignored. MPI will spawn the largest number of processes it can, consistent with some number in the set. The order in which triplets are given is not significant.

By Fortran-90 triplets, we mean:

1. `a` means a
2. `a:b` means $a, a + 1, a + 2, \dots, b$
3. `a:b:c` means $a, a + c, a + 2c, \dots, a + ck$, where for $c > 0$, k is the largest integer for which $a + ck \leq b$ and for $c < 0$, k is the largest integer for which $a + ck \geq b$. If $b > a$ then c must be positive. If $b < a$ then c must be negative.

Examples:

1. `a:b` gives a range between a and b
2. `0:N` gives full “soft” functionality
3. `1,2,4,8,16,32,64,128,256,512,1024,2048,4096` allows power-of-two number of processes.
4. `2:10000:2` allows even number of processes.
5. `2:10:2,7` allows 2, 4, 6, 7, 8, or 10 processes.

5.3.5 Spawn Example

Manager-worker Example, Using `MPI_SPAWN`.

```

/* manager */
#include "mpi.h"
int main(int argc, char *argv[])
{
    int world_size, universe_size, *universe_sizep, flag;
    MPI_Comm everyone;          /* intercommunicator */
    char worker_program[100];

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

```

```

    if (world_size != 1)      error("Top heavy with management");
                                1
                                2
    MPI_Attr_get(MPI_COMM_WORLD, MPI_UNIVERSE_SIZE,
                                3
                                &universe_sizep, &flag);
                                4
    if (!flag) {
                                5
        printf("This MPI does not support UNIVERSE_SIZE. How many\n\
processes total?");
                                6
        scanf("%d", &universe_size);
                                7
    } else universe_size = *universe_sizep;
                                8
    if (universe_size == 1) error("No room to start workers");
                                9
                                10
                                11
    /*
                                12
    * Now spawn the workers. Note that there is a run-time determination
                                13
    * of what type of worker to spawn, and presumably this calculation must
                                14
    * be done at run time and cannot be calculated before starting
                                15
    * the program. If everything is known when the application is
                                16
    * first started, it is generally better to start them all at once
                                17
    * in a single MPI_COMM_WORLD.
                                18
    */
                                19
                                20
    choose_worker_program(worker_program);
                                21
    MPI_Comm_spawn(worker_program, MPI_ARGV_NULL, universe_size-1,
                                22
                    MPI_INFO_NULL, 0, MPI_COMM_SELF, &everyone,
                                23
                    MPI_ERRCODES_IGNORE);
                                24
    /*
                                25
    * Parallel code here. The communicator "everyone" can be used
                                26
    * to communicate with the spawned processes, which have ranks 0,..
                                27
    * MPI_UNIVERSE_SIZE-1 in the remote group of the intercommunicator
                                28
    * "everyone".
                                29
    */
                                30
                                31
    MPI_Finalize();
                                32
    return 0;
                                33
}
                                34
                                35
/* worker */
                                36
                                37
#include "mpi.h"
                                38
int main(int argc, char *argv[])
                                39
{
                                40
    int size;
                                41
    MPI_Comm parent;
                                42
    MPI_Init(&argc, &argv);
                                43
    MPI_Comm_get_parent(&parent);
                                44
    if (parent == MPI_COMM_NULL) error("No parent!");
                                45
    MPI_Comm_remote_size(parent, &size);
                                46
    if (size != 1) error("Something's wrong with the parent");
                                47
                                48

```

```

1      /*
2      * Parallel code here.
3      * The manager is represented as the process with rank 0 in (the remote
4      * group of) MPI_COMM_PARENT. If the workers need to communicate among
5      * themselves, they can use MPI_COMM_WORLD.
6      */
7
8      MPI_Finalize();
9      return 0;
10     }

```

5.4 Establishing Communication

This section provides functions that establish communication between two sets of MPI processes that do not share a communicator.

Some situations in which these functions are useful are:

1. Two parts of an application that are started independently need to communicate.
2. A visualization tool wants to attach to a running process.
3. A server wants to accept connections from multiple clients. Both clients and server may be parallel programs.

In each of these situations, MPI must establish communication channels where none existed before, and there is no parent/child relationship. The routines described in this section establish communication between the two sets of processes by creating an MPI intercommunicator, where the two groups of the intercommunicator are the original sets of processes.

Establishing contact between two groups of processes that do not share an existing communicator is a collective but asymmetric process. One group of processes indicates its willingness to accept connections from other groups of processes. We will call this group the (parallel) *server*, even if this is not a client/server type of application. The other group connects to the server; we will call it the *client*.

Advice to users. While the names *client* and *server* are used throughout this section, MPI does not guarantee the traditional robustness of client server systems. The functionality described in this section is intended to allow two cooperating parts of the same application to communicate with one another. For instance, a client that gets a segmentation fault and dies, or one that doesn't participate in a collective operation may cause a server to crash or hang. (*End of advice to users.*)

5.4.1 Names, Addresses, Ports, and All That

Almost all of the complexity in MPI client/server routines addresses the question “how does the client find out how to contact the server?” The difficulty, of course, is that there is no existing communication channel between them, yet they must somehow agree on a rendezvous point where they will establish communication — Catch 22.

Agreeing on a rendezvous point always involves a third party. The third party may itself provide the rendezvous point or may communicate rendezvous information from server to client. Complicating matters might be the fact that a client doesn't really care what server it contacts, only that it be able to get in touch with one that can handle its request.

Ideally, MPI can accommodate a wide variety of run-time systems while retaining the ability to write simple portable code. The following should be compatible with MPI:

- The server resides at a well-known internet address `host:port`.
- The server prints out an address to the terminal, the user gives this address to the client program.
- The server places the address information on a nameserver, where it can be retrieved with an agreed-upon name.
- The server to which the client connects is actually a broker, acting as a middleman between the client and the real server.

MPI does not require a nameserver, so not all implementations will be able to support all of the above scenarios. However, MPI provides an optional nameserver interface, and is compatible with external name servers.

A **port_name** is a *system-supplied* string that encodes a low-level network address at which a server can be contacted. Typically this is an IP address and a port number, but an implementation is free to use any protocol. The server establishes a **port_name** with the `MPI_OPEN_PORT` routine. It accepts a connection to a given port with `MPI_COMM_ACCEPT`. A client uses **port_name** to connect to the server.

By itself, the **port_name** mechanism is completely portable, but it may be clumsy to use because of the necessity to communicate **port_name** to the client. It would be more convenient if a server could specify that it be known by an *application-supplied* **service_name** so that the client could connect to that **service_name** without knowing the **port_name**.

An MPI implementation may allow the server to publish a (**port_name**, **service_name**) pair with `MPI_PUBLISH_NAME` and the client to retrieve the port name from the service name with `MPI_LOOKUP_NAME`. This allows three levels of portability, with increasing levels of functionality.

1. Applications that do not rely on the ability to publish names are the most portable. Typically the **port_name** must be transferred "by hand" from server to client.
2. Applications that use the `MPI_PUBLISH_NAME` mechanism are completely portable among implementations that provide this service. To be portable among all implementations, these applications should have a fall-back mechanism that can be used when names are not published.
3. Applications may ignore MPI's name publishing functionality and use their own mechanism (possibly system-supplied) to publish names. This allows arbitrary flexibility but is not portable.

5.4.2 Server Routines

A server makes itself available with two routines. First it must call `MPI_OPEN_PORT` to establish a **port** at which it may be contacted. Secondly it must call `MPI_COMM_ACCEPT` to accept connections from clients.

```
1 MPI_OPEN_PORT(info, port_name)
```

```
2     IN          info          implementation-specific information on how to estab-
3                               lish an address (handle)
```

```
4     OUT        port_name      newly established port (string)
```

```
6
7 int MPI_Open_port(MPI_Info info, char *port_name)
```

```
8 MPI_OPEN_PORT(INFO, PORT_NAME, IERROR)
```

```
9     CHARACTER*(*) PORT_NAME
```

```
10    INTEGER INFO, IERROR
```

```
11
12 void MPI::Open_port(const MPI::Info& info, char* port_name)
```

13 This function establishes a network address, encoded in the **port_name** string, at which
 14 the server will be able to accept connections from clients. **port_name** is supplied by the
 15 system, possibly using information in the **info** argument.

16 MPI copies a system-supplied port name into **port_name**. **port_name** identifies the newly
 17 opened port and can be used by a client to contact the server. The maximum size string
 18 that may be supplied by the system is **MPI_MAX_PORT_NAME**.

19 *Advice to users.* The system copies the port name into **port_name**. The application
 20 must pass a buffer of sufficient size to hold this value. (*End of advice to users.*)

21 **port_name** is essentially a network address. It is unique within the communication
 22 universe to which it belongs (determined by the implementation), and may be used by any
 23 client within that communication universe. For instance, if it is an internet (host:port)
 24 address, it will be unique on the internet. If it is a low level switch address on an IBM SP,
 25 it will be unique to that SP.

26 *Advice to implementors.* These examples are not meant to constrain implementa-
 27 tions. A **port_name** could, for instance, contain a user name or the name of a batch
 28 job, as long as it is unique within some well-defined communication domain. The
 29 larger the communication domain, the more useful MPI's client/server functionality
 30 will be. (*End of advice to implementors.*)

31 The precise form of the address is implementation-defined. For instance, an internet address
 32 may be a host name or IP address, or anything that the implementation can decode into
 33 an IP address. A port name may be reused after it is freed with **MPI_CLOSE_PORT** and
 34 released by the system.

35 *Advice to implementors.* Since the user may type in **port_name** by hand, it is useful
 36 to choose a form that is easily readable and does not have embedded spaces. (*End of*
 37 *advice to implementors.*)

38 **info** may be used to tell the implementation how to establish the address. It may, and
 39 usually will, be **MPI_INFO_NULL** in order to get the implementation defaults.

```
43
44 MPI_CLOSE_PORT(port_name)
```

```
45     IN          port_name      a port (string)
```

```
46
47
48 int MPI_Close_port(char *port_name)
```

```

MPI_CLOSE_PORT(PORT_NAME, IERROR)
    CHARACTER*(*) PORT_NAME
    INTEGER IERROR

```

```

void MPI::Close_port(const char* port_name)

```

This function releases the network address represented by `port_name`.

```

MPI_COMM_ACCEPT(port_name, info, root, comm, newcomm)

```

IN	port_name	port name (string, used only on root)
IN	info	implementation-dependent information (handle, used only on root)
IN	root	rank in <code>comm</code> of root node (integer)
IN	comm	intracommunicator over which call is collective (handle)
OUT	newcomm	intercommunicator with client as remote group (handle)

```

int MPI_Comm_accept(char *port_name, MPI_Info info, int root, MPI_Comm comm,
    MPI_Comm *newcomm)

```

```

MPI_COMM_ACCEPT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
    CHARACTER*(*) PORT_NAME
    INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR

```

```

MPI::Intercomm MPI::Intracomm::Accept(const char* port_name,
    const MPI::Info& info, int root) const

```

`MPI_COMM_ACCEPT` establishes communication with a client. It is collective over the calling communicator. It returns an intercommunicator that allows communication with the client.

The `port_name` must have been established through a call to `MPI_OPEN_PORT`.

`info` is a implementation-defined string that may allow fine control over the `ACCEPT` call.

5.4.3 Client Routines

There is only one routine on the client side.

```

1  MPI_COMM_CONNECT(port_name, info, root, comm, newcomm)
2      IN          port_name          network address (string, used only on root)
3      IN          info               implementation-dependent information (handle, used
4                                     only on root)
5
6      IN          root               rank in comm of root node (integer)
7      IN          comm              intracommunicator over which call is collective (han-
8                                     dle)
9
10     OUT         newcomm            intercommunicator with server as remote group (han-
11                                     dle)

```

```

12
13  int MPI_Comm_connect(char *port_name, MPI_Info info, int root,
14                      MPI_Comm comm, MPI_Comm *newcomm)

```

```

15  MPI_COMM_CONNECT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
16      CHARACTER*(*) PORT_NAME
17      INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR

```

```

18
19  MPI::Intercomm MPI::Intracomm::Connect(const char* port_name,
20                                         const MPI::Info& info, int root) const

```

This routine establishes communication with a server specified by **port_name**. It is collective over the calling communicator and returns an intercommunicator in which the remote group participated in an **MPI_COMM_ACCEPT**.

If the named port does not exist (or has been closed), **MPI_COMM_CONNECT** raises an error of class **MPLERR_PORT**.

If the port exists, but does not have a pending **MPI_COMM_ACCEPT**, the connection attempt will eventually time out after an implementation-defined time, or succeed when the server calls **MPI_COMM_ACCEPT**. In the case of a time out, **MPI_COMM_CONNECT** raises an error of class **MPLERR_PORT**.

Advice to implementors. The time out period may be arbitrarily short or long. However, a high quality implementation will try to queue connection attempts so that a server can handle simultaneous requests from several clients. A high quality implementation may also provide a mechanism, through the **info** arguments to **MPI_OPEN_PORT**, **MPI_COMM_ACCEPT** and/or **MPI_COMM_CONNECT**, for the user to control timeout and queuing behavior. (*End of advice to implementors.*)

MPI provides no guarantee of fairness in servicing connection attempts. That is, connection attempts are not necessarily satisfied in the order they were initiated and competition from other connection attempts may prevent a particular connection attempt from being satisfied.

port_name is the address of the server. It must be the same as the name returned by **MPI_OPEN_PORT** on the server. Some freedom is allowed here. If there are equivalent forms of **port_name**, an implementation may accept them as well. For instance, if **port_name** is (hostname:port), an implementation may accept (ip_address:port) as well.

5.4.4 Name Publishing

The routines in this section provide a mechanism for publishing names. A (**service_name**, **port_name**) pair is published by the server, and may be retrieved by a client using the **service_name** only. An MPI implementation defines the *scope* of the **service_name**, that is, the domain over which the **service_name** can be retrieved. If the domain is the empty set, that is, if no client can retrieve the information, then we say that name publishing is not supported. Implementations should document how the scope is determined. High quality implementations will give some control to users through the **info** arguments to name publishing functions. Examples are given in the descriptions of individual functions.

MPI_PUBLISH_NAME(**service_name**, **info**, **port_name**)

IN	service_name	a service name to associate with the port (string)
IN	info	implementation-specific information (handle)
IN	port_name	a port name (string)

```
int MPI_Publish_name(char *service_name, MPI_Info info, char *port_name)
```

```
MPI_PUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
```

```
    INTEGER INFO, IERROR
```

```
    CHARACTER*(*) SERVICE_NAME, PORT_NAME
```

```
void MPI::Publish_name(const char* service_name, const MPI::Info& info,  
    const char* port_name)
```

This routine publishes the pair (**port_name**, **service_name**) so that an application may retrieve a system-supplied **port_name** using a well-known **service_name**.

The implementation must define the *scope* of a published service name, that is, the domain over which the service name is unique, and conversely, the domain over which the (port name, service name) pair may be retrieved. For instance, a service name may be unique to a job (where job is defined by a distributed operating system or batch scheduler), unique to a machine, or unique to a Kerberos realm. The scope may depend on the **info** argument to **MPI_PUBLISH_NAME**.

MPI permits publishing more than one **service_name** for a single **port_name**. On the other hand, if **service_name** has already been published within the scope determined by **info**, the behavior of **MPI_PUBLISH_NAME** is undefined. An MPI implementation may, through a mechanism in the **info** argument to **MPI_PUBLISH_NAME**, provide a way to allow multiple servers with the same service in the same scope. In this case, an implementation-defined policy will determine which of several port names is returned by **MPI_LOOKUP_NAME**.

Note that while **service_name** has a limited scope, determined by the implementation, **port_name** always has global scope within the communication universe used by the implementation (i.e., it is globally unique).

port_name should be the name of a port established by **MPI_OPEN_PORT** and not yet deleted by **MPI_CLOSE_PORT**. If it is not, the result is undefined.

Advice to implementors. In some cases, an MPI implementation may use a name service that a user can also access directly. In this case, a name published by MPI could easily conflict with a name published by a user. In order to avoid such conflicts,

MPI implementations should mangle service names so that they are unlikely to conflict with user code that makes use of the same service. Such name mangling will of course be completely transparent to the user.

The following situation is problematic but unavoidable, if we want to allow implementations to use nameservers. Suppose there are multiple instances of “ocean” running on a machine. If the scope of a service name is confined to a job, then multiple oceans can coexist. If an implementation provides site-wide scope, however, multiple instances are not possible as all calls to `MPI_PUBLISH_NAME` after the first may fail. There is no universal solution to this.

To handle these situations, a high quality implementation should make it possible to limit the domain over which names are published. (*End of advice to implementors.*)

```
MPI_UNPUBLISH_NAME(service_name, info, port_name)
```

IN	service_name	a service name (string)
IN	info	implementation-specific information (handle)
IN	port_name	a port name (string)

```
int MPI_Unpublish_name(char *service_name, MPI_Info info, char *port_name)
```

```
MPI_UNPUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
```

```
INTEGER INFO, IERROR
```

```
CHARACTER*(*) SERVICE_NAME, PORT_NAME
```

```
void MPI::Unpublish_name(const char* service_name, const MPI::Info& info,  
                        const char* port_name)
```

This routine unpublishes a service name that has been previously published. Attempting to unpublish a name that has not been published or has already been unpublished is erroneous and is indicated by the error class `MPI_ERR_SERVICE`.

All published names must be unpublished before the corresponding port is closed and before the publishing process exits. The behavior of `MPI_UNPUBLISH_NAME` is implementation dependent when a process tries to unpublish a name that it did not publish.

If the `info` argument was used with `MPI_PUBLISH_NAME` to tell the implementation how to publish names, the implementation may require that `info` passed to `MPI_UNPUBLISH_NAME` contain information to tell the implementation how to unpublish a name.

```
MPI_LOOKUP_NAME(service_name, info, port_name)
```

IN	service_name	a service name (string)
IN	info	implementation-specific information (handle)
OUT	port_name	a port name (string)

```
int MPI_Lookup_name(char *service_name, MPI_Info info, char *port_name)
```

```
MPI_LOOKUP_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
```

```

    CHARACTER*(*) SERVICE_NAME, PORT_NAME
    INTEGER INFO, IERROR

void MPI::Lookup_name(const char* service_name, const MPI::Info& info,
                     char* port_name)

```

This function retrieves a **port_name** published by **MPI_PUBLISH_NAME** with **service_name**. If **service_name** has not been published, it raises an error in the error class **MPI_ERR_NAME**. The application must supply a **port_name** buffer large enough to hold the largest possible port name (see discussion above under **MPI_OPEN_PORT**).

If an implementation allows multiple entries with the same **service_name** within the same scope, a particular **port_name** is chosen in a way determined by the implementation.

If the **info** argument was used with **MPI_PUBLISH_NAME** to tell the implementation how to publish names, a similar **info** argument may be required for **MPI_LOOKUP_NAME**.

5.4.5 Reserved Key Values

The following key values are reserved. An implementation is not required to interpret these key values, but if it does interpret the key value, it must provide the functionality described.

ip_port Value contains IP port number at which to establish a **port**. (Reserved for **MPI_OPEN_PORT** only).

ip_address Value contains IP address at which to establish a **port**. If the address is not a valid IP address of the host on which the **MPI_OPEN_PORT** call is made, the results are undefined. (Reserved for **MPI_OPEN_PORT** only).

5.4.6 Client/Server Examples

Simplest Example — Completely Portable.

The following example shows the simplest way to use the client/server interface. It does not use service names at all.

On the server side:

```

char myport[MPI_MAX_PORT_NAME];
MPI_Comm intercomm;
/* ... */
MPI_Open_port(MPI_INFO_NULL, myport);
printf("port name is: %s\n", myport);

MPI_Comm_accept(myport, MPI_INFO_NULL, 0, MPI_COMM_SELF, &intercomm);
/* do something with intercomm */

```

The server prints out the port name to the terminal and the user must type it in when starting up the client (assuming the MPI implementation supports stdin such that this works). On the client side:

```

MPI_Comm intercomm;
char name[MPI_MAX_PORT_NAME];
printf("enter port name: ");

```



```

        MPI_ANY_SOURCE, MPI_ANY_TAG, client, &status );
switch (status.MPI_TAG) {
    case 0: MPI_Comm_free( &client );
            MPI_Close_port(port_name);
            MPI_Finalize();
            return 0;
    case 1: MPI_Comm_disconnect( &client );
            again = 0;
            break;
    case 2: /* do something */
    ...
    default:
        /* Unexpected message type */
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
}
}
}

```

Here is the client.

```

#include "mpi.h"
int main( int argc, char **argv )
{
    MPI_Comm server;
    double buf[MAX_DATA];
    char port_name[MPI_MAX_PORT_NAME];

    MPI_Init( &argc, &argv );
    strcpy(port_name, argv[1] );/* assume server's name is cmd-line arg */

    MPI_Comm_connect( port_name, MPI_INFO_NULL, 0, MPI_COMM_WORLD,
                      &server );

    while (!done) {
        tag = 2; /* Action to perform */
        MPI_Send( buf, n, MPI_DOUBLE, 0, tag, server );
        /* etc */
    }
    MPI_Send( buf, 0, MPI_DOUBLE, 0, 1, server );
    MPI_Comm_disconnect( &server );
    MPI_Finalize();
    return 0;
}

```

5.5 Other Functionality

5.5.1 Universe Size

Many “dynamic” MPI applications are expected to exist in a static runtime environment, in which resources have been allocated before the application is run. When a user (or possibly a batch system) runs one of these quasi-static applications, she will usually specify a number of processes to start and a total number of processes that are expected. An application simply needs to know how many slots there are, i.e., how many processes it should spawn.

MPI provides an attribute on `MPI_COMM_WORLD`, `MPI_UNIVERSE_SIZE`, that allows the application to obtain this information in a portable manner. This attribute indicates the total number of processes that are expected. In Fortran, the attribute is the integer value. In C, the attribute is a pointer to the integer value. An application typically subtracts the size of `MPI_COMM_WORLD` from `MPI_UNIVERSE_SIZE` to find out how many processes it should spawn. `MPI_UNIVERSE_SIZE` is initialized in `MPI_INIT` and is not changed by MPI. If defined, it has the same value on all processes of `MPI_COMM_WORLD`. `MPI_UNIVERSE_SIZE` is determined by the application startup mechanism in a way not specified by MPI. (The size of `MPI_COMM_WORLD` is another example of such a parameter.)

Possibilities for how `MPI_UNIVERSE_SIZE` might be set include

- A `-universe_size` argument to a program that starts MPI processes.
- Automatic interaction with a batch scheduler to figure out how many processors have been allocated to an application.
- An environment variable set by the user.
- Extra information passed to `MPI_COMM_SPAWN` through the `info` argument.

An implementation must document how `MPI_UNIVERSE_SIZE` is set. An implementation may not support the ability to set `MPI_UNIVERSE_SIZE`, in which case the attribute `MPI_UNIVERSE_SIZE` is not set.

`MPI_UNIVERSE_SIZE` is a recommendation, not necessarily a hard limit. For instance, some implementations may allow an application to spawn 50 processes per processor, if they are requested. However, it is likely that the user only wants to spawn one process per processor.

`MPI_UNIVERSE_SIZE` is assumed to have been specified when an application was started, and is in essence a portable mechanism to allow the user to pass to the application (through the MPI process startup mechanism, such as `mpirexec`) a piece of critical runtime information. Note that no interaction with the runtime environment is required. If the runtime environment changes size while an application is running, `MPI_UNIVERSE_SIZE` is not updated, and the application must find out about the change through direct communication with the runtime system.

5.5.2 Singleton MPI_INIT

A high-quality implementation will allow any process (including those not started with a “parallel application” mechanism) to become an MPI process by calling `MPI_INIT`. Such a process can then connect to other MPI processes using the `MPI_COMM_ACCEPT` and

`MPI_COMM_CONNECT` routines, or spawn other MPI processes. MPI does not mandate this behavior, but strongly encourages it where technically feasible.

Advice to implementors. To start an MPI-1 application with more than one process requires some special coordination. The processes must be started at the “same” time, they must have a mechanism to establish communication, etc. Either the user or the operating system must take special steps beyond simply starting processes.

When an application enters `MPI_INIT`, clearly it must be able to determine if these special steps were taken. MPI-1 does not say what happens if these special steps were not taken — presumably this is treated as an error in starting the MPI application. MPI-2 recommends the following behavior.

If a process enters `MPI_INIT` and determines that no special steps were taken (i.e., it has not been given the information to form an `MPI_COMM_WORLD` with other processes) it succeeds and forms a singleton MPI program, that is, one in which `MPI_COMM_WORLD` has size 1.

In some implementations, MPI may not be able to function without an “MPI environment.” For example, MPI may require that daemons be running or MPI may not be able to work at all on the front-end of an MPP. In this case, an MPI implementation may either

1. Create the environment (e.g., start a daemon) or
2. Raise an error if it cannot create the environment and the environment has not been started independently.

A high quality implementation will try to create a singleton MPI process and not raise an error.

(End of advice to implementors.)

5.5.3 `MPI_APPNUM`

There is a predefined attribute `MPI_APPNUM` of `MPI_COMM_WORLD`. In Fortran, the attribute is an integer value. In C, the attribute is a pointer to an integer value. If a process was spawned with `MPI_COMM_SPAWN_MULTIPLE`, `MPI_APPNUM` is the command number that generated the current process. Numbering starts from zero. If a process was spawned with `MPI_COMM_SPAWN`, it will have `MPI_APPNUM` equal to zero.

Additionally, if the process was not started by a spawn call, but by an implementation-specific startup mechanism that can handle multiple process specifications, `MPI_APPNUM` should be set to the number of the corresponding process specification. In particular, if it is started with

```
mpiexec spec0 [: spec1 : spec2 : ...]
```

`MPI_APPNUM` should be set to the number of the corresponding specification.

If an application was not spawned with `MPI_COMM_SPAWN` or `MPI_COMM_SPAWN_MULTIPLE`, and `MPI_APPNUM` doesn’t make sense in the context of the implementation-specific startup mechanism, `MPI_APPNUM` is not set.

MPI implementations may optionally provide a mechanism to override the value of `MPI_APPNUM` through the `info` argument. MPI reserves the following key for all `SPAWN` calls.

1 **appnum** Value contains an integer that overrides the default value for **MPI_APPNUM** in the
 2 child.

3
 4 *Rationale.* When a single application is started, it is able to figure out how many pro-
 5 cesses there are by looking at the size of **MPI_COMM_WORLD**. An application consisting
 6 of multiple SPMD sub-applications has no way to find out how many sub-applications
 7 there are and to which sub-application the process belongs. While there are ways to
 8 figure it out in special cases, there is no general mechanism. **MPI_APPNUM** provides
 9 such a general mechanism. (*End of rationale.*)

10 5.5.4 Releasing Connections

11
 12 Before a client and server connect, they are independent MPI applications. An error in one
 13 does not affect the other. After establishing a connection with **MPI_COMM_CONNECT** and
 14 **MPI_COMM_ACCEPT**, an error in one may affect the other. It is desirable for a client and
 15 server to be able to disconnect, so that an error in one will not affect the other. Similarly,
 16 it might be desirable for a parent and child to disconnect, so that errors in the child do not
 17 affect the parent, or vice-versa.

- 18
 19 • Two processes are **connected** if there is a communication path (direct or indirect)
 20 between them. More precisely:
 - 21 1. Two processes are connected if
 - 22 (a) they both belong to the same communicator (inter- or intra-, including
 - 23 **MPI_COMM_WORLD**) *or*
 - 24 (b) they have previously belonged to a communicator that was freed with
 - 25 **MPI_COMM_FREE** instead of **MPI_COMM_DISCONNECT** *or*
 - 26 (c) they both belong to the group of the same window or filehandle.
 - 27
 - 28 2. If A is connected to B and B to C, then A is connected to C.
- 29
 30 • Two processes are **disconnected** (also **independent**) if they are not connected.
- 31
 32 • By the above definitions, connectivity is a transitive property, and divides the uni-
 33 verse of MPI processes into disconnected (independent) sets (equivalence classes) of
 34 processes.
- 35
 36 • Processes which are connected, but don't share the same **MPI_COMM_WORLD** may
 37 become disconnected (independent) if the communication path between them is bro-
 38 ken by using **MPI_COMM_DISCONNECT**.

39 The following additional rules apply to MPI-1 functions:

- 40
 41 • **MPI_FINALIZE** is collective over a set of connected processes.
- 42
 43 • **MPI_ABORT** does not abort independent processes. As in MPI-1, it may abort all
 44 processes in **MPI_COMM_WORLD** (ignoring its **comm** argument). Additionally, it
 45 may abort connected processes as well, though it makes a “best attempt” to abort
 46 only the processes in **comm**.
- 47
 48 • If a process terminates without calling **MPI_FINALIZE**, independent processes are not
 affected but the effect on connected processes is not defined.

`MPI_COMM_DISCONNECT(comm)`

INOUT `comm` communicator (handle)

`int MPI_Comm_disconnect(MPI_Comm *comm)`

`MPI_COMM_DISCONNECT(COMM, IERROR)`

INTEGER `COMM`, `IERROR`

`void MPI::Comm::Disconnect()`

This function waits for all pending communication on `comm` to complete internally, deallocates the communicator object, and sets the handle to `MPI_COMM_NULL`. It is a collective operation.

It may not be called with the communicator `MPI_COMM_WORLD` or `MPI_COMM_SELF`.

`MPI_COMM_DISCONNECT` may be called only if all communication is complete and matched, so that buffered data can be delivered to its destination. This requirement is the same as for `MPI_FINALIZE`.

`MPI_COMM_DISCONNECT` has the same action as `MPI_COMM_FREE`, except that it waits for pending communication to finish internally and enables the guarantee about the behavior of disconnected processes.

Advice to users. To disconnect two processes you may need to call `MPI_COMM_DISCONNECT`, `MPI_WIN_FREE` and `MPI_FILE_CLOSE` to remove all communication paths between the two processes. Notes that it may be necessary to disconnect several communicators (or to free several windows or files) before two processes are completely independent. (*End of advice to users.*)

Rationale. It would be nice to be able to use `MPI_COMM_FREE` instead, but that function explicitly does not wait for pending communication to complete. (*End of rationale.*)

5.5.5 Another Way to Establish MPI Communication

`MPI_COMM_JOIN(fd, intercomm)`

IN `fd` socket file descriptor

OUT `intercomm` new intercommunicator (handle)

`int MPI_Comm_join(int fd, MPI_Comm *intercomm)`

`MPI_COMM_JOIN(FD, INTERCOMM, IERROR)`

INTEGER `FD`, `INTERCOMM`, `IERROR`

`static MPI::Intercomm MPI::Comm::Join(const int fd)`

`MPI_COMM_JOIN` is intended for MPI implementations that exist in an environment supporting the Berkeley Socket interface [14, 17]. Implementations that exist in an environment not supporting Berkeley Sockets should provide the entry point for `MPI_COMM_JOIN` and should return `MPI_COMM_NULL`.

This call creates an intercommunicator from the union of two MPI processes which are connected by a socket. `MPI_COMM_JOIN` should normally succeed if the local and remote processes have access to the same implementation-defined MPI communication universe.

Advice to users. An MPI implementation may require a specific communication medium for MPI communication, such as a shared memory segment or a special switch. In this case, it may not be possible for two processes to successfully join even if there is a socket connecting them and they are using the same MPI implementation. (*End of advice to users.*)

Advice to implementors. A high quality implementation will attempt to establish communication over a slow medium if its preferred one is not available. If implementations do not do this, they must document why they cannot do MPI communication over the medium used by the socket (especially if the socket is a TCP connection). (*End of advice to implementors.*)

`fd` is a file descriptor representing a socket of type `SOCK_STREAM` (a two-way reliable byte-stream connection). Non-blocking I/O and asynchronous notification via `SIGIO` must not be enabled for the socket. The socket must be in a connected state. The socket must be quiescent when `MPI_COMM_JOIN` is called (see below). It is the responsibility of the application to create the socket using standard socket API calls.

`MPI_COMM_JOIN` must be called by the process at each end of the socket. It does not return until both processes have called `MPI_COMM_JOIN`. The two processes are referred to as the local and remote processes.

MPI uses the socket to bootstrap creation of the intercommunicator, and for nothing else. Upon return from `MPI_COMM_JOIN`, the file descriptor will be open and quiescent (see below).

If MPI is unable to create an intercommunicator, but is able to leave the socket in its original state, with no pending communication, it succeeds and sets `intercomm` to `MPI_COMM_NULL`.

The socket must be quiescent before `MPI_COMM_JOIN` is called and after `MPI_COMM_JOIN` returns. More specifically, on entry to `MPI_COMM_JOIN`, a `read` on the socket will not read any data that was written to the socket before the remote process called `MPI_COMM_JOIN`. On exit from `MPI_COMM_JOIN`, a `read` will not read any data that was written to the socket before the remote process returned from `MPI_COMM_JOIN`. It is the responsibility of the application to ensure the first condition, and the responsibility of the MPI implementation to ensure the second. In a multithreaded application, the application must ensure that one thread does not access the socket while another is calling `MPI_COMM_JOIN`, or call `MPI_COMM_JOIN` concurrently.

Advice to implementors. MPI is free to use any available communication path(s) for MPI messages in the new communicator; the socket is only used for the initial handshaking. (*End of advice to implementors.*)

`MPI_COMM_JOIN` uses non-MPI communication to do its work. The interaction of non-MPI communication with pending MPI communication is not defined. Therefore, the result of calling `MPI_COMM_JOIN` on two connected processes (see Section 5.5.4 on page 106 for the definition of connected) is undefined.

The returned communicator may be used to establish MPI communication with additional processes, through the usual MPI communicator creation mechanisms.

Chapter 6

One-Sided Communications

6.1 Introduction

Remote Memory Access (RMA) extends the communication mechanisms of MPI by allowing one process to specify all communication parameters, both for the sending side and for the receiving side. This mode of communication facilitates the coding of some applications with dynamically changing data access patterns where the data distribution is fixed or slowly changing. In such a case, each process can compute what data it needs to access or update at other processes. However, processes may not know which data in their own memory need to be accessed or updated by remote processes, and may not even know the identity of these processes. Thus, the transfer parameters are all available only on one side. Regular send/receive communication requires matching operations by sender and receiver. In order to issue the matching operations, an application needs to distribute the transfer parameters. This may require all processes to participate in a time consuming global computation, or to periodically poll for potential communication requests to receive and act upon. The use of RMA communication mechanisms avoids the need for global computations or explicit polling. A generic example of this nature is the execution of an assignment of the form $\mathbf{A} = \mathbf{B}(\text{map})$, where **map** is a permutation vector, and **A**, **B** and **map** are distributed in the same manner.

Message-passing communication achieves two effects: *communication* of data from sender to receiver; and *synchronization* of sender with receiver. The RMA design separates these two functions. Three communication calls are provided: **MPI_PUT** (remote write), **MPI_GET** (remote read) and **MPI_ACCUMULATE** (remote update). A larger number of synchronization calls are provided that support different synchronization styles. The design is similar to that of weakly coherent memory systems: correct ordering of memory accesses has to be imposed by the user, using synchronization calls; the implementation can delay communication operations until the synchronization calls occur, for efficiency.

The design of the RMA functions allows implementors to take advantage, in many cases, of fast communication mechanisms provided by various platforms, such as coherent or noncoherent shared memory, DMA engines, hardware-supported put/get operations, communication coprocessors, etc. The most frequently used RMA communication mechanisms can be layered on top of message passing. However, support for asynchronous communication agents (handlers, threads, etc.) is needed, for certain RMA functions, in a distributed memory environment.

We shall denote by **origin** the process that performs the call, and by **target** the

process in which the memory is accessed. Thus, in a put operation, source=origin and destination=target; in a get operation, source=target and destination=origin.

6.2 Initialization

6.2.1 Window Creation

The initialization operation allows each process in an intracommunicator group to specify, in a collective operation, a “window” in its memory that is made accessible to accesses by remote processes. The call returns an opaque object that represents the group of processes that own and access the set of windows, and the attributes of each window, as specified by the initialization call.

MPI_WIN_CREATE(base, size, disp_unit, info, comm, win)

IN	base	initial address of window (choice)
IN	size	size of window in bytes (nonnegative integer)
IN	disp_unit	local unit size for displacements, in bytes (positive integer)
IN	info	info argument (handle)
IN	comm	communicator (handle)
OUT	win	window object returned by the call (handle)

```
int MPI_Win_create(void *base, MPI_Aint size, int disp_unit, MPI_Info info,
                  MPI_Comm comm, MPI_Win *win)
```

```
MPI_WIN_CREATE(BASE, SIZE, DISP_UNIT, INFO, COMM, WIN, IERROR)
```

```
<type> BASE(*)
```

```
INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
```

```
INTEGER DISP_UNIT, INFO, COMM, WIN, IERROR
```

```
static MPI::Win MPI::Win::Create(const void* base, MPI::Aint size, int
                                disp_unit, const MPI::Info& info, const MPI::Intracomm& comm)
```

This is a collective call executed by all processes in the group of **comm**. It returns a window object that can be used by these processes to perform RMA operations. Each process specifies a window of existing memory that it exposes to RMA accesses by the processes in the group of **comm**. The window consists of **size** bytes, starting at address **base**. A process may elect to expose no memory by specifying **size** = 0.

The displacement unit argument is provided to facilitate address arithmetic in RMA operations: the target displacement argument of an RMA operation is scaled by the factor **disp_unit** specified by the target process, at window creation.

Rationale. The window size is specified using an address sized integer, so as to allow windows that span more than 4 GB of address space. (Even if the physical memory size is less than 4 GB, the address range may be larger than 4 GB, if addresses are not contiguous.) (*End of rationale.*)

Advice to users. Common choices for `disp_unit` are 1 (no scaling), and (in C syntax) `sizeof(type)`, for a window that consists of an array of elements of type `type`. The later choice will allow one to use array indices in RMA calls, and have those scaled correctly to byte displacements, even in a heterogeneous environment. (*End of advice to users.*)

The `info` argument provides optimization hints to the runtime about the expected usage pattern of the window. The following `info` key is predefined:

`no_locks` — if set to true, then the implementation may assume that the local window is never locked (by a call to `MPI_WIN_LOCK`). This implies that this window is not used for 3-party communication, and RMA can be implemented with no (less) asynchronous agent activity at this process.

The various processes in the group of `comm` may specify completely different target windows, in location, size, displacement units and info arguments. As long as all the get, put and accumulate accesses to a particular process fit their specific target window this should pose no problem. The same area in memory may appear in multiple windows, each associated with a different window object. However, concurrent communications to distinct, overlapping windows may lead to erroneous results.

Advice to users. A window can be created in any part of the process memory. However, on some systems, the performance of windows in memory allocated by `MPI_ALLOC_MEM` (Section 4.11, page 47) will be better. Also, on some systems, performance is improved when window boundaries are aligned at “natural” boundaries (word, double-word, cache line, page frame, etc.). (*End of advice to users.*)

Advice to implementors. In cases where RMA operations use different mechanisms in different memory areas (e.g., load/store in a shared memory segment, and an asynchronous handler in private memory), the `MPI_WIN_CREATE` call needs to figure out which type of memory is used for the window. To do so, MPI maintains, internally, the list of memory segments allocated by `MPI_ALLOC_MEM`, or by other, implementation specific, mechanisms, together with information on the type of memory segment allocated. When a call to `MPI_WIN_CREATE` occurs, then MPI checks which segment contains each window, and decides, accordingly, which mechanism to use for RMA operations.

Vendors may provide additional, implementation-specific mechanisms to allow “good” memory to be used for static variables.

Implementors should document any performance impact of window alignment. (*End of advice to implementors.*)

MPI_WIN_FREE(win)

INOUT	win	window object (handle)
-------	-----	------------------------

```
int MPI_Win_free(MPI_Win *win)
```

```
MPI_WIN_FREE(WIN, IERROR)
```

```

1      INTEGER WIN, IERROR
2
3      void MPI::Win::Free()

```

Frees the window object `win` and returns a null handle (equal to `MPI_WIN_NULL`). This is a collective call executed by all processes in the group associated with `win`. `MPI_WIN_FREE(win)` can be invoked by a process only after it has completed its involvement in RMA communications on window `win`: i.e., the process has called `MPI_WIN_FENCE`, or called `MPI_WIN_WAIT` to match a previous call to `MPI_WIN_POST` or called `MPI_WIN_COMPLETE` to match a previous call to `MPI_WIN_START` or called `MPI_WIN_UNLOCK` to match a previous call to `MPI_WIN_LOCK`. When the call returns, the window memory can be freed.

Advice to implementors. `MPI_WIN_FREE` requires a barrier synchronization: no process can return from free until all processes in the group of `win` called free. This, to ensure that no process will attempt to access a remote window (e.g., with lock/unlock) after it was freed. (*End of advice to implementors.*)

6.2.2 Window Attributes

The following three attributes are cached with a window, when the window is created.

<code>MPI_WIN_BASE</code>	window base address.
<code>MPI_WIN_SIZE</code>	window size, in bytes.
<code>MPI_WIN_DISP_UNIT</code>	displacement unit associated with the window.

In C, calls to `MPI_Win_get_attr(win, MPI_WIN_BASE, &base, &flag)`, `MPI_Win_get_attr(win, MPI_WIN_SIZE, &size, &flag)` and `MPI_Win_get_attr(win, MPI_WIN_DISP_UNIT, &disp_unit, &flag)` will return in `base` a pointer to the start of the window `win`, and will return in `size` and `disp_unit` pointers to the size and displacement unit of the window, respectively. And similarly, in C++.

In Fortran, calls to `MPI_WIN_GET_ATTR(win, MPI_WIN_BASE, base, flag, ierror)`, `MPI_WIN_GET_ATTR(win, MPI_WIN_SIZE, size, flag, ierror)` and `MPI_WIN_GET_ATTR(win, MPI_WIN_DISP_UNIT, disp_unit, flag, ierror)` will return in `base`, `size` and `disp_unit` the (integer representation of) the base address, the size and the displacement unit of the window `win`, respectively. (The window attribute access functions are defined in Section 8.8, page 198.)

The other “window attribute,” namely the group of processes attached to the window, can be retrieved using the call below.

```

39      MPI_WIN_GET_GROUP(win, group)

```

IN	<code>win</code>	window object (handle)
OUT	<code>group</code>	group of processes which share access to the window (handle)

```

45      int MPI_Win_get_group(MPI_Win win, MPI_Group *group)

```

```

46      MPI_WIN_GET_GROUP(WIN, GROUP, IERROR)
47
48      INTEGER WIN, GROUP, IERROR

```

`MPI::Group MPI::Win::Get_group() const`

`MPI_WIN_GET_GROUP` returns a duplicate of the group of the communicator used to create the window. associated with `win`. The group is returned in `group`.

6.3 Communication Calls

MPI supports three RMA communication calls: `MPI_PUT` transfers data from the caller memory (origin) to the target memory; `MPI_GET` transfers data from the target memory to the caller memory; and `MPI_ACCUMULATE` updates locations in the target memory, e.g. by adding to these locations values sent from the caller memory. These operations are *nonblocking*: the call initiates the transfer, but the transfer may continue after the call returns. The transfer is completed, both at the origin and at the target, when a subsequent *synchronization* call is issued by the caller on the involved window object. These synchronization calls are described in Section 6.4, page 121.

The local communication buffer of an RMA call should not be updated, and the local communication buffer of a get call should not be accessed after the RMA call, until the subsequent synchronization call completes.

Rationale. The rule above is more lenient than for message passing, where we do not allow two concurrent sends, with overlapping send buffers. Here, we allow two concurrent puts with overlapping send buffers. The reasons for this relaxation are

1. Users do not like that restriction, which is not very natural (it prohibits concurrent reads).
2. Weakening the rule does not prevent efficient implementation, as far as we know.
3. Weakening the rule is important for performance of RMA: we want to associate one synchronization call with as many RMA operations is possible. If puts from overlapping buffers cannot be concurrent, then we need to needlessly add synchronization points in the code.

(End of rationale.)

It is erroneous to have concurrent conflicting accesses to the same memory location in a window; if a location is updated by a put or accumulate operation, then this location cannot be accessed by a load or another RMA operation until the updating operation has completed at the target. There is one exception to this rule; namely, the same location can be updated by several concurrent accumulate calls, the outcome being as if these updates occurred in some order. In addition, a window cannot concurrently be updated by a put or accumulate operation and by a local store operation. This, even if these two updates access different locations in the window. The last restriction enables more efficient implementations of RMA operations on many systems. These restrictions are described in more detail in Section 6.7, page 137.

The calls use general datatype arguments to specify communication buffers at the origin and at the target. Thus, a transfer operation may also gather data at the source and scatter it at the destination. However, all arguments specifying both communication buffers are provided by the caller.

For all three calls, the target process may be identical with the origin process; i.e., a process may use an RMA operation to move data in its memory.

Rationale. The choice of supporting “self-communication” is the same as for message passing. It simplifies some coding, and is very useful with accumulate operations, to allow atomic updates of local variables. (*End of rationale.*)

6.3.1 Put

The execution of a put operation is similar to the execution of a send by the origin process and a matching receive by the target process. The obvious difference is that all arguments are provided by one call — the call executed by the origin process.

MPI_PUT(*origin_addr*, *origin_count*, *origin_datatype*, *target_rank*, *target_disp*, *target_count*, *target_datatype*, *win*)

IN	origin_addr	initial address of origin buffer (choice)
IN	origin_count	number of entries in origin buffer (nonnegative integer)
IN	origin_datatype	datatype of each entry in origin buffer (handle)
IN	target_rank	rank of target (nonnegative integer)
IN	target_disp	displacement from start of window to target buffer (nonnegative integer)
IN	target_count	number of entries in target buffer (nonnegative integer)
IN	target_datatype	datatype of each entry in target buffer (handle)
IN	win	window object used for communication (handle)

```

int MPI_Put(void *origin_addr, int origin_count, MPI_Datatype
            origin_datatype, int target_rank, MPI_Aint target_disp, int
            target_count, MPI_Datatype target_datatype, MPI_Win win)

MPI_PUT(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP,
        TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR)
    <type> ORIGIN_ADDR(*)
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
    TARGET_DATATYPE, WIN, IERROR

void MPI::Win::Put(const void* origin_addr, int origin_count, const
                  MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
                  target_disp, int target_count, const MPI::Datatype&
                  target_datatype) const

```

Transfers *origin_count* successive entries of the type specified by the *origin_datatype*, starting at address *origin_addr* on the origin node to the target node specified by the *win*, *target_rank* pair. The data are written in the target buffer at address *target_addr* = *window_base* + *target_disp* × *disp_unit*, where *window_base* and *disp_unit* are the base address and window displacement unit specified at window initialization, by the target process.

The target buffer is specified by the arguments *target_count* and *target_datatype*.

The data transfer is the same as that which would occur if the origin process executed a send operation with arguments `origin_addr`, `origin_count`, `origin_datatype`, `target_rank`, `tag`, `comm`, and the target process executed a receive operation with arguments `target_addr`, `target_count`, `target_datatype`, `source`, `tag`, `comm`, where `target_addr` is the target buffer address computed as explained above, and `comm` is a communicator for the group of `win`.

The communication must satisfy the same constraints as for a similar message-passing communication. The `target_datatype` may not specify overlapping entries in the target buffer. The message sent must fit, without truncation, in the target buffer. Furthermore, the target buffer must fit in the target window.

The `target_datatype` argument is a handle to a datatype object defined at the origin process. However, this object is interpreted at the target process: the outcome is as if the target datatype object was defined at the target process, by the same sequence of calls used to define it at the origin process. The target datatype must contain only relative displacements, not absolute addresses. The same holds for `get` and `accumulate`.

Advice to users. The `target_datatype` argument is a handle to a datatype object that is defined at the origin process, even though it defines a data layout in the target process memory. This causes no problems in a homogeneous environment, or in a heterogeneous environment, if only portable datatypes are used (portable datatypes are defined in Section 2.4, page 7).

The performance of a `put` transfer can be significantly affected, on some systems, from the choice of window location and the shape and location of the origin and target buffer: transfers to a target window in memory allocated by `MPI_ALLOC_MEM` may be much faster on shared memory systems; transfers from contiguous buffers will be faster on most, if not all, systems; the alignment of the communication buffers may also impact performance. (*End of advice to users.*)

Advice to implementors. A high quality implementation will attempt to prevent remote accesses to memory outside the window that was exposed by the process. This, both for debugging purposes, and for protection with client-server codes that use RMA. I.e., a high-quality implementation will check, if possible, window bounds on each RMA call, and raise an MPI exception at the origin call if an out-of-bound situation occurred. Note that the condition can be checked at the origin. Of course, the added safety achieved by such checks has to be weighed against the added cost of such checks. (*End of advice to implementors.*)

6.3.2 Get

MPI_GET(*origin_addr*, *origin_count*, *origin_datatype*, *target_rank*, *target_disp*, *target_count*, *target_datatype*, *win*)

OUT	origin_addr	initial address of origin buffer (choice)
IN	origin_count	number of entries in origin buffer (nonnegative integer)
IN	origin_datatype	datatype of each entry in origin buffer (handle)
IN	target_rank	rank of target (nonnegative integer)
IN	target_disp	displacement from window start to the beginning of the target buffer (nonnegative integer)
IN	target_count	number of entries in target buffer (nonnegative integer)
IN	target_datatype	datatype of each entry in target buffer (handle)
IN	win	window object used for communication (handle)

```

int MPI_Get(void *origin_addr, int origin_count, MPI_Datatype
            origin_datatype, int target_rank, MPI_Aint target_disp, int
            target_count, MPI_Datatype target_datatype, MPI_Win win)

MPI_GET(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP,
        TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR)
    <type> ORIGIN_ADDR(*)
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
    TARGET_DATATYPE, WIN, IERROR

void MPI::Win::Get(const void *origin_addr, int origin_count, const
                  MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
                  target_disp, int target_count, const MPI::Datatype&
                  target_datatype) const

```

Similar to **MPI_PUT**, except that the direction of data transfer is reversed. Data are copied from the target memory to the origin. The **origin_datatype** may not specify overlapping entries in the origin buffer. The target buffer must be contained within the target window, and the copied data must fit, without truncation, in the origin buffer.

6.3.3 Examples

Example 6.1 We show how to implement the generic indirect assignment **A = B(map)**, where **A**, **B** and **map** have the same distribution, and **map** is a permutation. To simplify, we assume a block distribution with equal size blocks.

```

SUBROUTINE MAPVALS(A, B, map, m, comm, p)
USE MPI
INTEGER m, map(m), comm, p

```

```

REAL A(m), B(m)
1
2
INTEGER otype(p), oindex(m),    & ! used to construct origin datatypes
3
    ttype(p), tindex(m),        & ! used to construct target datatypes
4
    count(p), total(p),         &
5
    sizeofreal, win, ierr
6
7
! This part does the work that depends on the locations of B.
8
! Can be reused while this does not change
9
10
CALL MPI_TYPE_EXTENT(MPI_REAL, sizeofreal, ierr)
11
CALL MPI_WIN_CREATE(B, m*sizeofreal, sizeofreal, MPI_INFO_NULL,    &
12
    comm, win, ierr)
13
14
! This part does the work that depends on the value of map and
15
! the locations of the arrays.
16
! Can be reused while these do not change
17
18
! Compute number of entries to be received from each process
19
20
DO i=1,p
21
    count(i) = 0
22
END DO
23
DO i=1,m
24
    j = map(i)/m+1
25
    count(j) = count(j)+1
26
END DO
27
28
total(1) = 0
29
DO i=2,p
30
    total(i) = total(i-1) + count(i-1)
31
END DO
32
33
DO i=1,p
34
    count(i) = 0
35
END DO
36
37
! compute origin and target indices of entries.
38
! entry i at current process is received from location
39
! k at process (j-1), where map(i) = (j-1)*m + (k-1),
40
! j = 1..p and k = 1..m
41
42
DO i=1,m
43
    j = map(i)/m+1
44
    k = MOD(map(i),m)+1
45
    count(j) = count(j)+1
46
    oindex(total(j) + count(j)) = i
47
    tindex(total(j) + count(j)) = k
48

```

```

1  END DO
2
3  ! create origin and target datatypes for each get operation
4  DO i=1,p
5      CALL MPI_TYPE_INDEXED_BLOCK(count(i), 1, oindex(total(i)+1), &
6                                  MPI_REAL, otype(i), ierr)
7      CALL MPI_TYPE_COMMIT(otype(i), ierr)
8      CALL MPI_TYPE_INDEXED_BLOCK(count(i), 1, tindex(total(i)+1), &
9                                  MPI_REAL, ttype(i), ierr)
10     CALL MPI_TYPE_COMMIT(ttype(i), ierr)
11 END DO
12
13 ! this part does the assignment itself
14 CALL MPI_WIN_FENCE(0, win, ierr)
15 DO i=1,p
16     CALL MPI_GET(A, 1, otype(i), i-1, 0, 1, ttype(i), win, ierr)
17 END DO
18 CALL MPI_WIN_FENCE(0, win, ierr)
19
20 CALL MPI_WIN_FREE(win, ierr)
21 DO i=1,p
22     CALL MPI_TYPE_FREE(otype(i), ierr)
23     CALL MPI_TYPE_FREE(ttype(i), ierr)
24 END DO
25 RETURN
26 END

```

Example 6.2 A simpler version can be written that does not require that a datatype be built for the target buffer. But, one then needs a separate get call for each entry, as illustrated below. This code is much simpler, but usually much less efficient, for large arrays.

```

31 SUBROUTINE MAPVALS(A, B, map, m, comm, p)
32 USE MPI
33 INTEGER m, map(m), comm, p
34 REAL A(m), B(m)
35 INTEGER sizeofreal, win, ierr
36
37 CALL MPI_TYPE_EXTENT(MPI_REAL, sizeofreal, ierr)
38 CALL MPI_WIN_CREATE(B, m*sizeofreal, sizeofreal, MPI_INFO_NULL, &
39                     comm, win, ierr)
40
41 CALL MPI_WIN_FENCE(0, win, ierr)
42 DO i=1,m
43     j = map(i)/p
44     k = MOD(map(i),p)
45     CALL MPI_GET(A(i), 1, MPI_REAL, j, k, 1, MPI_REAL, win, ierr)
46 END DO
47 CALL MPI_WIN_FENCE(0, win, ierr)
48

```

```
CALL MPI_WIN_FREE(win, ierr)
RETURN
END
```

6.3.4 Accumulate Functions

It is often useful in a put operation to combine the data moved to the target process with the data that resides at that process, rather than replacing the data there. This will allow, for example, the accumulation of a sum by having all involved processes add their contribution to the sum variable in the memory of one process.

MPI_ACCUMULATE(*origin_addr*, *origin_count*, *origin_datatype*, *target_rank*, *target_disp*, *target_count*, *target_datatype*, *op*, *win*)

IN	<i>origin_addr</i>	initial address of buffer (choice)
IN	<i>origin_count</i>	number of entries in buffer (nonnegative integer)
IN	<i>origin_datatype</i>	datatype of each buffer entry (handle)
IN	<i>target_rank</i>	rank of target (nonnegative integer)
IN	<i>target_disp</i>	displacement from start of window to beginning of target buffer (nonnegative integer)
IN	<i>target_count</i>	number of entries in target buffer (nonnegative integer)
IN	<i>target_datatype</i>	datatype of each entry in target buffer (handle)
IN	<i>op</i>	reduce operation (handle)
IN	<i>win</i>	window object (handle)

```
int MPI_Accumulate(void *origin_addr, int origin_count,
                  MPI_Datatype origin_datatype, int target_rank,
                  MPI_Aint target_disp, int target_count,
                  MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)

MPI_ACCUMULATE(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK,
               TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR)
    <type> ORIGIN_ADDR(*)
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
    TARGET_DATATYPE, OP, WIN, IERROR

void MPI::Win::Accumulate(const void* origin_addr, int origin_count, const
                        MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
                        target_disp, int target_count, const MPI::Datatype&
                        target_datatype, const MPI::Op& op) const
```

Accumulate the contents of the origin buffer (as defined by *origin_addr*, *origin_count* and *origin_datatype*) to the buffer specified by arguments *target_count* and *target_datatype*, at offset *target_disp*, in the target window specified by *target_rank* and *win*, using the operation

op. This is like `MPI_PUT` except that data is combined into the target area instead of overwriting it.

Any of the predefined operations for `MPI_REDUCE` can be used. User-defined functions cannot be used. For example, if **op** is `MPI_SUM`, each element of the origin buffer is added to the corresponding element in the target, replacing the former value in the target.

Each datatype argument must be a predefined datatype or a derived datatype, where all basic components are of the same predefined datatype. Both datatype arguments must be constructed from the same predefined datatype. The operation **op** applies to elements of that predefined type. **target_datatype** must not specify overlapping entries, and the target buffer must fit in the target window.

A new predefined operation, `MPI_REPLACE`, is defined. It corresponds to the associative function $f(a, b) = b$; i.e., the current value in the target memory is replaced by the value supplied by the origin.

Advice to users. `MPI_PUT` is a special case of `MPI_ACCUMULATE`, with the operation `MPI_REPLACE`. Note, however, that `MPI_PUT` and `MPI_ACCUMULATE` have different constraints on concurrent updates. (*End of advice to users.*)

Example 6.3 We want to compute $B(j) = \sum_{\text{map}(i)=j} A(i)$. The arrays `A`, `B` and `map` are distributed in the same manner. We write the simple version.

```

SUBROUTINE SUM(A, B, map, m, comm, p)
USE MPI
INTEGER m, map(m), comm, p, sizeofreal, win, ierr
REAL A(m), B(m)

CALL MPI_TYPE_EXTENT(MPI_REAL, sizeofreal, ierr)
CALL MPI_WIN_CREATE(B, m*sizeofreal, sizeofreal, MPI_INFO_NULL, &
                    comm, win, ierr)

CALL MPI_WIN_FENCE(0, win, ierr)
DO i=1,m
  j = map(i)/p
  k = MOD(map(i),p)
  CALL MPI_ACCUMULATE(A(i), 1, MPI_REAL, j, k, 1, MPI_REAL, &
                     MPI_SUM, win, ierr)
END DO
CALL MPI_WIN_FENCE(0, win, ierr)

CALL MPI_WIN_FREE(win, ierr)
RETURN
END

```

This code is identical to the code in Example 6.2, page 118, except that a call to `get` has been replaced by a call to `accumulate`. (Note that, if `map` is one-to-one, then the code computes $B = A(\text{map}^{-1})$, which is the reverse assignment to the one computed in that previous example.) In a similar manner, we can replace in Example 6.1, page 116, the call to `get` by a call to `accumulate`, thus performing the computation with only one communication between any two processes.

6.4 Synchronization Calls

RMA communications fall in two categories:

- **active target** communication, where data is moved from the memory of one process to the memory of another, and both are explicitly involved in the communication. This communication pattern is similar to message passing, except that all the data transfer arguments are provided by one process, and the second process only participates in the synchronization.
- **passive target** communication, where data is moved from the memory of one process to the memory of another, and only the origin process is explicitly involved in the transfer. Thus, two origin processes may communicate by accessing the same location in a target window. The process that owns the target window may be distinct from the two communicating processes, in which case it does not participate explicitly in the communication. This communication paradigm is closest to a shared memory model, where shared data can be accessed by all processes, irrespective of location.

RMA communication calls with argument **win** must occur at a process only within an **access epoch** for **win**. Such an epoch starts with an RMA synchronization call on **win**; it proceeds with zero or more RMA communication calls (**MPI_PUT**, **MPI_GET** or **MPI_ACCUMULATE**) on **win**; it completes with another synchronization call on **win**. This allows users to amortize one synchronization with multiple data transfers and provide implementors more flexibility in the implementation of RMA operations.

Distinct access epochs for **win** at the same process must be disjoint. On the other hand, epochs pertaining to different **win** arguments may overlap. Local operations or other MPI calls may also occur during an epoch.

In active target communication, a target window can be accessed by RMA operations only within an **exposure epoch**. Such an epoch is started and completed by RMA synchronization calls executed by the target process. Distinct exposure epochs at a process on the same window must be disjoint, but such an exposure epoch may overlap with exposure epochs on other windows or with access epochs for the same or other **win** arguments. There is a one-to-one matching between access epochs at origin processes and exposure epochs on target processes: RMA operations issued by an origin process for a target window will access that target window during the same exposure epoch if and only if they were issued during the same access epoch.

In passive target communication the target process does not execute RMA synchronization calls, and there is no concept of an exposure epoch.

MPI provides three synchronization mechanisms:

1. The **MPI_WIN_FENCE** collective synchronization call supports a simple synchronization pattern that is often used in parallel computations: namely a loosely-synchronous model, where global computation phases alternate with global communication phases. This mechanism is most useful for loosely synchronous algorithms where the graph of communicating processes changes very frequently, or where each process communicates with many others.

This call is used for active target communication. An access epoch at an origin process or an exposure epoch at a target process are started and completed by calls to **MPI_WIN_FENCE**. A process can access windows at all processes in the group of **win**

during such an access epoch, and the local window can be accessed by all processes in the group of `win` during such an exposure epoch.

2. The four functions `MPI_WIN_START`, `MPI_WIN_COMPLETE`, `MPI_WIN_POST` and `MPI_WIN_WAIT` can be used to restrict synchronization to the minimum: only pairs of communicating processes synchronize, and they do so only when a synchronization is needed to order correctly RMA accesses to a window with respect to local accesses to that same window. This mechanism may be more efficient when each process communicates with few (logical) neighbors, and the communication graph is fixed or changes infrequently.

These calls are used for active target communication. An access epoch is started at the origin process by a call to `MPI_WIN_START` and is terminated by a call to `MPI_WIN_COMPLETE`. The start call has a group argument that specifies the group of target processes for that epoch. An exposure epoch is started at the target process by a call to `MPI_WIN_POST` and is completed by a call to `MPI_WIN_WAIT`. The post call has a group argument that specifies the set of origin processes for that epoch.

3. Finally, shared and exclusive locks are provided by the two functions `MPI_WIN_LOCK` and `MPI_WIN_UNLOCK`. Lock synchronization is useful for MPI applications that emulate a shared memory model via MPI calls; e.g., in a “billboard” model, where processes can, at random times, access or update different parts of the billboard.

These two calls provide passive target communication. An access epoch is started by a call to `MPI_WIN_LOCK` and terminated by a call to `MPI_WIN_UNLOCK`. Only one target window can be accessed during that epoch with `win`.

Figure 6.1 illustrates the general synchronization pattern for active target communication. The synchronization between `post` and `start` ensures that the put call of the origin process does not start until the target process exposes the window (with the `post` call); the target process will expose the window only after preceding local accesses to the window have completed. The synchronization between `complete` and `wait` ensures that the put call of the origin process completes before the window is unexposed (with the `wait` call). The target process will execute following local accesses to the target window only after the `wait` returned.

Figure 6.1 shows operations occurring in the natural temporal order implied by the synchronizations: the `post` occurs before the matching `start`, and `complete` occurs before the matching `wait`. However, such **strong** synchronization is more than needed for correct ordering of window accesses. The semantics of MPI calls allow **weak** synchronization, as illustrated in Figure 6.2. The access to the target window is delayed until the window is exposed, after the `post`. However the `start` may complete earlier; the `put` and `complete` may also terminate earlier, if put data is buffered by the implementation. The synchronization calls order correctly window accesses, but do not necessarily synchronize other operations. This weaker synchronization semantic allows for more efficient implementations.

Figure 6.3 illustrates the general synchronization pattern for passive target communication. The first origin process communicates data to the second origin process, through the memory of the target process; the target process is not explicitly involved in the communication. The `lock` and `unlock` calls ensure that the two RMA accesses do not occur concurrently. However, they do *not* ensure that the `put` by origin 1 will precede the `get` by origin 2.

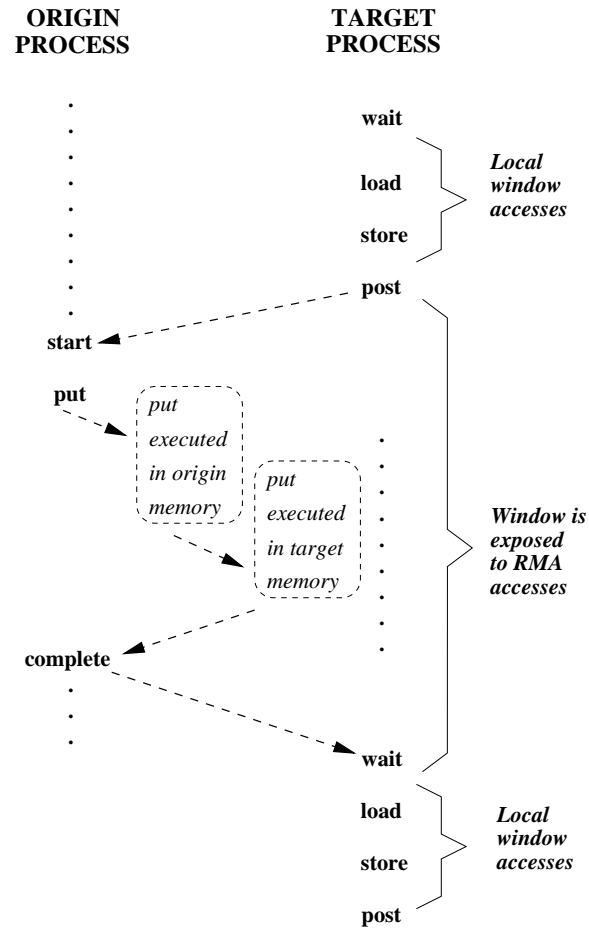


Figure 6.1: active target communication. Dashed arrows represent synchronizations (ordering of events).

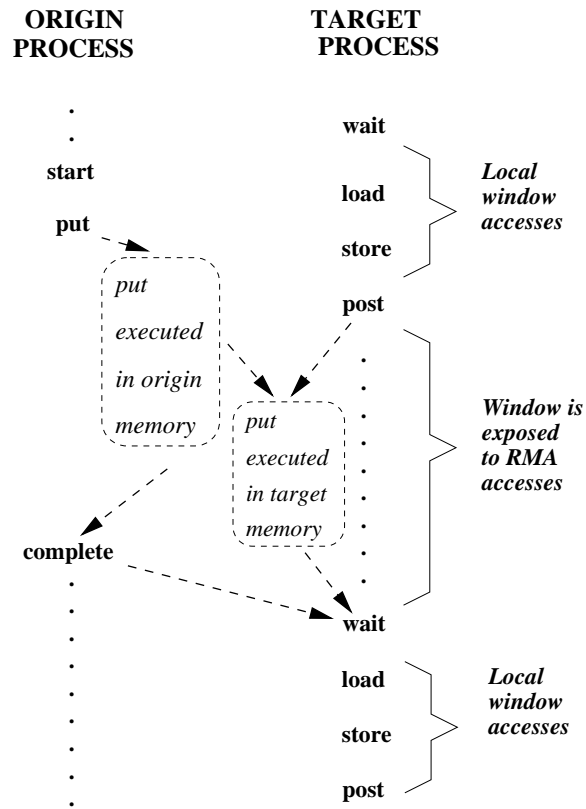


Figure 6.2: active target communication, with weak synchronization. Dashed arrows represent synchronizations (ordering of events)

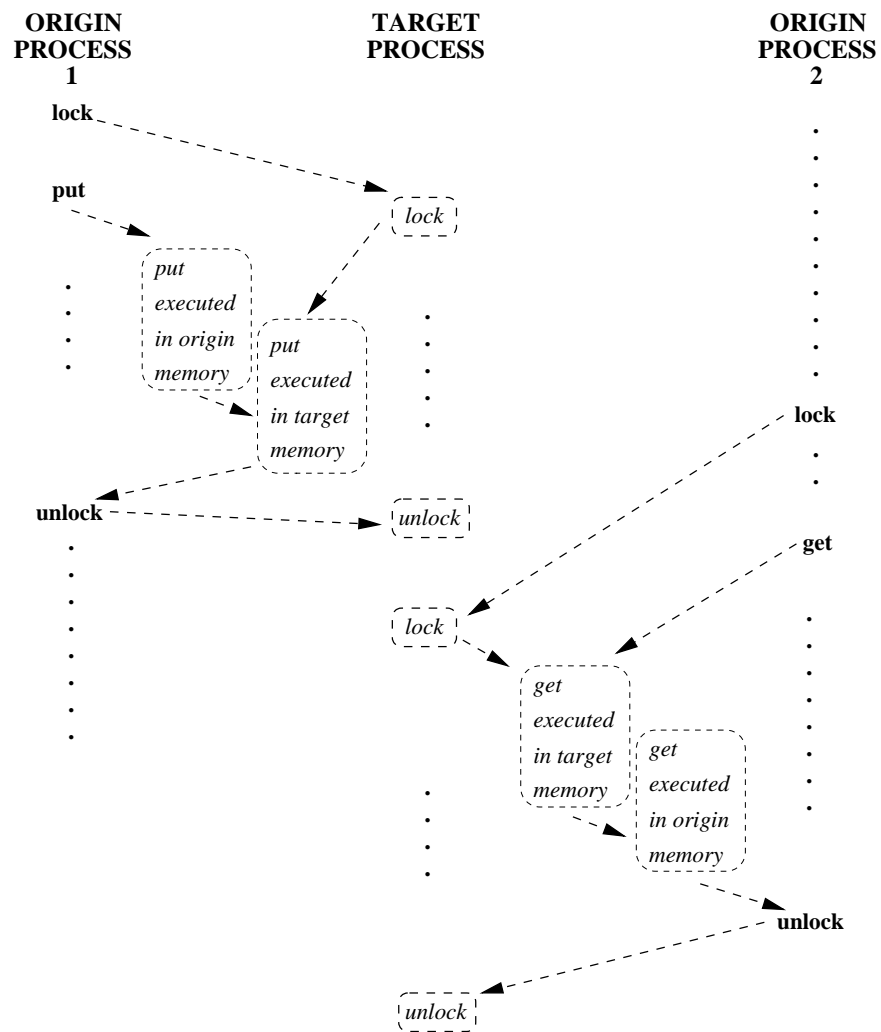


Figure 6.3: passive target communication. Dashed arrows represent synchronizations (ordering of events).

6.4.1 Fence

`MPI_WIN_FENCE(assert, win)`

IN	assert	program assertion (integer)
IN	win	window object (handle)

`int MPI_Win_fence(int assert, MPI_Win win)`

`MPI_WIN_FENCE(ASSERT, WIN, IERROR)`

INTEGER ASSERT, WIN, IERROR

`void MPI::Win::Fence(int assert) const`

The MPI call `MPI_WIN_FENCE(assert, win)` synchronizes RMA calls on `win`. The call is collective on the group of `win`. All RMA operations on `win` originating at a given process and started before the fence call will complete at that process before the fence call returns. They will be completed at their target before the fence call returns at the target. RMA operations on `win` started by a process after the fence call returns will access their target window only after `MPI_WIN_FENCE` has been called by the target process.

The call completes an RMA access epoch if it was preceded by another fence call and the local process issued RMA communication calls on `win` between these two calls. The call completes an RMA exposure epoch if it was preceded by another fence call and the local window was the target of RMA accesses between these two calls. The call starts an RMA access epoch if it is followed by another fence call and by RMA communication calls issued between these two fence calls. The call starts an exposure epoch if it is followed by another fence call and the local window is the target of RMA accesses between these two fence calls. Thus, the fence call is equivalent to calls to a subset of `post`, `start`, `complete`, `wait`.

A fence call usually entails a barrier synchronization: a process completes a call to `MPI_WIN_FENCE` only after all other processes in the group entered their matching call. However, a call to `MPI_WIN_FENCE` that is known not to end any epoch (in particular, a call with `assert = MPI_MODE_NOPRECEDE`) does not necessarily act as a barrier.

The `assert` argument is used to provide assertions on the context of the call that may be used for various optimizations. This is described in Section 6.4.4. A value of `assert = 0` is always valid.

Advice to users. Calls to `MPI_WIN_FENCE` should both precede and follow calls to put, get or accumulate that are synchronized with fence calls. (*End of advice to users.*)

6.4.2 General Active Target Synchronization

MPI_WIN_START(group, assert, win)

IN	group	group of target processes (handle)
IN	assert	program assertion (integer)
IN	win	window object (handle)

```
int MPI_Win_start(MPI_Group group, int assert, MPI_Win win)
```

```
MPI_WIN_START(GROUP, ASSERT, WIN, IERROR)
    INTEGER GROUP, ASSERT, WIN, IERROR
```

```
void MPI::Win::Start(const MPI::Group& group, int assert) const
```

Starts an RMA access epoch for **win**. RMA calls issued on **win** during this epoch must access only windows at processes in **group**. Each process in **group** must issue a matching call to **MPI_WIN_POST**. RMA accesses to each target window will be delayed, if necessary, until the target process executed the matching call to **MPI_WIN_POST**. **MPI_WIN_START** is allowed to block until the corresponding **MPI_WIN_POST** calls are executed, but is not required to.

The **assert** argument is used to provide assertions on the context of the call that may be used for various optimizations. This is described in Section 6.4.4. A value of **assert = 0** is always valid.

MPI_WIN_COMPLETE(win)

IN	win	window object (handle)
----	------------	------------------------

```
int MPI_Win_complete(MPI_Win win)
```

```
MPI_WIN_COMPLETE(WIN, IERROR)
    INTEGER WIN, IERROR
```

```
void MPI::Win::Complete() const
```

Completes an RMA access epoch on **win** started by a call to **MPI_WIN_START**. All RMA communication calls issued on **win** during this epoch will have completed at the origin when the call returns.

MPI_WIN_COMPLETE enforces completion of preceding RMA calls at the origin, but not at the target. A put or accumulate call may not have completed at the target when it has completed at the origin.

Consider the sequence of calls in the example below.

Example 6.4

```
MPI_Win_start(group, flag, win);
MPI_Put(...,win);
MPI_Win_complete(win);
```

The call to `MPI_WIN_COMPLETE` does not return until the put call has completed at the origin; and the target window will be accessed by the put operation only after the call to `MPI_WIN_START` has matched a call to `MPI_WIN_POST` by the target process. This still leaves much choice to implementors. The call to `MPI_WIN_START` can block until the matching call to `MPI_WIN_POST` occurs at all target processes. One can also have implementations where the call to `MPI_WIN_START` is nonblocking, but the call to `MPI_PUT` blocks until the matching call to `MPI_WIN_POST` occurred; or implementations where the first two calls are nonblocking, but the call to `MPI_WIN_COMPLETE` blocks until the call to `MPI_WIN_POST` occurred; or even implementations where all three calls can complete before any target process called `MPI_WIN_POST` — the data put must be buffered, in this last case, so as to allow the put to complete at the origin ahead of its completion at the target. However, once the call to `MPI_WIN_POST` is issued, the sequence above must complete, without further dependencies.

```
MPI_WIN_POST(group, assert, win)
```

IN	group	group of origin processes (handle)
IN	assert	program assertion (integer)
IN	win	window object (handle)

```
int MPI_Win_post(MPI_Group group, int assert, MPI_Win win)
```

```
MPI_WIN_POST(GROUP, ASSERT, WIN, IERROR)
    INTEGER GROUP, ASSERT, WIN, IERROR
```

```
void MPI::Win::Post(const MPI::Group& group, int assert) const
```

Starts an RMA exposure epoch for the local window associated with `win`. Only processes in `group` should access the window with RMA calls on `win` during this epoch. Each process in `group` must issue a matching call to `MPI_WIN_START`. `MPI_WIN_POST` does not block.

```
MPI_WIN_WAIT(win)
```

IN	win	window object (handle)
----	------------	------------------------

```
int MPI_Win_wait(MPI_Win win)
```

```
MPI_WIN_WAIT(WIN, IERROR)
    INTEGER WIN, IERROR
```

```
void MPI::Win::Wait() const
```

Completes an RMA exposure epoch started by a call to `MPI_WIN_POST` on `win`. This call matches calls to `MPI_WIN_COMPLETE(win)` issued by each of the origin processes that were granted access to the window during this epoch. The call to `MPI_WIN_WAIT` will block until all matching calls to `MPI_WIN_COMPLETE` have occurred. This guarantees that all these origin processes have completed their RMA accesses to the local window. When the call returns, all these RMA accesses will have completed at the target window.

Figure 6.4 illustrates the use of these four functions. Process 0 puts data in the windows

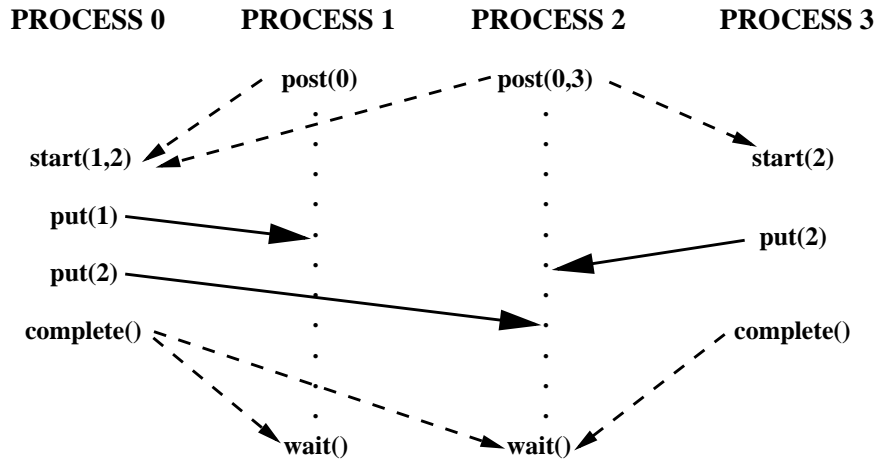


Figure 6.4: active target communication. Dashed arrows represent synchronizations and solid arrows represent data transfer.

of processes 1 and 2 and process 3 puts data in the window of process 2. Each start call lists the ranks of the processes whose windows will be accessed; each post call lists the ranks of the processes that access the local window. The figure illustrates a possible timing for the events, assuming strong synchronization; in a weak synchronization, the start, put or complete calls may occur ahead of the matching post calls.

MPI_WIN_TEST(win, flag)

IN	win	window object (handle)
OUT	flag	success flag (logical)

```
int MPI_Win_test(MPI_Win win, int *flag)
```

```
MPI_WIN_TEST(WIN, FLAG, IERROR)
```

```
    INTEGER WIN, IERROR
```

```
    LOGICAL FLAG
```

```
bool MPI::Win::Test() const
```

This is the nonblocking version of MPI_WIN_WAIT. It returns **flag = true** if MPI_WIN_WAIT would return, **flag = false**, otherwise. The effect of return of MPI_WIN_TEST with **flag = true** is the same as the effect of a return of MPI_WIN_WAIT. If **flag = false** is returned, then the call has no visible effect.

MPI_WIN_TEST should be invoked only where MPI_WIN_WAIT can be invoked. Once the call has returned **flag = true**, it must not be invoked anew, until the window is posted anew.

Assume that window **win** is associated with a “hidden” communicator **wincomm**, used for communication by the processes of **win**. The rules for matching of post and start calls and for matching complete and wait call can be derived from the rules for matching sends and receives, by considering the following (partial) model implementation.

`MPI_WIN_POST(group,0,win)` initiate a nonblocking send with tag **tag0** to each process in **group**, using **wincomm**. No need to wait for the completion of these sends.

`MPI_WIN_START(group,0,win)` initiate a nonblocking receive with tag **tag0** from each process in **group**, using **wincomm**. An RMA access to a window in target process *i* is delayed until the receive from *i* is completed.

`MPI_WIN_COMPLETE(win)` initiate a nonblocking send with tag **tag1** to each process in the group of the preceding start call. No need to wait for the completion of these sends.

`MPI_WIN_WAIT(win)` initiate a nonblocking receive with tag **tag1** from each process in the group of the preceding post call. Wait for the completion of all receives.

No races can occur in a correct program: each of the sends matches a unique receive, and vice-versa.

Rationale. The design for general active target synchronization requires the user to provide complete information on the communication pattern, at each end of a communication link: each origin specifies a list of targets, and each target specifies a list of origins. This provides maximum flexibility (hence, efficiency) for the implementor: each synchronization can be initiated by either side, since each “knows” the identity of the other. This also provides maximum protection from possible races. On the other hand, the design requires more information than RMA needs, in general: in general, it is sufficient for the origin to know the rank of the target, but not vice versa. Users that want more “anonymous” communication will be required to use the fence or lock mechanisms. (*End of rationale.*)

Advice to users. Assume a communication pattern that is represented by a directed graph $G = \langle V, E \rangle$, where $V = \{0, \dots, n-1\}$ and $ij \in E$ if origin process *i* accesses the window at target process *j*. Then each process *i* issues a call to `MPI_WIN_POST(ingroupi, ...)`, followed by a call to `MPI_WIN_START(outgroupi, ...)`, where $outgroup_i = \{j : ij \in E\}$ and $ingroup_i = \{j : ji \in E\}$. A call is a noop, and can be skipped, if the group argument is empty. After the communications calls, each process that issued a start will issue a complete. Finally, each process that issued a post will issue a wait.

Note that each process may call with a group argument that has different members. (*End of advice to users.*)

6.4.3 Lock

`MPI_WIN_LOCK(lock_type, rank, assert, win)`

IN	lock_type	either <code>MPI_LOCK_EXCLUSIVE</code> or <code>MPI_LOCK_SHARED</code> (state)
IN	rank	rank of locked window (nonnegative integer)
IN	assert	program assertion (integer)
IN	win	window object (handle)

`int MPI_Win_lock(int lock_type, int rank, int assert, MPI_Win win)`


```
MPI_WIN_LOCK(LOCK_TYPE, RANK, ASSERT, WIN, IERROR)
```

```
    INTEGER LOCK_TYPE, RANK, ASSERT, WIN, IERROR
```

```
void MPI::Win::Lock(int lock_type, int rank, int assert) const
```

Starts an RMA access epoch. Only the window at the process with rank `rank` can be accessed by RMA operations on `win` during that epoch.

```
MPI_WIN_UNLOCK(rank, win)
```

```
    IN          rank          rank of window (nonnegative integer)
```

```
    IN          win          window object (handle)
```

```
int MPI_Win_unlock(int rank, MPI_Win win)
```

```
MPI_WIN_UNLOCK(RANK, WIN, IERROR)
```

```
    INTEGER RANK, WIN, IERROR
```

```
void MPI::Win::Unlock(int rank) const
```

Completes an RMA access epoch started by a call to `MPI_WIN_LOCK(...,win)`. RMA operations issued during this period will have completed both at the origin and at the target when the call returns.

Locks are used to protect accesses to the locked target window effected by RMA calls issued between the lock and unlock call, and to protect local load/store accesses to a locked local window executed between the lock and unlock call. Accesses that are protected by an exclusive lock will not be concurrent at the window site with other accesses to the same window that are lock protected. Accesses that are protected by a shared lock will not be concurrent at the window site with accesses protected by an exclusive lock to the same window.

It is erroneous to have a window locked and exposed (in an exposure epoch) concurrently. I.e., a process may not call `MPI_WIN_LOCK` to lock a target window if the target process has called `MPI_WIN_POST` and has not yet called `MPI_WIN_WAIT`; it is erroneous to call `MPI_WIN_POST` while the local window is locked.

Rationale. An alternative is to require MPI to enforce mutual exclusion between exposure epochs and locking periods. But this would entail additional overheads when locks or active target synchronization do not interact in support of those rare interactions between the two mechanisms. The programming style that we encourage here is that a set of windows is used with only one synchronization mechanism at a time, with shifts from one mechanism to another being rare and involving global synchronization. (*End of rationale.*)

Advice to users. Users need to use explicit synchronization code in order to enforce mutual exclusion between locking periods and exposure epochs on a window. (*End of advice to users.*)

Implementors may restrict the use of RMA communication that is synchronized by lock calls to windows in memory allocated by `MPI_ALLOC_MEM` (Section 4.11, page 47). Locks can be used portably only in such memory.

Rationale. The implementation of passive target communication when memory is not shared requires an asynchronous agent. Such an agent can be implemented more easily, and can achieve better performance, if restricted to specially allocated memory. It can be avoided altogether if shared memory is used. It seems natural to impose restrictions that allows one to use shared memory for 3-rd party communication in shared memory machines.

The downside of this decision is that passive target communication cannot be used without taking advantage of nonstandard Fortran features: namely, the availability of C-like pointers; these are not supported by some Fortran compilers (g77 and Windows/NT compilers, at the time of writing). Also, passive target communication cannot be portably targeted to `COMMON` blocks, or other statically declared Fortran arrays. (*End of rationale.*)

Consider the sequence of calls in the example below.

Example 6.5

```
MPI_Win_lock(MPI_LOCK_EXCLUSIVE, rank, assert, win)
MPI_Put(..., rank, ..., win)
MPI_Win_unlock(rank, win)
```

The call to `MPI_WIN_UNLOCK` will not return until the put transfer has completed at the origin and at the target. This still leaves much freedom to implementors. The call to `MPI_WIN_LOCK` may block until an exclusive lock on the window is acquired; or, the call `MPI_WIN_LOCK` may not block, while the call to `MPI_PUT` blocks until a lock is acquired; or, the first two calls may not block, while `MPI_WIN_UNLOCK` blocks until a lock is acquired — the update of the target window is then postponed until the call to `MPI_WIN_UNLOCK` occurs. However, if the call to `MPI_WIN_LOCK` is used to lock a local window, then the call must block until the lock is acquired, since the lock may protect local load/store accesses to the window issued after the lock call returns.

6.4.4 Assertions

The `assert` argument in the calls `MPI_WIN_POST`, `MPI_WIN_START`, `MPI_WIN_FENCE` and `MPI_WIN_LOCK` is used to provide assertions on the context of the call that may be used to optimize performance. The `assert` argument does not change program semantics if it provides correct information on the program — it is erroneous to provides incorrect information. Users may always provide `assert = 0` to indicate a general case, where no guarantees are made.

Advice to users. Many implementations may not take advantage of the information in `assert`; some of the information is relevant only for noncoherent, shared memory machines. Users should consult their implementation manual to find which information is useful on each system. On the other hand, applications that provide correct assertions whenever applicable are portable and will take advantage of assertion specific optimizations, whenever available. (*End of advice to users.*)

Advice to implementors. Implementations can always ignore the `assert` argument. Implementors should document which `assert` values are significant on their implementation. (*End of advice to implementors.*)

assert is the bit-vector OR of zero or more of the following integer constants:
MPI_MODE_NOCHECK, **MPI_MODE_NOSTORE**, **MPI_MODE_NOPUT**, **MPI_MODE_NOPRECEDE** and
MPI_MODE_NOSUCCEED. The significant options are listed below, for each call.

Advice to users. C/C++ users can use bit vector or (|) to combine these constants;
 Fortran 90 users can use the bit-vector **IOR** intrinsic. Fortran 77 users can use (non-
 portably) bit vector **IOR** on systems that support it. Alternatively, Fortran users can
 portably use integer addition to OR the constants (each constant should appear at
 most once in the addition!). (*End of advice to users.*)

MPI_WIN_START:

MPI_MODE_NOCHECK — the matching calls to **MPI_WIN_POST** have already com-
 pleted on all target processes when the call to **MPI_WIN_START** is made. The
 nocheck option can be specified in a start call if and only if it is specified in
 each matching post call. This is similar to the optimization of “ready-send” that
 may save a handshake when the handshake is implicit in the code. (However,
 ready-send is matched by a regular receive, whereas both start and post must
 specify the nocheck option.)

MPI_WIN_POST:

MPI_MODE_NOCHECK — the matching calls to **MPI_WIN_START** have not yet oc-
 curred on any origin processes when the call to **MPI_WIN_POST** is made. The
 nocheck option can be specified by a post call if and only if it is specified by each
 matching start call.

MPI_MODE_NOSTORE — the local window was not updated by local stores (or local
 get or receive calls) since last synchronization. This may avoid the need for cache
 synchronization at the post call.

MPI_MODE_NOPUT — the local window will not be updated by put or accumulate
 calls after the post call, until the ensuing (wait) synchronization. This may avoid
 the need for cache synchronization at the wait call.

MPI_WIN_FENCE:

MPI_MODE_NOSTORE — the local window was not updated by local stores (or local
 get or receive calls) since last synchronization.

MPI_MODE_NOPUT — the local window will not be updated by put or accumulate
 calls after the fence call, until the ensuing (fence) synchronization.

MPI_MODE_NOPRECEDE — the fence does not complete any sequence of locally issued
 RMA calls. If this assertion is given by any process in the window group, then it
 must be given by all processes in the group.

MPI_MODE_NOSUCCEED — the fence does not start any sequence of locally issued
 RMA calls. If the assertion is given by any process in the window group, then it
 must be given by all processes in the group.

MPI_WIN_LOCK:

`MPI_MODE_NOCHECK` — no other process holds, or will attempt to acquire a conflicting lock, while the caller holds the window lock. This is useful when mutual exclusion is achieved by other means, but the coherence operations that may be attached to the lock and unlock calls are still required.

Advice to users. Note that the `nostore` and `noprecede` flags provide information on what happened *before* the call; the `noput` and `nosucceed` flags provide information on what will happen *after* the call. (*End of advice to users.*)

6.4.5 Miscellaneous Clarifications

Once an RMA routine completes, it is safe to free any opaque objects passed as argument to that routine. For example, the `datatype` argument of a `MPI_PUT` call can be freed as soon as the call returns, even though the communication may not be complete.

As in message passing, datatypes must be committed before they can be used in RMA communication.

6.5 Examples

Example 6.6 The following example shows a generic loosely synchronous, iterative code, using fence synchronization. The window at each process consists of array `A`, which contains the origin and target buffers of the put calls.

```
...
while(!converged(A)){
    update(A);
    MPI_Win_fence(MPI_MODE_NOPRECEDE, win);
    for(i=0; i < toneighbors; i++)
        MPI_Put(&frombuf[i], 1, fromtype[i], toneighbor[i],
                todisp[i], 1, totype[i], win);
    MPI_Win_fence((MPI_MODE_NOSTORE | MPI_MODE_NOSUCCEED), win);
}
```

The same code could be written with `get`, rather than `put`. Note that, during the communication phase, each window is concurrently read (as origin buffer of puts) and written (as target buffer of puts). This is OK, provided that there is no overlap between the target buffer of a put and another communication buffer.

Example 6.7 Same generic example, with more computation/communication overlap. We assume that the update phase is broken in two subphases: the first, where the “boundary,” which is involved in communication, is updated, and the second, where the “core,” which neither use nor provide communicated data, is updated.

```
...
while(!converged(A)){
    update_boundary(A);
    MPI_Win_fence((MPI_MODE_NOPUT | MPI_MODE_NOPRECEDE), win);
    for(i=0; i < fromneighbors; i++)
        MPI_Get(&tobuf[i], 1, totype[i], fromneighbor[i],
```

```

        fromdisp[i], 1, fromtype[i], win);
    update_core(A);
    MPI_Win_fence(MPI_MODE_NOSUCCEED, win);
}

```

The get communication can be concurrent with the core update, since they do not access the same locations, and the local update of the origin buffer by the get call can be concurrent with the local update of the core by the `update_core` call. In order to get similar overlap with put communication we would need to use separate windows for the core and for the boundary. This is required because we do not allow local stores to be concurrent with puts on the same, or on overlapping, windows.

Example 6.8 Same code as in Example 6.6, rewritten using post-start-complete-wait.

```

...
while(!converged(A)){
    update(A);
    MPI_Win_post(fromgroup, 0, win);
    MPI_Win_start(togroup, 0, win);
    for(i=0; i < toneighbors; i++)
        MPI_Put(&frombuf[i], 1, fromtype[i], toneighbor[i],
                todisp[i], 1, totype[i], win);
    MPI_Win_complete(win);
    MPI_Win_wait(win);
}

```

Example 6.9 Same example, with split phases, as in Example 6.7.

```

...
while(!converged(A)){
    update_boundary(A);
    MPI_Win_post(togroup, MPI_MODE_NOPUT, win);
    MPI_Win_start(fromgroup, 0, win);
    for(i=0; i < fromneighbors; i++)
        MPI_Get(&tobuf[i], 1, totype[i], fromneighbor[i],
                fromdisp[i], 1, fromtype[i], win);
    update_core(A);
    MPI_Win_complete(win);
    MPI_Win_wait(win);
}

```

Example 6.10 A checkerboard, or double buffer communication pattern, that allows more computation/communication overlap. Array `A0` is updated using values of array `A1`, and vice versa. We assume that communication is symmetric: if process A gets data from process B, then process B gets data from process A. Window `wini` consists of array `Ai`.

```

...
if (!converged(A0,A1))
    MPI_Win_post(neighbors, (MPI_MODE_NOCHECK | MPI_MODE_NOPUT), win0);
MPI_Barrier(comm0);

```

```

1  /* the barrier is needed because the start call inside the
2  loop uses the nocheck option */
3  while(!converged(A0, A1)){
4      /* communication on A0 and computation on A1 */
5      update2(A1, A0); /* local update of A1 that depends on A0 (and A1) */
6      MPI_Win_start(neighbors, MPI_MODE_NOCHECK, win0);
7      for(i=0; i < neighbors; i++)
8          MPI_Get(&tobuf0[i], 1, totype0[i], neighbor[i],
9                  fromdisp0[i], 1, fromtype0[i], win0);
10     update1(A1); /* local update of A1 that is
11                  concurrent with communication that updates A0 */
12     MPI_Win_post(neighbors, (MPI_MODE_NOCHECK | MPI_MODE_NOPUT), win1);
13     MPI_Win_complete(win0);
14     MPI_Win_wait(win0);
15
16     /* communication on A1 and computation on A0 */
17     update2(A0, A1); /* local update of A0 that depends on A1 (and A0)*/
18     MPI_Win_start(neighbors, MPI_MODE_NOCHECK, win1);
19     for(i=0; i < neighbors; i++)
20         MPI_Get(&tobuf1[i], 1, totype1[i], neighbor[i],
21                 fromdisp1[i], 1, fromtype1[i], win1);
22     update1(A0); /* local update of A0 that depends on A0 only,
23                  concurrent with communication that updates A1 */
24     if (!converged(A0,A1))
25         MPI_Win_post(neighbors, (MPI_MODE_NOCHECK | MPI_MODE_NOPUT), win0);
26     MPI_Win_complete(win1);
27     MPI_Win_wait(win1);
28 }

```

A process posts the local window associated with `win0` before it completes RMA accesses to the remote windows associated with `win1`. When the `wait(win1)` call returns, then all neighbors of the calling process have posted the windows associated with `win0`. Conversely, when the `wait(win0)` call returns, then all neighbors of the calling process have posted the windows associated with `win1`. Therefore, the `nocheck` option can be used with the calls to `MPI_WIN_START`.

Put calls can be used, instead of get calls, if the area of array `A0` (resp. `A1`) used by the `update(A1, A0)` (resp. `update(A0, A1)`) call is disjoint from the area modified by the RMA communication. On some systems, a put call may be more efficient than a get call, as it requires information exchange only in one direction.

6.6 Error Handling

6.6.1 Error Handlers

Errors occurring during calls to `MPI_WIN_CREATE(...,comm,...)` cause the error handler currently associated with `comm` to be invoked. All other RMA calls have an input `win` argument. When an error occurs during such a call, the error handler currently associated with `win` is invoked.




Figure 6.5: Schematic description of window

The default error handler associated with `win` is `MPI_ERRORS_FATAL`. Users may change this default by explicitly associating a new error handler with `win` (see Section 4.13, page 61).

6.6.2 Error Classes

The following new error classes are defined

<code>MPI_ERR_WIN</code>	invalid <code>win</code> argument
<code>MPI_ERR_BASE</code>	invalid <code>base</code> argument
<code>MPI_ERR_SIZE</code>	invalid <code>size</code> argument
<code>MPI_ERR_DISP</code>	invalid <code>disp</code> argument
<code>MPI_ERR_LOCKTYPE</code>	invalid <code>locktype</code> argument
<code>MPI_ERR_ASSERT</code>	invalid <code>assert</code> argument
<code>MPI_ERR_RMA_CONFLICT</code>	conflicting accesses to window
<code>MPI_ERR_RMA_SYNC</code>	wrong synchronization of RMA calls

6.7 Semantics and Correctness

The semantics of RMA operations is best understood by assuming that the system maintains a separate *public* copy of each window, in addition to the original location in process memory (the *private* window copy). There is only one instance of each variable in process memory, but a distinct *public* copy of the variable for each window that contains it. A load accesses the instance in process memory (this includes `MPI` sends). A store accesses and updates the instance in process memory (this includes `MPI` receives), but the update may affect other public copies of the same locations. A get on a window accesses the public copy of that window. A put or accumulate on a window accesses and updates the public copy of that window, but the update may affect the private copy of the same locations in process memory, and public copies of other overlapping windows. This is illustrated in Figure 6.5.

The following rules specify the latest time at which an operation must complete at the origin or the target. The update performed by a get call in the origin process memory is visible when the get operation is complete at the origin (or earlier); the update performed by a put or accumulate call in the public copy of the target window is visible when the put or accumulate has completed at the target (or earlier). The rules also specifies the latest time at which an update of one window copy becomes visible in another overlapping copy.

1. An RMA operation is completed at the origin by the ensuing call to `MPI_WIN_COMPLETE`, `MPI_WIN_FENCE` or `MPI_WIN_UNLOCK` that synchronizes this access at the origin.
2. If an RMA operation is completed at the origin by a call to `MPI_WIN_FENCE` then the operation is completed at the target by the matching call to `MPI_WIN_FENCE` by the target process.
3. If an RMA operation is completed at the origin by a call to `MPI_WIN_COMPLETE` then the operation is completed at the target by the matching call to `MPI_WIN_WAIT` by the target process.
4. If an RMA operation is completed at the origin by a call to `MPI_WIN_UNLOCK` then the operation is completed at the target by that same call to `MPI_WIN_UNLOCK`.
5. An update of a location in a private window copy in process memory becomes visible in the public window copy at latest when an ensuing call to `MPI_WIN_POST`, `MPI_WIN_FENCE`, or `MPI_WIN_UNLOCK` is executed on that window by the window owner.
6. An update by a put or accumulate call to a public window copy becomes visible in the private copy in process memory at latest when an ensuing call to `MPI_WIN_WAIT`, `MPI_WIN_FENCE`, or `MPI_WIN_LOCK` is executed on that window by the window owner.

The `MPI_WIN_FENCE` or `MPI_WIN_WAIT` call that completes the transfer from public copy to private copy (6) is the same call that completes the put or accumulate operation in the window copy (2, 3). If a put or accumulate access was synchronized with a lock, then the update of the public window copy is complete as soon as the updating process executed `MPI_WIN_UNLOCK`. On the other hand, the update of private copy in the process memory may be delayed until the target process executes a synchronization call on that window (6). Thus, updates to process memory can always be delayed until the process executes a suitable synchronization call. Updates to a public window copy can also be delayed until the window owner executes a synchronization call, if fences or post-start-complete-wait synchronization is used. Only when lock synchronization is used does it becomes necessary to update the public window copy, even if the window owner does not execute any related synchronization call.

The rules above also define, by implication, when an update to a public window copy becomes visible in another overlapping public window copy. Consider, for example, two overlapping windows, `win1` and `win2`. A call to `MPI_WIN_FENCE(0, win1)` by the window owner makes visible in the process memory previous updates to window `win1` by remote processes. A subsequent call to `MPI_WIN_FENCE(0, win2)` makes these updates visible in the public copy of `win2`.

A correct program must obey the following rules.

1. A location in a window must not be accessed locally once an update to that location has started, until the update becomes visible in the private window copy in process memory.
2. A location in a window must not be accessed as a target of an RMA operation once an update to that location has started, until the update becomes visible in the public window copy. There is one exception to this rule, in the case where the same variable is updated by two concurrent accumulates that use the same operation, with the same predefined datatype, on the same window.
3. A put or accumulate must not access a target window once a local update or a put or accumulate update to another (overlapping) target window have started on a location in the target window, until the update becomes visible in the public copy of the window. Conversely, a local update in process memory to a location in a window must not start once a put or accumulate update to that target window has started, until the put or accumulate update becomes visible in process memory. In both cases, the restriction applies to operations even if they access disjoint locations in the window.

A program is erroneous if it violates these rules.

Rationale. The last constraint on correct RMA accesses may seem unduly restrictive, as it forbids concurrent accesses to nonoverlapping locations in a window. The reason for this constraint is that, on some architectures, explicit coherence restoring operations may be needed at synchronization points. A different operation may be needed for locations that were locally updated by stores and for locations that were remotely updated by put or accumulate operations. Without this constraint, the MPI library will have to track precisely which locations in a window were updated by a put or accumulate call. The additional overhead of maintaining such information is considered prohibitive. (*End of rationale.*)

Advice to users. A user can write correct programs by following the following rules:

fence: During each period between fence calls, each window is either updated by put or accumulate calls, or updated by local stores, but not both. Locations updated by put or accumulate calls should not be accessed during the same period (with the exception of concurrent updates to the same location by accumulate calls). Locations accessed by get calls should not be updated during the same period.

post-start-complete-wait: A window should not be updated locally while being posted, if it is being updated by put or accumulate calls. Locations updated by put or accumulate calls should not be accessed while the window is posted (with the exception of concurrent updates to the same location by accumulate calls). Locations accessed by get calls should not be updated while the window is posted.

With the post-start synchronization, the target process can tell the origin process that its window is now ready for RMA access; with the complete-wait synchronization, the origin process can tell the target process that it has finished its RMA accesses to the window.

lock: Updates to the window are protected by exclusive locks if they may conflict. Nonconflicting accesses (such as read-only accesses or accumulate accesses) are protected by shared locks, both for local accesses and for RMA accesses.

changing window or synchronization mode: One can change synchronization mode, or change the window used to access a location that belongs to two overlapping windows, when the process memory and the window copy are guaranteed to have the same values. This is true after a local call to `MPI_WIN_FENCE`, if RMA accesses to the window are synchronized with fences; after a local call to `MPI_WIN_WAIT`, if the accesses are synchronized with post-start-complete-wait; after the call at the origin (local or remote) to `MPI_WIN_UNLOCK` if the accesses are synchronized with locks.

In addition, a process should not access the local buffer of a get operation until the operation is complete, and should not update the local buffer of a put or accumulate operation until that operation is complete. (*End of advice to users.*)

6.7.1 Atomicity

The outcome of concurrent accumulates to the same location, with the same operation and predefined datatype, is as if the accumulates were done at that location in some serial order. On the other hand, if two locations are both updated by two accumulate calls, then the updates may occur in reverse order at the two locations. Thus, there is no guarantee that the entire call to `MPI_ACCUMULATE` is executed atomically. The effect of this lack of atomicity is limited: The previous correctness conditions imply that a location updated by a call to `MPI_ACCUMULATE`, cannot be accessed by load or an RMA call other than accumulate, until the `MPI_ACCUMULATE` call has completed (at the target). Different interleavings can lead to different results only to the extent that computer arithmetics are not truly associative or commutative.

6.7.2 Progress

One-sided communication has the same progress requirements as point-to-point communication: once a communication is enabled, then it is guaranteed to complete. RMA calls must have local semantics, except when required for synchronization with other RMA calls.

There is some fuzziness in the definition of the time when a RMA communication becomes enabled. This fuzziness provides to the implementor more flexibility than with point-to-point communication. Access to a target window becomes enabled once the corresponding synchronization (such as `MPI_WIN_FENCE` or `MPI_WIN_POST`) has executed. On the origin process, an RMA communication may become enabled as soon as the corresponding put, get or accumulate call has executed, or as late as when the ensuing synchronization call is issued. Once the communication is enabled both at the origin and at the target, the communication must complete.

Consider the code fragment in Example 6.4, on page 127. Some of the calls may block if the target window is not posted. However, if the target window is posted, then the code fragment must complete. The data transfer may start as soon as the put call occur, but may be delayed until the ensuing complete call occurs.

Consider the code fragment in Example 6.5, on page 132. Some of the calls may block if another process holds a conflicting lock. However, if no conflicting lock is held, then the code fragment must complete.

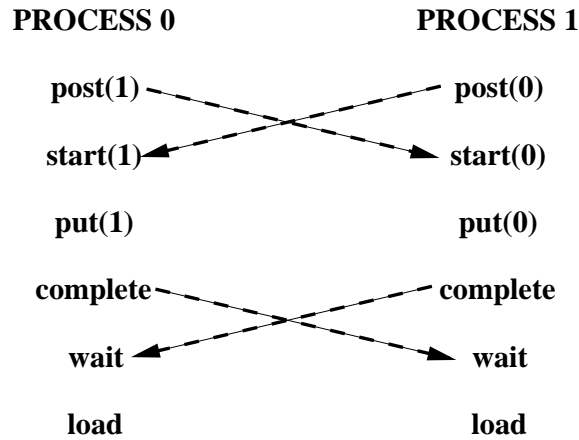


Figure 6.6: Symmetric communication




Figure 6.7: Deadlock situation

Consider the code illustrated in Figure 6.6. Each process updates the window of the other process using a put operation, then accesses its own window. The post calls are nonblocking, and should complete. Once the post calls occur, RMA access to the windows is enabled, so that each process should complete the sequence of calls start-put-complete. Once these are done, the wait calls should complete at both processes. Thus, this communication should not deadlock, irrespective of the amount of data transferred.

Assume, in the last example, that the order of the post and start calls is reversed, at each process. Then, the code may deadlock, as each process may block on the start call, waiting for the matching post to occur. Similarly, the program will deadlock, if the order of the complete and wait calls is reversed, at each process.

The following two examples illustrate the fact that the synchronization between complete and wait is not symmetric: the wait call blocks until the complete executes, but not vice-versa. Consider the code illustrated in Figure 6.7. This code will deadlock: the wait of process 1 blocks until process 0 calls complete, and the receive of process 0 blocks until process 1 calls send. Consider, on the other hand, the code illustrated in Figure 6.8. This code will not deadlock. Once process 1 calls post, then the sequence start, put, complete on process 0 can proceed to completion. Process 0 will reach the send call, allowing the receive call of process 1 to complete.

Rationale. MPI implementations must guarantee that a process makes progress on all enabled communications it participates in, while blocked on an MPI call. This is true for send-receive communication and applies to RMA communication as well. Thus, in the example in Figure 6.8, the put and complete calls of process 0 should complete

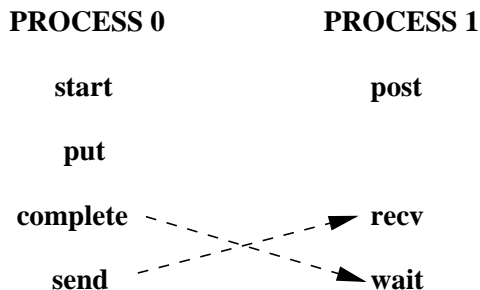


Figure 6.8: No deadlock

while process 1 is blocked on the receive call. This may require the involvement of process 1, e.g., to transfer the data put, while it is blocked on the receive call.

A similar issue is whether such progress must occur while a process is busy computing, or blocked in a non-MPI call. Suppose that in the last example the send-receive pair is replaced by a write-to-socket/read-from-socket pair. Then MPI does not specify whether deadlock is avoided. Suppose that the blocking receive of process 1 is replaced by a very long compute loop. Then, according to one interpretation of the MPI standard, process 0 must return from the complete call after a bounded delay, even if process 1 does not reach any MPI call in this period of time. According to another interpretation, the complete call may block until process 1 reaches the wait call, or reaches another MPI call. The qualitative behavior is the same, under both interpretations, unless a process is caught in an infinite compute loop, in which case the difference may not matter. However, the quantitative expectations are different. Different MPI implementations reflect these different interpretations. While this ambiguity is unfortunate, it does not seem to affect many real codes. The MPI forum decided not to decide which interpretation of the standard is the correct one, since the issue is very contentious, and a decision would have much impact on implementors but less impact on users. (*End of rationale.*)

6.7.3 Registers and Compiler Optimizations

Advice to users. All the material in this section is an advice to users. (*End of advice to users.*)

A coherence problem exists between variables kept in registers and the memory value of these variables. An RMA call may access a variable in memory (or cache), while the up-to-date value of this variable is in register. A get will not return the latest variable value, and a put may be overwritten when the register is stored back in memory.

The problem is illustrated by the following code:

Source of Process 1	Source of Process 2	Executed in Process 2
bbbb = 777	buff = 999	reg_A:=999
call MPI_WIN_FENCE	call MPI_WIN_FENCE	
call MPI_PUT(bbbb		stop appl.thread
into buff of process 2)		buff:=777 in PUT handler
		continue appl.thread
call MPI_WIN_FENCE	call MPI_WIN_FENCE	

```
ccc = buff
```

```
ccc:=reg_A
```

In this example, variable `buff` is allocated in the register `reg_A` and therefore `ccc` will have the old value of `buff` and not the new value `777`.

This problem, which also afflicts in some cases send/receive communication, is discussed more at length in Section 10.2.2.

MPI implementations will avoid this problem for standard conforming C programs. Many Fortran compilers will avoid this problem, without disabling compiler optimizations. However, in order to avoid register coherence problems in a completely portable manner, users should restrict their use of RMA windows to variables stored in `COMMON` blocks, or to variables that were declared `VOLATILE` (while `VOLATILE` is not a standard Fortran declaration, it is supported by many Fortran compilers). Details and an additional solution are discussed in Section 10.2.2, “A Problem with Register Optimization,” on page 289. See also, “Problems Due to Data Copying and Sequence Association,” on page 286, for additional Fortran problems.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 7

Extended Collective Operations

7.1 Introduction

MPI-1 defined collective communication for intracommunicators and two routines, `MPI_INTERCOMM_CREATE` and `MPI_COMM_DUP`, for creating new intercommunicators. In addition, in order to avoid argument aliasing problems with Fortran, MPI-1 requires separate send and receive buffers for collective operations. MPI-2 introduces extensions of many of the MPI-1 collective routines to intercommunicators, additional routines for creating intercommunicators, and two new collective routines: a generalized all-to-all and an exclusive scan. In addition, a way to specify “in place” buffers is provided for many of the intracommunicator collective operations.

7.2 Intercommunicator Constructors

The current MPI interface provides only two intercommunicator construction routines:

- `MPI_INTERCOMM_CREATE`, creates an intercommunicator from two intracommunicators,
- `MPI_COMM_DUP`, duplicates an existing intercommunicator (or intracommunicator).

The other communicator constructors, `MPI_COMM_CREATE` and `MPI_COMM_SPLIT`, currently apply only to intracommunicators. These operations in fact have well-defined semantics for intercommunicators [20].

In the following discussions, the two groups in an intercommunicator are called the *left* and *right* groups. A process in an intercommunicator is a member of either the left or the right group. From the point of view of that process, the group that the process is a member of is called the *local* group; the other group (relative to that process) is the *remote* group. The left and right group labels give us a way to describe the two groups in an intercommunicator that is not relative to any particular process (as the local and remote groups are).

In addition, the specification of collective operations (Section 4.1 of MPI-1) requires that all collective routines are called with matching arguments. For the intercommunicator extensions, this is weakened to matching for all members of the same local group.

```
1 MPI_COMM_CREATE(comm_in, group, comm_out)
```

```
2   IN      comm_in      original communicator (handle)
```

```
3   IN      group        group of processes to be in new communicator (handle)
```

```
4   OUT     comm_out     new communicator (handle)
```

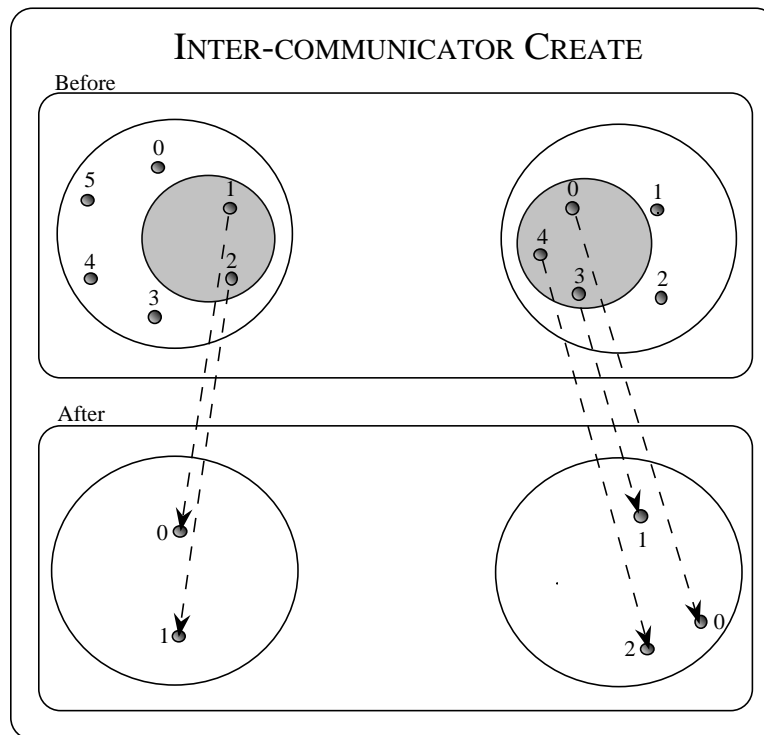
```
5
6
7
8 MPI::Intercomm MPI::Intercomm::Create(const Group& group) const
```

```
9
10 MPI::Intracomm MPI::Intracomm::Create(const Group& group) const
```

11 The C and Fortran language bindings are identical to those in MPI-1, so are omitted here.

12 If **comm_in** is an intercommunicator, then the output communicator is also an intercommunicator where the local group consists only of those processes contained in **group** (see Figure 7.1). The **group** argument should only contain those processes in the local group of the input intercommunicator that are to be a part of **comm_out**. If either **group** does not specify at least one process in the local group of the intercommunicator, or if the calling process is not included in the **group**, MPI_COMM_NULL is returned.

13 *Rationale.* In the case where either the left or right group is empty, a null communicator is returned instead of an intercommunicator with MPI_GROUP_EMPTY because the side with the empty group must return MPI_COMM_NULL. (*End of rationale.*)



46 Figure 7.1: Intercommunicator create using MPI_COMM_CREATE extended to intercommu-
47 nicators. The input groups are those in the grey circle.
48

Example 7.1 The following example illustrates how the first node in the left side of an intercommunicator could be joined with all members on the right side of an intercommunicator to form a new intercommunicator.

```

MPI_Comm  inter_comm, new_inter_comm;
MPI_Group local_group, group;
int        rank = 0; /* rank on left side to include in
                      new inter-comm */

/* Construct the original intercommunicator: "inter_comm" */
...

/* Construct the group of processes to be in new
   intercommunicator */
if (/* I'm on the left side of the intercommunicator */) {
    MPI_Comm_group ( inter_comm, &local_group );
    MPI_Group_incl ( local_group, 1, &rank, &group );
    MPI_Group_free ( &local_group );
}
else
    MPI_Comm_group ( inter_comm, &group );

MPI_Comm_create ( inter_comm, group, &new_inter_comm );
MPI_Group_free( &group );

```

MPI_COMM_SPLIT(comm_in, color, key, comm_out)

IN	comm_in	original communicator (handle)
IN	color	control of subset assignment (integer)
IN	key	control of rank assignment (integer)
OUT	comm_out	new communicator (handle)

MPI::Intercomm MPI::Intercomm::Split(int color, int key) const

MPI::Intracomm MPI::Intracomm::Split(int color, int key) const

The C and Fortran language bindings are identical to those in **MPI-1**, so are omitted here.

The result of **MPI_COMM_SPLIT** on an intercommunicator is that those processes on the left with the same **color** as those processes on the right combine to create a new intercommunicator. The **key** argument describes the relative rank of processes on each side of the intercommunicator (see Figure 7.2). For those colors that are specified only on one side of the intercommunicator, **MPI_COMM_NULL** is returned. **MPI_COMM_NULL** is also returned to those processes that specify **MPI_UNDEFINED** as the color.

Example 7.2 (Parallel client-server model). The following client code illustrates how clients on the left side of an intercommunicator could be assigned to a single server from a pool of servers on the right side of an intercommunicator.

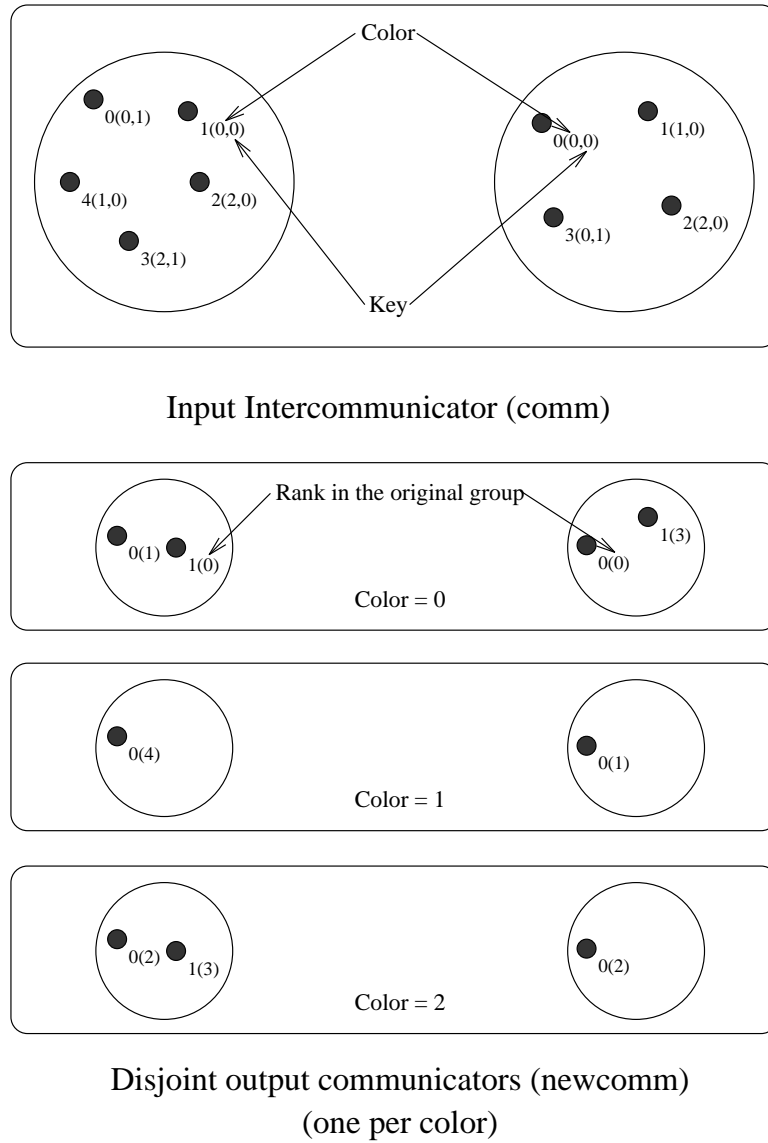


Figure 7.2: Intercommunicator construction achieved by splitting an existing intercommunicator with `MPI_COMM_SPLIT` extended to intercommunicators.

```

/* Client code */
MPI_Comm multiple_server_comm;
MPI_Comm single_server_comm;
int      color, rank, num_servers;

/* Create intercommunicator with clients and servers:
   multiple_server_comm */
...

/* Find out the number of servers available */
MPI_Comm_remote_size ( multiple_server_comm, &num_servers );

/* Determine my color */
MPI_Comm_rank ( multiple_server_comm, &rank );
color = rank % num_servers;

/* Split the intercommunicator */
MPI_Comm_split ( multiple_server_comm, color, rank,
                  &single_server_comm );

```

The following is the corresponding server code:

```

/* Server code */
MPI_Comm multiple_client_comm;
MPI_Comm single_server_comm;
int      rank;

/* Create intercommunicator with clients and servers:
   multiple_client_comm */
...

/* Split the intercommunicator for a single server per group
   of clients */
MPI_Comm_rank ( multiple_client_comm, &rank );
MPI_Comm_split ( multiple_client_comm, rank, 0,
                  &single_server_comm );

```

7.3 Extended Collective Operations

7.3.1 Intercommunicator Collective Operations

In the MPI-1 standard (Section 4.2), collective operations only apply to intracommunicators; however, most MPI collective operations can be generalized to intercommunicators. To understand how MPI can be extended, we can view most MPI intracommunicator collective operations as fitting one of the following categories (see, for instance, [20]):

All-To-All All processes contribute to the result. All processes receive the result.

- MPI_Allgather, MPI_Allgatherv

- `MPI_Alltoall`, `MPI_Alltoallv`
- `MPI_Allreduce`, `MPI_Reduce_scatter`

All-To-One All processes contribute to the result. One process receives the result.

- `MPI_Gather`, `MPI_Gatherv`
- `MPI_Reduce`

One-To-All One process contributes to the result. All processes receive the result.

- `MPI_Bcast`
- `MPI_Scatter`, `MPI_Scatterv`

Other Collective operations that do not fit into one of the above categories.

- `MPI_Scan`
- `MPI_Barrier`

The `MPI_Barrier` operation does not fit into this classification since no data is being moved (other than the implicit fact that a barrier has been called). The data movement pattern of `MPI_Scan` does not fit this taxonomy.

The extension of collective communication from intracommunicators to intercommunicators is best described in terms of the left and right groups. For example, an all-to-all `MPI_Allgather` operation can be described as collecting data from all members of one group with the result appearing in all members of the other group (see Figure 7.3). As another example, a one-to-all `MPI_Bcast` operation sends data from one member of one group to all members of the other group. Collective computation operations such as `MPI_REDUCE_SCATTER` have a similar interpretation (see Figure 7.4). For intracommunicators, these two groups are the same. For intercommunicators, these two groups are distinct. For the all-to-all operations, each such operation is described in two phases, so that it has a symmetric, full-duplex behavior.

For MPI-2, the following intracommunicator collective operations also apply to intercommunicators:

- `MPI_BCAST`,
- `MPI_GATHER`, `MPI_GATHERV`,
- `MPI_SCATTER`, `MPI_SCATTERV`,
- `MPI_ALLGATHER`, `MPI_ALLGATHERV`,
- `MPI_ALLTOALL`, `MPI_ALLTOALLV`, `MPI_ALLTOALLW`
- `MPI_REDUCE`, `MPI_ALLREDUCE`,
- `MPI_REDUCE_SCATTER`,
- `MPI_BARRIER`.

(`MPI_ALLTOALLW` is a new function described in Section 7.3.5.)

These functions use exactly the same argument list as their `MPI-1` counterparts and also work on intracommunicators, as expected. No new language bindings are consequently needed for Fortran or C. However, in C++, the bindings have been “relaxed”; these member functions have been moved from the `MPI::Intercomm` class to the `MPI::Comm` class. But since the collective operations do not make sense on a C++ `MPI::Comm` (since it is neither an intercommunicator nor an intracommunicator), the functions are all pure virtual. In an `MPI-2` implementation, the bindings in this chapter supersede the corresponding bindings for `MPI-1.2`.




Figure 7.3: Intercommunicator allgather. The focus of data to one process is represented, not mandated by the semantics. The two phases do allgathers in both directions.

7.3.2 Operations that Move Data

Two additions are made to many collective communication calls:

- Collective communication can occur “in place” for intracommunicators, with the output buffer being identical to the input buffer. This is specified by providing a special argument value, `MPI_IN_PLACE`, instead of the send buffer or the receive buffer argument.

Rationale. The “in place” operations are provided to reduce unnecessary memory motion by both the `MPI` implementation and by the user. Note that while the simple check of testing whether the send and receive buffers have the same address will work for some cases (e.g., `MPI_ALLREDUCE`), they are inadequate in others (e.g., `MPI_GATHER`, with root not equal to zero). Further, Fortran explicitly prohibits aliasing of arguments; the approach of using a special value to denote “in place” operation eliminates that difficulty. (*End of rationale.*)




Figure 7.4: Intercommunicator reduce-scatter. The focus of data to one process is represented, not mandated by the semantics. The two phases do reduce-scatters in both directions.

Advice to users. By allowing the “in place” option, the receive buffer in many of the collective calls becomes a send-and-receive buffer. For this reason, a Fortran binding that includes `INTENT` must mark these as `INOUT`, not `OUT`.

Note that `MPI_IN_PLACE` is a special kind of value; it has the same restrictions on its use that `MPI_BOTTOM` has.

Some intracommunicator collective operations do not support the “in place” option (e.g., `MPI_ALLTOALLV`). (*End of advice to users.*)

- Collective communication applies to intercommunicators. If the operation is rooted (e.g., broadcast, gather, scatter), then the transfer is unidirectional. The direction of the transfer is indicated by a special value of the root argument. In this case, for the group containing the root process, all processes in the group must call the routine using a special argument for the root. The root process uses the special root value `MPI_ROOT`; all other processes in the same group as the root use `MPI_PROC_NULL`. All processes in the other group (the group that is the remote group relative to the root process) must call the collective routine and provide the rank of the root. If the operation is unrooted (e.g., `alltoall`), then the transfer is bidirectional.

Note that the “in place” option for intracommunicators does not apply to intercommunicators since in the intercommunicator case there is no communication from a process to itself.

Rationale. Rooted operations are unidirectional by nature, and there is a clear way of specifying direction. Non-rooted operations, such as `all-to-all`, will often occur as part of an exchange, where it makes sense to communicate in both directions at once. (*End of rationale.*)

In the following, the definitions of the collective routines are provided to enhance the readability and understanding of the associated text. They do not change the definitions of the argument lists from **MPI-1**. The C and Fortran language bindings for these routines are unchanged from **MPI-1**, and are not repeated here. Since new C++ bindings for the intercommunicator versions are required, they are included. The text provided for each routine is appended to the definition of the routine in **MPI-1**.

Broadcast

MPI_BCAST(buffer, count, datatype, root, comm)

INOUT	buffer	starting address of buffer (choice)
IN	count	number of entries in buffer (integer)
IN	datatype	data type of buffer (handle)
IN	root	rank of broadcast root (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Bcast(void* buffer, int count,
    const MPI::Datatype& datatype, int root) const = 0
```

The “in place” option is not meaningful here.

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI_ROOT** in **root**. All other processes in group A pass the value **MPI_PROC_NULL** in **root**. Data is broadcast from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer argument of the root.

Gather

MPI_GATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	recvcount	number of elements for any single receive (integer, significant only at root)
IN	recvtype	data type of recv buffer elements (handle, significant only at root)
IN	root	rank of receiving process (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Gather(const void* sendbuf, int sendcount, const
    MPI::Datatype& sendtype, void* recvbuf, int recvcount,
    const MPI::Datatype& recvtype, int root) const = 0
```

The “in place” option for intracommunicators is specified by passing **MPI_IN_PLACE** as the value of **sendbuf** at the root. In such a case, **sendcount** and **sendtype** are ignored, and the contribution of the root to the gathered vector is assumed to be already in the correct place in the receive buffer

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI_ROOT** in **root**. All other processes in group A pass the value **MPI_PROC_NULL** in **root**. Data is gathered from all processes in group B to the root. The send buffer arguments of the processes in group B must be consistent with the receive buffer argument of the root.

MPI_GATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtype, root, comm)

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	recvcunts	integer array (of length group size) containing the number of elements that are received from each process (significant only at root)
IN	displs	integer array (of length group size). Entry i specifies the displacement relative to recvbuf at which to place the incoming data from process i (significant only at root)
IN	recvtype	data type of recv buffer elements (handle, significant only at root)
IN	root	rank of receiving process (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Gatherv(const void* sendbuf, int sendcount, const
    MPI::Datatype& sendtype, void* recvbuf,
    const int recvcunts[], const int displs[],
    const MPI::Datatype& recvtype, int root) const = 0
```

The “in place” option for intracommunicators is specified by passing **MPI_IN_PLACE** as the value of **sendbuf** at the root. In such a case, **sendcount** and **sendtype** are ignored, and the contribution of the root to the gathered vector is assumed to be already in the correct place in the receive buffer

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI_ROOT** in **root**. All other processes in group A pass the value **MPI_PROC_NULL** in **root**. Data is gathered from all processes in group B to the root. The send buffer arguments of the processes in group B must be consistent with the receive buffer argument of the root.

Scatter

MPI_SCATTER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)

IN	sendbuf	address of send buffer (choice, significant only at root)
IN	sendcount	number of elements sent to each process (integer, significant only at root)
IN	sendtype	data type of send buffer elements (handle, significant only at root)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements in receive buffer (integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	root	rank of sending process (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Scatter(const void* sendbuf, int sendcount, const
                        MPI::Datatype& sendtype, void* recvbuf, int recvcount,
                        const MPI::Datatype& recvtype, int root) const = 0
```

The “in place” option for intracommunicators is specified by passing `MPI_IN_PLACE` as the value of **recvbuf** at the root. In such case, **recvcount** and **recvtype** are ignored, and root “sends” no data to itself. The scattered vector is still assumed to contain n segments, where n is the group size; the *root*-th segment, which root should “send to itself,” is not moved.

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value `MPI_ROOT` in **root**. All other processes in group A pass the value `MPI_PROC_NULL` in **root**. Data is scattered from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer argument of the root.

MPI_SCATTERV(sendbuf, sendcounts, displs, sendtype, recvbuf, recvcount, recvtype, root, comm)

IN	sendbuf	address of send buffer (choice, significant only at root)
IN	sendcounts	integer array (of length group size) specifying the number of elements to send to each processor
IN	displs	integer array (of length group size). Entry i specifies the displacement (relative to sendbuf from which to take the outgoing data to process i)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements in receive buffer (integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	root	rank of sending process (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Scatterv(const void* sendbuf, const int sendcounts[],
                        const int displs[], const MPI::Datatype& sendtype,
                        void* recvbuf, int recvcount, const MPI::Datatype& recvtype,
                        int root) const = 0
```

The “in place” option for intracommunicators is specified by passing **MPI_IN_PLACE** as the value of **recvbuf** at the root. In such case, **recvcount** and **recvtype** are ignored, and root “sends” no data to itself. The scattered vector is still assumed to contain n segments, where n is the group size; the *root*-th segment, which root should “send to itself,” is not moved.

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI_ROOT** in **root**. All other processes in group A pass the value **MPI_PROC_NULL** in **root**. Data is scattered from the root to all processes in group B. The receive buffer arguments of the processes in group B must be consistent with the send buffer argument of the root.

“All” Forms and All-to-all

```
MPI_ALLGATHER(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)
```

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcount	number of elements received from any process (integer)
IN	recvtype	data type of receive buffer elements (handle)
IN	comm	communicator (handle)

```
void MPI::Comm::Allgather(const void* sendbuf, int sendcount, const
    MPI::Datatype& sendtype, void* recvbuf, int recvcount,
    const MPI::Datatype& recvtype) const = 0
```

The “in place” option for intracommunicators is specified by passing the value `MPI_IN_PLACE` to the argument **sendbuf** at all processes. **sendcount** and **sendtype** are ignored. Then the input data of each process is assumed to be in the area where that process would receive its own contribution to the receive buffer. Specifically, the outcome of a call to `MPI_ALLGATHER` in the “in place” case is as if all processes executed n calls to

```
MPI_GATHER( MPI_IN_PLACE, 0, MPI_DATATYPE_NULL, recvbuf, recvcount,
    recvtype, root, comm )
```

for $\text{root} = 0, \dots, n - 1$.

If **comm** is an intercommunicator, then each process in group A contributes a data item; these items are concatenated and the result is stored at each process in group B. Conversely the concatenation of the contributions of the processes in group B is stored at each process in group A. The send buffer arguments in group A must be consistent with the receive buffer arguments in group B, and vice versa.

Advice to users. The communication pattern of `MPI_ALLGATHER` executed on an intercommunication domain need not be symmetric. The number of items sent by processes in group A (as specified by the arguments **sendcount**, **sendtype** in group A and the arguments **recvcount**, **recvtype** in group B), need not equal the number of items sent by processes in group B (as specified by the arguments **sendcount**, **sendtype** in group B and the arguments **recvcount**, **recvtype** in group A). In particular, one can move data in only one direction by specifying **sendcount** = 0 for the communication in the reverse direction.

(End of advice to users.)

MPI_ALLGATHERV(sendbuf, sendcount, sendtype, recvbuf, recvcoun-
ts, displs, recvtype, comm)

IN	sendbuf	starting address of send buffer (choice)
IN	sendcount	number of elements in send buffer (integer)
IN	sendtype	data type of send buffer elements (handle)
OUT	recvbuf	address of receive buffer (choice)
IN	recvcoun- ts	integer array (of length group size) containing the num- ber of elements that are received from each process
IN	displs	integer array (of length group size). Entry i specifies the displacement (relative to recvbuf) at which to place the incoming data from process i
IN	recvtype	data type of receive buffer elements (handle)
IN	comm	communicator (handle)

```
void MPI::Comm::Allgatherv(const void* sendbuf, int sendcount, const
    MPI::Datatype& sendtype, void* recvbuf,
    const int recvcoun-
    ts[], const int displs[],
    const MPI::Datatype& recvtype) const = 0
```

The “in place” option for intracommunicators is specified by passing the value **MPI_IN_PLACE** to the argument **sendbuf** at all processes. **sendcount** and **sendtype** are ignored. Then the input data of each process is assumed to be in the area where that process would receive its own contribution to the receive buffer. Specifically, the outcome of a call to **MPI_ALLGATHER** in the “in place” case is as if all processes executed n calls to

```
MPI_GATHERV( MPI_IN_PLACE, 0, MPI_DATATYPE_NULL, recvbuf, recvcoun-
    ts, displs, recvtype, root, comm )
```

for $\text{root} = 0, \dots, n - 1$.

If **comm** is an intercommunicator, then each process in group A contributes a data item; these items are concatenated and the result is stored at each process in group B. Conversely the concatenation of the contributions of the processes in group B is stored at each process in group A. The send buffer arguments in group A must be consistent with the receive buffer arguments in group B, and vice versa.

```

1 MPI_ALLTOALL(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)
2     IN      sendbuf      starting address of send buffer (choice)
3     IN      sendcount    number of elements sent to each process (integer)
4     IN      sendtype     data type of send buffer elements (handle)
5     OUT     recvbuf      address of receive buffer (choice)
6     IN      recvcount    number of elements received from any process (inte-
7                          ger)
8     IN      recvtype     data type of receive buffer elements (handle)
9     IN      comm         communicator (handle)
10
11 void MPI::Comm::Alltoall(const void* sendbuf, int sendcount, const
12                          MPI::Datatype& sendtype, void* recvbuf, int recvcount,
13                          const MPI::Datatype& recvtype) const = 0
14
15 No “in place” option is supported.
16
17 If comm is an intercommunicator, then the outcome is as if each process in group A
18 sends a message to each process in group B, and vice versa. The  $j$ -th send buffer of process
19  $i$  in group A should be consistent with the  $i$ -th receive buffer of process  $j$  in group B, and
20 vice versa.
21
22
23 Advice to users. When all-to-all is executed on an intercommunication domain, then
24 the number of data items sent from processes in group A to processes in group B need
25 not equal the number of items sent in the reverse direction. In particular, one can have
26 unidirectional communication by specifying sendcount = 0 in the reverse direction.
27
28 (End of advice to users.)
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

MPI_ALLTOALLV(sendbuf, sendcounts, sdispls, sendtype, recvbuf, recvcoun-
ts, rdispls, recvtype, comm)

IN	sendbuf	starting address of send buffer (choice)	
IN	sendcounts	integer array equal to the group size specifying the number of elements to send to each processor	
IN	sdispls	integer array (of length group size). Entry <i>j</i> specifies the displacement (relative to sendbuf) from which to take the outgoing data destined for process <i>j</i>	
IN	sendtype	data type of send buffer elements (handle)	
OUT	recvbuf	address of receive buffer (choice)	
IN	recvcoun- ts	integer array equal to the group size specifying the number of elements that can be received from each processor	
IN	rdispls	integer array (of length group size). Entry <i>i</i> specifies the displacement (relative to recvbuf) at which to place the incoming data from process <i>i</i>	
IN	recvtype	data type of receive buffer elements (handle)	
IN	comm	communicator (handle)	

```
void MPI::Comm::Alltoallv(const void* sendbuf, const int sendcounts[],
                          const int sdispls[], const MPI::Datatype& sendtype,
                          void* recvbuf, const int recvcoun-
                          ts[], const int rdispls[],
                          const MPI::Datatype& recvtype) const = 0
```

No “in place” option is supported.

If **comm** is an intercommunicator, then the outcome is as if each process in group A sends a message to each process in group B, and vice versa. The *j*-th send buffer of process *i* in group A should be consistent with the *i*-th receive buffer of process *j* in group B, and vice versa.

7.3.3 Reductions

MPI_REDUCE(sendbuf, recvbuf, count, datatype, op, root, comm)

IN	sendbuf	address of send buffer (choice)
OUT	recvbuf	address of receive buffer (choice, significant only at root)
IN	count	number of elements in send buffer (integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	reduce operation (handle)
IN	root	rank of root process (integer)
IN	comm	communicator (handle)

```
void MPI::Comm::Reduce(const void* sendbuf, void* recvbuf, int count,
                      const MPI::Datatype& datatype, const MPI::Op& op, int root)
{
    const = 0
}
```

The “in place” option for intracommunicators is specified by passing the value **MPI_IN_PLACE** to the argument **sendbuf** at the root. In such case, the input data is taken at the root from the receive buffer, where it will be replaced by the output data.

If **comm** is an intercommunicator, then the call involves all processes in the intercommunicator, but with one group (group A) defining the root process. All processes in the other group (group B) pass the same value in argument **root**, which is the rank of the root in group A. The root passes the value **MPI_ROOT** in **root**. All other processes in group A pass the value **MPI_PROC_NULL** in **root**. Only send buffer arguments are significant in group B and only receive buffer arguments are significant at the root.

MPI_ALLREDUCE(sendbuf, recvbuf, count, datatype, op, comm)

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in send buffer (integer)
IN	datatype	data type of elements of send buffer (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)

```
void MPI::Comm::Allreduce(const void* sendbuf, void* recvbuf, int count,
                        const MPI::Datatype& datatype, const MPI::Op& op) const = 0
```

The “in place” option for intracommunicators is specified by passing the value **MPI_IN_PLACE** to the argument **sendbuf** at the root. In such case, the input data is taken at each process from the receive buffer, where it will be replaced by the output data.

If **comm** is an intercommunicator, then the result of the reduction of the data provided by processes in group A is stored at each process in group B, and vice versa. Both groups should provide the same **count** value.

MPI_REDUCE_SCATTER(sendbuf, recvbuf, recvcunts, datatype, op, comm)

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	recvcunts	integer array specifying the number of elements in result distributed to each process. Array must be identical on all calling processes.
IN	datatype	data type of elements of input buffer (handle)
IN	op	operation (handle)
IN	comm	communicator (handle)

```
void MPI::Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,
                              int recvcunts[], const MPI::Datatype& datatype,
                              const MPI::Op& op) const = 0
```

The “in place” option for intracommunicators is specified by passing **MPI_IN_PLACE** in the **sendbuf** argument. In this case, the input data is taken from the top of the receive buffer. Note that the area occupied by the input data may be either longer or shorter than the data filled by the output data.

If **comm** is an intercommunicator, then the result of the reduction of the data provided by processes in group A is scattered among processes in group B, and vice versa. Within each group, all processes provide the same **recvcunts** argument, and the sum of the **recvcunts** entries should be the same for the two groups.

Rationale. The last restriction is needed so that the length of the send buffer can be determined by the sum of the local **recvcunts** entries. Otherwise, a communication is needed to figure out how many elements are reduced. (*End of rationale.*)

7.3.4 Other Operations

MPI_BARRIER(comm)

IN	comm	communicator (handle)
----	-------------	-----------------------

```
void MPI::Comm::Barrier() const = 0
```

For **MPI-2**, **comm** may be an intercommunicator or an intracommunicator. If **comm** is an intercommunicator, the barrier is performed across all processes in the intercommunicator. In this case, all processes in the local group of the intercommunicator may exit the barrier when all of the processes in the remote group have entered the barrier.

1 **MPI_SCAN**(sendbuf, recvbuf, count, datatype, op, comm)

2	IN	sendbuf	starting address of send buffer (choice)
3	OUT	recvbuf	starting address of receive buffer (choice)
4			
5	IN	count	number of elements in input buffer (integer)
6	IN	datatype	data type of elements of input buffer (handle)
7			
8	IN	op	operation (handle)
9	IN	comm	communicator (handle)

10
11 The “in place” option for intracommunicators is specified by passing **MPI_IN_PLACE** in
12 the **sendbuf** argument. In this case, the input data is taken from the receive buffer, and
13 replaced by the output data.

14 This operation is illegal for intercommunicators.

15 7.3.5 Generalized All-to-all Function

16
17 One of the basic data movement operations needed in parallel signal processing is the 2-D
18 matrix transpose. This operation has motivated a generalization of the **MPI_ALLTOALLV**
19 function. This new collective operation is **MPI_ALLTOALLW**; the “W” indicates that it is
20 an extension to **MPI_ALLTOALLV**.

21 The following function is the most general form of **All-to-all**. Like
22 **MPI_TYPE_CREATE_STRUCT**, the most general type constructor, **MPI_ALLTOALLW** allows
23 separate specification of count, displacement and datatype. In addition, to allow maximum
24 flexibility, the displacement of blocks within the send and receive buffers is specified in
25 bytes.
26

27
28 *Rationale.* The **MPI_ALLTOALLW** function generalizes several MPI functions by care-
29 fully selecting the input arguments. For example, by making all but one process have
30 **sendcounts[i] = 0**, this achieves an **MPI_SCATTERW** function. (*End of rationale.*)

31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

MPI_ALLTOALLW(sendbuf, sendcounts, sdispls, sendtypes, recvbuf, recvcoun-
ts, rdispls, recv-
types, comm)

IN	sendbuf	starting address of send buffer (choice)	1
			2
IN	sendcounts	integer array equal to the group size specifying the number of elements to send to each processor (integer)	3
			4
IN	sdispls	integer array (of length group size). Entry <i>j</i> specifies the displacement in bytes (relative to sendbuf) from which to take the outgoing data destined for process <i>j</i>	5
			6
IN	sendtypes	array of datatypes (of length group size). Entry <i>j</i> spec- ifies the type of data to send to process <i>j</i> (handle)	7
			8
OUT	recvbuf	address of receive buffer (choice)	9
			10
IN	recvcoun-	integer array equal to the group size specifying the number of elements that can be received from each processor (integer)	11
	ts		12
IN	rdispls	integer array (of length group size). Entry <i>i</i> specifies the displacement in bytes (relative to recvbuf) at which to place the incoming data from process <i>i</i>	13
			14
IN	recvtypes	array of datatypes (of length group size). Entry <i>i</i> spec- ifies the type of data received from process <i>i</i> (handle)	15
			16
IN	comm	communicator (handle)	17
			18
			19
			20
			21
			22
			23
			24

```

int MPI_Alltoallw(void *sendbuf, int sendcounts[], int sdispls[],
                  MPI_Datatype sendtypes[], void *recvbuf, int recvcoun-
                  ts[], int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm)
MPI_ALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF, RECVCOUNTS,
               RDISPLS, RECVTYPES, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), REVCOUNTS(*),
RDISPLS(*), RECVTYPES(*), COMM, IERROR

void MPI::Comm::Alltoallw(const void* sendbuf, const int sendcounts[],
                          const int sdispls[], const MPI::Datatype sendtypes[], void*
                          recvbuf, const int recvcoun-
                          ts[], const int rdispls[], const
                          MPI::Datatype recvtypes[]) const = 0

```

No “in place” option is supported.

The *j*-th block sent from process *i* is received by process *j* and is placed in the *i*-th block of **recvbuf**. These blocks need not all have the same size.

The type signature associated with **sendcounts**[*j*], **sendtypes**[*j*] at process *i* must be equal to the type signature associated with **recvcoun-**ts[*i*], **recvtypes**[*i*] at process *j*. This implies that the amount of data sent must be equal to the amount of data received, pairwise between every pair of processes. Distinct type maps between sender and receiver are still allowed.

The outcome is as if each process sent a message to every other process with

```
MPI_Send(sendbuf + sdispls[i], sendcounts[i], sendtypes[i], i, ...),
```

and received a message from every other process with a call to

```
MPI_Recv(recvbuf + rdispls[i], recvcunts[i], recvtypes[i], i, ...).
```

All arguments on all processes are significant. The argument **comm** must describe the same communicator on all processes.

If **comm** is an intercommunicator, then the outcome is as if each process in group A sends a message to each process in group B, and vice versa. The j -th send buffer of process i in group A should be consistent with the i -th receive buffer of process j in group B, and vice versa.

7.3.6 Exclusive Scan

MPI-1 provides an inclusive scan operation. The exclusive scan is described here.

```
MPI_EXSCAN(sendbuf, recvbuf, count, datatype, op, comm)
```

IN	sendbuf	starting address of send buffer (choice)
OUT	recvbuf	starting address of receive buffer (choice)
IN	count	number of elements in input buffer (integer)
IN	datatype	data type of elements of input buffer (handle)
IN	op	operation (handle)
IN	comm	intracommunicator (handle)

```
int MPI_Exscan(void *sendbuf, void *recvbuf, int count,
               MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

```
MPI_EXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
<type> SENDBUF(*), RECVBUF(*)
INTEGER COUNT, DATATYPE, OP, COMM, IERROR
```

```
void MPI::Intracomm::Exscan(const void* sendbuf, void* recvbuf, int count,
                           const MPI::Datatype& datatype, const MPI::Op& op) const
```

MPI_EXSCAN is used to perform a prefix reduction on data distributed across the group. The value in **recvbuf** on the process with rank 0 is undefined, and **recvbuf** is not significant on process 0. The value in **recvbuf** on the process with rank 1 is defined as the value in **sendbuf** on the process with rank 0. For processes with rank $i > 1$, the operation returns, in the receive buffer of the process with rank i , the reduction of the values in the send buffers of processes with ranks $0, \dots, i - 1$ (inclusive). The type of operations supported, their semantics, and the constraints on send and receive buffers, are as for MPI_REDUCE.

No “in place” option is supported.

Advice to users. As for MPI_SCAN, MPI does not specify which processes may call the operation, only that the result be correctly computed. In particular, note that the process with rank 1 need not call the MPI_Op, since all it needs to do is to receive the value from the process with rank 0. However, all processes, even the processes with ranks zero and one, must provide the same **op**. (*End of advice to users.*)

Rationale. The exclusive scan is more general than the inclusive scan provided in MPI-1 as `MPI_SCAN`. Any inclusive scan operation can be achieved by using the exclusive scan and then locally combining the local contribution. Note that for non-invertable operations such as `MPI_MAX`, the exclusive scan cannot be computed with the inclusive scan.

The reason that MPI-1 chose the inclusive scan is that the definition of behavior on processes zero and one was thought to offer too many complexities in definition, particularly for user-defined operations. (*End of rationale.*)

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 8

External Interfaces

8.1 Introduction

This chapter begins with calls used to create **generalized requests**. The objective of this MPI-2 addition is to allow users of MPI to be able to create new nonblocking operations with an interface similar to what is present in MPI. This can be used to layer new functionality on top of MPI. Next, Section 8.3 deals with setting the information found in **status**. This is needed for generalized requests.

Section 8.4 allows users to associate names with communicators, windows, and datatypes. This will allow debuggers and profilers to identify communicators, windows, and datatypes with more useful labels. Section 8.5 allows users to add error codes, classes, and strings to MPI. With users being able to layer functionality on top of MPI, it is desirable for them to use the same error mechanisms found in MPI.

Section 8.6 deals with decoding datatypes. The opaque datatype object has found a number of uses outside MPI. Furthermore, a number of tools wish to display internal information about a datatype. To achieve this, datatype decoding functions are provided.

The chapter continues, in Section 8.7, with a discussion of how threads are to be handled in MPI-2. Although thread compliance is not required, the standard specifies how threads are to work if they are provided. Section 8.8 has information on caching on communicators, datatypes, and windows. Finally, Section 8.9 discusses duplicating a datatype.

8.2 Generalized Requests

The goal of this MPI-2 extension is to allow users to define new nonblocking operations. Such an outstanding nonblocking operation is represented by a (generalized) request. A fundamental property of nonblocking operations is that progress toward the completion of this operation occurs asynchronously, i.e., concurrently with normal program execution. Typically, this requires execution of code concurrently with the execution of the user code, e.g., in a separate thread or in a signal handler. Operating systems provide a variety of mechanisms in support of concurrent execution. MPI does not attempt to standardize or replace these mechanisms: it is assumed programmers who wish to define new asynchronous operations will use the mechanisms provided by the underlying operating system. Thus, the calls in this section only provide a means for defining the effect of MPI calls such as MPI_WAIT or MPI_CANCEL when they apply to generalized requests, and for signaling to MPI the completion of a generalized operation.

Rationale. It is tempting to also define an MPI standard mechanism for achieving concurrent execution of user-defined nonblocking operations. However, it is very difficult to define such a mechanism without consideration of the specific mechanisms used in the operating system. The Forum feels that concurrency mechanisms are a proper part of the underlying operating system and should not be standardized by MPI; the MPI standard should only deal with the interaction of such mechanisms with MPI. (*End of rationale.*)

For a regular request, the operation associated with the request is performed by the MPI implementation, and the operation completes without intervention by the application. For a generalized request, the operation associated with the request is performed by the application; therefore, the application must notify MPI when the operation completes. This is done by making a call to `MPI_GREQUEST_COMPLETE`. MPI maintains the “completion” status of generalized requests. Any other request state has to be maintained by the user.

A new generalized request is started with

`MPI_GREQUEST_START(query_fn, free_fn, cancel_fn, extra_state, request)`

IN	<code>query_fn</code>	callback function invoked when request status is queried (function)
IN	<code>free_fn</code>	callback function invoked when request is freed (function)
IN	<code>cancel_fn</code>	callback function invoked when request is cancelled (function)
IN	<code>extra_state</code>	extra state
OUT	<code>request</code>	generalized request (handle)

```
int MPI_Grequest_start(MPI_Grequest_query_function *query_fn,
                      MPI_Grequest_free_function *free_fn,
                      MPI_Grequest_cancel_function *cancel_fn, void *extra_state,
                      MPI_Request *request)
```

```
MPI_GREQUEST_START(QUERY_FN, FREE_FN, CANCEL_FN, EXTRA_STATE, REQUEST,
                   IERROR)
INTEGER REQUEST, IERROR
EXTERNAL QUERY_FN, FREE_FN, CANCEL_FN
INTEGER (KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
static MPI::Grequest
    MPI::Grequest::Start(const MPI::Grequest::Query_function
                        query_fn, const MPI::Grequest::Free_function free_fn,
                        const MPI::Grequest::Cancel_function cancel_fn,
                        void *extra_state)
```

Advice to users. Note that a generalized request belongs, in C++, to the class `MPI::Grequest`, which is a derived class of `MPI::Request`. It is of the same type as regular requests, in C and Fortran. (*End of advice to users.*)

The call starts a generalized request and returns a handle to it in **request**.

The syntax and meaning of the callback functions are listed below. All callback functions are passed the **extra_state** argument that was associated with the request by the starting call **MPI_GREQUEST_START**. This can be used to maintain user-defined state for the request. In C, the query function is

```
typedef int MPI_Grequest_query_function(void *extra_state,
                                       MPI_Status *status);
```

in Fortran

```
SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
  INTEGER STATUS(MPI_STATUS_SIZE), IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
typedef int MPI::Grequest::Query_function(void* extra_state,
                                       MPI::Status& status);
```

query_fn function computes the status that should be returned for the generalized request. The status also includes information about successful/unsuccessful cancellation of the request (result to be returned by **MPI_TEST_CANCELLED**).

query_fn callback is invoked by the **MPI_{WAIT|TEST}{ANY|SOME|ALL}** call that completed the generalized request associated with this callback. The callback function is also invoked by calls to **MPI_REQUEST_GET_STATUS**, if the request is complete when the call occurs. In both cases, the callback is passed a reference to the corresponding status variable passed by the user to the MPI call; the status set by the callback function is returned by the MPI call. If the user provided **MPI_STATUS_IGNORE** or **MPI_STATUSES_IGNORE** to the MPI function that causes **query_fn** to be called, then MPI will pass a valid status object to **query_fn**, and this status will be ignored upon return of the callback function. Note that **query_fn** is invoked only after **MPI_GREQUEST_COMPLETE** is called on the request; it may be invoked several times for the same generalized request, e.g., if the user calls **MPI_REQUEST_GET_STATUS** several times for this request. Note also that a call to **MPI_{WAIT|TEST}{SOME|ALL}** may cause multiple invocations of **query_fn** callback functions, one for each generalized request that is completed by the MPI call. The order of these invocations is not specified by MPI.

In C, the free function is

```
typedef int MPI_Grequest_free_function(void *extra_state);
```

and in Fortran

```
SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

and in C++

```
typedef int MPI::Grequest::Free_function(void* extra_state);
```

free_fn function is invoked to clean up user-allocated resources when the generalized request is freed.

free_fn callback is invoked by the **MPI_{WAIT|TEST}{ANY|SOME|ALL}** call that completed the generalized request associated with this callback. **free_fn** is invoked after the call

to `query_fn` for the same request. However, if the MPI call completed multiple generalized requests, the order in which `free_fn` callback functions are invoked is not specified by MPI.

`free_fn` callback is also invoked for generalized requests that are freed by a call to `MPI_REQUEST_FREE` (no call to `WAIT_{WAIT|TEST}_{ANY|SOME|ALL}` will occur for such a request). In this case, the callback function will be called either in the MPI call `MPI_REQUEST_FREE(request)`, or in the MPI call `MPI_GREQUEST_COMPLETE(request)`, whichever happens last. I.e., in this case the actual freeing code is executed as soon as both calls `MPI_REQUEST_FREE` and `MPI_GREQUEST_COMPLETE` have occurred. The `request` is not deallocated until after `free_fn` completes. Note that `free_fn` will be invoked only once per request by a correct program.

Advice to users. Calling `MPI_REQUEST_FREE(request)` will cause the `request` handle to be set to `MPI_REQUEST_NULL`. This handle to the generalized request is no longer valid. However, user copies of this handle are valid until after `free_fn` completes since MPI does not deallocate the object until then. Since `free_fn` is not called until after `MPI_GREQUEST_COMPLETE`, the user copy of the handle can be used to make this call. Users should note that MPI will deallocate the object after `free_fn` executes. At this point, user copies of the `request` handle no longer point to a valid request. MPI will not set user copies to `MPI_REQUEST_NULL` in this case, so it is up to the user to avoid accessing this stale handle. This is a special case where MPI defers deallocating the object until a later time that is known by the user. (*End of advice to users.*)

In C, the cancel function is

```
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);
```

in Fortran

```
SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
  INTEGER IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
  LOGICAL COMPLETE
```

and in C++

```
typedef int MPI::Grequest::Cancel_function(void* extra_state,
      bool complete);
```

`cancel_fn` function is invoked to start the cancellation of a generalized request. It is called by `MPI_REQUEST_CANCEL(request)`. MPI passes to the callback function `complete=true` if `MPI_GREQUEST_COMPLETE` was already called on the request, and `complete=false` otherwise.

All callback functions return an error code. The code is passed back and dealt with as appropriate for the error code by the MPI function that invoked the callback function. For example, if error codes are returned then the error code returned by the callback function will be returned by the MPI function that invoked the callback function. In the case of `MPI_{WAIT|TEST}_{ANY}` call that invokes both `query_fn` and `free_fn`, the MPI call will return the error code returned by the last callback, namely `free_fn`. If one or more of the requests in a call to `MPI_{WAIT|TEST}_{SOME|ALL}` failed, then the MPI call will return `MPI_ERR_IN_STATUS`. In such a case, if the MPI call was passed an array of statuses, then MPI will return in each of the statuses that correspond to a completed generalized request

Advice to users. `query_fn` must **not** set the error field of `status` since `query_fn` may be called by `MPI_WAIT` or `MPI_TEST`, in which case the error field of `status` should not change. The MPI library knows the “context” in which `query_fn` is invoked and can decide correctly when to put in the error field of `status` the returned error code. (*End of advice to users.*)

```
typedef struct {
    MPI_Comm comm;
    int tag;
}
```

```

1      int root;
2      int valin;
3      int *valout;
4      MPI_Request request;
5      } ARGS;
6
7
8      int myreduce(MPI_Comm comm, int tag, int root,
9                  int valin, int *valout, MPI_Request *request)
10     {
11         ARGS *args;
12         pthread_t thread;
13
14         /* start request */
15         MPI_Grequest_start(query_fn, free_fn, cancel_fn, NULL, request);
16
17         args = (ARGS*)malloc(sizeof(ARGS));
18         args->comm = comm;
19         args->tag = tag;
20         args->root = root;
21         args->valin = valin;
22         args->valout = valout;
23         args->request = *request;
24
25         /* spawn thread to handle request */
26         /* The availability of the pthread_create call is system dependent */
27         pthread_create(&thread, NULL, reduce_thread, args);
28
29         return MPI_SUCCESS;
30     }
31
32
33     /* thread code */
34     void reduce_thread(void *ptr)
35     {
36         int lchild, rchild, parent, lval, rval, val;
37         MPI_Request req[2];
38         ARGS *args;
39
40         args = (ARGS*)ptr;
41
42         /* compute left,right child and parent in tree; set
43            to MPI_PROC_NULL if does not exist */
44         /* code not shown */
45         ...
46
47         MPI_Irecv(&lval, 1, MPI_INT, lchild, args->tag, args->comm, &req[0]);
48         MPI_Irecv(&rval, 1, MPI_INT, rchild, args->tag, args->comm, &req[1]);

```

```

MPI_Waitall(2, req, MPI_STATUSES_IGNORE);           1
val = lval + args->valin + rval;                     2
MPI_Send( &val, 1, MPI_INT, parent, args->tag, args->comm ); 3
if (parent == MPI_PROC_NULL) *(args->valout) = val;    4
MPI_Grequest_complete((args->request));               5
free(ptr);                                           6
return;                                              7
}                                                    8

                                                    9
int query_fn(void *extra_state, MPI_Status *status)   10
{                                                    11
/* always send just one int */                      12
MPI_Status_set_elements(status, MPI_INT, 1);         13
/* can never cancel so always true */                14
MPI_Status_set_cancelled(status, 0);                 15
/* choose not to return a value for this */           16
status->MPI_SOURCE = MPI_UNDEFINED;                   17
/* tag has not meaning for this generalized request */ 18
status->MPI_TAG = MPI_UNDEFINED;                      19
/* this generalized request never fails */            20
return MPI_SUCCESS;                                  21
}                                                    22

                                                    23
                                                    24
int free_fn(void *extra_state)                       25
{                                                    26
/* this generalized request does not need to do any freeing */ 27
/* as a result it never fails here */                28
return MPI_SUCCESS;                                  29
}                                                    30

                                                    31
                                                    32
int cancel_fn(void *extra_state, int complete)        33
{                                                    34
/* This generalized request does not support cancelling.    35
   Abort if not already done. If done then treat as if cancel failed. */ 36
if (!complete) {                                       37
    fprintf(stderr, "Cannot cancel generalized request - aborting program\n"); 38
    MPI_Abort(MPI_COMM_WORLD, 99);                   39
}                                                      40
return MPI_SUCCESS;                                   41
}                                                      42
}                                                      43

```

8.3 Associating Information with Status

In MPI-1, requests were associated with point-to-point operations. In MPI-2 there are several different types of requests. These range from new MPI calls for I/O to generalized requests.

It is desirable to allow these calls use the same request mechanism. This allows one to wait or test on different types of requests. However, `MPI_{TEST|WAIT}_{ANY|SOME|ALL}` returns a status with information about the request. With the generalization of requests, one needs to define what information will be returned in the status object.

In MPI-2, each call fills in the appropriate fields in the status object. Any unused fields will have undefined values. A call to `MPI_{TEST|WAIT}_{ANY|SOME|ALL}` can modify any of the fields in the status object. Specifically, it can modify fields that are undefined. The fields with meaningful value for a given request are defined in the sections with the new request.

Generalized requests raise additional considerations. Here, the user provides the functions to deal with the request. Unlike other MPI calls, the user needs to provide the information to be returned in status. The status argument is provided directly to the callback function where the status needs to be set. Users can directly set the values in 3 of the 5 status values. The count and cancel fields are opaque. To overcome this, new calls are provided:

MPI_STATUS_SET_ELEMENTS(status, datatype, count)

INOUT	status	status to associate count with (Status)
IN	datatype	datatype associated with count (handle)
IN	count	number of elements to associate with status (integer)

```
int MPI_Status_set_elements(MPI_Status *status, MPI_Datatype datatype,
                           int count)
```

```
MPI_STATUS_SET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)
```

```
INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
```

```
void MPI::Status::Set_elements(const MPI::Datatype& datatype, int count)
```

This call modifies the opaque part of **status** so that a call to `MPI_GET_ELEMENTS` will return **count**. `MPI_GET_COUNT` will return a compatible value.

Rationale. The number of elements is set instead of the count because the former can deal with nonintegral number of datatypes. (*End of rationale.*)

A subsequent call to `MPI_GET_COUNT(status, datatype, count)` or to `MPI_GET_ELEMENTS(status, datatype, count)` must use a **datatype** argument that has the same type signature as the **datatype** argument that was used in the call to `MPI_STATUS_SET_ELEMENTS`.

Rationale. This is similar to the restriction that holds when when **count** is set by a receive operation: in that case, the calls to `MPI_GET_COUNT` and `MPI_GET_ELEMENTS` must use a **datatype** with the same signature as the datatype used in the receive call. (*End of rationale.*)

```
MPI_STATUS_SET_CANCELLED(status, flag)
```

```
INOUT  status          status to associate cancel flag with (Status)
IN      flag            if true indicates request was cancelled (logical)
```

```
int MPI_Status_set_cancelled(MPI_Status *status, int flag)
```

```
MPI_STATUS_SET_CANCELLED(STATUS, FLAG, IERROR)
    INTEGER STATUS(MPI_STATUS_SIZE), IERROR
    LOGICAL FLAG
```

```
void MPI::Status::Set_cancelled(bool flag)
```

If **flag** is set to **true** then a subsequent call to **MPI_TEST_CANCELLED(status, flag)** will also return **flag = true**, otherwise it will return **false**.

Advice to users. Users are advised not to reuse the status fields for values other than those for which they were intended. Doing so may lead to unexpected results when using the status object. For example, calling **MPI_GET_ELEMENTS** may cause an error if the value is out of range or it may be impossible to detect such an error. The **extra_state** argument provided with a generalized request can be used to return information that does not logically belong in status. Furthermore, modifying the values in a status set internally by **MPI**, e.g., **MPI_RECV**, may lead to unpredictable results and is strongly discouraged. (*End of advice to users.*)

8.4 Naming Objects

There are many occasions on which it would be useful to allow a user to associate a printable identifier with an **MPI** communicator, window, or datatype, for instance error reporting, debugging, and profiling. The names attached to opaque objects do not propagate when the object is duplicated or copied by **MPI** routines. For communicators this can be achieved using the following two functions.

```
MPI_COMM_SET_NAME(comm, comm_name)
```

```
INOUT  comm            communicator whose identifier is to be set (handle)
IN      comm_name       the character string which is remembered as the name
                        (string)
```

```
int MPI_Comm_set_name(MPI_Comm comm, char *comm_name)
```

```
MPI_COMM_SET_NAME(COMM, COMM_NAME, IERROR)
    INTEGER COMM, IERROR
    CHARACTER*(*) COMM_NAME
```

```
void MPI::Comm::Set_name(const char* comm_name)
```

MPI_COMM_SET_NAME allows a user to associate a name string with a communicator. The character string which is passed to **MPI_COMM_SET_NAME** will be saved inside the

MPI library (so it can be freed by the caller immediately after the call, or allocated on the stack). Leading spaces in **name** are significant but trailing ones are not.

MPI_COMM_SET_NAME is a local (non-collective) operation, which only affects the name of the communicator as seen in the process which made the **MPI_COMM_SET_NAME** call. There is no requirement that the same (or any) name be assigned to a communicator in every process where it exists.

Advice to users. Since **MPI_COMM_SET_NAME** is provided to help debug code, it is sensible to give the same name to a communicator in all of the processes where it exists, to avoid confusion. (*End of advice to users.*)

The length of the name which can be stored is limited to the value of **MPI_MAX_OBJECT_NAME** in Fortran and **MPI_MAX_OBJECT_NAME-1** in C and C++ to allow for the null terminator. Attempts to put names longer than this will result in truncation of the name. **MPI_MAX_OBJECT_NAME** must have a value of at least 64.

Advice to users. Under circumstances of store exhaustion an attempt to put a name of any length could fail, therefore the value of **MPI_MAX_OBJECT_NAME** should be viewed only as a strict upper bound on the name length, not a guarantee that setting names of less than this length will always succeed. (*End of advice to users.*)

Advice to implementors. Implementations which pre-allocate a fixed size space for a name should use the length of that allocation as the value of **MPI_MAX_OBJECT_NAME**. Implementations which allocate space for the name from the heap should still define **MPI_MAX_OBJECT_NAME** to be a relatively small value, since the user has to allocate space for a string of up to this size when calling **MPI_COMM_GET_NAME**. (*End of advice to implementors.*)

MPI_COMM_GET_NAME (comm, comm_name, resultlen)

IN	comm	communicator whose name is to be returned (handle)
OUT	comm_name	the name previously stored on the communicator, or an empty string if no such name exists (string)
OUT	resultlen	length of returned name (integer)

```
int MPI_Comm_get_name(MPI_Comm comm, char *comm_name, int *resultlen)
```

```
MPI_COMM_GET_NAME(COMM, COMM_NAME, RESULTLEN, IERROR)
```

```
INTEGER COMM, RESULTLEN, IERROR
```

```
CHARACTER*(*) COMM_NAME
```

```
void MPI::Comm::Get_name(char* comm_name, int& resultlen) const
```

MPI_COMM_GET_NAME returns the last name which has previously been associated with the given communicator. The name may be set and got from any language. The same name will be returned independent of the language used. **name** should be allocated so that it can hold a resulting string of length **MPI_MAX_OBJECT_NAME** characters.

MPI_COMM_GET_NAME returns a copy of the set name in **name**.

If the user has not associated a name with a communicator, or an error occurs, `MPI_COMM_GET_NAME` will return an empty string (all spaces in Fortran, "" in C and C++). The three predefined communicators will have predefined names associated with them. Thus, the names of `MPI_COMM_WORLD`, `MPI_COMM_SELF`, and `MPI_COMM_PARENT` will have the default of `MPI_COMM_WORLD`, `MPI_COMM_SELF`, and `MPI_COMM_PARENT`. The fact that the system may have chosen to give a default name to a communicator does not prevent the user from setting a name on the same communicator; doing this removes the old name and assigns the new one.

Rationale. We provide separate functions for setting and getting the name of a communicator, rather than simply providing a predefined attribute key for the following reasons:

- It is not, in general, possible to store a string as an attribute from Fortran.
- It is not easy to set up the delete function for a string attribute unless it is known to have been allocated from the heap.
- To make the attribute key useful additional code to call `strdup` is necessary. If this is not standardized then users have to write it. This is extra unneeded work which we can easily eliminate.
- The Fortran binding is not trivial to write (it will depend on details of the Fortran compilation system), and will not be portable. Therefore it should be in the library rather than in user code.

(End of rationale.)

Advice to users. The above definition means that it is safe simply to print the string returned by `MPI_COMM_GET_NAME`, as it is always a valid string even if there was no name.

Note that associating a name with a communicator has no effect on the semantics of an MPI program, and will (necessarily) increase the store requirement of the program, since the names must be saved. Therefore there is no requirement that users use these functions to associate names with communicators. However debugging and profiling MPI applications may be made easier if names are associated with communicators, since the debugger or profiler should then be able to present information in a less cryptic manner. *(End of advice to users.)*

The following functions are used for setting and getting names of datatypes.

`MPI_TYPE_SET_NAME (type, type_name)`

INOUT	<code>type</code>	datatype whose identifier is to be set (handle)
IN	<code>type_name</code>	the character string which is remembered as the name (string)

`int MPI_Type_set_name(MPI_Datatype type, char *type_name)`

`MPI_TYPE_SET_NAME(TYPE, TYPE_NAME, IERROR)`

INTEGER TYPE, IERROR

```

1      CHARACTER*(*) TYPE_NAME
2
3  void MPI::Datatype::Set_name(const char* type_name)
4
5
6  MPI_TYPE_GET_NAME (type, type_name, resultlen)
7
8      IN          type          datatype whose name is to be returned (handle)
9
10     OUT         type_name      the name previously stored on the datatype, or a empty
11                                string if no such name exists (string)
12
13     OUT         resultlen      length of returned name (integer)
14
15 int MPI_Type_get_name(MPI_Datatype type, char *type_name, int *resultlen)
16
17 MPI_TYPE_GET_NAME(TYPE, TYPE_NAME, RESULTLEN, IERROR)
18     INTEGER TYPE, RESULTLEN, IERROR
19     CHARACTER*(*) TYPE_NAME
20
21 void MPI::Datatype::Get_name(char* type_name, int& resultlen) const
22
23     Named predefined datatypes have the default names of the datatype name. For exam-
24     ple, MPI_WCHAR has the default name of MPI_WCHAR.
25     The following functions are used for setting and getting names of windows.
26
27
28 MPI_WIN_SET_NAME (win, win_name)
29
30     INOUT       win            window whose identifier is to be set (handle)
31
32     IN          win_name       the character string which is remembered as the name
33                                (string)
34
35 int MPI_Win_set_name(MPI_Win win, char *win_name)
36
37 MPI_WIN_SET_NAME(WIN, WIN_NAME, IERROR)
38     INTEGER WIN, IERROR
39     CHARACTER*(*) WIN_NAME
40
41 void MPI::Win::Set_name(const char* win_name)
42
43
44 MPI_WIN_GET_NAME (win, win_name, resultlen)
45
46     IN          win            window whose name is to be returned (handle)
47
48     OUT         win_name       the name previously stored on the window, or a empty
49                                string if no such name exists (string)
50
51     OUT         resultlen      length of returned name (integer)
52
53 int MPI_Win_get_name(MPI_Win win, char *win_name, int *resultlen)
54
55 MPI_WIN_GET_NAME(WIN, WIN_NAME, RESULTLEN, IERROR)
56     INTEGER WIN, RESULTLEN, IERROR

```

```

    CHARACTER*(*) WIN_NAME
void MPI::Win::Get_name(char* win_name, int& resultlen) const

```

8.5 Error Classes, Error Codes, and Error Handlers

Users may want to write a layered library on top of an existing MPI implementation, and this library may have its own set of error codes and classes. An example of such a library is an I/O library based on the I/O chapter in MPI-2. For this purpose, functions are needed to:

1. add a new error class to the ones an MPI implementation already knows.
2. associate error codes with this error class, so that `MPI_ERROR_CLASS` works.
3. associate strings with these error codes, so that `MPI_ERROR_STRING` works.
4. invoke the error handler associated with a communicator, window, or object.

Several new functions are provided to do this. They are all local. No functions are provided to free error handlers or error classes: it is not expected that an application will generate them in significant numbers.

`MPI_ADD_ERROR_CLASS(errorclass)`

OUT **errorclass** value for the new error class (integer)

```
int MPI_Add_error_class(int *errorclass)
```

```
MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)
```

```
    INTEGER ERRORCLASS, IERROR
```

```
int MPI::Add_error_class()
```

Creates a new error class and returns the value for it.

Rationale. To avoid conflicts with existing error codes and classes, the value is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high quality implementation will return the value for a new **errorclass** in the same deterministic way on all processes. (*End of advice to implementors.*)

Advice to users. Since a call to `MPI_ADD_ERROR_CLASS` is local, the same **errorclass** may not be returned on all processes that make this call. Thus, it is not safe to assume that registering a new error on a set of processes at the same time will yield the same **errorclass** on all of the processes. However, if an implementation returns the new **errorclass** in a deterministic way, and they are always generated in the same order on the same set of processes (for example, all processes), then the value will be the same. However, even if a deterministic algorithm is used, the value can vary

across processes. This can happen, for example, if different but overlapping groups of processes make a series of calls. As a result of these issues, getting the “same” error on multiple processes may not cause the same value of error code to be generated. (*End of advice to users.*)

The value of `MPI_ERR_LASTCODE` is not affected by new user-defined error codes and classes. As in `MPI-1`, it is a constant value. Instead, a predefined attribute key `MPI_LASTUSEDCLASS` is associated with `MPI_COMM_WORLD`. The attribute value corresponding to this key is the current maximum error class including the user-defined ones. This is a local value and may be different on different processes. The value returned by this key is always greater than or equal to `MPI_ERR_LASTCODE`.

Advice to users. The value returned by the key `MPI_LASTUSEDCLASS` will not change unless the user calls a function to explicitly add an error class/code. In a multi-threaded environment, the user must take extra care in assuming this value has not changed. Note that error codes and error classes are not necessarily dense. A user may not assume that each error class below `MPI_LASTUSEDCLASS` is valid. (*End of advice to users.*)

`MPI_ADD_ERROR_CODE(errorclass, errorcode)`

IN	errorclass	error class (integer)
OUT	errorcode	new error code to associated with errorclass (integer)

```
int MPI_Add_error_code(int errorclass, int *errorcode)
```

```
MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)
```

```
INTEGER ERRORCLASS, ERRORCODE, IERROR
```

```
int MPI::Add_error_code(int errorclass)
```

Creates new error code associated with **errorclass** and returns its value in **errorcode**.

Rationale. To avoid conflicts with existing error codes and classes, the value of the new error code is set by the implementation and not by the user. (*End of rationale.*)

Advice to implementors. A high quality implementation will return the value for a new **errorcode** in the same deterministic way on all processes. (*End of advice to implementors.*)

`MPI_ADD_ERROR_STRING(errorcode, string)`

IN	errorcode	error code or class (integer)
IN	string	text corresponding to errorcode (string)

```
int MPI_Add_error_string(int errorcode, char *string)
```

```
MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)
```

```
INTEGER ERRORCODE, IERROR
CHARACTER*(*) STRING
```

```
void MPI::Add_error_string(int errorcode, const char* string)
```

Associates an error string with an error code or class. The string must be no more than `MPI_MAX_ERROR_STRING` characters long. The length of the string is as defined in the calling language. The length of the string does not include the null terminator in C or C++. Trailing blanks will be stripped in Fortran. Calling `MPI_ADD_ERROR_STRING` for an `errorcode` that already has a string will replace the old string with the new string. It is erroneous to call `MPI_ADD_ERROR_STRING` for an error code or class with a value \leq `MPI_ERR_LASTCODE`.

If `MPI_ERROR_STRING` is called when no string has been set, it will return a empty string (all spaces in Fortran, "" in C and C++).

Section 4.13 on page 61 describes the methods for creating and associating error handlers with communicators, files, and windows.

```
MPI_COMM_CALL_ERRHANDLER (comm, errorcode)
```

```
IN      comm      communicator with error handler (handle)
```

```
IN      errorcode  error code (integer)
```

```
int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)
```

```
MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)
```

```
INTEGER COMM, ERRORCODE, IERROR
```

```
void MPI::Comm::Call_errhandler(int errorcode) const
```

This function invokes the error handler assigned to the communicator with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Users should note that the default error handler is `MPI_ERRORS_ARE_FATAL`. Thus, calling `MPI_COMM_CALL_ERRHANDLER` will abort the `comm` processes if the default error handler has not been changed for this communicator or on the parent before the communicator was created. (*End of advice to users.*)

```
MPI_WIN_CALL_ERRHANDLER (win, errorcode)
```

```
IN      win      window with error handler (handle)
```

```
IN      errorcode  error code (integer)
```

```
int MPI_Win_call_errhandler(MPI_Win win, int errorcode)
```

```
MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)
```

```

1      INTEGER WIN, ERRORCODE, IERROR
2
3      void MPI::Win::Call_errhandler(int errorcode) const

```

This function invokes the error handler assigned to the window with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. As with communicators, the default error handler for windows is `MPI_ERRORS_ARE_FATAL`. (*End of advice to users.*)

```

13     MPI_FILE_CALL_ERRHANDLER (fh, errorcode)
14
15     IN          fh                      file with error handler (handle)
16     IN          errorcode              error code (integer)

```

```

18     int MPI_File_call_errhandler(MPI_File fh, int errorcode)

```

```

19     MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)
20     INTEGER FH, ERRORCODE, IERROR

```

```

22     void MPI::File::Call_errhandler(int errorcode) const

```

This function invokes the error handler assigned to the file with the error code supplied. This function returns `MPI_SUCCESS` in C and C++ and the same value in `IERROR` if the error handler was successfully called (assuming the process is not aborted and the error handler returns).

Advice to users. Unlike errors on communicators and windows, the default behavior for files is to have `MPI_ERRORS_RETURN`. (*End of advice to users.*)

Advice to users. Users are warned that handlers should not be called recursively with `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER`. Doing this can create a situation where an infinite recursion is created. This can occur if `MPI_COMM_CALL_ERRHANDLER`, `MPI_FILE_CALL_ERRHANDLER`, or `MPI_WIN_CALL_ERRHANDLER` is called inside an error handler.

Error codes and classes are associated with a process. As a result, they may be used in any error handler. Error handlers should be prepared to deal with any error code it is given. Furthermore, it is good practice to only call an error handler with the appropriate error codes. For example, file errors would normally be sent to the file error handler. (*End of advice to users.*)

8.6 Decoding a Datatype

MPI-1 provides datatype objects, which allow users to specify an arbitrary layout of data in memory. The layout information, once put in a datatype, could not be decoded from

the datatype. There are several cases, however, where accessing the layout information in opaque datatype objects would be useful.

The two functions in this section are used together to decode datatypes to recreate the calling sequence used in their initial definition. These can be used to allow a user to determine the type map and type signature of a datatype.

MPI_TYPE_GET_ENVELOPE(datatype, num_integers, num_addresses, num_datatypes, combiner)

IN	datatype	datatype to access (handle)
OUT	num_integers	number of input integers used in the call constructing combiner (nonnegative integer)
OUT	num_addresses	number of input addresses used in the call constructing combiner (nonnegative integer)
OUT	num_datatypes	number of input datatypes used in the call constructing combiner (nonnegative integer)
OUT	combiner	combiner (state)

```
int MPI_Type_get_envelope(MPI_Datatype datatype, int *num_integers,
                          int *num_addresses, int *num_datatypes, int *combiner)
```

```
MPI_TYPE_GET_ENVELOPE(DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES,
                       COMBINER, IERROR)
    INTEGER DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES, COMBINER,
    IERROR
```

```
void MPI::Datatype::Get_envelope(int& num_integers, int& num_addresses,
                                int& num_datatypes, int& combiner) const
```

For the given **datatype**, **MPI_TYPE_GET_ENVELOPE** returns information on the number and type of input arguments used in the call that created the **datatype**. The number-of-arguments values returned can be used to provide sufficiently large arrays in the decoding routine **MPI_TYPE_GET_CONTENTS**. This call and the meaning of the returned values is described below. The **combiner** reflects the MPI datatype constructor call that was used in creating **datatype**.

Rationale. By requiring that the **combiner** reflect the constructor used in the creation of the **datatype**, the decoded information can be used to effectively recreate the calling sequence used in the original creation. One call is effectively the same as another when the information obtained from **MPI_TYPE_GET_CONTENTS** may be used with either to produce the same outcome. C calls **MPI_Type_hindexed** and **MPI_Type_create_hindexed** are always effectively the same while the Fortran call **MPI_TYPE_HINDEXED** will be different than either of these in some MPI implementations. This is the most useful information and was felt to be reasonable even though it constrains implementations to remember the original constructor sequence even if the internal representation is different.

The decoded information keeps track of datatype duplications. This is important as one needs to distinguish between a predefined datatype and a dup of a predefined

datatype. The former is a constant object that cannot be freed, while the latter is a derived datatype that can be freed. (*End of rationale.*)

The list below has the values that can be returned in **combiner** on the left and the call associated with them on the right.

<code>MPI_COMBINER_NAMED</code>	a named predefined datatype
<code>MPI_COMBINER_DUP</code>	<code>MPI_TYPE_DUP</code>
<code>MPI_COMBINER_CONTIGUOUS</code>	<code>MPI_TYPE_CONTIGUOUS</code>
<code>MPI_COMBINER_VECTOR</code>	<code>MPI_TYPE_VECTOR</code>
<code>MPI_COMBINER_HVECTOR_INTEGER</code>	<code>MPI_TYPE_HVECTOR</code> from Fortran
<code>MPI_COMBINER_HVECTOR</code>	<code>MPI_TYPE_HVECTOR</code> from C or C++ and in some case Fortran or <code>MPI_TYPE_CREATE_HVECTOR</code>
<code>MPI_COMBINER_INDEXED</code>	<code>MPI_TYPE_INDEXED</code>
<code>MPI_COMBINER_HINDEXED_INTEGER</code>	<code>MPI_TYPE_HINDEXED</code> from Fortran
<code>MPI_COMBINER_HINDEXED</code>	<code>MPI_TYPE_HINDEXED</code> from C or C++ and in some case Fortran or <code>MPI_TYPE_CREATE_HINDEXED</code>
<code>MPI_COMBINER_INDEXED_BLOCK</code>	<code>MPI_TYPE_CREATE_INDEXED_BLOCK</code>
<code>MPI_COMBINER_STRUCT_INTEGER</code>	<code>MPI_TYPE_STRUCT</code> from Fortran
<code>MPI_COMBINER_STRUCT</code>	<code>MPI_TYPE_STRUCT</code> from C or C++ and in some case Fortran or <code>MPI_TYPE_CREATE_STRUCT</code>
<code>MPI_COMBINER_SUBARRAY</code>	<code>MPI_TYPE_CREATE_SUBARRAY</code>
<code>MPI_COMBINER_DARRAY</code>	<code>MPI_TYPE_CREATE_DARRAY</code>
<code>MPI_COMBINER_F90_REAL</code>	<code>MPI_TYPE_CREATE_F90_REAL</code>
<code>MPI_COMBINER_F90_COMPLEX</code>	<code>MPI_TYPE_CREATE_F90_COMPLEX</code>
<code>MPI_COMBINER_F90_INTEGER</code>	<code>MPI_TYPE_CREATE_F90_INTEGER</code>
<code>MPI_COMBINER_RESIZED</code>	<code>MPI_TYPE_CREATE_RESIZED</code>

If **combiner** is `MPI_COMBINER_NAMED` then **datatype** is a named predefined datatype.

For calls with address arguments, we sometimes need to differentiate whether the call used an integer or an address size argument. For example, there are two combiners for `hvector`: `MPI_COMBINER_HVECTOR_INTEGER` and `MPI_COMBINER_HVECTOR`. The former is used if it was the MPI-1 call from Fortran, and the latter is used if it was the MPI-1 call from C or C++. However, on systems where `MPI_ADDRESS_KIND = MPI_INTEGER_KIND` (i.e., where integer arguments and address size arguments are the same), the combiner `MPI_COMBINER_HVECTOR` may be returned for a datatype constructed by a call to `MPI_TYPE_HVECTOR` from Fortran. Similarly, `MPI_COMBINER_HINDEXED` may be returned for a datatype constructed by a call to `MPI_TYPE_HINDEXED` from Fortran, and `MPI_COMBINER_STRUCT` may be returned for a datatype constructed by a call to `MPI_TYPE_STRUCT` from Fortran. On such systems, one need not differentiate constructors that take address size arguments from constructors that take integer arguments, since these are the same. The new MPI-2 calls all use address sized arguments.

Rationale. For recreating the original call, it is important to know if address information may have been truncated. The MPI-1 calls from Fortran for a few routines could be subject to truncation in the case where the default `INTEGER` size is smaller than the size of an address. (*End of rationale.*)

The actual arguments used in the creation call for a **datatype** can be obtained from the call:

MPI_TYPE_GET_CONTENTS(**datatype**, **max_integers**, **max_addresses**, **max_datatypes**, **array_of_integers**, **array_of_addresses**, **array_of_datatypes**)

IN	datatype	datatype to access (handle)
IN	max_integers	number of elements in array_of_integers (non-negative integer)
IN	max_addresses	number of elements in array_of_addresses (non-negative integer)
IN	max_datatypes	number of elements in array_of_datatypes (non-negative integer)
OUT	array_of_integers	contains integer arguments used in constructing datatype (array of integers)
OUT	array_of_addresses	contains address arguments used in constructing datatype (array of integers)
OUT	array_of_datatypes	contains datatype arguments used in constructing datatype (array of handles)

```
int MPI_Type_get_contents(MPI_Datatype datatype, int max_integers,
    int max_addresses, int max_datatypes, int array_of_integers[],
    MPI_Aint array_of_addresses[],
    MPI_Datatype array_of_datatypes[])
```

```
MPI_TYPE_GET_CONTENTS(DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
    ARRAY_OF_INTEGERS, ARRAY_OF_ADDRESSES, ARRAY_OF_DATATYPES,
    IERROR)
    INTEGER DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
    ARRAY_OF_INTEGERS(*), ARRAY_OF_DATATYPES(*), IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_ADDRESSES(*)
```

```
void MPI::Datatype::Get_contents(int max_integers, int max_addresses,
    int max_datatypes, int array_of_integers[],
    MPI::Aint array_of_addresses[],
    MPI::Datatype array_of_datatypes[]) const
```

datatype must be a predefined unnamed or a derived datatype; the call is erroneous if **datatype** is a predefined named datatype.

The values given for **max_integers**, **max_addresses**, and **max_datatypes** must be at least as large as the value returned in **num_integers**, **num_addresses**, and **num_datatypes**, respectively, in the call **MPI_TYPE_GET_ENVELOPE** for the same **datatype** argument.

Rationale. The arguments **max_integers**, **max_addresses**, and **max_datatypes** allow for error checking in the call. This is analogous to the topology calls in MPI-1. (*End of rationale.*)

The datatypes returned in `array_of_datatypes` are handles to datatype objects that are equivalent to the datatypes used in the original construction call. If these were derived datatypes, then the returned datatypes are new datatype objects, and the user is responsible for freeing these datatypes with `MPI_TYPE_FREE`. If these were predefined datatypes, then the returned datatype is equal to that (constant) predefined datatype and cannot be freed.

The committed state of returned derived datatypes is undefined, i.e., the datatypes may or may not be committed. Furthermore, the content of attributes of returned datatypes is undefined.

Note that `MPI_TYPE_GET_CONTENTS` can be invoked with a `datatype` argument that was constructed using `MPI_TYPE_CREATE_F90_REAL`, `MPI_TYPE_CREATE_F90_INTEGER`, or `MPI_TYPE_CREATE_F90_COMPLEX` (an unnamed predefined datatype). In such a case, an empty `array_of_datatypes` is returned.

Rationale. The definition of datatype equivalence implies that equivalent predefined datatypes are equal. By requiring the same handle for named predefined datatypes, it is possible to use the `==` or `.EQ.` comparison operator to determine the datatype involved. (*End of rationale.*)

Advice to implementors. The datatypes returned in `array_of_datatypes` must appear to the user as if each is an equivalent copy of the datatype used in the type constructor call. Whether this is done by creating a new datatype or via another mechanism such as a reference count mechanism is up to the implementation as long as the semantics are preserved. (*End of advice to implementors.*)

Rationale. The committed state and attributes of the returned datatype is deliberately left vague. The datatype used in the original construction may have been modified since its use in the constructor call. Attributes can be added, removed, or modified as well as having the datatype committed. The semantics given allow for a reference count implementation without having to track these changes. (*End of rationale.*)

In the MPI-1 datatype constructor calls, the address arguments in Fortran are of type `INTEGER`. In the new MPI-2 calls, the address arguments are of type `INTEGER(KIND=MPI_ADDRESS_KIND)`. The call `MPI_TYPE_GET_CONTENTS` returns all addresses in an argument of type `INTEGER(KIND=MPI_ADDRESS_KIND)`. This is true even if the old MPI-1 calls were used. Thus, the location of values returned can be thought of as being returned by the C bindings. It can also be determined by examining the new MPI-2 calls for datatype constructors for the deprecated MPI-1 calls that involve addresses.

Rationale. By having all address arguments returned in the `array_of_addresses` argument, the result from a C and Fortran decoding of a `datatype` gives the result in the same argument. It is assumed that an integer of type `INTEGER(KIND=MPI_ADDRESS_KIND)` will be at least as large as the `INTEGER` argument used in datatype construction with the old MPI-1 calls so no loss of information will occur. (*End of rationale.*)

The following defines what values are placed in each entry of the returned arrays depending on the datatype constructor used for `datatype`. It also specifies the size of the arrays needed which is the values returned by `MPI_TYPE_GET_ENVELOPE`. In Fortran, the following calls were made:

```

PARAMETER (LARGE = 1000)
INTEGER TYPE, NI, NA, ND, COMBINER, I(LARGE), D(LARGE), IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) A(LARGE)
! CONSTRUCT DATATYPE TYPE (NOT SHOWN)
CALL MPI_TYPE_GET_ENVELOPE(TYPE, NI, NA, ND, COMBINER, IERROR)
IF ((NI .GT. LARGE) .OR. (NA .GT. LARGE) .OR. (ND .GT. LARGE)) THEN
    WRITE (*, *) "NI, NA, OR ND = ", NI, NA, ND, &
    " RETURNED BY MPI_TYPE_GET_ENVELOPE IS LARGER THAN LARGE = ", LARGE
    CALL MPI_ABORT(MPI_COMM_WORLD, 99)
ENDIF
CALL MPI_TYPE_GET_CONTENTS(TYPE, NI, NA, ND, I, A, D, IERROR)

```

or in C the analogous calls of:

```

#define LARGE 1000
int ni, na, nd, combiner, i[LARGE];
MPI_Aint a[LARGE];
MPI_Datatype type, d[LARGE];
/* construct datatype type (not shown) */
MPI_Type_get_envelope(type, &ni, &na, &nd, &combiner);
if ((ni > LARGE) || (na > LARGE) || (nd > LARGE)) {
    fprintf(stderr, "ni, na, or nd = %d %d %d returned by ", ni, na, nd);
    fprintf(stderr, "MPI_Type_get_envelope is larger than LARGE = %d\n",
            LARGE);
    MPI_Abort(MPI_COMM_WORLD, 99);
};
MPI_Type_get_contents(type, ni, na, nd, i, a, d);

```

The C++ code is in analogy to the C code above with the same values returned.

In the descriptions that follow, the lower case name of arguments is used.

If combiner is `MPI_COMBINER_NAMED` then it is erroneous to call

`MPI_TYPE_GET_CONTENTS`.

If combiner is `MPI_COMBINER_DUP` then

Constructor argument	C & C++ location	Fortran location
oldtype	d[0]	D(1)

and `ni = 0, na = 0, nd = 1`.

If combiner is `MPI_COMBINER_CONTIGUOUS` then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
oldtype	d[0]	D(1)

and `ni = 1, na = 0, nd = 1`.

If combiner is `MPI_COMBINER_VECTOR` then

Constructor argument	C & C++ location	Fortran location
count	i[0]	I(1)
blocklength	i[1]	I(2)
stride	i[2]	I(3)
oldtype	d[0]	D(1)

and $ni = 3$, $na = 0$, $nd = 1$.

If combiner is `MPI_COMBINER_HVECTOR_INTEGER` or `MPI_COMBINER_HVECTOR` then

Constructor argument	C & C++ location	Fortran location
count	<code>i[0]</code>	<code>I(1)</code>
blocklength	<code>i[1]</code>	<code>I(2)</code>
stride	<code>a[0]</code>	<code>A(1)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = 2$, $na = 1$, $nd = 1$.

If combiner is `MPI_COMBINER_INDEXED` then

Constructor argument	C & C++ location	Fortran location
count	<code>i[0]</code>	<code>I(1)</code>
array_of_blocklengths	<code>i[1]</code> to <code>i[i[0]]</code>	<code>I(2)</code> to <code>I(I(1)+1)</code>
array_of_displacements	<code>i[i[0]+1]</code> to <code>i[2*i[0]]</code>	<code>I(I(1)+2)</code> to <code>I(2*I(1)+1)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = 2*count+1$, $na = 0$, $nd = 1$.

If combiner is `MPI_COMBINER_HINDEXED_INTEGER` or `MPI_COMBINER_HINDEXED` then

Constructor argument	C & C++ location	Fortran location
count	<code>i[0]</code>	<code>I(1)</code>
array_of_blocklengths	<code>i[1]</code> to <code>i[i[0]]</code>	<code>I(2)</code> to <code>I(I(1)+1)</code>
array_of_displacements	<code>a[0]</code> to <code>a[i[0]-1]</code>	<code>A(1)</code> to <code>A(I(1))</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = count+1$, $na = count$, $nd = 1$.

If combiner is `MPI_COMBINER_INDEXED_BLOCK` then

Constructor argument	C & C++ location	Fortran location
count	<code>i[0]</code>	<code>I(1)</code>
blocklength	<code>i[1]</code>	<code>I(2)</code>
array_of_displacements	<code>i[2]</code> to <code>i[i[0]+1]</code>	<code>I(3)</code> to <code>I(I(1)+2)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = count+2$, $na = 0$, $nd = 1$.

If combiner is `MPI_COMBINER_STRUCT_INTEGER` or `MPI_COMBINER_STRUCT` then

Constructor argument	C & C++ location	Fortran location
count	<code>i[0]</code>	<code>I(1)</code>
array_of_blocklengths	<code>i[1]</code> to <code>i[i[0]]</code>	<code>I(2)</code> to <code>I(I(1)+1)</code>
array_of_displacements	<code>a[0]</code> to <code>a[i[0]-1]</code>	<code>A(1)</code> to <code>A(I(1))</code>
array_of_types	<code>d[0]</code> to <code>d[i[0]-1]</code>	<code>D(1)</code> to <code>D(I(1))</code>

and $ni = count+1$, $na = count$, $nd = count$.

If combiner is `MPI_COMBINER_SUBARRAY` then

Constructor argument	C & C++ location	Fortran location
ndims	<code>i[0]</code>	<code>I(1)</code>
array_of_sizes	<code>i[1]</code> to <code>i[i[0]]</code>	<code>I(2)</code> to <code>I(I(1)+1)</code>
array_of_subsizes	<code>i[i[0]+1]</code> to <code>i[2*i[0]]</code>	<code>I(I(1)+2)</code> to <code>I(2*I(1)+1)</code>
array_of_starts	<code>i[2*i[0]+1]</code> to <code>i[3*i[0]]</code>	<code>I(2*I(1)+2)</code> to <code>I(3*I(1)+1)</code>
order	<code>i[3*i[0]+1]</code>	<code>I(3*I(1)+2)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = 3*ndims+2$, $na = 0$, $nd = 1$.

If combiner is `MPI_COMBINER_DARRAY` then

Constructor argument	C & C++ location	Fortran location
size	<code>i[0]</code>	<code>I(1)</code>
rank	<code>i[1]</code>	<code>I(2)</code>
ndims	<code>i[2]</code>	<code>I(3)</code>
array_of_gsizes	<code>i[3]</code> to <code>i[i[2]+2]</code>	<code>I(4)</code> to <code>I(I(3)+3)</code>
array_of_distribs	<code>i[i[2]+3]</code> to <code>i[2*i[2]+2]</code>	<code>I(I(3)+4)</code> to <code>I(2*I(3)+3)</code>
array_of_dargs	<code>i[2*i[2]+3]</code> to <code>i[3*i[2]+2]</code>	<code>I(2*I(3)+4)</code> to <code>I(3*I(3)+3)</code>
array_of_psize	<code>i[3*i[2]+3]</code> to <code>i[4*i[2]+2]</code>	<code>I(3*I(3)+4)</code> to <code>I(4*I(3)+3)</code>
order	<code>i[4*i[2]+3]</code>	<code>I(4*I(3)+4)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = 4*ndims+4$, $na = 0$, $nd = 1$.

If combiner is `MPI_COMBINER_F90_REAL` then

Constructor argument	C & C++ location	Fortran location
p	<code>i[0]</code>	<code>I(1)</code>
r	<code>i[1]</code>	<code>I(2)</code>

and $ni = 2$, $na = 0$, $nd = 0$.

If combiner is `MPI_COMBINER_F90_COMPLEX` then

Constructor argument	C & C++ location	Fortran location
p	<code>i[0]</code>	<code>I(1)</code>
r	<code>i[1]</code>	<code>I(2)</code>

and $ni = 2$, $na = 0$, $nd = 0$.

If combiner is `MPI_COMBINER_F90_INTEGER` then

Constructor argument	C & C++ location	Fortran location
r	<code>i[0]</code>	<code>I(1)</code>

and $ni = 1$, $na = 0$, $nd = 0$.

If combiner is `MPI_COMBINER_RESIZED` then

Constructor argument	C & C++ location	Fortran location
lb	<code>a[0]</code>	<code>A(1)</code>
extent	<code>a[1]</code>	<code>A(2)</code>
oldtype	<code>d[0]</code>	<code>D(1)</code>

and $ni = 0$, $na = 2$, $nd = 1$.

Example 8.2 This example shows how a datatype can be decoded. The routine `printdatatype` prints out the elements of the datatype. Note the use of `MPI_Type_free` for datatypes that are not predefined.

```

/*
  Example of decoding a datatype.

  Returns 0 if the datatype is predefined, 1 otherwise
*/
#include <stdio.h>
#include <stdlib.h>

```

```

1  #include "mpi.h"
2  int printdatatype( MPI_Datatype datatype )
3  {
4      int *array_of_ints;
5      MPI_Aint *array_of_adds;
6      MPI_Datatype *array_of_dtypes;
7      int num_ints, num_adds, num_dtypes, combiner;
8      int i;
9
10     MPI_Type_get_envelope( datatype,
11                           &num_ints, &num_adds, &num_dtypes, &combiner );
12     switch (combiner) {
13     case MPI_COMBINER_NAMED:
14         printf( "Datatype is named:" );
15         /* To print the specific type, we can match against the
16            predefined forms. We can NOT use a switch statement here
17            We could also use MPI_TYPE_GET_NAME if we preferred to use
18            names that the user may have changed.
19            */
20         if (datatype == MPI_INT) printf( "MPI_INT\n" );
21         else if (datatype == MPI_DOUBLE) printf( "MPI_DOUBLE\n" );
22         ... else test for other types ...
23         return 0;
24         break;
25     case MPI_COMBINER_STRUCT:
26     case MPI_COMBINER_STRUCT_INTEGER:
27         printf( "Datatype is struct containing" );
28         array_of_ints = (int *)malloc( num_ints * sizeof(int) );
29         array_of_adds =
30             (MPI_Aint *) malloc( num_adds * sizeof(MPI_Aint) );
31         array_of_dtypes = (MPI_Datatype *)
32             malloc( num_dtypes * sizeof(MPI_Datatype) );
33         MPI_Type_get_contents( datatype, num_ints, num_adds, num_dtypes,
34                               array_of_ints, array_of_adds, array_of_dtypes );
35         printf( " %d datatypes:\n", array_of_ints[0] );
36         for (i=0; i<array_of_ints[0]; i++) {
37             printf( "blocklength %d, displacement %ld, type:\n",
38                   array_of_ints[i+1], array_of_adds[i] );
39             if (printdatatype( array_of_dtypes[i] )) {
40                 /* Note that we free the type ONLY if it
41                    is not predefined */
42                 MPI_Type_free( &array_of_dtypes[i] );
43             }
44         }
45         free( array_of_ints );
46         free( array_of_adds );
47         free( array_of_dtypes );
48         break;

```

```

        ... other combiner values ...
default:
    printf( "Unrecognized combiner type\n" );
}
return 1;
}

```

8.7 MPI and Threads

This section specifies the interaction between MPI calls and threads. The section lists minimal requirements for **thread compliant** MPI implementations and defines functions that can be used for initializing the thread environment. MPI may be implemented in environments where threads are not supported or perform poorly. Therefore, it is not required that all MPI implementations fulfill all the requirements specified in this section.

This section generally assumes a thread package similar to POSIX threads [11], but the syntax and semantics of thread calls are not specified here — these are beyond the scope of this document.

8.7.1 General

In a thread-compliant implementation, an MPI process is a process that may be multi-threaded. Each thread can issue MPI calls; however, threads are not separately addressable: a rank in a send or receive call identifies a process, not a thread. A message sent to a process can be received by any thread in this process.

Rationale. This model corresponds to the POSIX model of interprocess communication: the fact that a process is multi-threaded, rather than single-threaded, does not affect the external interface of this process. MPI implementations where MPI ‘processes’ are POSIX threads inside a single POSIX process are not thread-compliant by this definition (indeed, their “processes” are single-threaded). (*End of rationale.*)

Advice to users. It is the user’s responsibility to prevent races when threads within the same application post conflicting communication calls. The user can make sure that two threads in the same process will not issue conflicting communication calls by using distinct communicators at each thread. (*End of advice to users.*)

The two main requirements for a thread-compliant implementation are listed below.

1. All MPI calls are *thread-safe*. I.e., two concurrently running threads may make MPI calls and the outcome will be as if the calls executed in some order, even if their execution is interleaved.
2. Blocking MPI calls will block the calling thread only, allowing another thread to execute, if available. The calling thread will be blocked until the event on which it is waiting occurs. Once the blocked communication is enabled and can proceed, then the call will complete and the thread will be marked runnable, within a finite time. A blocked thread will not prevent progress of other runnable threads on the same process, and will not prevent them from executing MPI calls.

Example 8.3 Process 0 consists of two threads. The first thread executes a blocking send call `MPI_Send(buff1, count, type, 0, 0, comm)`, whereas the second thread executes a blocking receive call `MPI_Recv(buff2, count, type, 0, 0, comm, &status)`. I.e., the first thread sends a message that is received by the second thread. This communication should always succeed. According to the first requirement, the execution will correspond to some interleaving of the two calls. According to the second requirement, a call can only block the calling thread and cannot prevent progress of the other thread. If the send call went ahead of the receive call, then the sending thread may block, but this will not prevent the receiving thread from executing. Thus, the receive call will occur. Once both calls occur, the communication is enabled and both calls will complete. On the other hand, a single-threaded process that posts a send, followed by a matching receive, may deadlock. The progress requirement for multithreaded implementations is stronger, as a blocked call cannot prevent progress in other threads.

Advice to implementors. MPI calls can be made thread-safe by executing only one at a time, e.g., by protecting MPI code with one process-global lock. However, blocked operations cannot hold the lock, as this would prevent progress of other threads in the process. The lock is held only for the duration of an atomic, locally-completing suboperation such as posting a send or completing a send, and is released in between. Finer locks can provide more concurrency, at the expense of higher locking overheads. Concurrency can also be achieved by having some of the MPI protocol executed by separate server threads. (*End of advice to implementors.*)

8.7.2 Clarifications

Initialization and Completion The call to `MPI_FINALIZE` should occur on the same thread that initialized MPI. We call this thread the **main thread**. The call should occur only after all the process threads have completed their MPI calls, and have no pending communications or I/O operations.

Rationale. This constraint simplifies implementation. (*End of rationale.*)

Multiple threads completing the same request. A program where two threads block, waiting on the same request, is erroneous. Similarly, the same request cannot appear in the array of requests of two concurrent `MPI_WAIT{ANY|SOME|ALL}` calls. In MPI, a request can only be completed once. Any combination of wait or test which violates this rule is erroneous.

Rationale. This is consistent with the view that a multithreaded execution corresponds to an interleaving of the MPI calls. In a single threaded implementation, once a wait is posted on a request the request handle will be nullified before it is possible to post a second wait on the same handle. With threads, an `MPI_WAIT{ANY|SOME|ALL}` may be blocked without having nullified its request(s) so it becomes the user's responsibility to avoid using the same request in an `MPI_WAIT` on another thread. This constraint also simplifies implementation, as only one thread will be blocked on any communication or I/O event. (*End of rationale.*)

Probe A receive call that uses source and tag values returned by a preceding call to `MPI_PROBE` or `MPI_IREQ_PROBE` will receive the message matched by the probe call only if there

was no other matching receive after the probe and before that receive. In a multithreaded environment, it is up to the user to enforce this condition using suitable mutual exclusion logic. This can be enforced by making sure that each communicator is used by only one thread on each process.

Collective calls Matching of collective calls on a communicator, window, or file handle is done according to the order in which the calls are issued at each process. If concurrent threads issue such calls on the same communicator, window or file handle, it is up to the user to make sure the calls are correctly ordered, using interthread synchronization.

Exception handlers An exception handler does not necessarily execute in the context of the thread that made the exception-raising MPI call; the exception handler may be executed by a thread that is distinct from the thread that will return the error code.

Rationale. The MPI implementation may be multithreaded, so that part of the communication protocol may execute on a thread that is distinct from the thread that made the MPI call. The design allows the exception handler to be executed on the thread where the exception occurred. (*End of rationale.*)

Interaction with signals and cancellations The outcome is undefined if a thread that executes an MPI call is cancelled (by another thread), or if a thread catches a signal while executing an MPI call. However, a thread of an MPI process may terminate, and may catch signals or be cancelled by another thread when not executing MPI calls.

Rationale. Few C library functions are signal safe, and many have cancellation points — points where the thread executing them may be cancelled. The above restriction simplifies implementation (no need for the MPI library to be “async-cancel-safe” or “async-signal-safe.”) (*End of rationale.*)

Advice to users. Users can catch signals in separate, non-MPI threads (e.g., by masking signals on MPI calling threads, and unmasking them in one or more non-MPI threads). A good programming practice is to have a distinct thread blocked in a call to `sigwait` for each user expected signal that may occur. Users must not catch signals used by the MPI implementation; as each MPI implementation is required to document the signals used internally, users can avoid these signals. (*End of advice to users.*)

Advice to implementors. The MPI library should not invoke library calls that are not thread safe, if multiple threads execute. (*End of advice to implementors.*)

8.7.3 Initialization

The following function may be used to initialize MPI, and initialize the MPI thread environment, instead of `MPI_INIT`.

`MPI_INIT_THREAD(required, provided)`

IN	required	desired level of thread support (integer)
OUT	provided	provided level of thread support (integer)

```
int MPI_Init_thread(int *argc, char *((*argv)[]), int required,
                    int *provided)
```

```
MPI_INIT_THREAD(REQUIRED, PROVIDED, IERROR)
    INTEGER REQUIRED, PROVIDED, IERROR
```

```
int MPI::Init_thread(int& argc, char**& argv, int required)
```

```
int MPI::Init_thread(int required)
```

Advice to users. In C and C++, the passing of **argc** and **argv** is optional. In C, this is accomplished by passing the appropriate null pointer. In C++, this is accomplished with two separate bindings to cover these two cases. This is as with `MPI_INIT` as discussed in Section 4.2. (*End of advice to users.*)

This call initializes `MPI` in the same way that a call to `MPI_INIT` would. In addition, it initializes the thread environment. The argument **required** is used to specify the desired level of thread support. The possible values are listed in increasing order of thread support.

`MPI_THREAD_SINGLE` Only one thread will execute.

`MPI_THREAD_FUNNELED` The process may be multi-threaded, but only the main thread will make `MPI` calls (all `MPI` calls are “funneled” to the main thread).

`MPI_THREAD_SERIALIZED` The process may be multi-threaded, and multiple threads may make `MPI` calls, but only one at a time: `MPI` calls are not made concurrently from two distinct threads (all `MPI` calls are “serialized”).

`MPI_THREAD_MULTIPLE` Multiple threads may call `MPI`, with no restrictions.

These values are monotonic; i.e., `MPI_THREAD_SINGLE` < `MPI_THREAD_FUNNELED` < `MPI_THREAD_SERIALIZED` < `MPI_THREAD_MULTIPLE`.

Different processes in `MPI_COMM_WORLD` may require different levels of thread support.

The call returns in **provided** information about the actual level of thread support that will be provided by `MPI`. It can be one of the four values listed above.

The level(s) of thread support that can be provided by `MPI_INIT_THREAD` will depend on the implementation, and may depend on information provided by the user before the program started to execute (e.g., with arguments to `mpiexec`). If possible, the call will return **provided** = **required**. Failing this, the call will return the least supported level such that **provided** > **required** (thus providing a stronger level of support than required by the user). Finally, if the user requirement cannot be satisfied, then the call will return in **provided** the highest supported level.

A **thread compliant** `MPI` implementation will be able to return **provided** = `MPI_THREAD_MULTIPLE`. Such an implementation may always return **provided** = `MPI_THREAD_MULTIPLE`, irrespective of the value of **required**. At the other extreme,

an MPI library that is not thread compliant may always return
provided = **MPI_THREAD_SINGLE**, irrespective of the value of **required**.

A call to **MPI_INIT** has the same effect as a call to **MPI_INIT_THREAD** with a **required** = **MPI_THREAD_SINGLE**.

Vendors may provide (implementation dependent) means to specify the level(s) of thread support available when the MPI program is started, e.g., with arguments to **mpiexec**. This will affect the outcome of calls to **MPI_INIT** and **MPI_INIT_THREAD**. Suppose, for example, that an MPI program has been started so that only **MPI_THREAD_MULTIPLE** is available. Then **MPI_INIT_THREAD** will return **provided** = **MPI_THREAD_MULTIPLE**, irrespective of the value of **required**; a call to **MPI_INIT** will also initialize the MPI thread support level to **MPI_THREAD_MULTIPLE**. Suppose, on the other hand, that an MPI program has been started so that all four levels of thread support are available. Then, a call to **MPI_INIT_THREAD** will return **provided** = **required**; on the other hand, a call to **MPI_INIT** will initialize the MPI thread support level to **MPI_THREAD_SINGLE**.

Rationale. Various optimizations are possible when MPI code is executed single-threaded, or is executed on multiple threads, but not concurrently: mutual exclusion code may be omitted. Furthermore, if only one thread executes, then the MPI library can use library functions that are not thread safe, without risking conflicts with user threads. Also, the model of one communication thread, multiple computation threads fits well many applications. E.g., if the process code is a sequential Fortran/C/C++ program with MPI calls that has been parallelized by a compiler for execution on an SMP node, in a cluster of SMPs, then the process computation is multi-threaded, but MPI calls will likely execute on a single thread.

The design accommodates a static specification of the thread support level, for environments that require static binding of libraries, and for compatibility for current multi-threaded MPI codes. (*End of rationale.*)

Advice to implementors. If **provided** is not **MPI_THREAD_SINGLE** then the MPI library should not invoke C/ C++/Fortran library calls that are not thread safe, e.g., in an environment where **malloc** is not thread safe, then **malloc** should not be used by the MPI library.

Some implementors may want to use different MPI libraries for different levels of thread support. They can do so using dynamic linking and selecting which library will be linked when **MPI_INIT_THREAD** is invoked. If this is not possible, then optimizations for lower levels of thread support will occur only when the level of thread support required is specified at link time. (*End of advice to implementors.*)

The following function can be used to query the current level of thread support.

MPI_QUERY_THREAD(provided)

OUT **provided** provided level of thread support (integer)

int MPI_Query_thread(int *provided)

MPI_QUERY_THREAD(PROVIDED, IERROR)

INTEGER PROVIDED, IERROR

```
1 int MPI::Query_thread()
```

```
2     The call returns in provided the current level of thread support. This will be the value
3     returned in provided by MPI_INIT_THREAD, if MPI was initialized by a call to
4     MPI_INIT_THREAD().
5
```

```
6
7 MPI_IS_THREAD_MAIN(flag)
```

```
8     OUT      flag                true if calling thread is main thread, false otherwise
9                                     (logical)
10
```

```
11 int MPI_Is_thread_main(int *flag)
```

```
12 MPI_IS_THREAD_MAIN(FLAG, IERROR)
```

```
13     LOGICAL FLAG
```

```
14     INTEGER IERROR
```

```
15
16 bool MPI::Is_thread_main()
```

```
17
18     This function can be called by a thread to find out whether it is the main thread (the
19     thread that called MPI_INIT or MPI_INIT_THREAD).
```

```
20     All routines listed in this section must be supported by all MPI implementations.
21
```

```
22     Rationale. MPI libraries are required to provide these calls even if they do not support
23     threads, so that portable code that contains invocations to these functions be able to
24     link correctly. MPI_INIT continues to be supported so as to provide compatibility with
25     current MPI codes. (End of rationale.)
26
```

```
27     Advice to users. It is possible to spawn threads before MPI is initialized, but
28     no MPI call other than MPI_INITIALIZED should be executed by these threads, un-
29     til MPI_INIT_THREAD is invoked by one thread (which, thereby, becomes the main
30     thread). In particular, it is possible to enter the MPI execution with a multi-threaded
31     process.
```

```
32     The level of thread support provided is a global property of the MPI process that can
33     be specified only once, when MPI is initialized on that process (or before). Portable
34     third party libraries have to be written so as to accommodate any provided level of
35     thread support. Otherwise, their usage will be restricted to specific level(s) of thread
36     support. If such a library can run only with specific level(s) of thread support, e.g.,
37     only with MPI_THREAD_MULTIPLE, then MPI_QUERY_THREAD can be used to check
38     whether the user initialized MPI to the correct level of thread support and, if not,
39     raise an exception. (End of advice to users.)
40
```

41 8.8 New Attribute Caching Functions

```
42
43     Caching on communicators has been a very useful feature. In MPI-2 it is expanded to
44     include caching on windows and datatypes.
45
```

```
46     Rationale. In one extreme you can allow caching on all opaque handles. The other
47     extreme is to only allow it on communicators. Caching has a cost associated with it
48
```

and should only be allowed when it is clearly needed and the increased cost is modest. This is the reason that windows and datatypes were added but not other handles. (*End of rationale.*)

One difficulty in MPI-1 is the potential for size differences between Fortran integers and C pointers. To overcome this problem with attribute caching on communicators, new functions are also given for this case. The new functions to cache on datatypes and windows also address this issue. For a general discussion of the address size problem, see Section 4.12.6.

The MPI-1.2 clarification, described in Section 3.2.8 on page 26, about the effect of returning other than MPI_SUCCESS from attribute callbacks applies to these new versions as well.

8.8.1 Communicators

The new functions that are replacements for the MPI-1 functions for caching on communicators are:

MPI_COMM_CREATE_KEYVAL(comm_copy_attr_fn, comm_delete_attr_fn, comm_keyval, extra_state)

IN	comm_copy_attr_fn	copy callback function for comm_keyval (function)
IN	comm_delete_attr_fn	delete callback function for comm_keyval (function)
OUT	comm_keyval	key value for future access (integer)
IN	extra_state	extra state for callback functions

```
int MPI_Comm_create_keyval(MPI_Comm_copy_attr_function *comm_copy_attr_fn,
                           MPI_Comm_delete_attr_function *comm_delete_attr_fn,
                           int *comm_keyval, void *extra_state)
```

```
MPI_COMM_CREATE_KEYVAL(COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN, COMM_KEYVAL,
                        EXTRA_STATE, IERROR)
EXTERNAL COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN
INTEGER COMM_KEYVAL, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
static int MPI::Comm::Create_keyval(MPI::Comm::Copy_attr_function*
                                     comm_copy_attr_fn,
                                     MPI::Comm::Delete_attr_function* comm_delete_attr_fn,
                                     void* extra_state)
```

This function replaces MPI_KEYVAL_CREATE, whose use is deprecated. The C binding is identical. The Fortran binding differs in that **extra_state** is an address-sized integer. Also, the copy and delete callback functions have Fortran bindings that are consistent with address-sized attributes.

The argument **comm_copy_attr_fn** may be specified as **MPI_COMM_NULL_COPY_FN** or **MPI_COMM_DUP_FN** from either C, C++, or Fortran. **MPI_COMM_NULL_COPY_FN** is a function that does nothing other than returning **flag = 0** and **MPI_SUCCESS**. **MPI_COMM_DUP_FN** is a simple-minded copy function that sets **flag = 1**, returns the value

of `attribute_val_in` in `attribute_val_out`, and returns `MPI_SUCCESS`. These replace the MPI-1 predefined callbacks `MPI_NULL_COPY_FN` and `MPI_DUP_FN`, whose use is deprecated.

The argument `comm_delete_attr_fn` may be specified as `MPI_COMM_NULL_DELETE_FN` from either C, C++, or Fortran. `MPI_COMM_NULL_DELETE_FN` is a function that does nothing, other than returning `MPI_SUCCESS`. `MPI_COMM_NULL_DELETE_FN` replaces `MPI_NULL_DELETE_FN`, whose use is deprecated.

The C callback functions are:

```
typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm, int comm_keyval,
                                       void *extra_state, void *attribute_val_in,
                                       void *attribute_val_out, int *flag);
```

and

```
typedef int MPI_Comm_delete_attr_function(MPI_Comm comm, int comm_keyval,
                                         void *attribute_val, void *extra_state);
```

which are the same as the MPI-1.1 calls but with a new name. The old names are deprecated.

The Fortran callback functions are:

```
SUBROUTINE COMM_COPY_ATTR_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
                             ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDCOMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
```

and

```
SUBROUTINE COMM_DELETE_ATTR_FN(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
                               IERROR)
    INTEGER COMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

The C++ callbacks are:

```
typedef int MPI::Comm::Copy_attr_function(const MPI::Comm& oldcomm,
                                           int comm_keyval, void* extra_state, void* attribute_val_in,
                                           void* attribute_val_out, bool& flag);
```

and

```
typedef int MPI::Comm::Delete_attr_function(MPI::Comm& comm,
                                             int comm_keyval, void* attribute_val, void* extra_state);
```

```
MPI_COMM_FREE_KEYVAL(comm_keyval)
```

```
    INOUT    comm_keyval          key value (integer)
```

```
int MPI_Comm_free_keyval(int *comm_keyval)
```

```
MPI_COMM_FREE_KEYVAL(COMM_KEYVAL, IERROR)
```

```
    INTEGER COMM_KEYVAL, IERROR
```

```
static void MPI::Comm::Free_keyval(int& comm_keyval)
```

This call is identical to the MPI-1 call `MPI_KEYVAL_FREE` but is needed to match the new communicator-specific creation function. The use of `MPI_KEYVAL_FREE` is deprecated.

`MPI_COMM_SET_ATTR(comm, comm_keyval, attribute_val)`

INOUT	<code>comm</code>	communicator from which attribute will be attached (handle)
IN	<code>comm_keyval</code>	key value (integer)
IN	<code>attribute_val</code>	attribute value

`int MPI_Comm_set_attr(MPI_Comm comm, int comm_keyval, void *attribute_val)`

`MPI_COMM_SET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, IERROR)`

`INTEGER COMM, COMM_KEYVAL, IERROR`

`INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL`

`void MPI::Comm::Set_attr(int comm_keyval, const void* attribute_val) const`

This function replaces `MPI_ATTR_PUT`, whose use is deprecated. The C binding is identical. The Fortran binding differs in that `attribute_val` is an address-sized integer.

`MPI_COMM_GET_ATTR(comm, comm_keyval, attribute_val, flag)`

IN	<code>comm</code>	communicator to which the attribute is attached (handle)
IN	<code>comm_keyval</code>	key value (integer)
OUT	<code>attribute_val</code>	attribute value, unless <code>flag = false</code>
OUT	<code>flag</code>	false if no attribute is associated with the key (logical)

`int MPI_Comm_get_attr(MPI_Comm comm, int comm_keyval, void *attribute_val, int *flag)`

`MPI_COMM_GET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)`

`INTEGER COMM, COMM_KEYVAL, IERROR`

`INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL`

`LOGICAL FLAG`

`bool MPI::Comm::Get_attr(int comm_keyval, void* attribute_val) const`

This function replaces `MPI_ATTR_GET`, whose use is deprecated. The C binding is identical. The Fortran binding differs in that `attribute_val` is an address-sized integer.

```
1 MPI_COMM_DELETE_ATTR(comm, comm_keyval)
```

```
2     INOUT    comm                communicator from which the attribute is deleted (han-
3                                     dle)
```

```
4     IN       comm_keyval        key value (integer)
```

```
6
7 int MPI_Comm_delete_attr(MPI_Comm comm, int comm_keyval)
```

```
8 MPI_COMM_DELETE_ATTR(COMM, COMM_KEYVAL, IERROR)
```

```
9     INTEGER COMM, COMM_KEYVAL, IERROR
```

```
10
11 void MPI::Comm::Delete_attr(int comm_keyval)
```

```
12     This function is the same as MPI_ATTR_DELETE but is needed to match the new
13     communicator specific functions. The use of MPI_ATTR_DELETE is deprecated.
```

15 8.8.2 Windows

```
16
17 The new functions for caching on windows are:
```

```
18
19 MPI_WIN_CREATE_KEYVAL(win_copy_attr_fn, win_delete_attr_fn, win_keyval, extra_state)
```

```
20
21     IN       win_copy_attr_fn    copy callback function for win_keyval (function)
```

```
22     IN       win_delete_attr_fn  delete callback function for win_keyval (function)
```

```
23     OUT      win_keyval         key value for future access (integer)
```

```
24     IN       extra_state        extra state for callback functions
```

```
25
26
27 int MPI_Win_create_keyval(MPI_Win_copy_attr_function *win_copy_attr_fn,
28                           MPI_Win_delete_attr_function *win_delete_attr_fn,
29                           int *win_keyval, void *extra_state)
```

```
30
31 MPI_WIN_CREATE_KEYVAL(WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN, WIN_KEYVAL,
32                       EXTRA_STATE, IERROR)
```

```
33     EXTERNAL WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN
```

```
34     INTEGER WIN_KEYVAL, IERROR
```

```
35     INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
36
37 static int MPI::Win::Create_keyval(MPI::Win::Copy_attr_function*
38                                   win_copy_attr_fn,
39                                   MPI::Win::Delete_attr_function* win_delete_attr_fn,
40                                   void* extra_state)
```

```
41     The argument win_copy_attr_fn may be specified as MPI_WIN_NULL_COPY_FN or
42     MPI_WIN_DUP_FN from either C, C++, or Fortran. MPI_WIN_NULL_COPY_FN is a function
43     that does nothing other than returning flag = 0 and MPI_SUCCESS. MPI_WIN_DUP_FN is
44     a simple-minded copy function that sets flag = 1, returns the value of attribute_val_in in
45     attribute_val_out, and returns MPI_SUCCESS.
```

```
46     The argument win_delete_attr_fn may be specified as MPI_WIN_NULL_DELETE_FN from
47     either C, C++, or Fortran. MPI_WIN_NULL_DELETE_FN is a function that does nothing,
48     other than returning MPI_SUCCESS.
```


The C callback functions are:

```
typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
                                       void *extra_state, void *attribute_val_in,
                                       void *attribute_val_out, int *flag);
```

and

```
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
                                         void *attribute_val, void *extra_state);
```

The Fortran callback functions are:

```
SUBROUTINE WIN_COPY_ATTR_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
                           ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
  INTEGER OLDWIN, WIN_KEYVAL, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
    ATTRIBUTE_VAL_OUT
  LOGICAL FLAG
```

and

```
SUBROUTINE WIN_DELETE_ATTR_FN(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
                              IERROR)
  INTEGER WIN, WIN_KEYVAL, IERROR
  INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

The C++ callbacks are:

```
typedef int MPI::Win::Copy_attr_function(const MPI::Win& oldwin,
                                         int win_keyval, void* extra_state, void* attribute_val_in,
                                         void* attribute_val_out, bool& flag);
```

and

```
typedef int MPI::Win::Delete_attr_function(MPI::Win& win, int win_keyval,
                                           void* attribute_val, void* extra_state);
```

MPI_WIN_FREE_KEYVAL(win_keyval)

INOUT win_keyval key value (integer)

int MPI_Win_free_keyval(int *win_keyval)

MPI_WIN_FREE_KEYVAL(WIN_KEYVAL, IERROR)

INTEGER WIN_KEYVAL, IERROR

static void MPI::Win::Free_keyval(int& win_keyval)

```

1  MPI_WIN_SET_ATTR(win, win_keyval, attribute_val)
2      INOUT    win                window to which attribute will be attached (handle)
3      IN       win_keyval         key value (integer)
4      IN       attribute_val      attribute value
5
6
7  int MPI_Win_set_attr(MPI_Win win, int win_keyval, void *attribute_val)
8
9  MPI_WIN_SET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, IERROR)
10     INTEGER WIN, WIN_KEYVAL, IERROR
11     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
12
13  void MPI::Win::Set_attr(int win_keyval, const void* attribute_val)
14
15
16  MPI_WIN_GET_ATTR(win, win_keyval, attribute_val, flag)
17      IN       win                window to which the attribute is attached (handle)
18      IN       win_keyval         key value (integer)
19      OUT      attribute_val      attribute value, unless flag = false
20      OUT      flag               false if no attribute is associated with the key (logical)
21
22
23  int MPI_Win_get_attr(MPI_Win win, int win_keyval, void *attribute_val,
24                      int *flag)
25
26  MPI_WIN_GET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
27     INTEGER WIN, WIN_KEYVAL, IERROR
28     INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
29     LOGICAL FLAG
30
31  bool MPI::Win::Get_attr(const MPI::Win& win, int win_keyval,
32                        void* attribute_val) const
33
34
35  MPI_WIN_DELETE_ATTR(win, win_keyval)
36      INOUT    win                window from which the attribute is deleted (handle)
37      IN       win_keyval         key value (integer)
38
39
40  int MPI_Win_delete_attr(MPI_Win win, int win_keyval)
41
42  MPI_WIN_DELETE_ATTR(WIN, WIN_KEYVAL, IERROR)
43     INTEGER WIN, WIN_KEYVAL, IERROR
44
45  void MPI::Win::Delete_attr(int win_keyval)
46
47
48

```

8.8.3 Datatypes

The new functions for caching on datatypes are:

`MPI_TYPE_CREATE_KEYVAL(type_copy_attr_fn, type_delete_attr_fn, type_keyval, extra_state)`

IN	<code>type_copy_attr_fn</code>	copy callback function for <code>type_keyval</code> (function)
IN	<code>type_delete_attr_fn</code>	delete callback function for <code>type_keyval</code> (function)
OUT	<code>type_keyval</code>	key value for future access (integer)
IN	<code>extra_state</code>	extra state for callback functions

```
int MPI_Type_create_keyval(MPI_Type_copy_attr_function *type_copy_attr_fn,
                          MPI_Type_delete_attr_function *type_delete_attr_fn,
                          int *type_keyval, void *extra_state)
```

```
MPI_TYPE_CREATE_KEYVAL(TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN, TYPE_KEYVAL,
                      EXTRA_STATE, IERROR)
EXTERNAL TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN
INTEGER TYPE_KEYVAL, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
```

```
static int MPI::Datatype::Create_keyval(MPI::Datatype::Copy_attr_function*
                                         type_copy_attr_fn, MPI::Datatype::Delete_attr_function*
                                         type_delete_attr_fn, void* extra_state)
```

The argument `type_copy_attr_fn` may be specified as `MPI_TYPE_NULL_COPY_FN` or `MPI_TYPE_DUP_FN` from either C, C++, or Fortran. `MPI_TYPE_NULL_COPY_FN` is a function that does nothing other than returning `flag = 0` and `MPI_SUCCESS`. `MPI_TYPE_DUP_FN` is a simple-minded copy function that sets `flag = 1`, returns the value of `attribute_val_in` in `attribute_val_out`, and returns `MPI_SUCCESS`.

The argument `type_delete_attr_fn` may be specified as `MPI_TYPE_NULL_DELETE_FN` from either C, C++, or Fortran. `MPI_TYPE_NULL_DELETE_FN` is a function that does nothing, other than returning `MPI_SUCCESS`.

The C callback functions are:

```
typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
                                         int type_keyval, void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int *flag);

and

typedef int MPI_Type_delete_attr_function(MPI_Datatype type, int type_keyval,
                                         void *attribute_val, void *extra_state);
```

The Fortran callback functions are:

```
SUBROUTINE TYPE_COPY_ATTR_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE,
                             ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE,
ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT
LOGICAL FLAG

and

SUBROUTINE TYPE_DELETE_ATTR_FN(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
                              IERROR)
```

```
1    INTEGER TYPE, TYPE_KEYVAL, IERROR
```

```
2    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

```
3    The C++ callbacks are:
```

```
4    typedef int MPI::Datatype::Copy_attr_function(const MPI::Datatype& oldtype,
5            int type_keyval, void* extra_state,
6            const void* attribute_val_in, void* attribute_val_out,
7            bool& flag);
8
```

```
9    and
```

```
10   typedef int MPI::Datatype::Delete_attr_function(MPI::Datatype& type,
11           int type_keyval, void* attribute_val, void* extra_state);
12
```

```
13
14   MPI_TYPE_FREE_KEYVAL(type_keyval)
```

```
15       INOUT    type_keyval                key value (integer)
16
```

```
17   int MPI_Type_free_keyval(int *type_keyval)
18
```

```
19   MPI_TYPE_FREE_KEYVAL(TYPE_KEYVAL, IERROR)
```

```
20       INTEGER TYPE_KEYVAL, IERROR
21
```

```
22   static void MPI::Datatype::Free_keyval(int& type_keyval)
23
```

```
24
25   MPI_TYPE_SET_ATTR(type, type_keyval, attribute_val)
```

```
26       INOUT    type                datatype to which attribute will be attached (handle)
27
```

```
28       IN      type_keyval          key value (integer)
29
```

```
29       IN      attribute_val        attribute value
30
```

```
31   int MPI_Type_set_attr(MPI_Datatype type, int type_keyval,
32       void *attribute_val)
33
```

```
34   MPI_TYPE_SET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, IERROR)
```

```
35       INTEGER TYPE, TYPE_KEYVAL, IERROR
```

```
36       INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
37
```

```
37   void MPI::Datatype::Set_attr(int type_keyval, const void* attribute_val)
38
39
40
41
42
43
44
45
46
47
48
```

```

MPI_TYPE_GET_ATTR(type, type_keyval, attribute_val, flag)
    IN      type          datatype to which the attribute is attached (handle)
    IN      type_keyval    key value (integer)
    OUT     attribute_val   attribute value, unless flag = false
    OUT     flag            false if no attribute is associated with the key (logical)

int MPI_Type_get_attr(MPI_Datatype type, int type_keyval, void
    *attribute_val, int *flag)

MPI_TYPE_GET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
    LOGICAL FLAG

bool MPI::Datatype::Get_attr(int type_keyval, void* attribute_val) const

MPI_TYPE_DELETE_ATTR(type, type_keyval)
    INOUT   type          datatype from which the attribute is deleted (handle)
    IN      type_keyval    key value (integer)

int MPI_Type_delete_attr(MPI_Datatype type, int type_keyval)

MPI_TYPE_DELETE_ATTR(TYPE, TYPE_KEYVAL, IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR

void MPI::Datatype::Delete_attr(int type_keyval)

```

8.9 Duplicating a Datatype

```

MPI_TYPE_DUP(type, newtype)
    IN      type          datatype (handle)
    OUT     newtype        copy of type (handle)

int MPI_Type_dup(MPI_Datatype type, MPI_Datatype *newtype)

MPI_TYPE_DUP(TYPE, NEWTYPE, IERROR)
    INTEGER TYPE, NEWTYPE, IERROR

MPI::Datatype MPI::Datatype::Dup() const

```

`MPI_TYPE_DUP` is a new type constructor which duplicates the existing `type` with associated key values. For each key value, the respective copy callback function determines the attribute value associated with this key in the new communicator; one particular action that a copy callback may take is to delete the attribute from the new datatype. Returns

in **newtype** a new datatype with exactly the same properties as **type** and any copied cached information. The new datatype has identical upper bound and lower bound and yields the same net result when fully decoded with the functions in Section 8.6. The **newtype** has the same committed state as the old **type**.

Chapter 9

I/O

9.1 Introduction

POSIX provides a model of a widely portable file system, but the portability and optimization needed for parallel I/O cannot be achieved with the POSIX interface.

The significant optimizations required for efficiency (e.g., grouping [15], collective buffering [1, 2, 16, 19, 22], and disk-directed I/O [13]) can only be implemented if the parallel I/O system provides a high-level interface supporting partitioning of file data among processes and a collective interface supporting complete transfers of global data structures between process memories and files. In addition, further efficiencies can be gained via support for asynchronous I/O, strided accesses, and control over physical file layout on storage devices (disks). The I/O environment described in this chapter provides these facilities.

Instead of defining I/O access modes to express the common patterns for accessing a shared file (broadcast, reduction, scatter, gather), we chose another approach in which data partitioning is expressed using derived datatypes. Compared to a limited set of predefined access patterns, this approach has the advantage of added flexibility and expressiveness.

9.1.1 Definitions

file An MPI file is an ordered collection of typed data items. MPI supports random or sequential access to any integral set of these items. A file is opened collectively by a group of processes. All collective I/O calls on a file are collective over this group.

displacement A file *displacement* is an absolute byte position relative to the beginning of a file. The displacement defines the location where a *view* begins. Note that a “file displacement” is distinct from a “typemap displacement.”

etype An *etype* (*elementary datatype*) is the unit of data access and positioning. It can be any MPI predefined or derived datatype. Derived etypes can be constructed using any of the MPI datatype constructor routines, provided all resulting typemap displacements are nonnegative and monotonically nondecreasing. Data access is performed in etype units, reading or writing whole data items of type etype. Offsets are expressed as a count of etypes; file pointers point to the beginning of etypes. Depending on context, the term “etype” is used to describe one of three aspects of an elementary datatype: a particular MPI type, a data item of that type, or the extent of that type.

filetype A *filetype* is the basis for partitioning a file among processes and defines a template for accessing the file. A filetype is either a single etype or a derived MPI datatype constructed from multiple instances of the same etype. In addition, the extent of any hole in the filetype must be a multiple of the etype's extent. The displacements in the typemap of the filetype are not required to be distinct, but they must be nonnegative and monotonically nondecreasing.

view A *view* defines the current set of data visible and accessible from an open file as an ordered set of etypes. Each process has its own view of the file, defined by three quantities: a displacement, an etype, and a filetype. The pattern described by a filetype is repeated, beginning at the displacement, to define the view. The pattern of repetition is defined to be the same pattern that `MPI_TYPE_CONTIGUOUS` would produce if it were passed the filetype and an arbitrarily large count. Figure 9.1 shows how the tiling works; note that the filetype in this example must have explicit lower and upper bounds set in order for the initial and final holes to be repeated in the view. Views can be changed by the user during program execution. The default view is a linear byte stream (displacement is zero, etype and filetype equal to `MPI_BYTE`).

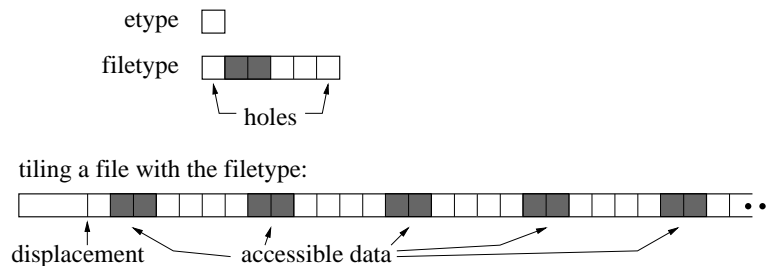


Figure 9.1: Etypes and filetypes

A group of processes can use complementary views to achieve a global data distribution such as a scatter/gather pattern (see Figure 9.2).




Figure 9.2: Partitioning a file among parallel processes

offset An *offset* is a position in the file relative to the current view, expressed as a count of etypes. Holes in the view's filetype are skipped when calculating this position. Offset 0 is the location of the first etype visible in the view (after skipping the displacement and any initial holes in the view). For example, an offset of 2 for process 1 in Figure 9.2 is the position of the 8th etype in the file after the displacement. An “explicit offset” is an offset that is used as a formal parameter in explicit data access routines.

file size and end of file The *size* of an MPI file is measured in bytes from the beginning of the file. A newly created file has a size of zero bytes. Using the size as an absolute displacement gives the position of the byte immediately following the last byte in the file. For any given view, the *end of file* is the offset of the first etype accessible in the current view starting after the last byte in the file.

file pointer A *file pointer* is an implicit offset maintained by MPI. “Individual file pointers” are file pointers that are local to each process that opened the file. A “shared file pointer” is a file pointer that is shared by the group of processes that opened the file.

file handle A *file handle* is an opaque object created by `MPI_FILE_OPEN` and freed by `MPI_FILE_CLOSE`. All operations on an open file reference the file through the file handle.

9.2 File Manipulation

9.2.1 Opening a File

`MPI_FILE_OPEN(comm, filename, amode, info, fh)`

IN	comm	communicator (handle)
IN	filename	name of file to open (string)
IN	amode	file access mode (integer)
IN	info	info object (handle)
OUT	fh	new file handle (handle)

```
int MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info,
                  MPI_File *fh)
```

```
MPI_FILE_OPEN(COMM, FILENAME, AMODE, INFO, FH, IERROR)
CHARACTER*(*) FILENAME
INTEGER COMM, AMODE, INFO, FH, IERROR
```

```
static MPI::File MPI::File::Open(const MPI::Intracomm& comm,
                                  const char* filename, int amode, const MPI::Info& info)
```

`MPI_FILE_OPEN` opens the file identified by the file name **filename** on all processes in the **comm** communicator group. `MPI_FILE_OPEN` is a collective routine: all processes must provide the same value for **amode**, and all processes must provide **filenames** that reference the same file. (Values for **info** may vary.) **comm** must be an intracommunicator; it is erroneous to pass an intercommunicator to `MPI_FILE_OPEN`. Errors in `MPI_FILE_OPEN` are raised using the default file error handler (see Section 9.7, page 265). A process can open a file independently of other processes by using the `MPI_COMM_SELF` communicator. The file handle returned, **fh**, can be subsequently used to access the file until the file is closed using `MPI_FILE_CLOSE`. Before calling `MPI_FINALIZE`, the user is required to close (via `MPI_FILE_CLOSE`) all files that were opened with `MPI_FILE_OPEN`. Note that the communicator **comm** is unaffected by `MPI_FILE_OPEN` and continues to be usable in all

MPI routines (e.g., `MPI_SEND`). Furthermore, the use of `comm` will not interfere with I/O behavior.

The format for specifying the file name in the `filename` argument is implementation dependent and must be documented by the implementation.

Advice to implementors. An implementation may require that `filename` include a string or strings specifying additional information about the file. Examples include the type of filesystem (e.g., a prefix of `ufs:`), a remote hostname (e.g., a prefix of `machine.univ.edu:`), or a file password (e.g., a suffix of `/PASSWORD=SECRET`). (*End of advice to implementors.*)

Advice to users. On some implementations of MPI, the file namespace may not be identical from all processes of all applications. For example, `"/tmp/foo"` may denote different files on different processes, or a single file may have many names, dependent on process location. The user is responsible for ensuring that a single file is referenced by the `filename` argument, as it may be impossible for an implementation to detect this type of namespace error. (*End of advice to users.*)

Initially, all processes view the file as a linear byte stream, and each process views data in its own native representation (no data representation conversion is performed). (POSIX files are linear byte streams in the native representation.) The file view can be changed via the `MPI_FILE_SET_VIEW` routine.

The following access modes are supported (specified in `amode`, a bit vector OR of the following integer constants):

- `MPI_MODE_RDONLY` — read only,
- `MPI_MODE_RDWR` — reading and writing,
- `MPI_MODE_WRONLY` — write only,
- `MPI_MODE_CREATE` — create the file if it does not exist,
- `MPI_MODE_EXCL` — error if creating file that already exists,
- `MPI_MODE_DELETE_ON_CLOSE` — delete file on close,
- `MPI_MODE_UNIQUE_OPEN` — file will not be concurrently opened elsewhere,
- `MPI_MODE_SEQUENTIAL` — file will only be accessed sequentially,
- `MPI_MODE_APPEND` — set initial position of all file pointers to end of file.

Advice to users. C/C++ users can use bit vector OR (`|`) to combine these constants; Fortran 90 users can use the bit vector `IOR` intrinsic. Fortran 77 users can use (non-portably) bit vector `IOR` on systems that support it. Alternatively, Fortran users can portably use integer addition to OR the constants (each constant should appear at most once in the addition). (*End of advice to users.*)

Advice to implementors. The values of these constants must be defined such that the bitwise OR and the sum of any distinct set of these constants is equivalent. (*End of advice to implementors.*)

The modes `MPI_MODE_RDONLY`, `MPI_MODE_RDWR`, `MPI_MODE_WRONLY`, `MPI_MODE_CREATE`, and `MPI_MODE_EXCL` have identical semantics to their POSIX counterparts [11]. Exactly one of `MPI_MODE_RDONLY`, `MPI_MODE_RDWR`, or `MPI_MODE_WRONLY`, must be specified. It is erroneous to specify `MPI_MODE_CREATE` or `MPI_MODE_EXCL` in conjunction with `MPI_MODE_RDONLY`; it is erroneous to specify `MPI_MODE_SEQUENTIAL` together with `MPI_MODE_RDWR`.

The `MPI_MODE_DELETE_ON_CLOSE` mode causes the file to be deleted (equivalent to performing an `MPI_FILE_DELETE`) when the file is closed.

The `MPI_MODE_UNIQUE_OPEN` mode allows an implementation to optimize access by eliminating the overhead of file locking. It is erroneous to open a file in this mode unless the file will not be concurrently opened elsewhere.

Advice to users. For `MPI_MODE_UNIQUE_OPEN`, *not opened elsewhere* includes both inside and outside the MPI environment. In particular, one needs to be aware of potential external events which may open files (e.g., automated backup facilities). When `MPI_MODE_UNIQUE_OPEN` is specified, the user is responsible for ensuring that no such external events take place. (*End of advice to users.*)

The `MPI_MODE_SEQUENTIAL` mode allows an implementation to optimize access to some sequential devices (tapes and network streams). It is erroneous to attempt nonsequential access to a file that has been opened in this mode.

Specifying `MPI_MODE_APPEND` only guarantees that all shared and individual file pointers are positioned at the initial end of file when `MPI_FILE_OPEN` returns. Subsequent positioning of file pointers is application dependent. In particular, the implementation does not ensure that all writes are appended.

Errors related to the access mode are raised in the class `MPI_ERR_AMODE`.

The **info** argument is used to provide information regarding file access patterns and file system specifics (see Section 9.2.8, page 218). The constant `MPI_INFO_NULL` can be used when no info needs to be specified.

Advice to users. Some file attributes are inherently implementation dependent (e.g., file permissions). These attributes must be set using either the `info` argument or facilities outside the scope of `MPI`. (*End of advice to users.*)

Files are opened by default using nonatomic mode file consistency semantics (see Section 9.6.1, page 255). The more stringent atomic mode consistency semantics, required for atomicity of conflicting accesses, can be set using `MPI_FILE_SET_ATOMICITY`.

9.2.2 Closing a File

MPI_FILE_CLOSE(fh)

INOUT	fh	file handle (handle)
-------	----	----------------------

```
int MPI_File_close(MPI_File *fh)
```

```
MPI_FILE_CLOSE(FH, IERROR)
```

INTEGER FH, IERROR

```
1 void MPI::File::Close()
```

```
2     MPI_FILE_CLOSE first synchronizes file state (equivalent to performing an
3     MPI_FILE_SYNC), then closes the file associated with fh. The file is deleted if it was opened
4     with access mode MPI_MODE_DELETE_ON_CLOSE (equivalent to performing an
5     MPI_FILE_DELETE). MPI_FILE_CLOSE is a collective routine.
```

```
7     Advice to users. If the file is deleted on close, and there are other processes currently
8     accessing the file, the status of the file and the behavior of future accesses by these
9     processes are implementation dependent. (End of advice to users.)
```

```
11    The user is responsible for ensuring that all outstanding nonblocking requests and
12    split collective operations associated with fh made by a process have completed before that
13    process calls MPI_FILE_CLOSE.
```

```
14    The MPI_FILE_CLOSE routine deallocates the file handle object and sets fh to
15    MPI_FILE_NULL.
```

17 9.2.3 Deleting a File

```
20 MPI_FILE_DELETE(filename, info)
```

```
22     IN          filename          name of file to delete (string)
```

```
23     IN          info             info object (handle)
```

```
25 int MPI_File_delete(char *filename, MPI_Info info)
```

```
27 MPI_FILE_DELETE(FILENAME, INFO, IERROR)
```

```
28     CHARACTER*(*) FILENAME
```

```
29     INTEGER INFO, IERROR
```

```
30 static void MPI::File::Delete(const char* filename, const MPI::Info& info)
```

```
32     MPI_FILE_DELETE deletes the file identified by the file name filename. If the file does
33     not exist, MPI_FILE_DELETE raises an error in the class MPI_ERR_NO_SUCH_FILE.
```

```
34     The info argument can be used to provide information regarding file system specifics
35     (see Section 9.2.8, page 218). The constant MPI_INFO_NULL refers to the null info, and can
36     be used when no info needs to be specified.
```

```
37     If a process currently has the file open, the behavior of any access to the file (as well
38     as the behavior of any outstanding accesses) is implementation dependent. In addition,
39     whether an open file is deleted or not is also implementation dependent. If the file is not
40     deleted, an error in the class MPI_ERR_FILE_IN_USE or MPI_ERR_ACCESS will be raised. Errors
41     are raised using the default error handler (see Section 9.7, page 265).
```

9.2.4 Resizing a File

`MPI_FILE_SET_SIZE(fh, size)`

INOUT	fh	file handle (handle)
IN	size	size to truncate or expand file (integer)

`int MPI_File_set_size(MPI_File fh, MPI_Offset size)`

`MPI_FILE_SET_SIZE(FH, SIZE, IERROR)`

INTEGER FH, IERROR

INTEGER(KIND=MPI_OFFSET_KIND) SIZE

`void MPI::File::Set_size(MPI::Offset size)`

`MPI_FILE_SET_SIZE` resizes the file associated with the file handle `fh`. `size` is measured in bytes from the beginning of the file. `MPI_FILE_SET_SIZE` is collective; all processes in the group must pass identical values for `size`.

If `size` is smaller than the current file size, the file is truncated at the position defined by `size`. The implementation is free to deallocate file blocks located beyond this position.

If `size` is larger than the current file size, the file size becomes `size`. Regions of the file that have been previously written are unaffected. The values of data in the new regions in the file (those locations with displacements between old file size and `size`) are undefined. It is implementation dependent whether the `MPI_FILE_SET_SIZE` routine allocates file space—use `MPI_FILE_PREALLOCATE` to force file space to be reserved.

`MPI_FILE_SET_SIZE` does not affect the individual file pointers or the shared file pointer. If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call this routine.

Advice to users. It is possible for the file pointers to point beyond the end of file after a `MPI_FILE_SET_SIZE` operation truncates a file. This is legal, and equivalent to seeking beyond the current end of file. (*End of advice to users.*)

All nonblocking requests and split collective operations on `fh` must be completed before calling `MPI_FILE_SET_SIZE`. Otherwise, calling `MPI_FILE_SET_SIZE` is erroneous. As far as consistency semantics are concerned, `MPI_FILE_SET_SIZE` is a write operation that conflicts with operations that access bytes at displacements between the old and new file sizes (see Section 9.6.1, page 255).

9.2.5 Preallocating Space for a File

`MPI_FILE_PREALLOCATE(fh, size)`

INOUT	fh	file handle (handle)
IN	size	size to preallocate file (integer)

`int MPI_File_preallocate(MPI_File fh, MPI_Offset size)`

```

1 MPI_FILE_PREALLOCATE(FH, SIZE, IERROR)
2     INTEGER FH, IERROR
3     INTEGER(KIND=MPI_OFFSET_KIND) SIZE
4
5 void MPI::File::Preallocate(MPI::Offset size)

```

`MPI_FILE_PREALLOCATE` ensures that storage space is allocated for the first `size` bytes of the file associated with `fh`. `MPI_FILE_PREALLOCATE` is collective; all processes in the group must pass identical values for `size`. Regions of the file that have previously been written are unaffected. For newly allocated regions of the file, `MPI_FILE_PREALLOCATE` has the same effect as writing undefined data. If `size` is larger than the current file size, the file size increases to `size`. If `size` is less than or equal to the current file size, the file size is unchanged.

The treatment of file pointers, pending nonblocking accesses, and file consistency is the same as with `MPI_FILE_SET_SIZE`. If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call this routine.

Advice to users. In some implementations, file preallocation may be expensive. (*End of advice to users.*)

9.2.6 Querying the Size of a File

```

23 MPI_FILE_GET_SIZE(fh, size)
24     IN          fh          file handle (handle)
25     OUT         size        size of the file in bytes (integer)
26
27
28 int MPI_File_get_size(MPI_File fh, MPI_Offset *size)
29
30 MPI_FILE_GET_SIZE(FH, SIZE, IERROR)
31     INTEGER FH, IERROR
32     INTEGER(KIND=MPI_OFFSET_KIND) SIZE
33
34 MPI::Offset MPI::File::Get_size() const

```

`MPI_FILE_GET_SIZE` returns, in `size`, the current size in bytes of the file associated with the file handle `fh`. As far as consistency semantics are concerned, `MPI_FILE_GET_SIZE` is a data access operation (see Section 9.6.1, page 255).

9.2.7 Querying File Parameters

```

42 MPI_FILE_GET_GROUP(fh, group)
43     IN          fh          file handle (handle)
44     OUT         group       group which opened the file (handle)
45
46
47 int MPI_File_get_group(MPI_File fh, MPI_Group *group)
48

```

```
MPI_FILE_GET_GROUP(FH, GROUP, IERROR)
```

```
    INTEGER FH, GROUP, IERROR
```

```
MPI::Group MPI::File::Get_group() const
```

`MPI_FILE_GET_GROUP` returns a duplicate of the group of the communicator used to open the file associated with `fh`. The group is returned in `group`. The user is responsible for freeing `group`.

```
MPI_FILE_GET_AMODE(fh, amode)
```

```
    IN      fh                file handle (handle)
```

```
    OUT     amode             file access mode used to open the file (integer)
```

```
int MPI_File_get_amode(MPI_File fh, int *amode)
```

```
MPI_FILE_GET_AMODE(FH, AMODE, IERROR)
```

```
    INTEGER FH, AMODE, IERROR
```

```
int MPI::File::Get_amode() const
```

`MPI_FILE_GET_AMODE` returns, in `amode`, the access mode of the file associated with `fh`.

Example 9.1 In Fortran 77, decoding an `amode` bit vector will require a routine such as the following:

```

      SUBROUTINE BIT_QUERY(TEST_BIT, MAX_BIT, AMODE, BIT_FOUND)
      !
      !  TEST IF THE INPUT TEST_BIT IS SET IN THE INPUT AMODE
      !  IF SET, RETURN 1 IN BIT_FOUND, 0 OTHERWISE
      !
      INTEGER TEST_BIT, AMODE, BIT_FOUND, CP_AMODE, HIFOUND
      BIT_FOUND = 0
      CP_AMODE = AMODE
100  CONTINUE
      LBIT = 0
      HIFOUND = 0
      DO 20 L = MAX_BIT, 0, -1
          MATCHER = 2**L
          IF (CP_AMODE .GE. MATCHER .AND. HIFOUND .EQ. 0) THEN
              HIFOUND = 1
              LBIT = MATCHER
              CP_AMODE = CP_AMODE - MATCHER
          END IF
20  CONTINUE
      IF (HIFOUND .EQ. 1 .AND. LBIT .EQ. TEST_BIT) BIT_FOUND = 1
      IF (BIT_FOUND .EQ. 0 .AND. HIFOUND .EQ. 1 .AND. &
          CP_AMODE .GT. 0) GO TO 100
      END

```

This routine could be called successively to decode `amode`, one bit at a time. For example, the following code fragment would check for `MPI_MODE_RDONLY`.

```

CALL BIT_QUERY(MPI_MODE_RDONLY, 30, AMODE, BIT_FOUND)
IF (BIT_FOUND .EQ. 1) THEN
    PRINT *, ' FOUND READ-ONLY BIT IN AMODE=', AMODE
ELSE
    PRINT *, ' READ-ONLY BIT NOT FOUND IN AMODE=', AMODE
END IF

```

9.2.8 File Info

Hints specified via info (see Section 4.10, page 43) allow a user to provide information such as file access patterns and file system specifics to direct optimization. Providing hints may enable an implementation to deliver increased I/O performance or minimize the use of system resources. However, hints do not change the semantics of any of the I/O interfaces. In other words, an implementation is free to ignore all hints. Hints are specified on a per file basis, in `MPI_FILE_OPEN`, `MPI_FILE_DELETE`, `MPI_FILE_SET_VIEW`, and `MPI_FILE_SET_INFO`, via the opaque info object.

Advice to implementors. It may happen that a program is coded with hints for one system, and later executes on another system that does not support these hints. In general, unsupported hints should simply be ignored. Needless to say, no hint can be mandatory. However, for each hint used by a specific implementation, a default value must be provided when the user does not specify a value for this hint. (*End of advice to implementors.*)

`MPI_FILE_SET_INFO(fh, info)`

INOUT	fh	file handle (handle)
IN	info	info object (handle)

```
int MPI_File_set_info(MPI_File fh, MPI_Info info)
```

```
MPI_FILE_SET_INFO(FH, INFO, IERROR)
INTEGER FH, INFO, IERROR
```

```
void MPI::File::Set_info(const MPI::Info& info)
```

`MPI_FILE_SET_INFO` sets new values for the hints of the file associated with `fh`. `MPI_FILE_SET_INFO` is a collective routine. The info object may be different on each process, but any info entries that an implementation requires to be the same on all processes must appear with the same value in each process's info object.

Advice to users. Many info items that an implementation can use when it creates or opens a file cannot easily be changed once the file has been created or opened. Thus, an implementation may ignore hints issued in this call that it would have accepted in an open call. (*End of advice to users.*)

`MPI_FILE_GET_INFO(fh, info_used)`

IN	<code>fh</code>	file handle (handle)
OUT	<code>info_used</code>	new info object (handle)

`int MPI_File_get_info(MPI_File fh, MPI_Info *info_used)`

`MPI_FILE_GET_INFO(FH, INFO_USED, IERROR)`

INTEGER FH, INFO_USED, IERROR

`MPI::Info MPI::File::Get_info() const`

`MPI_FILE_GET_INFO` returns a new info object containing the hints of the file associated with `fh`. The current setting of all hints actually used by the system related to this open file is returned in `info_used`. The user is responsible for freeing `info_used` via `MPI_INFO_FREE`.

Advice to users. The info object returned in `info_used` will contain all hints currently active for this file. This set of hints may be greater or smaller than the set of hints passed in to `MPI_FILE_OPEN`, `MPI_FILE_SET_VIEW`, and `MPI_FILE_SET_INFO`, as the system may not recognize some hints set by the user, and may recognize other hints that the user has not set. (*End of advice to users.*)

Reserved File Hints

Some potentially useful hints (info key values) are outlined below. The following key values are reserved. An implementation is not required to interpret these key values, but if it does interpret the key value, it must provide the functionality described. (For more details on “info,” see Section 4.10, page 43.)

These hints mainly affect access patterns and the layout of data on parallel I/O devices. For each hint name introduced, we describe the purpose of the hint, and the type of the hint value. The “[**SAME**]” annotation specifies that the hint values provided by all participating processes must be identical; otherwise the program is erroneous. In addition, some hints are context dependent, and are only used by an implementation at specific times (e.g., `file_perm` is only useful during file creation).

access_style (comma separated list of strings): This hint specifies the manner in which the file will be accessed until the file is closed or until the `access_style` key value is altered. The hint value is a comma separated list of the following: `read_once`, `write_once`, `read_mostly`, `write_mostly`, `sequential`, `reverse_sequential`, and `random`.

collective_buffering (boolean) [SAME]: This hint specifies whether the application may benefit from collective buffering. Collective buffering is an optimization performed on collective accesses. Accesses to the file are performed on behalf of all processes in the group by a number of target nodes. These target nodes coalesce small requests into large disk accesses. Legal values for this key are `true` and `false`. Collective buffering parameters are further directed via additional hints: `cb_block_size`, `cb_buffer_size`, and `cb_nodes`.

cb_block_size (integer) [SAME]: This hint specifies the block size to be used for collective buffering file access. *Target nodes* access data in chunks of this size. The chunks are distributed among target nodes in a round-robin (CYCLIC) pattern.

cb_buffer_size (integer) [SAME]: This hint specifies the total buffer space that can be used for collective buffering on each target node, usually a multiple of **cb_block_size**.

cb_nodes (integer) [SAME]: This hint specifies the number of target nodes to be used for collective buffering.

chunked (comma separated list of integers) [SAME]: This hint specifies that the file consists of a multidimensional array that is often accessed by subarrays. The value for this hint is a comma separated list of array dimensions, starting from the most significant one (for an array stored in row-major order, as in C, the most significant dimension is the first one; for an array stored in column-major order, as in Fortran, the most significant dimension is the last one, and array dimensions should be reversed).

chunked_item (comma separated list of integers) [SAME]: This hint specifies the size of each array entry, in bytes.

chunked_size (comma separated list of integers) [SAME]: This hint specifies the dimensions of the subarrays. This is a comma separated list of array dimensions, starting from the most significant one.

filename (string): This hint specifies the file name used when the file was opened. If the implementation is capable of returning the file name of an open file, it will be returned using this key by **MPI_FILE_GET_INFO**. This key is ignored when passed to **MPI_FILE_OPEN**, **MPI_FILE_SET_VIEW**, **MPI_FILE_SET_INFO**, and **MPI_FILE_DELETE**.

file_perm (string) [SAME]: This hint specifies the file permissions to use for file creation. Setting this hint is only useful when passed to **MPI_FILE_OPEN** with an **amode** that includes **MPI_MODE_CREATE**. The set of legal values for this key is implementation dependent.

io_node_list (comma separated list of strings) [SAME]: This hint specifies the list of I/O devices that should be used to store the file. This hint is most relevant when the file is created.

nb_proc (integer) [SAME]: This hint specifies the number of parallel processes that will typically be assigned to run programs that access this file. This hint is most relevant when the file is created.

num_io_nodes (integer) [SAME]: This hint specifies the number of I/O devices in the system. This hint is most relevant when the file is created.

striping_factor (integer) [SAME]: This hint specifies the number of I/O devices that the file should be striped across, and is relevant only when the file is created.

striping_unit (integer) [SAME]: This hint specifies the suggested striping unit to be used for this file. The striping unit is the amount of consecutive data assigned to one I/O device before progressing to the next device, when striping across a number of devices. It is expressed in bytes. This hint is relevant only when the file is created.

9.3 File Views

```
MPI_FILE_SET_VIEW(fh, disp, etype, filetype, datarep, info)
```

INOUT	fh	file handle (handle)
IN	disp	displacement (integer)
IN	etype	elementary datatype (handle)
IN	filetype	filetype (handle)
IN	datarep	data representation (string)
IN	info	info object (handle)

```
int MPI_File_set_view(MPI_File fh, MPI_Offset disp, MPI_Datatype etype,
                     MPI_Datatype filetype, char *datarep, MPI_Info info)
```

```
MPI_FILE_SET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, INFO, IERROR)
    INTEGER FH, ETYPE, FILETYPE, INFO, IERROR
    CHARACTER*(*) DATAREP
    INTEGER(KIND=MPI_OFFSET_KIND) DISP
```

```
void MPI::File::Set_view(MPI::Offset disp, const MPI::Datatype& etype,
                        const MPI::Datatype& filetype, const char* datarep,
                        const MPI::Info& info)
```

The `MPI_FILE_SET_VIEW` routine changes the process's view of the data in the file. The start of the view is set to **disp**; the type of data is set to **etype**; the distribution of data to processes is set to **filetype**; and the representation of data in the file is set to **datarep**. In addition, `MPI_FILE_SET_VIEW` resets the individual file pointers and the shared file pointer to zero. `MPI_FILE_SET_VIEW` is collective; the values for **datarep** and the extents of **etype** in the file data representation must be identical on all processes in the group; values for **disp**, **filetype**, and **info** may vary. The datatypes passed in **etype** and **filetype** must be committed.

The **etype** always specifies the data layout in the file. If **etype** is a portable datatype (see Section 2.4, page 7), the extent of **etype** is computed by scaling any displacements in the datatype to match the file data representation. If **etype** is not a portable datatype, no scaling is done when computing the extent of **etype**. The user must be careful when using nonportable **etypes** in heterogeneous environments; see Section 9.5.1, page 248 for further details.

If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, the special displacement `MPI_DISPLACEMENT_CURRENT` must be passed in **disp**. This sets the displacement to the current position of the shared file pointer.

Rationale. For some sequential files, such as those corresponding to magnetic tapes or streaming network connections, the *displacement* may not be meaningful. `MPI_DISPLACEMENT_CURRENT` allows the view to be changed for these types of files. (*End of rationale.*)

Advice to implementors. It is expected that a call to `MPI_FILE_SET_VIEW` will immediately follow `MPI_FILE_OPEN` in numerous instances. A high quality implementation will ensure that this behavior is efficient. (*End of advice to implementors.*)

The **disp** displacement argument specifies the position (absolute offset in bytes from the beginning of the file) where the view begins.

Advice to users. **disp** can be used to skip headers or when the file includes a sequence of data segments that are to be accessed in different patterns (see Figure 9.3). Separate views, each using a different displacement and filetype, can be used to access each segment.




Figure 9.3: Displacements

(*End of advice to users.*)

An *etype* (*elementary datatype*) is the unit of data access and positioning. It can be any MPI predefined or derived datatype. Derived etypes can be constructed by using any of the MPI datatype constructor routines, provided all resulting typemap displacements are nonnegative and monotonically nondecreasing. Data access is performed in etype units, reading or writing whole data items of type etype. Offsets are expressed as a count of **etypes**; file pointers point to the beginning of etypes.

Advice to users. In order to ensure interoperability in a heterogeneous environment, additional restrictions must be observed when constructing the **etype** (see Section 9.5, page 246). (*End of advice to users.*)

A filetype is either a single etype or a derived MPI datatype constructed from multiple instances of the same etype. In addition, the extent of any hole in the filetype must be a multiple of the etype's extent. These displacements are not required to be distinct, but they cannot be negative, and they must be monotonically nondecreasing.

If the file is opened for writing, neither the **etype** nor the **filetype** is permitted to contain overlapping regions. This restriction is equivalent to the “datatype used in a receive cannot specify overlapping regions” restriction for communication. Note that **filetypes** from different processes may still overlap each other.

If **filetype** has holes in it, then the data in the holes is inaccessible to the calling process. However, the **disp**, **etype** and **filetype** arguments can be changed via future calls to **MPI_FILE_SET_VIEW** to access a different part of the file.

It is erroneous to use absolute addresses in the construction of the **etype** and **filetype**.

The **info** argument is used to provide information regarding file access patterns and file system specifics to direct optimization (see Section 9.2.8, page 218). The constant **MPI_INFO_NULL** refers to the null info and can be used when no info needs to be specified.

The **datarep** argument is a string that specifies the representation of data in the file. See the file interoperability section (Section 9.5, page 246) for details and a discussion of valid values.

The user is responsible for ensuring that all nonblocking requests and split collective operations on `fh` have been completed before calling `MPI_FILE_SET_VIEW`—otherwise, the call to `MPI_FILE_SET_VIEW` is erroneous.

`MPI_FILE_GET_VIEW(fh, disp, etype, filetype, datarep)`

IN	<code>fh</code>	file handle (handle)
OUT	<code>disp</code>	displacement (integer)
OUT	<code>etype</code>	elementary datatype (handle)
OUT	<code>filetype</code>	filetype (handle)
OUT	<code>datarep</code>	data representation (string)

```
int MPI_File_get_view(MPI_File fh, MPI_Offset *disp, MPI_Datatype *etype,
                     MPI_Datatype *filetype, char *datarep)
```

```
MPI_FILE_GET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, IERROR)
    INTEGER FH, ETYPE, FILETYPE, IERROR
    CHARACTER*(*) DATAREP, INTEGER(KIND=MPI_OFFSET_KIND) DISP
```

```
void MPI::File::Get_view(MPI::Offset& disp, MPI::Datatype& etype,
                        MPI::Datatype& filetype, char* datarep) const
```

`MPI_FILE_GET_VIEW` returns the process's view of the data in the file. The current value of the displacement is returned in `disp`. The `etype` and `filetype` are new datatypes with typemaps equal to the typemaps of the current `etype` and `filetype`, respectively.

The data representation is returned in `datarep`. The user is responsible for ensuring that `datarep` is large enough to hold the returned data representation string. The length of a data representation string is limited to the value of `MPI_MAX_DATAREP_STRING`.

In addition, if a portable datatype was used to set the current view, then the corresponding datatype returned by `MPI_FILE_GET_VIEW` is also a portable datatype. If `etype` or `filetype` are derived datatypes, the user is responsible for freeing them. The `etype` and `filetype` returned are both in a committed state.

9.4 Data Access

9.4.1 Data Access Routines

Data is moved between files and processes by issuing read and write calls. There are three orthogonal aspects to data access: positioning (explicit offset *vs.* implicit file pointer), synchronism (blocking *vs.* nonblocking and split collective), and coordination (noncollective *vs.* collective). The following combinations of these data access routines, including two types of file pointers (individual and shared) are provided:

positioning	synchronism	coordination	
		<i>noncollective</i>	<i>collective</i>
<i>explicit offsets</i>	<i>blocking</i>	MPI_FILE_READ_AT MPI_FILE_WRITE_AT	MPI_FILE_READ_AT_ALL MPI_FILE_WRITE_AT_ALL
	<i>nonblocking & split collective</i>	MPI_FILE_IREAD_AT MPI_FILE_IWRITE_AT	MPI_FILE_READ_AT_ALL_BEGIN MPI_FILE_READ_AT_ALL_END MPI_FILE_WRITE_AT_ALL_BEGIN MPI_FILE_WRITE_AT_ALL_END
<i>individual file pointers</i>	<i>blocking</i>	MPI_FILE_READ MPI_FILE_WRITE	MPI_FILE_READ_ALL MPI_FILE_WRITE_ALL
	<i>nonblocking & split collective</i>	MPI_FILE_IREAD MPI_FILE_IWRITE	MPI_FILE_READ_ALL_BEGIN MPI_FILE_READ_ALL_END MPI_FILE_WRITE_ALL_BEGIN MPI_FILE_WRITE_ALL_END
<i>shared file pointer</i>	<i>blocking</i>	MPI_FILE_READ_SHARED MPI_FILE_WRITE_SHARED	MPI_FILE_READ_ORDERED MPI_FILE_WRITE_ORDERED
	<i>nonblocking & split collective</i>	MPI_FILE_IREAD_SHARED MPI_FILE_IWRITE_SHARED	MPI_FILE_READ_ORDERED_BEGIN MPI_FILE_READ_ORDERED_END MPI_FILE_WRITE_ORDERED_BEGIN MPI_FILE_WRITE_ORDERED_END

POSIX `read()`/`fread()` and `write()`/`fwrite()` are blocking, noncollective operations and use individual file pointers. The MPI equivalents are `MPI_FILE_READ` and `MPI_FILE_WRITE`.

Implementations of data access routines may buffer data to improve performance. This does not affect reads, as the data is always available in the user's buffer after a read operation completes. For writes, however, the `MPI_FILE_SYNC` routine provides the only guarantee that data has been transferred to the storage device.

Positioning

MPI provides three types of positioning for data access routines: explicit offsets, individual file pointers, and shared file pointers. The different positioning methods may be mixed within the same program and do not affect each other.

The data access routines that accept explicit offsets contain `_AT` in their name (e.g., `MPI_FILE_WRITE_AT`). Explicit offset operations perform data access at the file position given directly as an argument—no file pointer is used nor updated. Note that this is not equivalent to an atomic seek-and-read or seek-and-write operation, as no “seek” is issued. Operations with explicit offsets are described in Section 9.4.2, page 226.

The names of the individual file pointer routines contain no positional qualifier (e.g., `MPI_FILE_WRITE`). Operations with individual file pointers are described in Section 9.4.3, page 230. The data access routines that use shared file pointers contain `_SHARED` or `_ORDERED` in their name (e.g., `MPI_FILE_WRITE_SHARED`). Operations with shared file pointers are described in Section 9.4.4, page 235.

The main semantic issues with MPI-maintained file pointers are how and when they are updated by I/O operations. In general, each I/O operation leaves the file pointer pointing to the next data item after the last one that is accessed by the operation. In a nonblocking or split collective operation, the pointer is updated by the call that initiates the I/O, possibly before the access completes.

More formally,

$$new_file_offset = old_file_offset + \frac{elements(datatype)}{elements(etype)} \times count$$

where *count* is the number of *datatype* items to be accessed, *elements(X)* is the number of predefined datatypes in the typemap of *X*, and *old_file_offset* is the value of the implicit offset before the call. The file position, *new_file_offset*, is in terms of a count of etypes relative to the current view.

Synchronism

MPI supports blocking and nonblocking I/O routines.

A *blocking* I/O call will not return until the I/O request is completed.

A *nonblocking* I/O call initiates an I/O operation, but does not wait for it to complete. Given suitable hardware, this allows the transfer of data out/in the user's buffer to proceed concurrently with computation. A separate *request complete* call (`MPI_WAIT`, `MPI_TEST`, or any of their variants) is needed to complete the I/O request, i.e., to confirm that the data has been read or written and that it is safe for the user to reuse the buffer. The nonblocking versions of the routines are named `MPI_FILE_XXX`, where the *l* stands for immediate.

It is erroneous to access the local buffer of a nonblocking data access operation, or to use that buffer as the source or target of other communications, between the initiation and completion of the operation.

The split collective routines support a restricted form of “nonblocking” operations for collective data access (see Section 9.4.5, page 240).

Coordination

Every noncollective data access routine `MPI_FILE_XXX` has a collective counterpart. For most routines, this counterpart is `MPI_FILE_XXX_ALL` or a pair of `MPI_FILE_XXX_BEGIN` and `MPI_FILE_XXX_END`. The counterparts to the `MPI_FILE_XXX_SHARED` routines are `MPI_FILE_XXX_ORDERED`.

The completion of a noncollective call only depends on the activity of the calling process. However, the completion of a collective call (which must be called by all members of the process group) may depend on the activity of the other processes participating in the collective call. See Section 9.6.4, page 259, for rules on semantics of collective calls.

Collective operations may perform much better than their noncollective counterparts, as global data accesses have significant potential for automatic optimization.

Data Access Conventions

Data is moved between files and processes by calling read and write routines. Read routines move data from a file into memory. Write routines move data from memory into a file. The file is designated by a file handle, *fh*. The location of the file data is specified by an offset into the current view. The data in memory is specified by a triple: *buf*, *count*, and *datatype*. Upon completion, the amount of data accessed by the calling process is returned in a *status*.

An offset designates the starting position in the file for an access. The offset is always in etype units relative to the current view. Explicit offset routines pass *offset* as an argument (negative values are erroneous). The file pointer routines use implicit offsets maintained by MPI.

A data access routine attempts to transfer (read or write) **count** data items of type **datatype** between the user's buffer **buf** and the file. The **datatype** passed to the routine must be a committed datatype. The layout of data in memory corresponding to **buf**, **count**, **datatype** is interpreted the same way as in MPI-1 communication functions; see Section 3.12.5 in [6]. The data is accessed from those parts of the file specified by the current view (Section 9.3, page 221). The type signature of **datatype** must match the type signature of some number of contiguous copies of the **etype** of the current view. As in a receive, it is erroneous to specify a **datatype** for reading that contains overlapping regions (areas of memory which would be stored into more than once).

The nonblocking data access routines indicate that MPI can start a data access and associate a request handle, **request**, with the I/O operation. Nonblocking operations are completed via **MPI_TEST**, **MPI_WAIT**, or any of their variants.

Data access operations, when completed, return the amount of data accessed in **status**.

Advice to users. To prevent problems with the argument copying and register optimization done by Fortran compilers, please note the hints in subsections “Problems Due to Data Copying and Sequence Association,” and “A Problem with Register Optimization” in Section 10.2.2, pages 286 and 289. (*End of advice to users.*)

For blocking routines, **status** is returned directly. For nonblocking routines and split collective routines, **status** is returned when the operation is completed. The number of **datatype** entries and predefined elements accessed by the calling process can be extracted from **status** by using **MPI_GET_COUNT** and **MPI_GET_ELEMENTS**, respectively. The interpretation of the **MPI_ERROR** field is the same as for other operations — normally undefined, but meaningful if an MPI routine returns **MPI_ERR_IN_STATUS**. The user can pass (in C and Fortran) **MPI_STATUS_IGNORE** in the **status** argument if the return value of this argument is not needed. In C++, the **status** argument is optional. The **status** can be passed to **MPI_TEST_CANCELLED** to determine if the operation was cancelled. All other fields of **status** are undefined.

When reading, a program can detect the end of file by noting that the amount of data read is less than the amount requested. Writing past the end of file increases the file size. The amount of data accessed will be the amount requested, unless an error is raised (or a read reaches the end of file).

9.4.2 Data Access with Explicit Offsets

If **MPI_MODE_SEQUENTIAL** mode was specified when the file was opened, it is erroneous to call the routines in this section.


```

MPI_FILE_READ_AT(fh, offset, buf, count, datatype, status)
    IN      fh                file handle (handle)
    IN      offset            file offset (integer)
    OUT     buf                initial address of buffer (choice)
    IN      count             number of elements in buffer (integer)
    IN      datatype           datatype of each buffer element (handle)
    OUT     status             status object (Status)

int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
    MPI_Datatype datatype, MPI_Status *status)

MPI_FILE_READ_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

void MPI::File::Read_at(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status)

void MPI::File::Read_at(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype)

    MPI_FILE_READ_AT reads a file beginning at the position specified by offset.

MPI_FILE_READ_AT_ALL(fh, offset, buf, count, datatype, status)
    IN      fh                file handle (handle)
    IN      offset            file offset (integer)
    OUT     buf                initial address of buffer (choice)
    IN      count             number of elements in buffer (integer)
    IN      datatype           datatype of each buffer element (handle)
    OUT     status             status object (Status)

int MPI_File_read_at_all(MPI_File fh, MPI_Offset offset, void *buf,
    int count, MPI_Datatype datatype, MPI_Status *status)

MPI_FILE_READ_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

void MPI::File::Read_at_all(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status)

void MPI::File::Read_at_all(MPI::Offset offset, void* buf, int count,
    const MPI::Datatype& datatype)

```



```
void MPI::File::Write_at_all(MPI::Offset offset, const void* buf,
                             int count, const MPI::Datatype& datatype)
```

MPI_FILE_WRITE_AT_ALL is a collective version of the blocking MPI_FILE_WRITE_AT interface.

```
MPI_FILE_READ_AT(fh, offset, buf, count, datatype, request)
```

IN	fh	file handle (handle)
IN	offset	file offset (integer)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```
int MPI_File_iread_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                     MPI_Datatype datatype, MPI_Request *request)
```

```
MPI_FILE_IREAD_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
MPI::Request MPI::File::Iread_at(MPI::Offset offset, void* buf, int count,
                                  const MPI::Datatype& datatype)
```

MPI_FILE_READ_AT is a nonblocking version of the MPI_FILE_READ_AT interface.

```
MPI_FILE_WRITE_AT(fh, offset, buf, count, datatype, request)
```

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```
int MPI_File_fwrite_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                     MPI_Datatype datatype, MPI_Request *request)
```

```
MPI_FILE_IWRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
MPI::Request MPI::File::Iwrite_at(MPI::Offset offset, const void* buf,
                                   int count, const MPI::Datatype& datatype)
```

`MPI_FILE_WRITE_AT` is a nonblocking version of the `MPI_FILE_WRITE_AT` interface.

9.4.3 Data Access with Individual File Pointers

MPI maintains one individual file pointer per process per file handle. The current value of this pointer implicitly specifies the offset in the data access routines described in this section. These routines only use and update the individual file pointers maintained by MPI. The shared file pointer is not used nor updated.

The individual file pointer routines have the same semantics as the data access with explicit offset routines described in Section 9.4.2, page 226, with the following modification:

- the **offset** is defined to be the current value of the MPI-maintained individual file pointer.

After an individual file pointer operation is initiated, the individual file pointer is updated to point to the next etype after the last one that will be accessed. The file pointer is updated relative to the current view of the file.

If `MPI_MODE_SEQUENTIAL` mode was specified when the file was opened, it is erroneous to call the routines in this section.

`MPI_FILE_READ(fh, buf, count, datatype, status)`

INOUT	<code>fh</code>	file handle (handle)
OUT	<code>buf</code>	initial address of buffer (choice)
IN	<code>count</code>	number of elements in buffer (integer)
IN	<code>datatype</code>	datatype of each buffer element (handle)
OUT	<code>status</code>	status object (Status)

```
int MPI_File_read(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
                  MPI_Status *status)
```

```
MPI_FILE_READ(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```
void MPI::File::Read(void* buf, int count, const MPI::Datatype& datatype,
                     MPI::Status& status)
```

```
void MPI::File::Read(void* buf, int count, const MPI::Datatype& datatype)
```

`MPI_FILE_READ` reads a file using the individual file pointer.

Example 9.2 The following Fortran code fragment is an example of reading a file until the end of file is reached:

```
!   Read a preexisting input file until all data has been read.
!   Call routine "process_input" if all requested data is read.
!   The Fortran 90 "exit" statement exits the loop.
```

```
integer    bufsize, numread, totprocessed, status(MPI_STATUS_SIZE)
```

```

parameter (bufsize=100)
real      localbuffer(bufsize)

call MPI_FILE_OPEN( MPI_COMM_WORLD, 'myoldfile', &
                    MPI_MODE_RDONLY, MPI_INFO_NULL, myfh, ierr )
call MPI_FILE_SET_VIEW( myfh, 0, MPI_REAL, MPI_REAL, 'native', &
                        MPI_INFO_NULL, ierr )

totprocessed = 0
do
    call MPI_FILE_READ( myfh, localbuffer, bufsize, MPI_REAL, &
                        status, ierr )
    call MPI_GET_COUNT( status, MPI_REAL, numread, ierr )
    call process_input( localbuffer, numread )
    totprocessed = totprocessed + numread
    if ( numread < bufsize ) exit
enddo

write(6,1001) numread, bufsize, totprocessed
1001 format( "No more data:  read", I3, "and expected", I3, &
            "Processed total of", I6, "before terminating job." )

call MPI_FILE_CLOSE( myfh, ierr )

```

MPI_FILE_READ_ALL(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```

int MPI_File_read_all(MPI_File fh, void *buf, int count,
                     MPI_Datatype datatype, MPI_Status *status)

```

```

MPI_FILE_READ_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
<type> BUF(*)
INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR

```

```

void MPI::File::Read_all(void* buf, int count,
                        const MPI::Datatype& datatype, MPI::Status& status)

```

```

void MPI::File::Read_all(void* buf, int count,
                        const MPI::Datatype& datatype)

```

MPI_FILE_READ_ALL is a collective version of the blocking MPI_FILE_READ interface.

1 MPI_FILE_WRITE(fh, buf, count, datatype, status)

2	INOUT	fh	file handle (handle)
3	IN	buf	initial address of buffer (choice)
4			
5	IN	count	number of elements in buffer (integer)
6	IN	datatype	datatype of each buffer element (handle)
7			
8	OUT	status	status object (Status)

9
10 int MPI_File_write(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
11 MPI_Status *status)

12 MPI_FILE_WRITE(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)

13 <type> BUF(*)

14 INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR

15
16 void MPI::File::Write(const void* buf, int count,
17 const MPI::Datatype& datatype, MPI::Status& status)

18 void MPI::File::Write(const void* buf, int count,
19 const MPI::Datatype& datatype)
20

21 MPI_FILE_WRITE writes a file using the individual file pointer.

22
23 MPI_FILE_WRITE_ALL(fh, buf, count, datatype, status)

24			
25	INOUT	fh	file handle (handle)
26	IN	buf	initial address of buffer (choice)
27			
28	IN	count	number of elements in buffer (integer)
29	IN	datatype	datatype of each buffer element (handle)
30	OUT	status	status object (Status)
31			

32
33 int MPI_File_write_all(MPI_File fh, void *buf, int count,
34 MPI_Datatype datatype, MPI_Status *status)

35 MPI_FILE_WRITE_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)

36 <type> BUF(*)

37 INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR

38
39 void MPI::File::Write_all(const void* buf, int count,
40 const MPI::Datatype& datatype, MPI::Status& status)

41 void MPI::File::Write_all(const void* buf, int count,
42 const MPI::Datatype& datatype)
43

44 MPI_FILE_WRITE_ALL is a collective version of the blocking MPI_FILE_WRITE interface.
45
46
47
48

```

MPI_FILE_IREAD(fh, buf, count, datatype, request)
    INOUT  fh                file handle (handle)
    OUT    buf               initial address of buffer (choice)
    IN     count             number of elements in buffer (integer)
    IN     datatype          datatype of each buffer element (handle)
    OUT    request           request object (handle)

int MPI_File_iread(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
                  MPI_Request *request)

MPI_FILE_IREAD(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR

MPI::Request MPI::File::Iread(void* buf, int count,
                              const MPI::Datatype& datatype)

    MPI_FILE_IREAD is a nonblocking version of the MPI_FILE_READ interface.

```

Example 9.3 The following Fortran code fragment illustrates file pointer update semantics:

```

!   Read the first twenty real words in a file into two local
!   buffers.  Note that when the first MPI_FILE_IREAD returns,
!   the file pointer has been updated to point to the
!   eleventh real word in the file.

    integer  bufsize, req1, req2
    integer, dimension(MPI_STATUS_SIZE) :: status1, status2
    parameter (bufsize=10)
    real     buf1(bufsize), buf2(bufsize)

    call MPI_FILE_OPEN( MPI_COMM_WORLD, 'myoldfile', &
                       MPI_MODE_RDONLY, MPI_INFO_NULL, myfh, ierr )
    call MPI_FILE_SET_VIEW( myfh, 0, MPI_REAL, MPI_REAL, 'native', &
                           MPI_INFO_NULL, ierr )
    call MPI_FILE_IREAD( myfh, buf1, bufsize, MPI_REAL, &
                       req1, ierr )
    call MPI_FILE_IREAD( myfh, buf2, bufsize, MPI_REAL, &
                       req2, ierr )

    call MPI_WAIT( req1, status1, ierr )
    call MPI_WAIT( req2, status2, ierr )

    call MPI_FILE_CLOSE( myfh, ierr )

```

MPI_FILE_IWRITE(fh, buf, count, datatype, request)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	request	request object (handle)

```
int MPI_File_iwrite(MPI_File fh, void *buf, int count,
                    MPI_Datatype datatype, MPI_Request *request)
```

```
MPI_FILE_IWRITE(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
```

```
<type> BUF(*)
```

```
INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
```

```
MPI::Request MPI::File::Iwrite(const void* buf, int count,
                               const MPI::Datatype& datatype)
```

MPI_FILE_IWRITE is a nonblocking version of the **MPI_FILE_WRITE** interface.

MPI_FILE_SEEK(fh, offset, whence)

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	whence	update mode (state)

```
int MPI_File_seek(MPI_File fh, MPI_Offset offset, int whence)
```

```
MPI_FILE_SEEK(FH, OFFSET, WHENCE, IERROR)
```

```
INTEGER FH, WHENCE, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
void MPI::File::Seek(MPI::Offset offset, int whence)
```

MPI_FILE_SEEK updates the individual file pointer according to **whence**, which has the following possible values:

- **MPI_SEEK_SET**: the pointer is set to **offset**
- **MPI_SEEK_CUR**: the pointer is set to the current pointer position plus **offset**
- **MPI_SEEK_END**: the pointer is set to the end of file plus **offset**

The **offset** can be negative, which allows seeking backwards. It is erroneous to seek to a negative position in the view.

MPI_FILE_GET_POSITION(fh, offset)

IN	fh	file handle (handle)
OUT	offset	offset of individual pointer (integer)

```
int MPI_File_get_position(MPI_File fh, MPI_Offset *offset)
```

```
MPI_FILE_GET_POSITION(FH, OFFSET, IERROR)
```

```
INTEGER FH, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
MPI::Offset MPI::File::Get_position() const
```

MPI_FILE_GET_POSITION returns, in **offset**, the current position of the individual file pointer in etype units relative to the current view.

Advice to users. The **offset** can be used in a future call to **MPI_FILE_SEEK** using **whence = MPI_SEEK_SET** to return to the current position. To set the displacement to the current file pointer position, first convert **offset** into an absolute byte position using **MPI_FILE_GET_BYTE_OFFSET**, then call **MPI_FILE_SET_VIEW** with the resulting displacement. (*End of advice to users.*)

MPI_FILE_GET_BYTE_OFFSET(fh, offset, disp)

IN	fh	file handle (handle)
IN	offset	offset (integer)
OUT	disp	absolute byte position of offset (integer)

```
int MPI_File_get_byte_offset(MPI_File fh, MPI_Offset offset,
                             MPI_Offset *disp)
```

```
MPI_FILE_GET_BYTE_OFFSET(FH, OFFSET, DISP, IERROR)
```

```
INTEGER FH, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET, DISP
```

```
MPI::Offset MPI::File::Get_byte_offset(const MPI::Offset disp) const
```

MPI_FILE_GET_BYTE_OFFSET converts a view-relative offset into an absolute byte position. The absolute byte position (from the beginning of the file) of **offset** relative to the current view of **fh** is returned in **disp**.

9.4.4 Data Access with Shared File Pointers

MPI maintains exactly one shared file pointer per collective **MPI_FILE_OPEN** (shared among processes in the communicator group). The current value of this pointer implicitly specifies the offset in the data access routines described in this section. These routines only use and update the shared file pointer maintained by MPI. The individual file pointers are not used nor updated.

The shared file pointer routines have the same semantics as the data access with explicit offset routines described in Section 9.4.2, page 226, with the following modifications:

- the **offset** is defined to be the current value of the MPI-maintained shared file pointer,
- the effect of multiple calls to shared file pointer routines is defined to behave as if the calls were serialized, and
- the use of shared file pointer routines is erroneous unless all processes use the same file view.

For the noncollective shared file pointer routines, the serialization ordering is not deterministic. The user needs to use other synchronization means to enforce a specific order.

After a shared file pointer operation is initiated, the shared file pointer is updated to point to the next etype after the last one that will be accessed. The file pointer is updated relative to the current view of the file.

Noncollective Operations

MPI_FILE_READ_SHARED(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_read_shared(MPI_File fh, void *buf, int count,
                        MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_READ_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```
void MPI::File::Read_shared(void* buf, int count,
                           const MPI::Datatype& datatype, MPI::Status& status)
```

```
void MPI::File::Read_shared(void* buf, int count,
                           const MPI::Datatype& datatype)
```

MPI_FILE_READ_SHARED reads a file using the shared file pointer.

MPI_FILE_WRITE_SHARED(fh, buf, count, datatype, status)

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_write_shared(MPI_File fh, void *buf, int count,
```

```

        MPI_Datatype datatype, MPI_Status *status)
MPI_FILE_WRITE_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
void MPI::File::Write_shared(const void* buf, int count,
    const MPI::Datatype& datatype, MPI::Status& status)
void MPI::File::Write_shared(const void* buf, int count,
    const MPI::Datatype& datatype)

```

MPI_FILE_WRITE_SHARED writes a file using the shared file pointer.

```

MPI_FILE_READ_SHARED(fh, buf, count, datatype, request)
    INOUT   fh                file handle (handle)
    OUT     buf                initial address of buffer (choice)
    IN      count              number of elements in buffer (integer)
    IN      datatype           datatype of each buffer element (handle)
    OUT     request            request object (handle)

```

```

int MPI_File_iread_shared(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Request *request)

```

```

MPI_FILE_IREAD_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR

```

```

MPI::Request MPI::File::Iread_shared(void* buf, int count,
    const MPI::Datatype& datatype)

```

MPI_FILE_READ_SHARED is a nonblocking version of the MPI_FILE_READ_SHARED interface.

```

MPI_FILE_WRITE_SHARED(fh, buf, count, datatype, request)
    INOUT   fh                file handle (handle)
    IN      buf                initial address of buffer (choice)
    IN      count              number of elements in buffer (integer)
    IN      datatype           datatype of each buffer element (handle)
    OUT     request            request object (handle)

```

```

int MPI_File_fwrite_shared(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype, MPI_Request *request)

```

```

MPI_FILE_IWRITE_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)

```

```

1      INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
2
3      MPI::Request MPI::File::Iwrite_shared(const void* buf, int count,
4          const MPI::Datatype& datatype)

```

MPI_FILE_WRITE_SHARED is a nonblocking version of the MPI_FILE_WRITE_SHARED interface.

Collective Operations

The semantics of a collective access using a shared file pointer is that the accesses to the file will be in the order determined by the ranks of the processes within the group. For each process, the location in the file at which data is accessed is the position at which the shared file pointer would be after all processes whose ranks within the group less than that of this process had accessed their data. In addition, in order to prevent subsequent shared offset accesses by the same processes from interfering with this collective access, the call might return only after all the processes within the group have initiated their accesses. When the call returns, the shared file pointer points to the next etype accessible, according to the file view used by all processes, after the last etype requested.

Advice to users. There may be some programs in which all processes in the group need to access the file using the shared file pointer, but the program may not *require* that data be accessed in order of process rank. In such programs, using the shared ordered routines (e.g., MPI_FILE_WRITE_ORDERED rather than MPI_FILE_WRITE_SHARED) may enable an implementation to optimize access, improving performance. (*End of advice to users.*)

Advice to implementors. Accesses to the data requested by all processes do not have to be serialized. Once all processes have issued their requests, locations within the file for all accesses can be computed, and accesses can proceed independently from each other, possibly in parallel. (*End of advice to implementors.*)

```

32 MPI_FILE_READ_ORDERED(fh, buf, count, datatype, status)

```

INOUT	fh	file handle (handle)
OUT	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```

41 int MPI_File_read_ordered(MPI_File fh, void *buf, int count,
42     MPI_Datatype datatype, MPI_Status *status)
43
44 MPI_FILE_READ_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
45     <type> BUF(*)
46     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
47
48 void MPI::File::Read_ordered(void* buf, int count,
49     const MPI::Datatype& datatype, MPI::Status& status)

```

```
void MPI::File::Read_ordered(void* buf, int count,
                             const MPI::Datatype& datatype)
```

MPI_FILE_READ_ORDERED is a collective version of the MPI_FILE_READ_SHARED interface.

```
MPI_FILE_WRITE_ORDERED(fh, buf, count, datatype, status)
```

INOUT	fh	file handle (handle)
IN	buf	initial address of buffer (choice)
IN	count	number of elements in buffer (integer)
IN	datatype	datatype of each buffer element (handle)
OUT	status	status object (Status)

```
int MPI_File_write_ordered(MPIFile fh, void *buf, int count,
                           MPI_Datatype datatype, MPI_Status *status)
```

```
MPI_FILE_WRITE_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
```

```
void MPI::File::Write_ordered(const void* buf, int count,
                              const MPI::Datatype& datatype, MPI::Status& status)
```

```
void MPI::File::Write_ordered(const void* buf, int count,
                              const MPI::Datatype& datatype)
```

MPI_FILE_WRITE_ORDERED is a collective version of the MPI_FILE_WRITE_SHARED interface.

Seek

If MPI_MODE_SEQUENTIAL mode was specified when the file was opened, it is erroneous to call the following two routines (MPI_FILE_SEEK_SHARED and MPI_FILE_GET_POSITION_SHARED).

```
MPI_FILE_SEEK_SHARED(fh, offset, whence)
```

INOUT	fh	file handle (handle)
IN	offset	file offset (integer)
IN	whence	update mode (state)

```
int MPI_File_seek_shared(MPIFile fh, MPI_Offset offset, int whence)
```

```
MPI_FILE_SEEK_SHARED(FH, OFFSET, WHENCE, IERROR)
    INTEGER FH, WHENCE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
void MPI::File::Seek_shared(MPI::Offset offset, int whence)
```

`MPI_FILE_SEEK_SHARED` updates the shared file pointer according to **whence**, which has the following possible values:

- `MPI_SEEK_SET`: the pointer is set to **offset**
- `MPI_SEEK_CUR`: the pointer is set to the current pointer position plus **offset**
- `MPI_SEEK_END`: the pointer is set to the end of file plus **offset**

`MPI_FILE_SEEK_SHARED` is collective; all the processes in the communicator group associated with the file handle **fh** must call `MPI_FILE_SEEK_SHARED` with the same values for **offset** and **whence**.

The **offset** can be negative, which allows seeking backwards. It is erroneous to seek to a negative position in the view.

`MPI_FILE_GET_POSITION_SHARED(fh, offset)`

IN	fh	file handle (handle)
OUT	offset	offset of shared pointer (integer)

```
int MPI_File_get_position_shared(MPI_File fh, MPI_Offset *offset)
```

```
MPI_FILE_GET_POSITION_SHARED(FH, OFFSET, IERROR)
```

```
INTEGER FH, IERROR
```

```
INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
```

```
MPI::Offset MPI::File::Get_position_shared() const
```

`MPI_FILE_GET_POSITION_SHARED` returns, in **offset**, the current position of the shared file pointer in etype units relative to the current view.

Advice to users. The **offset** can be used in a future call to `MPI_FILE_SEEK_SHARED` using **whence** = `MPI_SEEK_SET` to return to the current position. To set the displacement to the current file pointer position, first convert **offset** into an absolute byte position using `MPI_FILE_GET_BYTE_OFFSET`, then call `MPI_FILE_SET_VIEW` with the resulting displacement. (*End of advice to users.*)

9.4.5 Split Collective Data Access Routines

MPI provides a restricted form of “nonblocking collective” I/O operations for all data accesses using split collective data access routines. These routines are referred to as “split” collective routines because a single collective operation is split in two: a begin routine and an end routine. The begin routine begins the operation, much like a nonblocking data access (e.g., `MPI_FILE_READ`). The end routine completes the operation, much like the matching test or wait (e.g., `MPI_WAIT`). As with nonblocking data access operations, the user must not use the buffer passed to a begin routine while the routine is outstanding; the operation must be completed with an end routine before it is safe to free buffers, etc.

Split collective data access operations on a file handle **fh** are subject to the semantic rules given below.

- On any MPI process, each file handle may have at most one active split collective operation at any time.

- Begin calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls.
- End calls are collective over the group of processes that participated in the collective open and follow the ordering rules for collective calls. Each end call matches the preceding begin call for the same collective operation. When an “end” call is made, exactly one unmatched “begin” call for the same operation must precede it.
- An implementation is free to implement any split collective data access routine using the corresponding blocking collective routine when either the begin call (e.g., `MPI_FILE_READ_ALL_BEGIN`) or the end call (e.g., `MPI_FILE_READ_ALL_END`) is issued. The begin and end calls are provided to allow the user and MPI implementation to optimize the collective operation.
- Split collective operations do not match the corresponding regular collective operation. For example, in a single collective read operation, an `MPI_FILE_READ_ALL` on one process does not match an `MPI_FILE_READ_ALL_BEGIN`/`MPI_FILE_READ_ALL_END` pair on another process.
- Split collective routines must specify a buffer in both the begin and end routines. By specifying the buffer that receives data in the end routine, we can avoid many (though not all) of the problems described in “A Problem with Register Optimization,” Section 10.2.2, page 289.
- No collective I/O operations are permitted on a file handle concurrently with a split collective access on that file handle (i.e., between the begin and end of the access). That is

```

MPI_File_read_all_begin(fh, ...);
...
MPI_File_read_all(fh, ...);
...
MPI_File_read_all_end(fh, ...);

```

is erroneous.

- In a multithreaded implementation, any split collective begin and end operation called by a process must be called from the same thread. This restriction is made to simplify the implementation in the multithreaded case. (Note that we have already disallowed having two threads begin a split collective operation on the same file handle since only one split collective operation can be active on a file handle at any time.)

The arguments for these routines have the same meaning as for the equivalent collective versions (e.g., the argument definitions for `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END` are equivalent to the arguments for `MPI_FILE_READ_ALL`). The begin routine (e.g., `MPI_FILE_READ_ALL_BEGIN`) begins a split collective operation that, when completed with the matching end routine (i.e., `MPI_FILE_READ_ALL_END`) produces the result as defined for the equivalent collective routine (i.e., `MPI_FILE_READ_ALL`).

For the purpose of consistency semantics (Section 9.6.1, page 255), a matched pair of split collective data access operations (e.g., `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END`) compose a single data access.


```

MPI_FILE_WRITE_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET

void MPI::File::Write_at_all_begin(MPI::Offset offset, const void* buf,
    int count, const MPI::Datatype& datatype)

MPI_FILE_WRITE_AT_ALL_END(fh, buf, status)

    INOUT   fh                file handle (handle)
    IN      buf               initial address of buffer (choice)
    OUT     status            status object (Status)

int MPI_File_write_at_all_end(MPI_File fh, void *buf, MPI_Status *status)

MPI_FILE_WRITE_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR

void MPI::File::Write_at_all_end(const void* buf, MPI::Status& status)

void MPI::File::Write_at_all_end(const void* buf)

MPI_FILE_READ_ALL_BEGIN(fh, buf, count, datatype)

    INOUT   fh                file handle (handle)
    OUT     buf               initial address of buffer (choice)
    IN      count             number of elements in buffer (integer)
    IN      datatype          datatype of each buffer element (handle)

int MPI_File_read_all_begin(MPI_File fh, void *buf, int count,
    MPI_Datatype datatype)

MPI_FILE_READ_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR

void MPI::File::Read_all_begin(void* buf, int count,
    const MPI::Datatype& datatype)

```

```

1  MPI_FILE_READ_ALL_END(fh, buf, status)
2      INOUT   fh                file handle (handle)
3      OUT     buf                initial address of buffer (choice)
4
5      OUT     status             status object (Status)
6
7  int MPI_File_read_all_end(MPI_File fh, void *buf, MPI_Status *status)
8
9  MPI_FILE_READ_ALL_END(FH, BUF, STATUS, IERROR)
10     <type> BUF(*)
11     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
12
13 void MPI::File::Read_all_end(void* buf, MPI::Status& status)
14
15 void MPI::File::Read_all_end(void* buf)
16
17 MPI_FILE_WRITE_ALL_BEGIN(fh, buf, count, datatype)
18     INOUT   fh                file handle (handle)
19     IN       buf                initial address of buffer (choice)
20
21     IN       count              number of elements in buffer (integer)
22
23     IN       datatype            datatype of each buffer element (handle)
24
25 int MPI_File_write_all_begin(MPI_File fh, void *buf, int count,
26                               MPI_Datatype datatype)
27
28 MPI_FILE_WRITE_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
29     <type> BUF(*)
30     INTEGER FH, COUNT, DATATYPE, IERROR
31
32 void MPI::File::Write_all_begin(const void* buf, int count,
33                                 const MPI::Datatype& datatype)
34
35 MPI_FILE_WRITE_ALL_END(fh, buf, status)
36     INOUT   fh                file handle (handle)
37     IN       buf                initial address of buffer (choice)
38
39     OUT     status             status object (Status)
40
41 int MPI_File_write_all_end(MPI_File fh, void *buf, MPI_Status *status)
42
43 MPI_FILE_WRITE_ALL_END(FH, BUF, STATUS, IERROR)
44     <type> BUF(*)
45     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
46
47 void MPI::File::Write_all_end(const void* buf, MPI::Status& status)
48
49 void MPI::File::Write_all_end(const void* buf)

```

```

MPI_FILE_READ_ORDERED_BEGIN(fh, buf, count, datatype) 1
    INOUT fh file handle (handle) 2
    OUT buf initial address of buffer (choice) 3
    IN count number of elements in buffer (integer) 4
    IN datatype datatype of each buffer element (handle) 5
    6
    7
    8
int MPI_File_read_ordered_begin(MPI_File fh, void *buf, int count, 9
    MPI_Datatype datatype) 10
MPI_FILE_READ_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR) 11
    <type> BUF(*) 12
    INTEGER FH, COUNT, DATATYPE, IERROR 13
    14
void MPI::File::Read_ordered_begin(void* buf, int count, 15
    const MPI::Datatype& datatype) 16
    17
    18
MPI_FILE_READ_ORDERED_END(fh, buf, status) 19
    INOUT fh file handle (handle) 20
    OUT buf initial address of buffer (choice) 21
    OUT status status object (Status) 22
    23
    24
int MPI_File_read_ordered_end(MPI_File fh, void *buf, MPI_Status *status) 25
MPI_FILE_READ_ORDERED_END(FH, BUF, STATUS, IERROR) 26
    <type> BUF(*) 27
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR 28
    29
void MPI::File::Read_ordered_end(void* buf, MPI::Status& status) 30
    31
void MPI::File::Read_ordered_end(void* buf) 32
    33
    34
MPI_FILE_WRITE_ORDERED_BEGIN(fh, buf, count, datatype) 35
    INOUT fh file handle (handle) 36
    IN buf initial address of buffer (choice) 37
    IN count number of elements in buffer (integer) 38
    IN datatype datatype of each buffer element (handle) 39
    40
    41
    42
int MPI_File_write_ordered_begin(MPI_File fh, void *buf, int count, 43
    MPI_Datatype datatype) 44
MPI_FILE_WRITE_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR) 45
    <type> BUF(*) 46
    INTEGER FH, COUNT, DATATYPE, IERROR 47
    48

```

```

1 void MPI::File::Write_ordered_begin(const void* buf, int count,
2                                     const MPI::Datatype& datatype)
3
4
5 MPI_FILE_WRITE_ORDERED_END(fh, buf, status)
6
7     INOUT    fh                file handle (handle)
8
9     IN       buf              initial address of buffer (choice)
10
11     OUT      status           status object (Status)
12
13 int MPI_File_write_ordered_end(MPI_File fh, void *buf, MPI_Status *status)
14
15 MPI_FILE_WRITE_ORDERED_END(FH, BUF, STATUS, IERROR)
16     <type> BUF(*)
17     INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
18
19 void MPI::File::Write_ordered_end(const void* buf, MPI::Status& status)
20
21 void MPI::File::Write_ordered_end(const void* buf)

```

9.5 File Interoperability

At the most basic level, file interoperability is the ability to read the information previously written to a file—not just the bits of data, but the actual information the bits represent. MPI guarantees full interoperability within a single MPI environment, and supports increased interoperability outside that environment through the external data representation (Section 9.5.2, page 250) as well as the data conversion functions (Section 9.5.3, page 251).

Interoperability within a single MPI environment (which could be considered “operability”) ensures that file data written by one MPI process can be read by any other MPI process, subject to the consistency constraints (see Section 9.6.1, page 255), provided that it would have been possible to start the two processes simultaneously and have them reside in a single MPI_COMM_WORLD. Furthermore, both processes must see the same data values at every absolute byte offset in the file for which data was written.

This single environment file interoperability implies that file data is accessible regardless of the number of processes.

There are three aspects to file interoperability:

- transferring the bits,
- converting between different file structures, and
- converting between different machine representations.

The first two aspects of file interoperability are beyond the scope of this standard, as both are highly machine dependent. However, transferring the bits of a file into and out of the MPI environment (e.g., by writing a file to tape) is required to be supported by all MPI implementations. In particular, an implementation must specify how familiar operations similar to POSIX `cp`, `rm`, and `mv` can be performed on the file. Furthermore, it is expected that the facility provided maintains the correspondence between absolute byte

offsets (e.g., after possible file structure conversion, the data bits at byte offset 102 in the MPI environment are at byte offset 102 outside the MPI environment). As an example, a simple off-line conversion utility that transfers and converts files between the native file system and the MPI environment would suffice, provided it maintained the offset coherence mentioned above. In a high quality implementation of MPI, users will be able to manipulate MPI files using the same or similar tools that the native file system offers for manipulating its files.

The remaining aspect of file interoperability, converting between different machine representations, is supported by the typing information specified in the etype and filetype. This facility allows the information in files to be shared between any two applications, regardless of whether they use MPI, and regardless of the machine architectures on which they run.

MPI supports multiple data representations: “native,” “internal,” and “external32.” An implementation may support additional data representations. MPI also supports user-defined data representations (see Section 9.5.3, page 251). The native and internal data representations are implementation dependent, while the external32 representation is common to all MPI implementations and facilitates file interoperability. The data representation is specified in the *datarep* argument to `MPI_FILE_SET_VIEW`.

Advice to users. MPI is not guaranteed to retain knowledge of what data representation was used when a file is written. Therefore, to correctly retrieve file data, an MPI application is responsible for specifying the same data representation as was used to create the file. (*End of advice to users.*)

“native” Data in this representation is stored in a file exactly as it is in memory. The advantage of this data representation is that data precision and I/O performance are not lost in type conversions with a purely homogeneous environment. The disadvantage is the loss of transparent interoperability within a heterogeneous MPI environment.

Advice to users. This data representation should only be used in a homogeneous MPI environment, or when the MPI application is capable of performing the data type conversions itself. (*End of advice to users.*)

Advice to implementors. When implementing read and write operations on top of MPI message passing, the message data should be typed as `MPI_BYTE` to ensure that the message routines do not perform any type conversions on the data. (*End of advice to implementors.*)

“internal” This data representation can be used for I/O operations in a homogeneous or heterogeneous environment; the implementation will perform type conversions if necessary. The implementation is free to store data in any format of its choice, with the restriction that it will maintain constant extents for all predefined datatypes in any one file. The environment in which the resulting file can be reused is implementation-defined and must be documented by the implementation.

Rationale. This data representation allows the implementation to perform I/O efficiently in a heterogeneous environment, though with implementation-defined restrictions on how the file can be reused. (*End of rationale.*)

Advice to implementors. Since “external32” is a superset of the functionality provided by “internal,” an implementation may choose to implement “internal” as “external32.” (*End of advice to implementors.*)

“external32” This data representation states that read and write operations convert all data from and to the “external32” representation defined in Section 9.5.2, page 250. The data conversion rules for communication also apply to these conversions (see Section 3.3.2, page 25-27, of the MPI-1 document). The data on the storage medium is always in this canonical representation, and the data in memory is always in the local process’s native representation.

This data representation has several advantages. First, all processes reading the file in a heterogeneous MPI environment will automatically have the data converted to their respective native representations. Second, the file can be exported from one MPI environment and imported into any other MPI environment with the guarantee that the second environment will be able to read all the data in the file.

The disadvantage of this data representation is that data precision and I/O performance may be lost in data type conversions.

Advice to implementors. When implementing read and write operations on top of MPI message passing, the message data should be converted to and from the “external32” representation in the client, and sent as type `MPI_BYTE`. This will avoid possible double data type conversions and the associated further loss of precision and performance. (*End of advice to implementors.*)

9.5.1 Datatypes for File Interoperability

If the file data representation is other than “native,” care must be taken in constructing etypes and filetypes. Any of the datatype constructor functions may be used; however, for those functions that accept displacements in bytes, the displacements must be specified in terms of their values in the file for the file data representation being used. MPI will interpret these byte displacements as is; no scaling will be done. The function `MPI_FILE_GET_TYPE_EXTENT` can be used to calculate the extents of datatypes in the file. For etypes and filetypes that are portable datatypes (see Section 2.4, page 7), MPI will scale any displacements in the datatypes to match the file data representation. Datatypes passed as arguments to read/write routines specify the data layout in memory; therefore, they must always be constructed using displacements corresponding to displacements in memory.

Advice to users. One can logically think of the file as if it were stored in the memory of a file server. The **etype** and **filetype** are interpreted as if they were defined at this file server, by the same sequence of calls used to define them at the calling process. If the data representation is “native”, then this logical file server runs on the same architecture as the calling process, so that these types define the same data layout on the file as they would define in the memory of the calling process. If the **etype** and **filetype** are portable datatypes, then the data layout defined in the file is the same as would be defined in the calling process memory, up to a scaling factor. The routine `MPI_FILE_GET_FILE_EXTENT` can be used to calculate this scaling factor. Thus, two equivalent, portable datatypes will define the same data layout in the file,

even in a heterogeneous environment with “internal”, “external32”, or user defined data representations. Otherwise, the **etype** and **filetype** must be constructed so that their typemap and extent are the same on any architecture. This can be achieved if they have an explicit upper bound and lower bound (defined either using **MPI_LB** and **MPI_UB** markers, or using **MPI_TYPE_CREATE_RESIZED**). This condition must also be fulfilled by any datatype that is used in the construction of the **etype** and **filetype**, if this datatype is replicated contiguously, either explicitly, by a call to **MPI_TYPE_CONTIGUOUS**, or implicitly, by a blocklength argument that is greater than one. If an **etype** or **filetype** is not portable, and has a typemap or extent that is architecture dependent, then the data layout specified by it on a file is implementation dependent.

File data representations other than “native” may be different from corresponding data representations in memory. Therefore, for these file data representations, it is important not to use hardwired byte offsets for file positioning, including the initial displacement that specifies the view. When a portable datatype (see Section 2.4, page 7) is used in a data access operation, any holes in the datatype are scaled to match the data representation. However, note that this technique only works when all the processes that created the file view build their etypes from the same predefined datatypes. For example, if one process uses an etype built from **MPI_INT** and another uses an etype built from **MPI_FLOAT**, the resulting views may be nonportable because the relative sizes of these types may differ from one data representation to another. (*End of advice to users.*)

MPI_FILE_GET_TYPE_EXTENT(fh, datatype, extent)

IN	fh	file handle (handle)
IN	datatype	datatype (handle)
OUT	extent	datatype extent (integer)

```
int MPI_File_get_type_extent(MPI_File fh, MPI_Datatype datatype,
                             MPI_Aint *extent)
```

```
MPI_FILE_GET_TYPE_EXTENT(FH, DATATYPE, EXTENT, IERROR)
    INTEGER FH, DATATYPE, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT
```

```
MPI::Aint MPI::File::Get_type_extent(const MPI::Datatype& datatype) const
```

Returns the extent of **datatype** in the file **fh**. This extent will be the same for all processes accessing the file **fh**. If the current view uses a user-defined data representation (see Section 9.5.3, page 251), MPI uses the **dtype_file_extent_fn** callback to calculate the extent.

Advice to implementors. In the case of user-defined data representations, the extent of a derived datatype can be calculated by first determining the extents of the predefined datatypes in this derived datatype using **dtype_file_extent_fn** (see Section 9.5.3, page 251). (*End of advice to implementors.*)

9.5.2 External Data Representation: “external32”

All MPI implementations are required to support the data representation defined in this section. Support of optional datatypes (e.g., `MPI_INTEGER2`) is not required.

All floating point values are in big-endian IEEE format [9] of the appropriate size. Floating point values are represented by one of three IEEE formats. These are the IEEE “Single,” “Double,” and “Double Extended” formats, requiring 4, 8 and 16 bytes of storage, respectively. For the IEEE “Double Extended” formats, MPI specifies a Format Width of 16 bytes, with 15 exponent bits, bias = +10383, 112 fraction bits, and an encoding analogous to the “Double” format. All integral values are in two’s complement big-endian format. Big-endian means most significant byte at lowest address byte. For Fortran `LOGICAL` and C++ `bool`, 0 implies false and nonzero implies true. Fortran `COMPLEX` and `DOUBLE COMPLEX` are represented by a pair of floating point format values for the real and imaginary components. Characters are in ISO 8859-1 format [10]. Wide characters (of type `MPI_WCHAR`) are in Unicode format [23].

All signed numerals (e.g., `MPI_INT`, `MPI_REAL`) have the sign bit at the most significant bit. `MPI_COMPLEX` and `MPI_DOUBLE_COMPLEX` have the sign bit of the real and imaginary parts at the most significant bit of each part.

According to IEEE specifications [9], the “NaN” (not a number) is system dependent. It should not be interpreted within MPI as anything other than “NaN.”

Advice to implementors. The MPI treatment of “NaN” is similar to the approach used in XDR (see <ftp://ds.internic.net/rfc/rfc1832.txt>). (*End of advice to implementors.*)

All data is byte aligned, regardless of type. All data items are stored contiguously in the file.

Advice to implementors. All bytes of `LOGICAL` and `bool` must be checked to determine the value. (*End of advice to implementors.*)

Advice to users. The type `MPI_PACKED` is treated as bytes and is not converted. The user should be aware that `MPI_PACK` has the option of placing a header in the beginning of the pack buffer. (*End of advice to users.*)

Type	Length
-----	-----
<code>MPI_PACKED</code>	1
<code>MPI_BYTE</code>	1
<code>MPI_CHAR</code>	1
<code>MPI_UNSIGNED_CHAR</code>	1
<code>MPI_SIGNED_CHAR</code>	1
<code>MPI_WCHAR</code>	2
<code>MPI_SHORT</code>	2
<code>MPI_UNSIGNED_SHORT</code>	2
<code>MPI_INT</code>	4
<code>MPI_UNSIGNED</code>	4
<code>MPI_LONG</code>	4
<code>MPI_UNSIGNED_LONG</code>	4
<code>MPI_FLOAT</code>	4

MPI_DOUBLE	8	1
MPI_LONG_DOUBLE	16	2
		3
MPI_CHARACTER	1	4
MPI_LOGICAL	4	5
MPI_INTEGER	4	6
MPI_REAL	4	7
MPI_DOUBLE_PRECISION	8	8
MPI_COMPLEX	2*4	9
MPI_DOUBLE_COMPLEX	2*8	10
		11
Optional Type	Length	12
-----	-----	13
MPI_INTEGER1	1	14
MPI_INTEGER2	2	15
MPI_INTEGER4	4	16
MPI_INTEGER8	8	17
MPI_LONG_LONG	8	18
MPI_UNSIGNED_LONG_LONG	8	19
		20
MPI_REAL4	4	21
MPI_REAL8	8	22
MPI_REAL16	16	23
		24

The size of the predefined datatypes returned from `MPI_TYPE_CREATE_F90_REAL`, `MPI_TYPE_CREATE_F90_COMPLEX`, and `MPI_TYPE_CREATE_F90_INTEGER` are defined in Section 10.2.5, page 296.

Advice to implementors. When converting a larger size integer to a smaller size integer, only the less significant bytes are moved. Care must be taken to preserve the sign bit value. This allows no conversion errors if the data range is within the range of the smaller size integer. (*End of advice to implementors.*)

9.5.3 User-Defined Data Representations

There are two situations that cannot be handled by the required representations:

1. a user wants to write a file in a representation unknown to the implementation, and
2. a user wants to read a file written in a representation unknown to the implementation.

User-defined data representations allow the user to insert a third party converter into the I/O stream to do the data representation conversion.

```

1  MPI_REGISTER_DATAREP(datarep, read_conversion_fn, write_conversion_fn,
2  dtype_file_extent_fn, extra_state)
3
4      IN      datarep          data representation identifier (string)
5
6      IN      read_conversion_fn  function invoked to convert from file representation to
7                                  native representation (function)
8
9      IN      write_conversion_fn function invoked to convert from native representation
10                                 to file representation (function)
11
12     IN      dtype_file_extent_fn function invoked to get the extent of a datatype as
13                                 represented in the file (function)
14
15     IN      extra_state        extra state
16
17 int MPI_Register_datarep(char *datarep,
18     MPI_Datarep_conversion_function *read_conversion_fn,
19     MPI_Datarep_conversion_function *write_conversion_fn,
20     MPI_Datarep_extent_function *dtype_file_extent_fn,
21     void *extra_state)
22
23 MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
24     DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
25
26 CHARACTER*(*) DATAREP
27 EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
28 INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
29 INTEGER IERROR
30
31 void MPI::Register_datarep(const char* datarep,
32     MPI::Datarep_conversion_function* read_conversion_fn,
33     MPI::Datarep_conversion_function* write_conversion_fn,
34     MPI::Datarep_extent_function* dtype_file_extent_fn,
35     void* extra_state)
36
37
38
39
40
41
42
43
44
45
46
47
48

```

The call associates `read_conversion_fn`, `write_conversion_fn`, and `dtype_file_extent_fn` with the data representation identifier `datarep`. `datarep` can then be used as an argument to `MPI_FILE_SET_VIEW`, causing subsequent data access operations to call the conversion functions to convert all data items accessed between file data representation and native representation. `MPI_REGISTER_DATAREP` is a local operation and only registers the data representation for the calling MPI process. If `datarep` is already defined, an error in the error class `MPI_ERR_DUP_DATAREP` is raised using the default file error handler (see Section 9.7, page 265). The length of a data representation string is limited to the value of `MPI_MAX_DATAREP_STRING`. `MPI_MAX_DATAREP_STRING` must have a value of at least 64. No routines are provided to delete data representations and free the associated resources; it is not expected that an application will generate them in significant numbers.

Extent Callback

```

45 typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
46     MPI_Aint *file_extent, void *extra_state);
47
48 SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)

```

```

INTEGER DATATYPE, IERROR
INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE

typedef MPI::Datarep_extent_function(const MPI::Datatype& datatype,
                                     MPI::Aint& file_extent, void* extra_state);

```

The function `dtype_file_extent_fn` must return, in `file_extent`, the number of bytes required to store `datatype` in the file representation. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. MPI will only call this routine with predefined datatypes employed by the user.

Datarep Conversion Functions

```

typedef int MPI_Datarep_conversion_function(void *userbuf,
                                           MPI_Datatype datatype, int count, void *filebuf,
                                           MPI_Offset position, void *extra_state);

SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
                                       POSITION, EXTRA_STATE, IERROR)
  <TYPE> USERBUF(*), FILEBUF(*)
  INTEGER COUNT, DATATYPE, IERROR
  INTEGER(KIND=MPI_OFFSET_KIND) POSITION
  INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE

typedef MPI::Datarep_conversion_function(void* userbuf,
                                         MPI::Datatype& datatype, int count, void* filebuf,
                                         MPI::Offset position, void* extra_state);

```

The function `read_conversion_fn` must convert from file data representation to native representation. Before calling this routine, MPI allocates and fills `filebuf` with `count` contiguous data items. The type of each data item matches the corresponding entry for the predefined datatype in the type signature of `datatype`. The function is passed, in `extra_state`, the argument that was passed to the `MPI_REGISTER_DATAREP` call. The function must copy all `count` data items from `filebuf` to `userbuf` in the distribution described by `datatype`, converting each data item from file representation to native representation. `datatype` will be equivalent to the datatype that the user passed to the read or write function. If the size of `datatype` is less than the size of the `count` data items, the conversion function must treat `datatype` as being contiguously tiled over the `userbuf`. The conversion function must begin storing converted data at the location in `userbuf` specified by `position` into the (tiled) `datatype`.

Advice to users. Although the conversion functions have similarities to `MPI_PACK` and `MPI_UNPACK` in MPI-1, one should note the differences in the use of the arguments `count` and `position`. In the conversion functions, `count` is a count of data items (i.e., count of typemap entries of `datatype`), and `position` is an index into this typemap. In `MPI_PACK`, `incount` refers to the number of whole datatypes, and `position` is a number of bytes. (*End of advice to users.*)

Advice to implementors. A converted read operation could be implemented as follows:

1. Get file extent of all data items

2. Allocate a filebuf large enough to hold all count data items
3. Read data from file into filebuf
4. Call `read_conversion_fn` to convert data and place it into userbuf
5. Deallocate filebuf

(End of advice to implementors.)

If MPI cannot allocate a buffer large enough to hold all the data to be converted from a read operation, it may call the conversion function repeatedly using the same **datatype** and **userbuf**, and reading successive chunks of data to be converted in **filebuf**. For the first call (and in the case when all the data to be converted fits into **filebuf**), MPI will call the function with **position** set to zero. Data converted during this call will be stored in the **userbuf** according to the first **count** data items in **datatype**. Then in subsequent calls to the conversion function, MPI will increment the value in **position** by the **count** of items converted in the previous call.

Rationale. Passing the conversion function a position and one datatype for the transfer allows the conversion function to decode the datatype only once and cache an internal representation of it on the datatype. Then on subsequent calls, the conversion function can use the **position** to quickly find its place in the datatype and continue storing converted data where it left off at the end of the previous call. *(End of rationale.)*

Advice to users. Although the conversion function may usefully cache an internal representation on the datatype, it should not cache any state information specific to an ongoing conversion operation, since it is possible for the same datatype to be used concurrently in multiple conversion operations. *(End of advice to users.)*

The function `write_conversion_fn` must convert from native representation to file data representation. Before calling this routine, MPI allocates **filebuf** of a size large enough to hold **count** contiguous data items. The type of each data item matches the corresponding entry for the predefined datatype in the type signature of **datatype**. The function must copy **count** data items from **userbuf** in the distribution described by **datatype**, to a contiguous distribution in **filebuf**, converting each data item from native representation to file representation. If the size of **datatype** is less than the size of **count** data items, the conversion function must treat **datatype** as being contiguously tiled over the **userbuf**.

The function must begin copying at the location in **userbuf** specified by **position** into the (tiled) **datatype**. **datatype** will be equivalent to the datatype that the user passed to the read or write function. The function is passed, in **extra_state**, the argument that was passed to the `MPI_REGISTER_DATAREP` call.

The predefined constant `MPI_CONVERSION_FN_NULL` may be used as either `write_conversion_fn` or `read_conversion_fn`. In that case, MPI will not attempt to invoke `write_conversion_fn` or `read_conversion_fn`, respectively, but will perform the requested data access using the native data representation.

An MPI implementation must ensure that all data accessed is converted, either by using a filebuf large enough to hold all the requested data items or else by making repeated

calls to the conversion function with the same **datatype** argument and appropriate values for **position**.

An implementation will only invoke the callback routines in this section (**read_conversion_fn**, **write_conversion_fn**, and **dtype_file_extent_fn**) when one of the read or write routines in Section 9.4, page 223, or **MPI_FILE_GET_TYPE_EXTENT** is called by the user. **dtype_file_extent_fn** will only be passed predefined datatypes employed by the user. The conversion functions will only be passed datatypes equivalent to those that the user has passed to one of the routines noted above.

The conversion functions must be reentrant. User defined data representations are restricted to use byte alignment for all types. Furthermore, it is erroneous for the conversion functions to call any collective routines or to free **datatype**.

The conversion functions should return an error code. If the returned error code has a value other than **MPI_SUCCESS**, the implementation will raise an error in the class **MPI_ERR_CONVERSION**.

9.5.4 Matching Data Representations

It is the user's responsibility to ensure that the data representation used to read data from a file is *compatible* with the data representation that was used to write that data to the file.

In general, using the same data representation name when writing and reading a file does not guarantee that the representation is compatible. Similarly, using different representation names on two different implementations may yield compatible representations.

Compatibility can be obtained when "external32" representation is used, although precision may be lost and the performance may be less than when "native" representation is used. Compatibility is guaranteed using "external32" provided at least one of the following conditions is met.

- The data access routines directly use types enumerated in Section 9.5.2, page 250, that are supported by all implementations participating in the I/O. The predefined type used to write a data item must also be used to read a data item.
- In the case of Fortran 90 programs, the programs participating in the data accesses obtain compatible datatypes using MPI routines that specify precision and/or range (Section 10.2.5, page 292).
- For any given data item, the programs participating in the data accesses use compatible predefined types to write and read the data item.

User-defined data representations may be used to provide an implementation compatibility with another implementation's "native" or "internal" representation.

Advice to users. Section 10.2.5, page 292, defines routines that support the use of matching datatypes in heterogeneous environments and contains examples illustrating their use. (*End of advice to users.*)

9.6 Consistency and Semantics

9.6.1 File Consistency

Consistency semantics define the outcome of multiple accesses to a single file. All file accesses in MPI are relative to a specific file handle created from a collective open. MPI

provides three levels of consistency: sequential consistency among all accesses using a single file handle, sequential consistency among all accesses using file handles created from a single collective open with atomic mode enabled, and user-imposed consistency among accesses other than the above. Sequential consistency means the behavior of a set of operations will be as if the operations were performed in some serial order consistent with program order; each access appears atomic, although the exact ordering of accesses is unspecified. User-imposed consistency may be obtained using program order and calls to `MPI_FILE_SYNC`.

Let FH_1 be the set of file handles created from one particular collective open of the file FOO , and FH_2 be the set of file handles created from a different collective open of FOO . Note that nothing restrictive is said about FH_1 and FH_2 : the sizes of FH_1 and FH_2 may be different, the groups of processes used for each open may or may not intersect, the file handles in FH_1 may be destroyed before those in FH_2 are created, etc. Consider the following three cases: a single file handle (e.g., $fh_1 \in FH_1$), two file handles created from a single collective open (e.g., $fh_{1a} \in FH_1$ and $fh_{1b} \in FH_1$), and two file handles from different collective opens (e.g., $fh_1 \in FH_1$ and $fh_2 \in FH_2$).

For the purpose of consistency semantics, a matched pair (Section 9.4.5, page 240) of split collective data access operations (e.g., `MPI_FILE_READ_ALL_BEGIN` and `MPI_FILE_READ_ALL_END`) compose a single data access operation. Similarly, a nonblocking data access routine (e.g., `MPI_FILE_READ`) and the routine which completes the request (e.g., `MPI_WAIT`) also compose a single data access operation. For all cases below, these data access operations are subject to the same constraints as blocking data access operations.

Advice to users. For an `MPI_FILE_READ` and `MPI_WAIT` pair, the operation begins when `MPI_FILE_READ` is called and ends when `MPI_WAIT` returns. (*End of advice to users.*)

Assume that A_1 and A_2 are two data access operations. Let D_1 (D_2) be the set of absolute byte displacements of every byte accessed in A_1 (A_2). The two data accesses *overlap* if $D_1 \cap D_2 \neq \emptyset$. The two data accesses *conflict* if they overlap and at least one is a write access.

Let SEQ_{fh} be a sequence of file operations on a single file handle, bracketed by `MPI_FILE_SYNC`s on that file handle. (Both opening and closing a file implicitly perform an `MPI_FILE_SYNC`.) SEQ_{fh} is a “write sequence” if any of the data access operations in the sequence are writes or if any of the file manipulation operations in the sequence change the state of the file (e.g., `MPI_FILE_SET_SIZE` or `MPI_FILE_PREALLOCATE`). Given two sequences, SEQ_1 and SEQ_2 , we say they are not *concurrent* if one sequence is guaranteed to completely precede the other (temporally).

The requirements for guaranteeing sequential consistency among all accesses to a particular file are divided into the three cases given below. If any of these requirements are not met, then the value of all data in that file is implementation dependent.

Case 1: $fh_1 \in FH_1$ All operations on fh_1 are sequentially consistent if atomic mode is set. If nonatomic mode is set, then all operations on fh_1 are sequentially consistent if they are either nonconcurrent, nonconflicting, or both.

Case 2: $fh_{1a} \in FH_1$ and $fh_{1b} \in FH_1$ Assume A_1 is a data access operation using fh_{1a} , and A_2 is a data access operation using fh_{1b} . If for any access A_1 , there is no access A_2 that conflicts with A_1 , then MPI guarantees sequential consistency.

However, unlike POSIX semantics, the default MPI semantics for conflicting accesses do not guarantee sequential consistency. If A_1 and A_2 conflict, sequential consistency can be guaranteed by either enabling atomic mode via the `MPI_FILE_SET_ATOMICITY` routine, or meeting the condition described in Case 3 below.

Case 3: $fh_1 \in FH_1$ and $fh_2 \in FH_2$ Consider access to a single file using file handles from distinct collective opens. In order to guarantee sequential consistency, `MPI_FILE_SYNC` must be used (both opening and closing a file implicitly perform an `MPI_FILE_SYNC`).

Sequential consistency is guaranteed among accesses to a single file if for any write sequence SEQ_1 to the file, there is no sequence SEQ_2 to the file which is *concurrent* with SEQ_1 . To guarantee sequential consistency when there are write sequences, `MPI_FILE_SYNC` must be used together with a mechanism that guarantees nonconcurrency of the sequences.

See the examples in Section 9.6.10, page 261, for further clarification of some of these consistency semantics.

`MPI_FILE_SET_ATOMICITY(fh, flag)`

INOUT	<code>fh</code>	file handle (handle)
IN	<code>flag</code>	<code>true</code> to set atomic mode, <code>false</code> to set nonatomic mode (logical)

```
int MPI_File_set_atomicity(MPIFile fh, int flag)
```

```
MPI_FILE_SET_ATOMICITY(FH, FLAG, IERROR)
```

```
    INTEGER FH, IERROR
```

```
    LOGICAL FLAG
```

```
void MPI::File::Set_atomicity(bool flag)
```

Let FH be the set of file handles created by one collective open. The consistency semantics for data access operations using FH is set by collectively calling `MPI_FILE_SET_ATOMICITY` on FH . `MPI_FILE_SET_ATOMICITY` is collective; all processes in the group must pass identical values for `fh` and `flag`. If `flag` is `true`, atomic mode is set; if `flag` is `false`, nonatomic mode is set.

Changing the consistency semantics for an open file only affects new data accesses. All completed data accesses are guaranteed to abide by the consistency semantics in effect during their execution. Nonblocking data accesses and split collective operations that have not completed (e.g., via `MPI_WAIT`) are only guaranteed to abide by nonatomic mode consistency semantics.

Advice to implementors. Since the semantics guaranteed by atomic mode are stronger than those guaranteed by nonatomic mode, an implementation is free to adhere to the more stringent atomic mode semantics for outstanding requests. (*End of advice to implementors.*)

1 `MPI_FILE_GET_ATOMICITY(fh, flag)`

2 IN **fh** file handle (handle)

3 OUT **flag** true if atomic mode, false if nonatomic mode (logical)

5
6 `int MPI_File_get_atomicity(MPI_File fh, int *flag)`

7 `MPI_FILE_GET_ATOMICITY(FH, FLAG, IERROR)`

8 INTEGER FH, IERROR

9 LOGICAL FLAG

10
11 `bool MPI::File::Get_atomicity() const`

12 **MPI_FILE_GET_ATOMICITY** returns the current consistency semantics for data access
13 operations on the set of file handles created by one collective open. If **flag** is true, atomic
14 mode is enabled; if **flag** is false, nonatomic mode is enabled.

16
17 `MPI_FILE_SYNC(fh)`

18 INOUT **fh** file handle (handle)

20
21 `int MPI_File_sync(MPI_File fh)`

22 `MPI_FILE_SYNC(FH, IERROR)`

23 INTEGER FH, IERROR

24
25 `void MPI::File::Sync()`

26 Calling **MPI_FILE_SYNC** with **fh** causes all previous writes to **fh** by the calling process
27 to be transferred to the storage device. If other processes have made updates to the storage
28 device, then all such updates become visible to subsequent reads of **fh** by the calling process.
29 **MPI_FILE_SYNC** may be necessary to ensure sequential consistency in certain cases (see
30 above).

31 **MPI_FILE_SYNC** is a collective operation.

32 The user is responsible for ensuring that all nonblocking requests and split collective
33 operations on **fh** have been completed before calling **MPI_FILE_SYNC**—otherwise, the call
34 to **MPI_FILE_SYNC** is erroneous.

36 9.6.2 Random Access vs. Sequential Files

37
38 **MPI** distinguishes ordinary random access files from sequential stream files, such as pipes
39 and tape files. Sequential stream files must be opened with the **MPI_MODE_SEQUENTIAL**
40 flag set in the amode. For these files, the only permitted data access operations are shared
41 file pointer reads and writes. Filetypes and etypes with holes are erroneous. In addition,
42 the notion of file pointer is not meaningful; therefore, calls to **MPI_FILE_SEEK_SHARED** and
43 **MPI_FILE_GET_POSITION_SHARED** are erroneous, and the pointer update rules specified
44 for the data access routines do not apply. The amount of data accessed by a data access
45 operation will be the amount requested unless the end of file is reached or an error is raised.

46
47 *Rationale.* This implies that reading on a pipe will always wait until the requested
48 amount of data is available or until the process writing to the pipe has issued an end

of file. (*End of rationale.*)

Finally, for some sequential files, such as those corresponding to magnetic tapes or streaming network connections, writes to the file may be destructive. In other words, a write may act as a truncate (a `MPI_FILE_SET_SIZE` with `size` set to the current position) followed by the write.

9.6.3 Progress

The progress rules of `MPI` are both a promise to users and a set of constraints on implementors. In cases where the progress rules restrict possible implementation choices more than the interface specification alone, the progress rules take precedence.

All blocking routines must complete in finite time unless an exceptional condition (such as resource exhaustion) causes an error.

Nonblocking data access routines inherit the following progress rule from nonblocking point to point communication: a nonblocking write is equivalent to a nonblocking send for which a receive is eventually posted, and a nonblocking read is equivalent to a nonblocking receive for which a send is eventually posted.

Finally, an implementation is free to delay progress of collective routines until all processes in the group associated with the collective call have invoked the routine. Once all processes in the group have invoked the routine, the progress rule of the equivalent noncollective routine must be followed.

9.6.4 Collective File Operations

Collective file operations are subject to the same restrictions as collective communication operations. For a complete discussion, please refer to the semantics set forth in `MPI-1` [6], Section 4.12.

Collective file operations are collective over a dup of the communicator used to open the file—this duplicate communicator is implicitly specified via the file handle argument. Different processes can pass different values for other arguments of a collective routine unless specified otherwise.

9.6.5 Type Matching

The type matching rules for I/O mimic the type matching rules for communication with one exception: if `etype` is `MPI_BYTE`, then this matches any `datatype` in a data access operation. In general, the `etype` of data items written must match the `etype` used to read the items, and for each data access operation, the current `etype` must also match the type declaration of the data access buffer.

Advice to users. In most cases, use of `MPI_BYTE` as a wild card will defeat the file interoperability features of `MPI`. File interoperability can only perform automatic conversion between heterogeneous data representations when the exact datatypes accessed are explicitly specified. (*End of advice to users.*)

9.6.6 Miscellaneous Clarifications

Once an I/O routine completes, it is safe to free any opaque objects passed as arguments to that routine. For example, the `comm` and `info` used in an `MPI_FILE_OPEN`, or the `etype`

and `filetype` used in an `MPI_FILE_SET_VIEW`, can be freed without affecting access to the file. Note that for nonblocking routines and split collective operations, the operation must be completed before it is safe to reuse data buffers passed as arguments.

As in communication, datatypes must be committed before they can be used in file manipulation or data access operations. For example, the `etype` and `filetype` must be committed before calling `MPI_FILE_SET_VIEW`, and the `datatype` must be committed before calling `MPI_FILE_READ` or `MPI_FILE_WRITE`.

9.6.7 MPI_Offset Type

`MPI_Offset` is an integer type of size sufficient to represent the size (in bytes) of the largest file supported by MPI. Displacements and offsets are always specified as values of type `MPI_Offset`.

In Fortran, the corresponding integer is an integer of kind `MPI_OFFSET_KIND`, defined in `mpif.h` and the `mpi` module.

In Fortran 77 environments that do not support `KIND` parameters, `MPI_Offset` arguments should be declared as an `INTEGER` of suitable size. The language interoperability implications for `MPI_Offset` are similar to those for addresses (see Section 4.12, page 49).

9.6.8 Logical vs. Physical File Layout

MPI specifies how the data should be laid out in a virtual file structure (the view), not how that file structure is to be stored on one or more disks. Specification of the physical file structure was avoided because it is expected that the mapping of files to disks will be system specific, and any specific control over file layout would therefore restrict program portability. However, there are still cases where some information may be necessary to optimize file layout. This information can be provided as *hints* specified via *info* when a file is created (see Section 9.2.8, page 218).

9.6.9 File Size

The size of a file may be increased by writing to the file after the current end of file. The size may also be changed by calling MPI *size changing* routines, such as `MPI_FILE_SET_SIZE`. A call to a size changing routine does not necessarily change the file size. For example, calling `MPI_FILE_PREALLOCATE` with a size less than the current size does not change the size.

Consider a set of bytes that has been written to a file since the most recent call to a size changing routine, or since `MPI_FILE_OPEN` if no such routine has been called. Let the *high byte* be the byte in that set with the largest displacement. The file size is the larger of

- One plus the displacement of the high byte.
- The size immediately after the size changing routine, or `MPI_FILE_OPEN`, returned.

When applying consistency semantics, calls to `MPI_FILE_SET_SIZE` and `MPI_FILE_PREALLOCATE` are considered writes to the file (which conflict with operations that access bytes at displacements between the old and new file sizes), and `MPI_FILE_GET_SIZE` is considered a read of the file (which overlaps with all accesses to the file).

Advice to users. Any sequence of operations containing the collective routines `MPI_FILE_SET_SIZE` and `MPI_FILE_PREALLOCATE` is a write sequence. As such, sequential consistency in nonatomic mode is not guaranteed unless the conditions in Section 9.6.1, page 255, are satisfied. (*End of advice to users.*)

File pointer update semantics (i.e., file pointers are updated by the amount accessed) are only guaranteed if file size changes are sequentially consistent.

Advice to users. Consider the following example. Given two operations made by separate processes to a file containing 100 bytes: an `MPI_FILE_READ` of 10 bytes and an `MPI_FILE_SET_SIZE` to 0 bytes. If the user does not enforce sequential consistency between these two operations, the file pointer may be updated by the amount requested (10 bytes) even if the amount accessed is zero bytes. (*End of advice to users.*)

9.6.10 Examples

The examples in this section illustrate the application of the MPI consistency and semantics guarantees. These address

- conflicting accesses on file handles obtained from a single collective open, and
- all accesses on file handles obtained from two separate collective opens.

The simplest way to achieve consistency for conflicting accesses is to obtain sequential consistency by setting atomic mode. For the code below, process 1 will read either 0 or 10 integers. If the latter, every element of `b` will be 5. If nonatomic mode is set, the results of the read are undefined.

```
/* Process 0 */
int i, a[10] ;
int TRUE = 1;

for ( i=0;i<10;i++)
    a[i] = 5 ;

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_set_atomicity( fh0, TRUE ) ;
MPI_File_write_at(fh0, 0, a, 10, MPI_INT, &status) ;
/* MPI_Barrier( MPI_COMM_WORLD ) ; */

/* Process 1 */
int b[10] ;
int TRUE = 1;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_set_atomicity( fh1, TRUE ) ;
/* MPI_Barrier( MPI_COMM_WORLD ) ; */
MPI_File_read_at(fh1, 0, b, 10, MPI_INT, &status) ;
```

A user may guarantee that the write on process 0 precedes the read on process 1 by imposing temporal order with, for example, calls to `MPI_BARRIER`.

Advice to users. Routines other than `MPI_BARRIER` may be used to impose temporal order. In the example above, process 0 could use `MPI_SEND` to send a 0 byte message, received by process 1 using `MPI_RECV`. (*End of advice to users.*)

Alternatively, a user can impose consistency with nonatomic mode set:

```

/* Process 0 */
int i, a[10] ;
for ( i=0;i<10;i++)
    a[i] = 5 ;

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_write_at(fh0, 0, a, 10, MPI_INT, &status ) ;
MPI_File_sync( fh0 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh0 ) ;

/* Process 1 */
int b[10] ;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_sync( fh1 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh1 ) ;
MPI_File_read_at(fh1, 0, b, 10, MPI_INT, &status ) ;

```

The “sync-barrier-sync” construct is required because:

- The barrier ensures that the write on process 0 occurs before the read on process 1.
- The first sync guarantees that the data written by all processes is transferred to the storage device.
- The second sync guarantees that all data which has been transferred to the storage device is visible to all processes. (This does not affect process 0 in this example.)

The following program represents an erroneous attempt to achieve consistency by eliminating the apparently superfluous second “sync” call for each process.

```

/* ----- THIS EXAMPLE IS ERRONEOUS ----- */
/* Process 0 */
int i, a[10] ;
for ( i=0;i<10;i++)
    a[i] = 5 ;

```

```

MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh0 ) ;
MPI_File_set_view( fh0, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_File_write_at( fh0, 0, a, 10, MPI_INT, &status ) ;
MPI_File_sync( fh0 ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;

/* Process 1 */
int b[10] ;
MPI_File_open( MPI_COMM_WORLD, "workfile",
               MPI_MODE_RDWR | MPI_MODE_CREATE, MPI_INFO_NULL, &fh1 ) ;
MPI_File_set_view( fh1, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
MPI_Barrier( MPI_COMM_WORLD ) ;
MPI_File_sync( fh1 ) ;
MPI_File_read_at( fh1, 0, b, 10, MPI_INT, &status ) ;

/* ----- THIS EXAMPLE IS ERRONEOUS ----- */

```

The above program also violates the MPI rule against out-of-order collective operations and will deadlock for implementations in which MPI_FILE_SYNC blocks.

Advice to users. Some implementations may choose to implement MPI_FILE_SYNC as a temporally synchronizing function. When using such an implementation, the “sync-barrier-sync” construct above can be replaced by a single “sync.” The results of using such code with an implementation for which MPI_FILE_SYNC is not temporally synchronizing is undefined. (*End of advice to users.*)

Asynchronous I/O

The behavior of asynchronous I/O operations is determined by applying the rules specified above for synchronous I/O operations.

The following examples all access a preexisting file “myfile.” Word 10 in myfile initially contains the integer 2. Each example writes and reads word 10.

First consider the following code fragment:

```

int a = 4, b, TRUE=1;
MPI_File_open( MPI_COMM_WORLD, "myfile",
               MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
/* MPI_File_set_atomicity( fh, TRUE ) ; Use this to set atomic mode. */
MPI_File_irewrite_at( fh, 10, &a, 1, MPI_INT, &reqs[0] ) ;
MPI_File_iread_at( fh, 10, &b, 1, MPI_INT, &reqs[1] ) ;
MPI_Waitall(2, reqs, statuses) ;

```

For asynchronous data access operations, MPI specifies that the access occurs at any time between the call to the asynchronous data access routine and the return from the corresponding request complete routine. Thus, executing either the read before the write, or the write before the read is consistent with program order. If atomic mode is set, then MPI guarantees sequential consistency, and the program will read either 2 or 4 into b. If atomic

mode is not set, then sequential consistency is not guaranteed and the program may read something other than 2 or 4 due to the conflicting data access.

Similarly, the following code fragment does not order file accesses:

```

1  int a = 4, b;
2  MPI_File_open( MPI_COMM_WORLD, "myfile",
3                  MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
4
5  MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
6
7  /* MPI_File_set_atomicity( fh, TRUE ) ; Use this to set atomic mode. */
8
9  MPI_File_iread_at(fh, 10, &a, 1, MPI_INT, &reqs[0]) ;
10 MPI_File_iread_at(fh, 10, &b, 1, MPI_INT, &reqs[1]) ;
11 MPI_Wait(&reqs[0], &status) ;
12 MPI_Wait(&reqs[1], &status) ;
13

```

If atomic mode is set, either 2 or 4 will be read into b. Again, MPI does not guarantee sequential consistency in nonatomic mode.

On the other hand, the following code fragment:

```

14 int a = 4, b;
15 MPI_File_open( MPI_COMM_WORLD, "myfile",
16                 MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
17
18 MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
19
20 MPI_File_iread_at(fh, 10, &a, 1, MPI_INT, &reqs[0]) ;
21 MPI_Wait(&reqs[0], &status) ;
22 MPI_File_iread_at(fh, 10, &b, 1, MPI_INT, &reqs[1]) ;
23 MPI_Wait(&reqs[1], &status) ;
24
25

```

defines the same ordering as:

```

26 int a = 4, b;
27 MPI_File_open( MPI_COMM_WORLD, "myfile",
28                 MPI_MODE_RDWR, MPI_INFO_NULL, &fh ) ;
29
30 MPI_File_set_view( fh, 0, MPI_INT, MPI_INT, "native", MPI_INFO_NULL ) ;
31 MPI_File_write_at(fh, 10, &a, 1, MPI_INT, &status) ;
32 MPI_File_read_at(fh, 10, &b, 1, MPI_INT, &status) ;
33
34

```

Since

- nonconcurrent operations on a single file handle are sequentially consistent, and
- the program fragments specify an order for the operations,

MPI guarantees that both program fragments will read the value 4 into b. There is no need to set atomic mode for this example.

Similar considerations apply to conflicting accesses of the form:

```

41 MPI_File_write_all_begin(fh,...) ;
42 MPI_File_iread(fh,...) ;
43 MPI_Wait(fh,...) ;
44 MPI_File_write_all_end(fh,...) ;
45
46

```

Recall that constraints governing consistency and semantics are not relevant to the following:

```

MPI_File_write_all_begin(fh,...) ;
MPI_File_read_all_begin(fh,...) ;
MPI_File_read_all_end(fh,...) ;
MPI_File_write_all_end(fh,...) ;

```

since split collective operations on the same file handle may not overlap (see Section 9.4.5, page 240).

9.7 I/O Error Handling

By default, communication errors are fatal—`MPI_ERRORS_ARE_FATAL` is the default error handler associated with `MPI_COMM_WORLD`. I/O errors are usually less catastrophic (e.g., “file not found”) than communication errors, and common practice is to catch these errors and continue executing. For this reason, MPI provides additional error facilities for I/O.

Advice to users. MPI does not specify the state of a computation after an erroneous MPI call has occurred. A high quality implementation will support the I/O error handling facilities, allowing users to write programs using common practice for I/O. (*End of advice to users.*)

Like communicators, each file handle has an error handler associated with it. The MPI-2 I/O error handling routines are defined in Section 4.13, page 61.

When MPI calls a user-defined error handler resulting from an error on a particular file handle, the first two arguments passed to the file error handler are the file handle and the error code. For I/O errors that are not associated with a valid file handle (e.g., in `MPI_FILE_OPEN` or `MPI_FILE_DELETE`), the first argument passed to the error handler is `MPI_FILE_NULL`,

I/O error handling differs from communication error handling in another important aspect. By default, the predefined error handler for file handles is `MPI_ERRORS_RETURN`. The default file error handler has two purposes: when a new file handle is created (by `MPI_FILE_OPEN`), the error handler for the new file handle is initially set to the default error handler, and I/O routines that have no valid file handle on which to raise an error (e.g., `MPI_FILE_OPEN` or `MPI_FILE_DELETE`) use the default file error handler. The default file error handler can be changed by specifying `MPI_FILE_NULL` as the `fh` argument to `MPI_FILE_SET_ERRHANDLER`. The current value of the default file error handler can be determined by passing `MPI_FILE_NULL` as the `fh` argument to `MPI_FILE_GET_ERRHANDLER`.

Rationale. For communication, the default error handler is inherited from `MPI_COMM_WORLD`. In I/O, there is no analogous “root” file handle from which default properties can be inherited. Rather than invent a new global file handle, the default file error handler is manipulated as if it were attached to `MPI_FILE_NULL`. (*End of rationale.*)

9.8 I/O Error Classes

The implementation dependent error codes returned by the I/O routines can be converted into the following error classes. In addition, calls to routines in this chapter may raise errors in other MPI classes, such as `MPI_ERR_TYPE`.

1	<code>MPI_ERR_FILE</code>	Invalid file handle
2	<code>MPI_ERR_NOT_SAME</code>	Collective argument not identical on all
3		processes, or collective routines called in
4		a different order by different processes
5	<code>MPI_ERR_AMODE</code>	Error related to the <code>amode</code> passed to
6		<code>MPI_FILE_OPEN</code>
7	<code>MPI_ERR_UNSUPPORTED_DATAREP</code>	Unsupported <code>datarep</code> passed to
8		<code>MPI_FILE_SET_VIEW</code>
9	<code>MPI_ERR_UNSUPPORTED_OPERATION</code>	Unsupported operation, such as seeking on
10		a file which supports sequential access only
11	<code>MPI_ERR_NO_SUCH_FILE</code>	File does not exist
12	<code>MPI_ERR_FILE_EXISTS</code>	File exists
13	<code>MPI_ERR_BAD_FILE</code>	Invalid file name (e.g., path name too long)
14	<code>MPI_ERR_ACCESS</code>	Permission denied
15	<code>MPI_ERR_NO_SPACE</code>	Not enough space
16	<code>MPI_ERR_QUOTA</code>	Quota exceeded
17	<code>MPI_ERR_READ_ONLY</code>	Read-only file or file system
18	<code>MPI_ERR_FILE_IN_USE</code>	File operation could not be completed, as
19		the file is currently open by some process
20	<code>MPI_ERR_DUP_DATAREP</code>	Conversion functions could not be regis-
21		tered because a data representation identi-
22		fier that was already defined was passed to
23		<code>MPI_REGISTER_DATAREP</code>
24	<code>MPI_ERR_CONVERSION</code>	An error occurred in a user supplied data
25		conversion function.
26	<code>MPI_ERR_IO</code>	Other I/O error

9.9 Examples

9.9.1 Double Buffering with Split Collective I/O

This example shows how to overlap computation and output. The computation is performed by the function `compute_buffer()`.

```

35  /*=====
36  *
37  * Function:          double_buffer
38  *
39  * Synopsis:
40  *     void double_buffer(
41  *         MPI_File fh,                ** IN
42  *         MPI_Datatype buftype,      ** IN
43  *         int bufcount                ** IN
44  *     )
45  *
46  * Description:
47  *     Performs the steps to overlap computation with a collective write
48  */

```



```

1      by using a double-buffering technique.
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1
```

```

1      TOGGLE_PTR(write_buf_ptr);
2      MPI_File_write_all_begin(fh, write_buf_ptr, bufcount, buftype);
3  }
4
5  /* DOUBLE-BUFFER epilog:
6   *   wait for final write to complete.
7   */
8  MPI_File_write_all_end(fh, write_buf_ptr, &status);
9
10
11 /* buffer cleanup */
12 free(buffer1);
13 free(buffer2);
14 }

```

9.9.2 Subarray Filetype Constructor




Figure 9.4: Example array file layout




Figure 9.5: Example local array filetype for process 1

Assume we are writing out a 100x100 2D array of double precision floating point numbers that is distributed among 4 processes such that each process has a block of 25 columns (e.g.,

process 0 has columns 0-24, process 1 has columns 25-49, etc.; see Figure 9.4). To create the filetypes for each process one could use the following C program:

```
double subarray[100][25];
MPI_Datatype filetype;
int sizes[2], subsizes[2], starts[2];
int rank;

MPI_Comm_rank(MPI_COMM_WORLD, &rank);
sizes[0]=100; sizes[1]=100;
subsizes[0]=100; subsizes[1]=25;
starts[0]=0; starts[1]=rank*subsizes[1];

MPI_Type_create_subarray(2, sizes, subsizes, starts, MPI_ORDER_C,
                        MPI_DOUBLE, &filetype);
```

Or, equivalently in Fortran:

```
double precision subarray(100,25)
integer filetype, rank, ierror
integer sizes(2), subsizes(2), starts(2)

call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierror)
sizes(1)=100
sizes(2)=100
subsizes(1)=100
subsizes(2)=25
starts(1)=0
starts(2)=rank*subsizes(2)

call MPI_TYPE_CREATE_SUBARRAY(2, sizes, subsizes, starts, &
                             MPI_ORDER_FORTRAN, MPI_DOUBLE_PRECISION, &
                             filetype, ierror)
```

The generated filetype will then describe the portion of the file contained within the process's subarray with holes for the space taken by the other processes. Figure 9.5 shows the filetype created for process 1.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Chapter 10

Language Bindings

10.1 C++

10.1.1 Overview

This section presents a complete C++ language interface for MPI. There are some issues specific to C++ that must be considered in the design of this interface that go beyond the simple description of language bindings. In particular, in C++, we must be concerned with the design of objects and their interfaces, rather than just the design of a language-specific functional interface to MPI. Fortunately, the original design of MPI was based on the notion of objects, so a natural set of classes is already part of MPI.

Since the original design of MPI-1 did not include a C++ language interface, a complete list of C++ bindings for MPI-1 functions is provided in Annex B. MPI-2 includes C++ bindings as part of its function specifications. In some cases, MPI-2 provides new names for the C bindings of MPI-1 functions. In this case, the C++ binding matches the new C name — there is no binding for the deprecated name. As such, the C++ binding for the new name appears in Annex A, not Annex B.

10.1.2 Design

The C++ language interface for MPI is designed according to the following criteria:

1. The C++ language interface consists of a small set of classes with a lightweight functional interface to MPI. The classes are based upon the fundamental MPI object types (e.g., communicator, group, etc.).
2. The MPI C++ language bindings provide a semantically correct interface to MPI.
3. To the greatest extent possible, the C++ bindings for MPI functions are member functions of MPI classes.

Rationale. Providing a lightweight set of MPI objects that correspond to the basic MPI types is the best fit to MPI's implicit object-based design; methods can be supplied for these objects to realize MPI functionality. The existing C bindings can be used in C++ programs, but much of the expressive power of the C++ language is forfeited. On the other hand, while a comprehensive class library would make user programming more elegant, such a library it is not suitable as a language binding for MPI since a

binding must provide a direct and unambiguous mapping to the specified functionality of MPI. (*End of rationale.*)

10.1.3 C++ Classes for MPI

All MPI classes, constants, and functions are declared within the scope of an MPI `namespace`. Thus, instead of the `MPI_` prefix that is used in C and Fortran, MPI functions essentially have an `MPI::` prefix.

Advice to implementors. Although `namespace` is officially part of the draft ANSI C++ standard, as of this writing it not yet widely implemented in C++ compilers. Implementations using compilers without `namespace` may obtain the same scoping through the use of a non-instantiable MPI class. (To make the MPI class non-instantiable, all constructors must be `private`.) (*End of advice to implementors.*)

The members of the MPI namespace are those classes corresponding to objects implicitly used by MPI. An abbreviated definition of the MPI namespace for MPI-1 and its member classes is as follows:

```
namespace MPI {
    class Comm                {...};
    class Intracomm : public Comm    {...};
    class Graphcomm : public Intracomm {...};
    class Cartcomm  : public Intracomm {...};
    class Intercomm : public Comm    {...};
    class Datatype   {...};
    class Errhandler {...};
    class Exception  {...};
    class Group      {...};
    class Op          {...};
    class Request     {...};
    class Prequest    : public Request {...};
    class Status      {...};
};
```

Additionally, the following classes defined for MPI-2:

```
namespace MPI {
    class File                {...};
    class Grequest : public Request {...};
    class Info                {...};
    class Win                 {...};
};
```

Note that there are a small number of derived classes, and that virtual inheritance is *not* used.

10.1.4 Class Member Functions for MPI

Besides the member functions which constitute the C++ language bindings for MPI, the C++ language interface has additional functions (as required by the C++ language). In particular, the C++ language interface must provide a constructor and destructor, an assignment operator, and comparison operators.

The complete set of C++ language bindings for MPI-1 is presented in Annex B. The bindings take advantage of some important C++ features, such as references and `const`. Declarations (which apply to all MPI member classes) for construction, destruction, copying, assignment, comparison, and mixed-language operability are also provided. To maintain consistency with what has gone before, the binding definitions are given in the same order as given for the C bindings in [6].

Except where indicated, all non-static member functions (except for constructors and the assignment operator) of MPI member classes are virtual functions.

Rationale. Providing virtual member functions is an important part of design for inheritance. Virtual functions can be bound at run-time, which allows users of libraries to re-define the behavior of objects already contained in a library. There is a small performance penalty that must be paid (the virtual function must be looked up before it can be called). However, users concerned about this performance penalty can force compile-time function binding. (*End of rationale.*)

Example 10.1 Example showing a derived MPI class.

```
class foo_comm : public MPI::Intracomm {
public:
    void Send(void* buf, int count, const MPI::Datatype& type,
              int dest, int tag) const
    {
        // Class library functionality
        MPI::Intracomm::Send(buf, count, type, dest, tag);
        // More class library functionality
    }
};
```

Advice to implementors. Implementors must be careful to avoid unintended side effects from class libraries that use inheritance, especially in layered implementations. For example, if `MPI_BCAST` is implemented by repeated calls to `MPI_SEND` or `MPI_RECV`, the behavior of `MPI_BCAST` cannot be changed by derived communicator classes that might redefine `MPI_SEND` or `MPI_RECV`. The implementation of `MPI_BCAST` must explicitly use the `MPI_SEND` (or `MPI_RECV`) of the base `MPI::Comm` class. (*End of advice to implementors.*)

10.1.5 Semantics

The semantics of the member functions constituting the C++ language binding for MPI are specified by the MPI function description itself. Here, we specify the semantics for those portions of the C++ language interface that are not part of the language binding. In this subsection, functions are prototyped using the type `MPI::<CLASS>` rather than listing each function for every MPI class; the word `<CLASS>` can be replaced with any valid MPI class name (e.g., `Group`), except as noted.

Construction / Destruction The default constructor and destructor are prototyped as follows:

```
MPI::<CLASS>()
~MPI::<CLASS>()
```

In terms of construction and destruction, opaque MPI user level objects behave like handles. Default constructors for all MPI objects except `MPI::Status` create corresponding `MPI::*NULL` handles. That is, when an MPI object is instantiated, comparing it with its corresponding `MPI::*NULL` object will return `true`. The default constructors do not create new MPI opaque objects. Some classes have a member function `Create()` for this purpose.

Example 10.2 In the following code fragment, the test will return `true` and the message will be sent to `cout`.

```
void foo()
{
    MPI::Intracomm bar;

    if (bar == MPI::COMM_NULL)
        cout << "bar is MPI::COMM_NULL" << endl;
}
```

The destructor for each MPI user level object does *not* invoke the corresponding `MPI*_FREE` function (if it exists).

Rationale. `MPI*_FREE` functions are not automatically invoked for the following reasons:

1. Automatic destruction contradicts the shallow-copy semantics of the MPI classes.
2. The model put forth in MPI makes memory allocation and deallocation the responsibility of the user, not the implementation.
3. Calling `MPI*_FREE` upon destruction could have unintended side effects, including triggering collective operations (this also affects the copy, assignment, and construction semantics). In the following example, we would want neither `foo_comm` nor `bar_comm` to automatically invoke `MPI*_FREE` upon exit from the function.

```
void example_function()
{
    MPI::Intracomm foo_comm(MPI::COMM_WORLD), bar_comm;
    bar_comm = MPI::COMM_WORLD.Dup();
    // rest of function
}
```

(*End of rationale.*)

Copy / Assignment The copy constructor and assignment operator are prototyped as follows:

```
MPI::<CLASS>(const MPI::<CLASS>& data)
```

```
MPI::<CLASS>& MPI::<CLASS>::operator=(const MPI::<CLASS>& data)
```

In terms of copying and assignment, opaque MPI user level objects behave like handles. Copy constructors perform handle-based (shallow) copies. MPI::Status objects are exceptions to this rule. These objects perform deep copies for assignment and copy construction.

Advice to implementors. Each MPI user level object is likely to contain, by value or by reference, implementation-dependent state information. The assignment and copying of MPI object handles may simply copy this value (or reference). (*End of advice to implementors.*)

Example 10.3 Example using assignment operator. In this example, MPI::Intracomm::Dup() is *not* called for foo_comm. The object foo_comm is simply an alias for MPI::COMM_WORLD. But bar_comm is created with a call to MPI::Intracomm::Dup() and is therefore a different communicator than foo_comm (and thus different from MPI::COMM_WORLD). baz_comm becomes an alias for bar_comm. If one of bar_comm or baz_comm is freed with MPI_COMM_FREE it will be set to MPI::COMM_NULL. The state of the other handle will be undefined — it will be invalid, but not necessarily set to MPI::COMM_NULL.

```
MPI::Intracomm foo_comm, bar_comm, baz_comm;
```

```
foo_comm = MPI::COMM_WORLD;
```

```
bar_comm = MPI::COMM_WORLD.Dup();
```

```
baz_comm = bar_comm;
```

Comparison The comparison operators are prototyped as follows:

```
bool MPI::<CLASS>::operator==(const MPI::<CLASS>& data) const
```

```
bool MPI::<CLASS>::operator!=(const MPI::<CLASS>& data) const
```

The member function operator==() returns true only when the handles reference the same internal MPI object, false otherwise. operator!=() returns the boolean complement of operator==(). However, since the Status class is not a handle to an underlying MPI object, it does not make sense to compare Status instances. Therefore, the operator==() and operator!=() functions are not defined on the Status class.

Constants Constants are singleton objects and are declared const. Note that not all globally defined MPI objects are constant. For example, MPI::COMM_WORLD and MPI::COMM_SELF are not const.

10.1.6 C++ Datatypes

Table 10.1 lists all of the C++ predefined MPI datatypes and their corresponding C and C++ datatypes, Table 10.2 lists all of the Fortran predefined MPI datatypes and their

corresponding Fortran 77 datatypes. Table 10.3 lists the C++ names for all other MPI datatypes.

`MPI::BYTE` and `MPI::PACKED` conform to the same restrictions as `MPI_BYTE` and `MPI_PACKED`, listed in Sections 3.2.2 and 3.13 of MPI-1, respectively.

MPI datatype	C datatype	C++ datatype
<code>MPI::CHAR</code>	<code>char</code>	<code>char</code>
<code>MPI::WCHAR</code>	<code>wchar_t</code>	<code>wchar_t</code>
<code>MPI::SHORT</code>	<code>signed short</code>	<code>signed short</code>
<code>MPI::INT</code>	<code>signed int</code>	<code>signed int</code>
<code>MPI::LONG</code>	<code>signed long</code>	<code>signed long</code>
<code>MPI::SIGNED_CHAR</code>	<code>signed char</code>	<code>signed char</code>
<code>MPI::UNSIGNED_CHAR</code>	<code>unsigned char</code>	<code>unsigned char</code>
<code>MPI::UNSIGNED_SHORT</code>	<code>unsigned short</code>	<code>unsigned short</code>
<code>MPI::UNSIGNED</code>	<code>unsigned int</code>	<code>unsigned int</code>
<code>MPI::UNSIGNED_LONG</code>	<code>unsigned long</code>	<code>unsigned long int</code>
<code>MPI::FLOAT</code>	<code>float</code>	<code>float</code>
<code>MPI::DOUBLE</code>	<code>double</code>	<code>double</code>
<code>MPI::LONG_DOUBLE</code>	<code>long double</code>	<code>long double</code>
<code>MPI::BOOL</code>		<code>bool</code>
<code>MPI::COMPLEX</code>		<code>Complex<float></code>
<code>MPI::DOUBLE_COMPLEX</code>		<code>Complex<double></code>
<code>MPI::LONG_DOUBLE_COMPLEX</code>		<code>Complex<long double></code>
<code>MPI::BYTE</code>		
<code>MPI::PACKED</code>		

Table 10.1: C++ names for the MPI C and C++ predefined datatypes, and their corresponding C/C++ datatypes.

MPI datatype	Fortran datatype
<code>MPI::CHARACTER</code>	<code>CHARACTER(1)</code>
<code>MPI::INTEGER</code>	<code>INTEGER</code>
<code>MPI::REAL</code>	<code>REAL</code>
<code>MPI::DOUBLE_PRECISION</code>	<code>DOUBLE PRECISION</code>
<code>MPI::LOGICAL</code>	<code>LOGICAL</code>
<code>MPI::F_COMPLEX</code>	<code>COMPLEX</code>
<code>MPI::BYTE</code>	
<code>MPI::PACKED</code>	

Table 10.2: C++ names for the MPI Fortran predefined datatypes, and their corresponding Fortran 77 datatypes.

The following table defines groups of MPI predefined datatypes:

C integer: `MPI::INT`, `MPI::LONG`, `MPI::SHORT`,

MPI datatype	Description
MPI::FLOAT_INT	C/C++ reduction type
MPI::DOUBLE_INT	C/C++ reduction type
MPI::LONG_INT	C/C++ reduction type
MPI::TWO_INT	C/C++ reduction type
MPI::SHORT_INT	C/C++ reduction type
MPI::LONG_DOUBLE_INT	C/C++ reduction type
MPI::LONG_LONG	Optional C/C++ type
MPI::UNSIGNED_LONG_LONG	Optional C/C++ type
MPI::TWO_REAL	Fortran reduction type
MPI::TWO_DOUBLE_PRECISION	Fortran reduction type
MPI::TWO_INTEGER	Fortran reduction type
MPI::F_DOUBLE_COMPLEX	Optional Fortran type
MPI::INTEGER1	Explicit size type
MPI::INTEGER2	Explicit size type
MPI::INTEGER4	Explicit size type
MPI::INTEGER8	Explicit size type
MPI::REAL4	Explicit size type
MPI::REAL8	Explicit size type
MPI::REAL16	Explicit size type

Table 10.3: C++ names for other MPI datatypes. Implementations may also define other optional types (e.g., `MPI::INTEGER8`).

	MPI::UNSIGNED_SHORT, MPI::UNSIGNED,
	MPI::UNSIGNED_LONG, MPI::SIGNED_CHAR,
	MPI::UNSIGNED_CHAR
Fortran integer:	MPI::INTEGER
Floating point:	MPI::FLOAT, MPI::DOUBLE, MPI::REAL,
	MPI::DOUBLE_PRECISION,
	MPI::LONG_DOUBLE
Logical:	MPI::LOGICAL, MPI::BOOL
Complex:	MPI::F_COMPLEX, MPI::COMPLEX,
	MPI::F_DOUBLE_COMPLEX,
	MPI::DOUBLE_COMPLEX,
	MPI::LONG_DOUBLE_COMPLEX
Byte:	MPI::BYTE

Valid datatypes for each reduction operation is specified below in terms of the groups defined above.

Op	Allowed Types
MPI::MAX, MPI::MIN	C integer, Fortran integer, Floating point
MPI::SUM, MPI::PROD	C integer, Fortran integer, Floating point, Complex
MPI::LAND, MPI::LOR, MPI::LXOR	C integer, Logical
MPI::BAND, MPI::BOR, MPI::BXOR	C integer, Fortran integer, Byte

MPI::MINLOC and MPI::MAXLOC perform just as their C and Fortran counterparts; see Section 4.9.3 in MPI-1.

10.1.7 Communicators

The `MPI::Comm` class hierarchy makes explicit the different kinds of communicators implicitly defined by MPI and allows them to be strongly typed. Since the original design of MPI defined only one type of handle for all types of communicators, the following clarifications are provided for the C++ design.

Types of communicators There are five different types of communicators: `MPI::Comm`, `MPI::Intercomm`, `MPI::Intracomm`, `MPI::Cartcomm`, and `MPI::Graphcomm`. `MPI::Comm` is the abstract base communicator class, encapsulating the functionality common to all MPI communicators. `MPI::Intercomm` and `MPI::Intracomm` are derived from `MPI::Comm`. `MPI::Cartcomm` and `MPI::Graphcomm` are derived from `MPI::Intracomm`.

Advice to users. Initializing a derived class with an instance of a base class is not legal in C++. For instance, it is not legal to initialize a `Cartcomm` from an `Intracomm`. Moreover, because `MPI::Comm` is an abstract base class, it is non-instantiable, so that it is not possible to have an object of class `MPI::Comm`. However, it is possible to have a reference or a pointer to an `MPI::Comm`.

Example 10.4 The following code is erroneous.

```
Intracomm intra = MPI::COMM_WORLD.Dup();
Cartcomm cart(intra);           // This is erroneous
```

(End of advice to users.)

`MPI::COMM_NULL` The specific type of `MPI::COMM_NULL` is implementation dependent. `MPI::COMM_NULL` must be able to be used in comparisons and initializations with all types of communicators. `MPI::COMM_NULL` must also be able to be passed to a function that expects a communicator argument in the parameter list (provided that `MPI::COMM_NULL` is an allowed value for the communicator argument).

Rationale. There are several possibilities for implementation of `MPI::COMM_NULL`. Specifying its required behavior, rather than its realization, provides maximum flexibility to implementors. (*End of rationale.*)

Example 10.5 The following example demonstrates the behavior of assignment and comparison using `MPI::COMM_NULL`.

```
MPI::Intercomm comm;
comm = MPI::COMM_NULL;           // assign with COMM_NULL
if (comm == MPI::COMM_NULL)      // true
    cout << "comm is NULL" << endl;
if (MPI::COMM_NULL == comm)      // note -- a different function!
    cout << "comm is still NULL" << endl;
```

`Dup()` is not defined as a member function of `MPI::Comm`, but it is defined for the derived classes of `MPI::Comm`. `Dup()` is not virtual and it returns its OUT/ parameter by value.

`MPI::Comm::Clone()` The C++ language interface for MPI includes a new function `Clone()`. `MPI::Comm::Clone()` is a pure virtual function. For the derived communicator classes, `Clone()` behaves like `Dup()` except that it returns a new object by reference. The `Clone()` functions are prototyped as follows:

```
Comm& Comm::Clone() const = 0

Intracomm& Intracomm::Clone() const

Intercomm& Intercomm::Clone() const

Cartcomm& Cartcomm::Clone() const

Graphcomm& Graphcomm::Clone() const
```

Rationale. `Clone()` provides the “virtual dup” functionality that is expected by C++ programmers and library writers. Since `Clone()` returns a new object by reference, users are responsible for eventually deleting the object. A new name is introduced rather than changing the functionality of `Dup()`. (*End of rationale.*)

Advice to implementors. Within their class declarations, prototypes for `Clone()` and `Dup()` would look like the following:

```
namespace MPI {
    class Comm {
        virtual Comm& Clone() const = 0;
    };
}
```

```

1      class Intracomm : public Comm {
2          Intracomm Dup() const { ... };
3          virtual Intracomm& Clone() const { ... };
4      };
5      class Intercomm : public Comm {
6          Intercomm Dup() const { ... };
7          virtual Intercomm& Clone() const { ... };
8      };
9      // Cartcomm and Graphcomm are similarly defined
10     };

```

Compilers that do not support the variable return type feature of virtual functions may return a reference to `Comm`. Users can cast to the appropriate type as necessary. (*End of advice to implementors.*)

10.1.8 Exceptions

The C++ language interface for MPI includes the predefined error handler `MPI::ERRORS_THROW_EXCEPTIONS` for use with the `Set_errhandler()` member functions. `MPI::ERRORS_THROW_EXCEPTIONS` can only be set or retrieved by C++ functions. If a non-C++ program causes an error that invokes the `MPI::ERRORS_THROW_EXCEPTIONS` error handler, the exception will pass up the calling stack until C++ code can catch it. If there is no C++ code to catch it, the behavior is undefined. In a multi-threaded environment or if a non-blocking MPI call throws an exception while making progress in the background, the behavior is implementation dependent.

The error handler `MPI::ERRORS_THROW_EXCEPTIONS` causes an `MPI::Exception` to be thrown for any MPI result code other than `MPI::SUCCESS`. The public interface to `MPI::Exception` class is defined as follows:

```

28     namespace MPI {
29         class Exception {
30             public:
31
32                 Exception(int error_code);
33
34                 int Get_error_code() const;
35                 int Get_error_class() const;
36                 const char *Get_error_string() const;
37             };
38     };
39

```

Advice to implementors.

The exception will be thrown within the body of `MPI::ERRORS_THROW_EXCEPTIONS`. It is expected that control will be returned to the user when the exception is thrown. Some MPI functions specify certain return information in their parameters in the case of an error and `MPI_ERRORS_RETURN` is specified. The same type of return information must be provided when exceptions are thrown.

For example, `MPI_WAITALL` puts an error code for each request in the corresponding entry in the status array and returns `MPI_ERR_IN_STATUS`. When using

`MPI::ERRORS_THROW_EXCEPTIONS`, it is expected that the error codes in the status array will be set appropriately before the exception is thrown.

(End of advice to implementors.)

10.1.9 Mixed-Language Operability

The C++ language interface provides functions listed below for mixed-language operability. These functions provide for a seamless transition between C and C++. For the case where the C++ class corresponding to `<CLASS>` has derived classes, functions are also provided for converting between the derived classes and the C `MPI_<CLASS>`.

```
MPI::<CLASS>& MPI::<CLASS>::operator=(const MPI_<CLASS>& data)
```

```
MPI::<CLASS>(const MPI_<CLASS>& data)
```

```
MPI::<CLASS>::operator MPI_<CLASS>() const
```

These functions are discussed in Section 4.12.4.

10.1.10 Profiling

This section specifies the requirements of a C++ profiling interface to MPI.

Advice to implementors. Since the main goal of profiling is to intercept function calls from user code, it is the implementor's decision how to layer the underlying implementation to allow function calls to be intercepted and profiled. If an implementation of the MPI C++ bindings is layered on top of MPI bindings in another language (such as C), or if the C++ bindings are layered on top of a profiling interface in another language, no extra profiling interface is necessary because the underlying MPI implementation already meets the MPI profiling interface requirements.

Native C++ MPI implementations that do not have access to other profiling interfaces must implement an interface that meets the requirements outlined in this section.

High quality implementations can implement the interface outlined in this section in order to promote portable C++ profiling libraries. Implementors may wish to provide an option whether to build the C++ profiling interface or not; C++ implementations that are already layered on top of bindings in another language or another profiling interface will have to insert a third layer to implement the C++ profiling interface.

(End of advice to implementors.)

To meet the requirements of the C++ MPI profiling interface, an implementation of the MPI functions *must*:

1. Provide a mechanism through which all of the MPI defined functions may be accessed with a name shift. Thus all of the MPI functions (which normally start with the prefix "MPI:") should also be accessible with the prefix "PMPI:."
2. Ensure that those MPI functions which are not replaced may still be linked into an executable image without causing name clashes.

3. Document the implementation of different language bindings of the `MPI` interface if they are layered on top of each other, so that profiler developer knows whether they must implement the profile interface for each binding, or can economize by implementing it only for the lowest level routines.

4. Where the implementation of different language bindings is done through a layered approach (e.g., the C++ binding is a set of “wrapper” functions which call the C implementation), ensure that these wrapper functions are separable from the rest of the library.

This is necessary to allow a separate profiling library to be correctly implemented, since (at least with Unix linker semantics) the profiling library must contain these wrapper functions if it is to perform as expected. This requirement allows the author of the profiling library to extract these functions from the original `MPI` library and add them into the profiling library without bringing along any other unnecessary code.

5. Provide a no-op routine `MPI::Pcontrol` in the `MPI` library.

Advice to implementors. There are (at least) two apparent options for implementing the C++ profiling interface: inheritance or caching. An inheritance-based approach may not be attractive because it may require a virtual inheritance implementation of the communicator classes. Thus, it is most likely that implementors still cache `PMPI` objects on their corresponding `MPI` objects. The caching scheme is outlined below.

The “real” entry points to each routine can be provided within a `namespace PMPI`. The non-profiling version can then be provided within a `namespace MPI`.

Caching instances of `PMPI` objects in the `MPI` handles provides the “has a” relationship that is necessary to implement the profiling scheme.

Each instance of an `MPI` object simply “wraps up” an instance of a `PMPI` object. `MPI` objects can then perform profiling actions before invoking the corresponding function in their internal `PMPI` object.

The key to making the profiling work by simply re-linking programs is by having a header file that *declares* all the `MPI` functions. The functions must be *defined* elsewhere, and compiled into a library. `MPI` constants should be declared `extern` in the `MPI` namespace. For example, the following is an excerpt from a sample `mpi.h` file:

Example 10.6 Sample `mpi.h` file.

```
namespace PMPI {
    class Comm {
    public:
        int Get_size() const;
    };
    // etc.
};

namespace MPI {
public:
```



```

class Comm {
public:
    int Get_size() const;

private:
    PMPI::Comm pmpi_comm;
};

```

Note that all constructors, the assignment operator, and the destructor in the `MPI` class will need to initialize/destroy the internal `PMPI` object as appropriate.

The definitions of the functions must be in separate object files; the `PMPI` class member functions and the non-profiling versions of the `MPI` class member functions can be compiled into `libmpi.a`, while the profiling versions can be compiled into `libpmpi.a`. Note that the `PMPI` class member functions and the `MPI` constants must be in different object files than the non-profiling `MPI` class member functions in the `libmpi.a` library to prevent multiple definitions of `MPI` class member function names when linking both `libmpi.a` and `libpmpi.a`. For example:

Example 10.7 `pmapi.cc`, to be compiled into `libmpi.a`.

```

int PMPI::Comm::Get_size() const
{
    // Implementation of MPI_COMM_SIZE
}

```

Example 10.8 `constants.cc`, to be compiled into `libmpi.a`.

```

const MPI::Intracomm MPI::COMM_WORLD;

```

Example 10.9 `mpi_no_profile.cc`, to be compiled into `libmpi.a`.

```

int MPI::Comm::Get_size() const
{
    return pmpi_comm.Get_size();
}

```

Example 10.10 `mpi_profile.cc`, to be compiled into `libpmpi.a`.

```

int MPI::Comm::Get_size() const
{
    // Do profiling stuff
    int ret = pmpi_comm.Get_size();
    // More profiling stuff
    return ret;
}

```

(End of advice to implementors.)

10.2 Fortran Support

10.2.1 Overview

Fortran 90 is the current international Fortran standard. MPI-2 Fortran bindings are Fortran 90 bindings that in most cases are “Fortran 77 friendly.” That is, with few exceptions (e.g., `KIND`-parameterized types, and the `mpi` module, both of which can be avoided) Fortran 77 compilers should be able to compile MPI programs.

Rationale. Fortran 90 contains numerous features designed to make it a more “modern” language than Fortran 77. It seems natural that MPI should be able to take advantage of these new features with a set of bindings tailored to Fortran 90. MPI does not (yet) use many of these features because of a number of technical difficulties. (*End of rationale.*)

MPI defines two levels of Fortran support, described in Sections 10.2.3 and 10.2.4. A third level of Fortran support is envisioned, but is deferred to future standardization efforts. In the rest of this section, “Fortran” shall refer to Fortran 90 (or its successor) unless qualified.

1. **Basic Fortran Support** An implementation with this level of Fortran support provides the original Fortran bindings specified in MPI-1, with small additional requirements specified in Section 10.2.3.
2. **Extended Fortran Support** An implementation with this level of Fortran support provides Basic Fortran Support plus additional features that specifically support Fortran 90, as described in Section 10.2.4.

A compliant MPI-2 implementation providing a Fortran interface must provide Extended Fortran Support unless the target compiler does not support modules or `KIND`-parameterized types.

10.2.2 Problems With Fortran Bindings for MPI

This section discusses a number of problems that may arise when using MPI in a Fortran program. It is intended as advice to users, and clarifies how MPI interacts with Fortran. It does not add to the standard, but is intended to clarify the standard.

As noted in the original MPI specification, the interface violates the Fortran standard in several ways. While these cause few problems for Fortran 77 programs, they become more significant for Fortran 90 programs, so that users must exercise care when using new Fortran 90 features. The violations were originally adopted and have been retained because they are important for the usability of MPI. The rest of this section describes the potential problems in detail. It supersedes and replaces the discussion of Fortran bindings in the original MPI specification (for Fortran 90, not Fortran 77).

The following MPI features are inconsistent with Fortran 90.

1. An MPI subroutine with a choice argument may be called with different argument types.
2. An MPI subroutine with an assumed-size dummy argument may be passed an actual scalar argument.

3. Many MPI routines assume that actual arguments are passed by address and that arguments are not copied on entrance to or exit from the subroutine.
4. An MPI implementation may read or modify user data (e.g., communication buffers used by nonblocking communications) concurrently with a user program that is executing outside of MPI calls.
5. Several named “constants,” such as `MPI_BOTTOM`, `MPI_IN_PLACE`, `MPI_STATUS_IGNORE`, `MPI_STATUSES_IGNORE`, `MPI_ERRCODES_IGNORE`, `MPI_ARGV_NULL`, and `MPI_ARGVS_NULL` are not ordinary Fortran constants and require a special implementation. See Section 2.5.4 on page 10 for more information.
6. The memory allocation routine `MPI_ALLOC_MEM` can’t be usefully used in Fortran without a language extension that allows the allocated memory to be associated with a Fortran variable.

MPI-1 contained several routines that take address-sized information as input or return address-sized information as output. In C such arguments were of type `MPI_Aint` and in Fortran of type `INTEGER`. On machines where integers are smaller than addresses, these routines can lose information. In MPI-2 the use of these functions has been deprecated and they have been replaced by routines taking `INTEGER` arguments of `KIND=MPI_ADDRESS_KIND`. A number of new MPI-2 functions also take `INTEGER` arguments of non-default `KIND`. See Section 2.6 on page 11 and Section 4.14 on page 65 for more information.

Problems Due to Strong Typing

All MPI functions with choice arguments associate actual arguments of different Fortran datatypes with the same dummy argument. This is not allowed by Fortran 77, and in Fortran 90 is technically only allowed if the function is overloaded with a different function for each type. In C, the use of `void*` formal arguments avoids these problems.

The following code fragment is technically illegal and may generate a compile-time error.

```
integer i(5)
real    x(5)
...
call mpi_send(x, 5, MPI_REAL, ...)
call mpi_send(i, 5, MPI_INTEGER, ...)
```

In practice, it is rare for compilers to do more than issue a warning, though there is concern that Fortran 90 compilers are more likely to return errors.

It is also technically illegal in Fortran to pass a scalar actual argument to an array dummy argument. Thus the following code fragment may generate an error since the `buf` argument to `MPI_SEND` is declared as an assumed-size array `<type> buf(*)`.

```
integer a
call mpi_send(a, 1, MPI_INTEGER, ...)
```

Advice to users. In the event that you run into one of the problems related to type checking, you may be able to work around it by using a compiler flag, by compiling

separately, or by using an MPI implementation with Extended Fortran Support as described in Section 10.2.4. An alternative that will usually work with variables local to a routine but not with arguments to a function or subroutine is to use the **EQUIVALENCE** statement to create another variable with a type accepted by the compiler. (*End of advice to users.*)

Problems Due to Data Copying and Sequence Association

Implicit in MPI is the idea of a contiguous chunk of memory accessible through a linear address space. MPI copies data to and from this memory. An MPI program specifies the location of data by providing memory addresses and offsets. In the C language, sequence association rules plus pointers provide all the necessary low-level structure.

In Fortran 90, user data is not necessarily stored contiguously. For example, the array section **A(1:N:2)** involves only the elements of **A** with indices 1, 3, 5, The same is true for a pointer array whose target is such a section. Most compilers ensure that an array that is a dummy argument is held in contiguous memory if it is declared with an explicit shape (e.g., **B(N)**) or is of assumed size (e.g., **B(*)**). If necessary, they do this by making a copy of the array into contiguous memory. Both Fortran 77 and Fortran 90 are carefully worded to allow such copying to occur, but few Fortran 77 compilers do it.¹

Because MPI dummy buffer arguments are assumed-size arrays, this leads to a serious problem for a non-blocking call: the compiler copies the temporary array back on return but MPI continues to copy data to the memory that held it. For example, consider the following code fragment:

```
real a(100)
call MPI_IRECV(a(1:100:2), MPI_REAL, 50, ...)
```

Since the first dummy argument to **MPI_IRECV** is an assumed-size array (**<type> buf(*)**), the array section **a(1:100:2)** is copied to a temporary before being passed to **MPI_IRECV**, so that it is contiguous in memory. **MPI_IRECV** returns immediately, and data is copied from the temporary back into the array **a**. Sometime later, MPI may write to the address of the deallocated temporary. Copying is also a problem for **MPI_SEND** since the temporary array may be deallocated before the data has all been sent from it.

Most Fortran 90 compilers do not make a copy if the actual argument is the whole of an explicit-shape or assumed-size array or is a ‘simple’ section such as **A(1:N)** of such an array. (We define ‘simple’ more fully in the next paragraph.) Also, many compilers treat allocatable arrays the same as they treat explicit-shape arrays in this regard (though we know of one that does not). However, the same is not true for assumed-shape and pointer arrays; since they may be discontinuous, copying is often done. It is this copying that causes problems for MPI as described in the previous paragraph.

Our formal definition of a ‘simple’ array section is

```
name ( [:,]... [<subscript>]:<subscript> [,<subscript>]... )
```

That is, there are zero or more dimensions that are selected in full, then one dimension selected without a stride, then zero or more dimensions that are selected with a simple subscript. Examples are

```
A(1:N), A(:,N), A(:,1:N,1), A(1:6,N), A(:, :, 1:N)
```

¹Technically, the Fortran standards are worded to allow non-contiguous storage of any array data.

Because of Fortran's column-major ordering, where the first index varies fastest, a simple section of a contiguous array will also be contiguous.²

The same problem can occur with a scalar argument. Some compilers, even for Fortran 77, make a copy of some scalar dummy arguments within a called procedure. That this can cause a problem is illustrated by the example

```
call user1(a,rq)
call MPI_WAIT(rq,status,ierr)
write (*,*) a

subroutine user1(buf,request)
call MPI_IRECV(buf,...,request,...)
end
```

If `a` is copied, `MPI_IRECV` will alter the copy when it completes the communication and will not alter `a` itself.

Note that copying will almost certainly occur for an argument that is a non-trivial expression (one with at least one operator or function call), a section that does not select a contiguous part of its parent (e.g., `A(1:n:2)`), a pointer whose target is such a section, or an assumed-shape array that is (directly or indirectly) associated with such a section.

If there is a compiler option that inhibits copying of arguments, in either the calling or called procedure, this should be employed.

If a compiler makes copies in the calling procedure of arguments that are explicit-shape or assumed-size arrays, simple array sections of such arrays, or scalars, and if there is no compiler option to inhibit this, then the compiler cannot be used for applications that use `MPI_GET_ADDRESS`, or any non-blocking MPI routine. If a compiler copies scalar arguments in the called procedure and there is no compiler option to inhibit this, then this compiler cannot be used for applications that use memory references across subroutine calls as in the example above.

Special Constants

MPI requires a number of special "constants" that cannot be implemented as normal Fortran constants, including `MPI_BOTTOM`, `MPI_STATUS_IGNORE`, `MPI_IN_PLACE`, `MPI_STATUSES_IGNORE` and `MPI_ERRCODES_IGNORE`. In C, these are implemented as constant pointers, usually as `NULL` and are used where the function prototype calls for a pointer to a variable, not the variable itself.

In Fortran the implementation of these special constants may require the use of language constructs that are outside the Fortran standard. Using special values for the constants (e.g., by defining them through `parameter` statements) is not possible because an implementation cannot distinguish these values from legal data. Typically these constants are implemented as predefined static variables (e.g., a variable in an MPI-declared `COMMON` block), relying on the fact that the target compiler passes data by address. Inside the subroutine, this address can be extracted by some mechanism outside the Fortran standard (e.g., by Fortran extensions or by implementing the function in C).

²To keep the definition of 'simple' simple, we have chosen to require all but one of the section subscripts to be without bounds. A colon without bounds makes it obvious both to the compiler and to the reader that the whole of the dimension is selected. It would have been possible to allow cases where the whole dimension is selected with one or two bounds, but this means for the reader that the array declaration or most recent allocation has to be consulted and for the compiler that a run-time check may be required.

Fortran 90 Derived Types

MPI does not explicitly support passing Fortran 90 derived types to choice dummy arguments. Indeed, for MPI implementations that provide explicit interfaces through the `mpi` module a compiler will reject derived type actual arguments at compile time. Even when no explicit interfaces are given, users should be aware that Fortran 90 provides no guarantee of sequence association for derived types or arrays of derived types. For instance, an array of a derived type consisting of two elements may be implemented as an array of the first elements followed by an array of the second. Use of the `SEQUENCE` attribute may help here, somewhat.

The following code fragment shows one possible way to send a derived type in Fortran. The example assumes that all data is passed by address.

```

type mytype
  integer i
  real x
  double precision d
end type mytype

type(mytype) foo
integer blocklen(3), type(3)
integer(MPI_ADDRESS_KIND) disp(3), base

call MPI_GET_ADDRESS(foo%i, disp(1), ierr)
call MPI_GET_ADDRESS(foo%x, disp(2), ierr)
call MPI_GET_ADDRESS(foo%d, disp(3), ierr)

base = disp(1)
disp(1) = disp(1) - base
disp(2) = disp(2) - base
disp(3) = disp(3) - base

blocklen(1) = 1
blocklen(2) = 1
blocklen(3) = 1

type(1) = MPI_INTEGER
type(2) = MPI_REAL
type(3) = MPI_DOUBLE_PRECISION

call MPI_TYPE_CREATE_STRUCT(3, blocklen, disp, type, newtype, ierr)
call MPI_TYPE_COMMIT(newtype, ierr)

! unpleasant to send foo%i instead of foo, but it works for scalar
! entities of type mytype
call MPI_SEND(foo%i, 1, newtype, ...)

```

A Problem with Register Optimization

MPI provides operations that may be hidden from the user code and run concurrently with it, accessing the same memory as user code. Examples include the data transfer for an `MPI_RECV`. The optimizer of a compiler will assume that it can recognize periods when a copy of a variable can be kept in a register without reloading from or storing to memory. When the user code is working with a register copy of some variable while the hidden operation reads or writes the memory copy, problems occur. This section discusses register optimization pitfalls.

When a variable is local to a Fortran subroutine (i.e., not in a module or `COMMON block`), the compiler will assume that it cannot be modified by a called subroutine unless it is an actual argument of the call. In the most common linkage convention, the subroutine is expected to save and restore certain registers. Thus, the optimizer will assume that a register which held a valid copy of such a variable before the call will still hold a valid copy on return.

Normally users are not afflicted with this. But the user should pay attention to this section if in his/her program a buffer argument to an `MPI_SEND`, `MPI_RECV` etc., uses a name which hides the actual variables involved. `MPI_BOTTOM` with an `MPI_Datatype` containing absolute addresses is one example. Creating a datatype which uses one variable as an anchor and brings along others by using `MPI_GET_ADDRESS` to determine their offsets from the anchor is another. The anchor variable would be the only one mentioned in the call. Also attention must be paid if MPI operations are used that run in parallel with the user's application.

The following example shows what Fortran compilers are allowed to do.

This source ...	can be compiled as:
<code>call MPI_GET_ADDRESS(buf,bufaddr, ierror)</code>	<code>call MPI_GET_ADDRESS(buf,...)</code>
<code>call MPI_TYPE_CREATE_STRUCT(1,1, bufaddr, MPI_REAL,type,ierror)</code>	<code>call MPI_TYPE_CREATE_STRUCT(...)</code>
<code>call MPI_TYPE_COMMIT(type,ierror)</code>	<code>call MPI_TYPE_COMMIT(...)</code>
<code>val_old = buf</code>	<code>register = buf</code>
	<code>val_old = register</code>
<code>call MPI_RECV(MPI_BOTTOM,1,type,...)</code>	<code>call MPI_RECV(MPI_BOTTOM,...)</code>
<code>val_new = buf</code>	<code>val_new = register</code>

The compiler does not invalidate the register because it cannot see that `MPI_RECV` changes the value of `buf`. The access of `buf` is hidden by the use of `MPI_GET_ADDRESS` and `MPI_BOTTOM`.

The next example shows extreme, but allowed, possibilities.

Source	compiled as	or compiled as
<code>call MPI_RECV(buf,..req)</code>	<code>call MPI_RECV(buf,..req)</code>	<code>call MPI_RECV(buf,..req)</code>
	<code>register = buf</code>	<code>b1 = buf</code>
<code>call MPI_WAIT(req,..)</code>	<code>call MPI_WAIT(req,..)</code>	<code>call MPI_WAIT(req,..)</code>
<code>b1 = buf</code>	<code>b1 := register</code>	

`MPI_WAIT` on a concurrent thread modifies `buf` between the invocation of `MPI_RECV` and the finish of `MPI_WAIT`. But the compiler cannot see any possibility that `buf` can be changed after `MPI_RECV` has returned, and may schedule the load of `buf` earlier than typed in the source. It has no reason to avoid using a register to hold `buf` across the call to `MPI_WAIT`. It also may reorder the instructions as in the case on the right.

To prevent instruction reordering or the allocation of a buffer in a register there are two possibilities in portable Fortran code:

- The compiler may be prevented from moving a reference to a buffer across a call to an `MPI` subroutine by surrounding the call by calls to an external subroutine with the buffer as an actual argument. Note that if the intent is declared in the external subroutine, it must be `OUT` or `INOUT`. The subroutine itself may have an empty body, but the compiler does not know this and has to assume that the buffer may be altered. For example, the above call of `MPI_RECV` might be replaced by

```
call DD(buf)
call MPI_RECV(MPI_BOTTOM,...)
call DD(buf)
```

with the separately compiled

```
subroutine DD(buf)
  integer buf
end
```

(assuming that `buf` has type `INTEGER`). The compiler may be similarly prevented from moving a reference to a variable across a call to an `MPI` subroutine.

In the case of a non-blocking call, as in the above call of `MPI_WAIT`, no reference to the buffer is permitted until it has been verified that the transfer has been completed. Therefore, in this case, the extra call ahead of the `MPI` call is not necessary, i.e., the call of `MPI_WAIT` in the example might be replaced by

```
call MPI_WAIT(req,...)
call DD(buf)
```

- An alternative is to put the buffer or variable into a module or a common block and access it through a `USE` or `COMMON` statement in each scope where it is referenced, defined or appears as an actual argument in a call to an `MPI` routine. The compiler will then have to assume that the `MPI` procedure (`MPI_RECV` in the above example) may alter the buffer or variable, provided that the compiler cannot analyze that the `MPI` procedure does not reference the module or common block.

In the longer term, the attribute `VOLATILE` is under consideration for Fortran 2000 and would give the buffer or variable the properties needed, but it would inhibit optimization of any code containing the buffer or variable.

In C, subroutines which modify variables that are not in the argument list will not cause register optimization problems. This is because taking pointers to storage objects by using the `&` operator and later referencing the objects by way of the pointer is an integral part of the language. A C compiler understands the implications, so that the problem should not occur, in general. However, some compilers do offer optional aggressive optimization levels which may not be safe.

10.2.3 Basic Fortran Support

Because Fortran 90 is (for all practical purposes) a superset of Fortran 77, Fortran 90 (and future) programs can use the original Fortran interface. The following additional requirements are added:

1. Implementations are required to provide the file `mpif.h`, as described in the original MPI-1 specification.
2. `mpif.h` must be valid and equivalent for both fixed- and free- source form.

Advice to implementors. To make `mpif.h` compatible with both fixed- and free-source forms, to allow automatic inclusion by preprocessors, and to allow extended fixed-form line length, it is recommended that requirement two be met by constructing `mpif.h` without any continuation lines. This should be possible because `mpif.h` contains only declarations, and because common block declarations can be split among several lines. To support Fortran 77 as well as Fortran 90, it may be necessary to eliminate all comments from `mpif.h`. (*End of advice to implementors.*)

10.2.4 Extended Fortran Support

Implementations with Extended Fortran support must provide:

1. An `mpi` module
2. A new set of functions to provide additional support for Fortran intrinsic numeric types, including parameterized types: `MPI_SIZEOF`, `MPI_TYPE_MATCH_SIZE`, `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL` and `MPI_TYPE_CREATE_F90_COMPLEX`. Parameterized types are Fortran intrinsic types which are specified using `KIND` type parameters. These routines are described in detail in Section 10.2.5.

Additionally, high quality implementations should provide a mechanism to prevent fatal type mismatch errors for MPI routines with choice arguments.

The `mpi` Module

An MPI implementation must provide a module named `mpi` that can be `USED` in a Fortran 90 program. This module must:

- Define all named MPI constants
- Declare MPI functions that return a value.

An MPI implementation may provide in the `mpi` module other features that enhance the usability of MPI while maintaining adherence to the standard. For example, it may:

- Provide interfaces for all or for a subset of MPI routines.
- Provide `INTENT` information in these interface blocks.

Advice to implementors. The appropriate `INTENT` may be different from what is given in the MPI generic interface. Implementations must choose `INTENT` so that the function adheres to the MPI standard. (*End of advice to implementors.*)

Rationale. The intent given by the MPI generic interface is not precisely defined and does not in all cases correspond to the correct Fortran `INTENT`. For instance, receiving into a buffer specified by a datatype with absolute addresses may require associating `MPI_BOTTOM` with a dummy `OUT` argument. Moreover, “constants” such as `MPI_BOTTOM` and `MPI_STATUS_IGNORE` are not constants as defined by Fortran, but “special addresses” used in a nonstandard way. Finally, the MPI-1 generic intent is changed in several places by MPI-2. For instance, `MPI_IN_PLACE` changes the sense of an `OUT` argument to be `INOUT`. (*End of rationale.*)

Applications may use either the `mpi` module or the `mpif.h` include file. An implementation may require use of the module to prevent type mismatch errors (see below).

Advice to users. It is recommended to use the `mpi` module even if it is not necessary to use it to avoid type mismatch errors on a particular system. Using a module provides several potential advantages over using an include file. (*End of advice to users.*)

It must be possible to link together routines some of which `USE mpi` and others of which `INCLUDE mpif.h`.

No Type Mismatch Problems for Subroutines with Choice Arguments

A high quality MPI implementation should provide a mechanism to ensure that MPI choice arguments do not cause fatal compile-time or run-time errors due to type mismatch. An MPI implementation may require applications to use the `mpi` module, or require that it be compiled with a particular compiler flag, in order to avoid type mismatch problems.

Advice to implementors. In the case where the compiler does not generate errors, nothing needs to be done to the existing interface. In the case where the compiler may generate errors, a set of overloaded functions may be used. See the paper of M. Hennecke [8]. Even if the compiler does not generate errors, explicit interfaces for all routines would be useful for detecting errors in the argument list. Also, explicit interfaces which give `INTENT` information can reduce the amount of copying for `BUF(*)` arguments. (*End of advice to implementors.*)

10.2.5 Additional Support for Fortran Numeric Intrinsic Types

The routines in this section are part of Extended Fortran Support described in Section 10.2.4.

MPI-1 provides a small number of named datatypes that correspond to named intrinsic types supported by C and Fortran. These include `MPI_INTEGER`, `MPI_REAL`, `MPI_INT`, `MPI_DOUBLE`, etc., as well as the optional types `MPI_REAL4`, `MPI_REAL8`, etc. There is a one-to-one correspondence between language declarations and MPI types.

Fortran (starting with Fortran 90) provides so-called `KIND`-parameterized types. These types are declared using an intrinsic type (one of `INTEGER`, `REAL`, `COMPLEX`, `LOGICAL` and `CHARACTER`) with an optional integer `KIND` parameter that selects from among one or more variants. The specific meaning of different `KIND` values themselves are implementation dependent and not specified by the language. Fortran provides the `KIND` selection functions `selected_real_kind` for `REAL` and `COMPLEX` types, and `selected_int_kind` for `INTEGER` types that allow users to declare variables with a minimum precision or number of digits. These functions provide a portable way to declare `KIND`-parameterized `REAL`, `COMPLEX` and

INTEGER variables in Fortran. This scheme is backward compatible with Fortran 77. REAL and INTEGER Fortran variables have a default KIND if none is specified. Fortran DOUBLE PRECISION variables are of intrinsic type REAL with a non-default KIND. The following two declarations are equivalent:

```
double precision x
real(KIND(0.0d0)) x
```

MPI provides two orthogonal methods to communicate using numeric intrinsic types. The first method can be used when variables have been declared in a portable way — using default KIND or using KIND parameters obtained with the `selected_int_kind` or `selected_real_kind` functions. With this method, MPI automatically selects the correct data size (e.g., 4 or 8 bytes) and provides representation conversion in heterogeneous environments. The second method gives the user complete control over communication by exposing machine representations.

Parameterized Datatypes with Specified Precision and Exponent Range

MPI-1 provides named datatypes corresponding to standard Fortran 77 numeric types — `MPI_INTEGER`, `MPI_COMPLEX`, `MPI_REAL`, `MPI_DOUBLE_PRECISION` and `MPI_DOUBLE_COMPLEX`. MPI automatically selects the correct data size and provides representation conversion in heterogeneous environments. The mechanism described in this section extends this MPI-1 model to support portable parameterized numeric types.

The model for supporting portable parameterized types is as follows. Real variables are declared (perhaps indirectly) using `selected_real_kind(p, r)` to determine the KIND parameter, where `p` is decimal digits of precision and `r` is an exponent range. Implicitly MPI maintains a two-dimensional array of predefined MPI datatypes `D(p, r)`. `D(p, r)` is defined for each value of `(p, r)` supported by the compiler, including pairs for which one value is unspecified. Attempting to access an element of the array with an index `(p, r)` not supported by the compiler is erroneous. MPI implicitly maintains a similar array of `COMPLEX` datatypes. For integers, there is a similar implicit array related to `selected_int_kind` and indexed by the requested number of digits `r`. Note that the predefined datatypes contained in these implicit arrays are not the same as the named MPI datatypes `MPI_REAL`, etc., but a new set.

Advice to implementors. The above description is for explanatory purposes only. It is not expected that implementations will have such internal arrays. (*End of advice to implementors.*)

Advice to users. `selected_real_kind()` maps a large number of `(p,r)` pairs to a much smaller number of KIND parameters supported by the compiler. KIND parameters are not specified by the language and are not portable. From the language point of view intrinsic types of the same base type and KIND parameter are of the same type. In order to allow interoperability in a heterogeneous environment, MPI is more stringent. The corresponding MPI datatypes match if and only if they have the same `(p,r)` value (REAL and COMPLEX) or `r` value (INTEGER). Thus MPI has many more datatypes than there are fundamental language types. (*End of advice to users.*)

```
1 MPI_TYPE_CREATE_F90_REAL(p, r, newtype)
```

```
2     IN      p                precision, in decimal digits (integer)
3     IN      r                decimal exponent range (integer)
4
5     OUT     newtype          the requested MPI datatype (handle)
```

```
6
7 int MPI_Type_create_f90_real(int p, int r, MPI_Datatype *newtype)
```

```
8 MPI_TYPE_CREATE_F90_REAL(P, R, NEWTYPE, IERROR)
```

```
9     INTEGER P, R, NEWTYPE, IERROR
```

```
10
11 static MPI::Datatype MPI::Datatype::Create_f90_real(int p, int r)
```

This function returns a predefined MPI datatype that matches a **REAL** variable of **KIND selected_real_kind(p, r)**. In the model described above it returns a handle for the element **D(p, r)**. Either **p** or **r** may be omitted from calls to **selected_real_kind(p, r)** (but not both). Analogously, either **p** or **r** may be set to **MPI_UNDEFINED**. In communication, an MPI datatype **A** returned by **MPI_TYPE_CREATE_F90_REAL** matches a datatype **B** if and only if **B** was returned by **MPI_TYPE_CREATE_F90_REAL** called with the same values for **p** and **r** or **B** is a duplicate of such a datatype. Restrictions on using the returned datatype with the “external32” data representation are given on page 296.

It is erroneous to supply values for **p** and **r** not supported by the compiler.

```
22
23 MPI_TYPE_CREATE_F90_COMPLEX(p, r, newtype)
```

```
24
25     IN      p                precision, in decimal digits (integer)
26     IN      r                decimal exponent range (integer)
27
28     OUT     newtype          the requested MPI datatype (handle)
```

```
29
30 int MPI_Type_create_f90_complex(int p, int r, MPI_Datatype *newtype)
```

```
31 MPI_TYPE_CREATE_F90_COMPLEX(P, R, NEWTYPE, IERROR)
```

```
32     INTEGER P, R, NEWTYPE, IERROR
```

```
33
34 static MPI::Datatype MPI::Datatype::Create_f90_complex(int p, int r)
```

This function returns a predefined MPI datatype that matches a **COMPLEX** variable of **KIND selected_real_kind(p, r)**. Either **p** or **r** may be omitted from calls to **selected_real_kind(p, r)** (but not both). Analogously, either **p** or **r** may be set to **MPI_UNDEFINED**. Matching rules for datatypes created by this function are analogous to the matching rules for datatypes created by **MPI_TYPE_CREATE_F90_REAL**. Restrictions on using the returned datatype with the “external32” data representation are given on page 296.

It is erroneous to supply values for **p** and **r** not supported by the compiler.

```

MPI_TYPE_CREATE_F90_INTEGER(r, newtype)
IN      r                      decimal exponent range, i.e., number of decimal digits
                                (integer)
OUT     newtype                the requested MPI datatype (handle)

```

```
int MPI_Type_create_f90_integer(int r, MPI_Datatype *newtype)
```

```
MPI_TYPE_CREATE_F90_INTEGER(R, NEWTYPE, IERROR)
```

```
INTEGER R, NEWTYPE, IERROR
```

```
static MPI::Datatype MPI::Datatype::Create_f90_integer(int r)
```

This function returns a predefined MPI datatype that matches a `INTEGER` variable of `KIND selected_int_kind(r)`. Matching rules for datatypes created by this function are analogous to the matching rules for datatypes created by `MPI_TYPE_CREATE_F90_REAL`. Restrictions on using the returned datatype with the “external32” data representation are given on page 296.

It is erroneous to supply a value for `r` that is not supported by the compiler.

Example:

```

integer      longtype, quadtype
integer, parameter :: long = selected_int_kind(15)
integer(long) ii(10)
real(selected_real_kind(30)) x(10)
call MPI_TYPE_CREATE_F90_INTEGER(15, longtype, ierror)
call MPI_TYPE_CREATE_F90_REAL(30, MPI_UNDEFINED, quadtype, ierror)
...

call MPI_SEND(ii, 10, longtype, ...)
call MPI_SEND(x, 10, quadtype, ...)

```

Advice to users. The datatypes returned by the above functions are predefined datatypes. They cannot be freed; they do not need to be committed; they can be used with predefined reduction operations. There are two situations in which they behave differently syntactically, but not semantically, from the MPI named predefined datatypes.

1. `MPI_TYPE_GET_ENVELOPE` returns special combiners that allow a program to retrieve the values of `p` and `r`.
2. Because the datatypes are not named, they cannot be used as compile-time initializers or otherwise accessed before a call to one of the `MPI_TYPE_CREATE_F90_` routines.

If a variable was declared specifying a non-default `KIND` value that was not obtained with `selected_real_kind()` or `selected_int_kind()`, the only way to obtain a matching MPI datatype is to use the size-based mechanism described in the next section.

(End of advice to users.)

Rationale. The `MPI_TYPE_CREATE_F90_REAL/COMPLEX/INTEGER` interface needs as input the original range and precision values to be able to define useful and compiler-independent external (Section 9.5.2 on page 250) or user-defined (Section 9.5.3 on page 251) data representations, and in order to be able to perform automatic and efficient data conversions in a heterogeneous environment. (*End of rationale.*)

We now specify how the datatypes described in this section behave when used with the “external32” external data representation described in Section 9.5.2 on page 250.

The external32 representation specifies data formats for integer and floating point values. Integer values are represented in two’s complement big-endian format. Floating point values are represented by one of three IEEE formats. These are the IEEE “Single,” “Double” and “Double Extended” formats, requiring 4, 8 and 16 bytes of storage, respectively. For the IEEE “Double Extended” formats, MPI specifies a Format Width of 16 bytes, with 15 exponent bits, bias = +10383, 112 fraction bits, and an encoding analogous to the “Double” format.

The external32 representations of the datatypes returned by `MPI_TYPE_CREATE_F90_REAL/COMPLEX/INTEGER` are given by the following rules.

For `MPI_TYPE_CREATE_F90_REAL`:

```

if      (p > 33) or (r > 4931) then  external32 representation
                                     is undefined
else if (p > 15) or (r > 307) then  external32_size = 16
else if (p > 6) or (r > 37) then   external32_size = 8
else                                external32_size = 4

```

For `MPI_TYPE_CREATE_F90_COMPLEX`: twice the size as for `MPI_TYPE_CREATE_F90_REAL`.

For `MPI_TYPE_CREATE_F90_INTEGER`:

```

if      (r > 38) then  external32 representation is undefined
else if (r > 18) then  external32_size = 16
else if (r > 9) then   external32_size = 8
else if (r > 4) then   external32_size = 4
else if (r > 2) then   external32_size = 2
else                  external32_size = 1

```

If the external32 representation of a datatype is undefined, the result of using the datatype directly or indirectly (i.e., as part of another datatype or through a duplicated datatype) in operations that require the external32 representation is undefined. These operations include `MPI_PACK_EXTERNAL`, `MPI_UNPACK_EXTERNAL` and many `MPI_FILE` functions, when the “external32” data representation is used. The ranges for which the external32 representation is undefined are reserved for future standardization.

Support for Size-specific MPI Datatypes

MPI-1 provides named datatypes corresponding to optional Fortran 77 numeric types that contain explicit byte lengths — `MPI_REAL4`, `MPI_INTEGER8`, etc. This section describes a mechanism that generalizes this model to support all Fortran numeric intrinsic types.

We assume that for each **typeclass** (integer, real, complex) and each word size there is a unique machine representation. For every pair (**typeclass**, **n**) supported by a compiler, MPI must provide a named size-specific datatype. The name of this datatype is of the form

MPI_<TYPE>n in C and Fortran and of the form MPI::<TYPE>n in C++ where <TYPE> is one of REAL, INTEGER and COMPLEX, and n is the length in bytes of the machine representation. This datatype locally matches all variables of type (**typeclass**, n). The list of names for such types includes:

```
MPI_REAL4
MPI_REAL8
MPI_REAL16
MPI_COMPLEX8
MPI_COMPLEX16
MPI_COMPLEX32
MPI_INTEGER1
MPI_INTEGER2
MPI_INTEGER4
MPI_INTEGER8
MPI_INTEGER16
```

In MPI-1 these datatypes are all optional and correspond to the optional, nonstandard declarations supported by many Fortran compilers. In MPI-2, one datatype is required for each representation supported by the compiler. To be backward compatible with the interpretation of these types in MPI-1, we assume that the nonstandard declarations **REAL*n**, **INTEGER*n**, always create a variable whose representation is of size n. All these datatypes are predefined.

The following functions allow a user to obtain a size-specific MPI datatype for any intrinsic Fortran type.

MPI_SIZEOF(x, size)

IN	x	a Fortran variable of numeric intrinsic type (choice)
OUT	size	size of machine representation of that type (integer)

MPI_SIZEOF(X, SIZE, IERROR)

<type> X

INTEGER SIZE, IERROR

This function returns the size in bytes of the machine representation of the given variable. It is a generic Fortran routine and has a Fortran binding only.

Advice to users. This function is similar to the C and C++ *sizeof* operator but behaves slightly differently. If given an array argument, it returns the size of the base element, not the size of the whole array. (*End of advice to users.*)

Rationale. This function is not available in other languages because it would not be useful. (*End of rationale.*)

```
1 MPI_TYPE_MATCH_SIZE(typeclass, size, type)
```

```
2     IN      typeclass          generic type specifier (integer)
```

```
3     IN      size              size, in bytes, of representation (integer)
```

```
4     OUT     type              datatype with correct type, size (handle)
```

```
6
7 int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *type)
```

```
8 MPI_TYPE_MATCH_SIZE(TYPECLASS, SIZE, TYPE, IERROR)
```

```
9     INTEGER TYPECLASS, SIZE, TYPE, IERROR
```

```
10
11 static MPI::Datatype MPI::Datatype::Match_size(int typeclass, int size)
```

```
12
13     typeclass is one of MPI_TYPECLASS_REAL, MPI_TYPECLASS_INTEGER and
14 MPI_TYPECLASS_COMPLEX, corresponding to the desired typeclass. The function returns
15 an MPI datatype matching a local variable of type (typeclass, size).
```

```
16     This function returns a reference (handle) to one of the predefined named datatypes, not
17 a duplicate. This type cannot be freed. MPI_TYPE_MATCH_SIZE can be used to obtain a
18 size-specific type that matches a Fortran numeric intrinsic type by first calling MPI_SIZEOF
19 in order to compute the variable size, and then calling MPI_TYPE_MATCH_SIZE to find a
20 suitable datatype. In C and C++, one can use the C function sizeof(), instead of
21 MPI_SIZEOF. In addition, for variables of default kind the variable's size can be computed
22 by a call to MPI_TYPE_GET_EXTENT, if the typeclass is known. It is erroneous to specify
23 a size not supported by the compiler.
```

```
24
25     Rationale. This is a convenience function. Without it, it can be tedious to find the
26 correct named type. See note to implementors below. (End of rationale.)
```

```
27
28     Advice to implementors. This function could be implemented as a series of tests.
```

```
29
30 int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *rtype)
31 {
32     switch(typeclass) {
33         case MPI_TYPECLASS_REAL: switch(size) {
34             case 4: *rtype = MPI_REAL4; return MPI_SUCCESS;
35             case 8: *rtype = MPI_REAL8; return MPI_SUCCESS;
36             default: error(...);
37         }
38         case MPI_TYPECLASS_INTEGER: switch(size) {
39             case 4: *rtype = MPI_INTEGER4; return MPI_SUCCESS;
40             case 8: *rtype = MPI_INTEGER8; return MPI_SUCCESS;
41             default: error(...);          }
42         ... etc ...
43     }
44 }
```

```
45
46     (End of advice to implementors.)
47
48
```


Communication With Size-specific Types

The usual type matching rules apply to size-specific datatypes: a value sent with datatype `MPI_<TYPE>n` can be received with this same datatype on another process. Most modern computers use 2's complement for integers and IEEE format for floating point. Thus, communication using these size-specific datatypes will not entail loss of precision or truncation errors.

Advice to users. Care is required when communicating in a heterogeneous environment. Consider the following code:

```
real(selected_real_kind(5)) x(100)
call MPI_SIZEOF(x, size, ierror)
call MPI_TYPE_MATCH_SIZE(MPI_TYPECLASS_REAL, size, xtype, ierror)
if (myrank .eq. 0) then
    ... initialize x ...
    call MPI_SEND(x, xtype, 100, 1, ...)
else if (myrank .eq. 1) then
    call MPI_RECV(x, xtype, 100, 0, ...)
endif
```

This may not work in a heterogeneous environment if the value of `size` is not the same on process 1 and process 0. There should be no problem in a homogeneous environment. To communicate in a heterogeneous environment, there are at least four options. The first is to declare variables of default type and use the `MPI` datatypes for these types, e.g., declare a variable of type `REAL` and use `MPI_REAL`. The second is to use `selected_real_kind` or `selected_int_kind` and with the functions of the previous section. The third is to declare a variable that is known to be the same size on all architectures (e.g., `selected_real_kind(12)` on almost all compilers will result in an 8-byte representation). The fourth is to carefully check representation size before communication. This may require explicit conversion to a variable of size that can be communicated and handshaking between sender and receiver to agree on a size.

Note finally that using the “external32” representation for I/O requires explicit attention to the representation sizes. Consider the following code:

```
real(selected_real_kind(5)) x(100)
call MPI_SIZEOF(x, size, ierror)
call MPI_TYPE_MATCH_SIZE(MPI_TYPECLASS_REAL, size, xtype, ierror)

if (myrank .eq. 0) then
    call MPI_FILE_OPEN(MPI_COMM_SELF, 'foo',
                      MPI_MODE_CREATE+MPI_MODE_WRONLY,
                      MPI_INFO_NULL, fh, ierror)
    call MPI_FILE_SET_VIEW(fh, 0, xtype, xtype, 'external32',
                          MPI_INFO_NULL, ierror)
    call MPI_FILE_WRITE(fh, x, 100, xtype, status, ierror)
    call MPI_FILE_CLOSE(fh, ierror)
```

```
1      endif
2
3      call MPI_BARRIER(MPI_COMM_WORLD, ierror)
4
5      if (myrank .eq. 1) then
6          call MPI_FILE_OPEN(MPI_COMM_SELF, 'foo', MPI_MODE_RDONLY, &
7                               MPI_INFO_NULL, fh, ierror)
8          call MPI_FILE_SET_VIEW(fh, 0, xtype, xtype, 'external32', &
9                                   MPI_INFO_NULL, ierror)
10         call MPI_FILE_WRITE(fh, x, 100, xtype, status, ierror)
11         call MPI_FILE_CLOSE(fh, ierror)
12     endif
```

13
14
15 If processes 0 and 1 are on different machines, this code may not work as expected if
16 the **size** is different on the two machines. (*End of advice to users.*)
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

Bibliography

- [1] Rajesh Bordawekar, Juan Miguel del Rosario, and Alok Choudhary. Design and evaluation of primitives for parallel I/O. In *Proceedings of Supercomputing '93*, pages 452–461, 1993.
- [2] Juan Miguel del Rosario, Rajesh Bordawekar, and Alok Choudhary. Improved parallel I/O via a two-phase run-time access strategy. In *IPPS '93 Workshop on Input/Output in Parallel Computer Systems*, pages 56–70, 1993. Also published in *Computer Architecture News* 21(5), December 1993, pages 31–38.
- [3] Margaret A. Ellis and Bjarne Stroustrup. *The Annotated C++ Reference Manual*. Addison Wesley, 1990.
- [4] C++ Forum. Working paper for draft proposed international standard for information systems — programming language C++. Technical report, American National Standards Institute, 1995.
- [5] Message Passing Interface Forum. MPI: A Message-Passing Interface standard. *The International Journal of Supercomputer Applications and High Performance Computing*, 8, 1994.
- [6] Message Passing Interface Forum. MPI: A Message-Passing Interface standard (version 1.1). Technical report, 1995. <http://www.mpi-forum.org>.
- [7] Al Geist, Adam Beguelin, Jack Dongarra, Weicheng Jiang, Bob Manchek, and Vaidy Sunderam. *PVM: Parallel Virtual Machine—A User's Guide and Tutorial for Network Parallel Computing*. MIT Press, 1994.
- [8] Michael Hennecke. A Fortran 90 interface to MPI version 1.1. Technical Report Internal Report 63/96, Rechenzentrum, Universität Karlsruhe, D-76128 Karlsruhe, Germany, June 1996. Available via world wide web from http://www.uni-karlsruhe.de/~Michael.Hennecke/Publications/#MPI_F90.
- [9] Institute of Electrical and Electronics Engineers, New York. *IEEE Standard for Binary Floating-Point Arithmetic, ANSI/IEEE Standard 754-1985*, 1985.
- [10] International Organization for Standardization, Geneva. *Information processing — 8-bit single-byte coded graphic character sets — Part 1: Latin alphabet No. 1*, 1987.
- [11] International Organization for Standardization, Geneva. *Information technology — Portable Operating System Interface (POSIX) — Part 1: System Application Program Interface (API) [C Language]*, December 1996.

- [12] Charles H. Koelbel, David B. Loveman, Robert S. Schreiber, Guy L. Steele Jr., and Mary E. Zosel. *The High Performance Fortran Handbook*. MIT Press, 1993.
- [13] David Kotz. Disk-directed I/O for MIMD multiprocessors. In *Proceedings of the 1994 Symposium on Operating Systems Design and Implementation*, pages 61–74, November 1994. Updated as Dartmouth TR PCS-TR94-226 on November 8, 1994.
- [14] S. J. Lefflet, R. S. Fabry, W. N. Joy, P. Lapsley, S. Miller, and C. Torek. An advanced 4.4BSD interprocess communication tutorial, Unix programmer’s supplementary documents (PSD) 21. Technical report, Computer Systems Research Group, Department of Electrical Engineering and Computer Science, University of California, Berkeley, 1993. Also available at <http://www.netbsd.org/Documentation/lite2/psd/>.
- [15] Bill Nitzberg. Performance of the iPSC/860 Concurrent File System. Technical Report RND-92-020, NAS Systems Division, NASA Ames, December 1992.
- [16] William J. Nitzberg. *Collective Parallel I/O*. PhD thesis, Department of Computer and Information Science, University of Oregon, December 1995.
- [17] *4.4BSD Programmer’s Supplementary Documents (PSD)*. O’Reilly and Associates, 1994.
- [18] Perry Partow and Dennis Cattel. Scalable Programming Environment. Technical Report 1672, Naval Command Control and Ocean Surveillance Center (NRAD), September 1994.
- [19] K. E. Seamons, Y. Chen, P. Jones, J. Jozwiak, and M. Winslett. Server-directed collective I/O in Panda. In *Proceedings of Supercomputing ’95*, December 1995.
- [20] Anthony Skjellum, Nathan E. Doss, and Kishore Viswanathan. Inter-communicator extensions to MPI in the MPIX (MPI eXtension) Library. Technical Report MSU-940722, Mississippi State University — Dept. of Computer Science, April 1994. <http://www.erc.msstate.edu/mpi/mpix.html>.
- [21] Anthony Skjellum, Ziyang Lu, Purushotham V. Bangalore, and Nathan E. Doss. Explicit parallel programming in C++ based on the message-passing interface (MPI). In Gregory V. Wilson, editor, *Parallel Programming Using C++*, Engineering Computation Series. MIT Press, July 1996. ISBN 0-262-73118-5.
- [22] Rajeev Thakur and Alok Choudhary. An Extended Two-Phase Method for Accessing Sections of Out-of-Core Arrays. *Scientific Programming*, 5(4):301–317, Winter 1996.
- [23] *The Unicode Standard, Version 2.0*. Addison-Wesley, 1996. ISBN 0-201-48345-9.

Annex A

Language Binding

A.1 Introduction

This annex summarizes the specific bindings for Fortran, C, and C++. First the constants, error codes, info keys, and info values are presented. Second, the MPI-1.2 bindings are given. Third, the MPI-2 bindings are given.

A.2 Defined Values and Handles

A.2.1 Defined Constants

The C and Fortran name is listed in the left column and the C++ name is listed in the right column.

Return Codes

MPI_ERR_ACCESS	MPI::ERR_ACCESS
MPI_ERR_AMODE	MPI::ERR_AMODE
MPI_ERR_ASSERT	MPI::ERR_ASSERT
MPI_ERR_BAD_FILE	MPI::ERR_BAD_FILE
MPI_ERR_BASE	MPI::ERR_BASE
MPI_ERR_CONVERSION	MPI::ERR_CONVERSION
MPI_ERR_DISP	MPI::ERR_DISP
MPI_ERR_DUP_DATAREP	MPI::ERR_DUP_DATAREP
MPI_ERR_FILE_EXISTS	MPI::ERR_FILE_EXISTS
MPI_ERR_FILE_IN_USE	MPI::ERR_FILE_IN_USE
MPI_ERR_FILE	MPI::ERR_FILE
MPI_ERR_INFO_KEY	MPI::ERR_INFO_VALUE
MPI_ERR_INFO_NOKEY	MPI::ERR_INFO_NOKEY
MPI_ERR_INFO_VALUE	MPI::ERR_INFO_KEY
MPI_ERR_INFO	MPI::ERR_INFO
MPI_ERR_IO	MPI::ERR_IO
MPI_ERR_KEYVAL	MPI::ERR_KEYVAL
MPI_ERR_LOCKTYPE	MPI::ERR_LOCKTYPE
MPI_ERR_NAME	MPI::ERR_NAME
MPI_ERR_NO_MEM	MPI::ERR_NO_MEM
MPI_ERR_NOT_SAME	MPI::ERR_NOT_SAME
MPI_ERR_NO_SPACE	MPI::ERR_NO_SPACE
MPI_ERR_NO_SUCH_FILE	MPI::ERR_NO_SUCH_FILE
MPI_ERR_PORT	MPI::ERR_PORT
MPI_ERR_QUOTA	MPI::ERR_QUOTA
MPI_ERR_READ_ONLY	MPI::ERR_READ_ONLY
MPI_ERR_RMA_CONFLICT	MPI::ERR_RMA_CONFLICT
MPI_ERR_RMA_SYNC	MPI::ERR_RMA_SYNC
MPI_ERR_SERVICE	MPI::ERR_SERVICE
MPI_ERR_SIZE	MPI::ERR_SIZE
MPI_ERR_SPAWN	MPI::ERR_SPAWN
MPI_ERR_UNSUPPORTED_DATAREP	MPI::ERR_UNSUPPORTED_DATAREP
MPI_ERR_UNSUPPORTED_OPERATION	MPI::ERR_UNSUPPORTED_OPERATION
MPI_ERR_WIN	MPI::ERR_WIN

Assorted Constants

MPI_IN_PLACE	MPI::IN_PLACE
MPI_LOCK_EXCLUSIVE	MPI::LOCK_EXCLUSIVE
MPI_LOCK_SHARED	MPI::LOCK_SHARED
MPI_ROOT	MPI::ROOT

Variable Address Size (Fortran only)

MPI_ADDRESS_KIND	Not defined for C++
MPI_INTEGER_KIND	Not defined for C++
MPI_OFFSET_KIND	Not defined for C++

Maximum Sizes for Strings

MPI_MAX_DATAREP_STRING	MPI::MAX_DATAREP_STRING
MPI_MAX_INFO_KEY	MPI::MAX_INFO_KEY
MPI_MAX_INFO_VAL	MPI::MAX_INFO_VAL
MPI_MAX_OBJECT_NAME	MPI::MAX_OBJECT_NAME
MPI_MAX_PORT_NAME	MPI::MAX_PORT_NAME

Named Predefined Datatypes

MPI_WCHAR	MPI::WCHAR
-----------	------------

C and C++ (no Fortran) Named Predefined Datatypes

MPI_Fint	MPI::Fint
----------	-----------

Optional C and C++ (no Fortran) Named Predefined Datatypes

MPI_UNSIGNED_LONG_LONG	MPI::UNSIGNED_LONG_LONG
MPI_SIGNED_CHAR	MPI::SIGNED_CHAR

Predefined Attribute Keys

MPI_APPNUM	MPI::APPNUM
MPI_LASTUSED_CODE	MPI::LASTUSED_CODE
MPI_UNIVERSE_SIZE	MPI::UNIVERSE_SIZE
MPI_WIN_BASE	MPI::WIN_BASE
MPI_WIN_DISP_UNIT	MPI::WIN_DISP_UNIT
MPI_WIN_SIZE	MPI::WIN_SIZE

Collective Operations

MPI_REPLACE	MPI::REPLACE
-------------	--------------

Null Handles

MPI_FILE_NULL	MPI::FILE_NULL
MPI_INFO_NULL	MPI::INFO_NULL
MPI_WIN_NULL	MPI::WIN_NULL

Mode Constants

MPI_MODE_APPEND	MPI::MODE_APPEND
MPI_MODE_CREATE	MPI::MODE_CREATE
MPI_MODE_DELETE_ON_CLOSE	MPI::MODE_DELETE_ON_CLOSE
MPI_MODE_EXCL	MPI::MODE_EXCL
MPI_MODE_NOCHECK	MPI::MODE_NOCHECK
MPI_MODE_NOPRECEDE	MPI::MODE_NOPRECEDE
MPI_MODE_NOPUT	MPI::MODE_NOPUT
MPI_MODE_NOSTORE	MPI::MODE_NOSTORE
MPI_MODE_NOSUCCEED	MPI::MODE_NOSUCCEED
MPI_MODE_RDONLY	MPI::MODE_RDONLY
MPI_MODE_RDWR	MPI::MODE_RDWR
MPI_MODE_SEQUENTIAL	MPI::MODE_SEQUENTIAL
MPI_MODE_UNIQUE_OPEN	MPI::MODE_UNIQUE_OPEN
MPI_MODE_WRONLY	MPI::MODE_WRONLY

Datatype Decoding Constants

MPI_COMBINER_CONTIGUOUS	MPI::COMBINER_CONTIGUOUS
MPI_COMBINER_DARRAY	MPI::COMBINER_DARRAY
MPI_COMBINER_DUP	MPI::COMBINER_DUP
MPI_COMBINER_F90_COMPLEX	MPI::COMBINER_F90_COMPLEX
MPI_COMBINER_F90_INTEGER	MPI::COMBINER_F90_INTEGER
MPI_COMBINER_F90_REAL	MPI::COMBINER_F90_REAL
MPI_COMBINER_HINDEXED_INTEGER	MPI::COMBINER_HINDEXED_INTEGER
MPI_COMBINER_HINDEXED	MPI::COMBINER_HINDEXED
MPI_COMBINER_HVECTOR_INTEGER	MPI::COMBINER_HVECTOR_INTEGER
MPI_COMBINER_HVECTOR	MPI::COMBINER_HVECTOR
MPI_COMBINER_INDEXED_BLOCK	MPI::COMBINER_INDEXED_BLOCK
MPI_COMBINER_INDEXED	MPI::COMBINER_INDEXED
MPI_COMBINER_NAMED	MPI::COMBINER_NAMED
MPI_COMBINER_RESIZED	MPI::COMBINER_RESIZED
MPI_COMBINER_STRUCT_INTEGER	MPI::COMBINER_STRUCT_INTEGER
MPI_COMBINER_STRUCT	MPI::COMBINER_STRUCT
MPI_COMBINER_SUBARRAY	MPI::COMBINER_SUBARRAY
MPI_COMBINER_VECTOR	MPI::COMBINER_VECTOR

Threads Constants

MPI_THREAD_FUNNELED	MPI::THREAD_FUNNELED
MPI_THREAD_MULTIPLE	MPI::THREAD_MULTIPLE
MPI_THREAD_SERIALIZED	MPI::THREAD_SERIALIZED
MPI_THREAD_SINGLE	MPI::THREAD_SINGLE

File Operation Constants

MPI_DISPLACEMENT_CURRENT	MPI::DISPLACEMENT_CURRENT
MPI_DISTRIBUTE_BLOCK	MPI::DISTRIBUTE_BLOCK
MPI_DISTRIBUTE_CYCLIC	MPI::DISTRIBUTE_CYCLIC
MPI_DISTRIBUTE_DFLT_DARG	MPI::DISTRIBUTE_DFLT_DARG
MPI_DISTRIBUTE_NONE	MPI::DISTRIBUTE_NONE
MPI_ORDER_C	MPI::ORDER_C
MPI_ORDER_FORTRAN	MPI::ORDER_FORTRAN
MPI_SEEK_CUR	MPI::SEEK_CUR
MPI_SEEK_END	MPI::SEEK_END
MPI_SEEK_SET	MPI::SEEK_SET

F90 Datatype Matching Constants

MPI_TYPECLASS_COMPLEX	MPI::TYPECLASS_COMPLEX
MPI_TYPECLASS_INTEGER	MPI::TYPECLASS_INTEGER
MPI_TYPECLASS_REAL	MPI::TYPECLASS_REAL

Handles to Assorted Structures in C and C++ (no Fortran)

MPI_File	MPI::File
MPI_Info	MPI::Info
MPI_Win	MPI::Win

Constants Specifying Empty or Ignored Input

MPI_ARGVS_NULL	MPI::ARGVS_NULL
MPI_ARGV_NULL	MPI::ARGV_NULL
MPI_ERRCODES_IGNORE	Not defined for C++
MPI_STATUSES_IGNORE	Not defined for C++
MPI_STATUS_IGNORE	Not defined for C++

C Constants Specifying Ignored Input (no C++ or Fortran)

MPI_F_STATUSES_IGNORE	Not defined for C++
MPI_F_STATUS_IGNORE	Not defined for C++

C and C++ cpp Constants and Fortran Parameters

MPI_SUBVERSION
MPI_VERSION

A.2.2 Info Keys

access_style
appnum
arch
cb_block_size

cb_buffer_size
cb_nodes
chunked_item
chunked_size
chunked
collective_buffering
file_perm
filename
file
host
io_node_list
ip_address
ip_port
nb_proc
no_locks
num_io_nodes
path
soft
striping_factor
striping_unit
wdir

A.2.3 Info Values

false
random
read_mostly
read_once
reverse_sequential
sequential
true
write_mostly
write_once

A.3 MPI-1.2 C Bindings

```
int MPI_Get_version(int *version, int *subversion)
```

A.4 MPI-1.2 Fortran Bindings

```
MPI_GET_VERSION(VERSION, SUBVERSION, IERROR)  
    INTEGER VERSION, SUBVERSION, IERROR
```

A.5 MPI-1.2 C++ Bindings

See Section B.11.

A.6 MPI-2 C Bindings

A.6.1 Miscellany

```

int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)
MPI_Fint MPI_Comm_c2f(MPI_Comm comm)
int MPI_Comm_create_errhandler(MPI_Comm_errhandler_fn *function,
                               MPI_Errhandler *errhandler)
MPI_Comm MPI_Comm_f2c(MPI_Fint comm)
int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)
int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
MPI_Fint MPI_File_c2f(MPI_File file)
int MPI_File_create_errhandler(MPI_File_errhandler_fn *function,
                               MPI_Errhandler *errhandler)
MPI_File MPI_File_f2c(MPI_Fint file)
int MPI_File_get_errhandler(MPI_File file, MPI_Errhandler *errhandler)
int MPI_File_set_errhandler(MPI_File file, MPI_Errhandler errhandler)
int MPI_Finalized(int *flag)
int MPI_Free_mem(void *base)
int MPI_Get_address(void *location, MPI_Aint *address)
MPI_Fint MPI_Group_c2f(MPI_Group group)
MPI_Group MPI_Group_f2c(MPI_Fint group)
MPI_Fint MPI_Info_c2f(MPI_Info info)
int MPI_Info_create(MPI_Info *info)
int MPI_Info_delete(MPI_Info info, char *key)
int MPI_Info_dup(MPI_Info info, MPI_Info *newinfo)
MPI_Info MPI_Info_f2c(MPI_Fint info)
int MPI_Info_free(MPI_Info *info)
int MPI_Info_get(MPI_Info info, char *key, int valuelen, char *value,
                int *flag)
int MPI_Info_get_nkeys(MPI_Info info, int *nkeys)

```

```

1  int MPI_Info_get_nthkey(MPI_Info info, int n, char *key)
2
3  int MPI_Info_get_valuelen(MPI_Info info, char *key, int *valuelen,
4                          int *flag)
5
6  int MPI_Info_set(MPI_Info info, char *key, char *value)
7
8  MPI_Fint MPI_Op_c2f(MPI_Op op)
9
10 MPI_Op MPI_Op_f2c(MPI_Fint op)
11
12 int MPI_Pack_external(char *datarep, void *inbuf, int incout,
13                     MPI_Datatype datatype, void *outbuf, MPI_Aint outsize,
14                     MPI_Aint *position)
15
16 int MPI_Pack_external_size(char *datarep, int incout,
17                           MPI_Datatype datatype, MPI_Aint *size)
18
19 MPI_Fint MPI_Request_c2f(MPI_Request request)
20
21 MPI_Request MPI_Request_f2c(MPI_Fint request)
22
23 int MPI_Request_get_status(MPI_Request request, int *flag,
24                           MPI_Status *status)
25
26 int MPI_Status_c2f(MPI_Status *c_status, MPI_Fint *f_status)
27
28 int MPI_Status_f2c(MPI_Fint *f_status, MPI_Status *c_status)
29
30 MPI_Fint MPI_Type_c2f(MPI_Datatype datatype)
31
32 int MPI_Type_create_darray(int size, int rank, int ndims,
33                           int array_of_gsizes[], int array_of_distribs[], int
34                           array_of_dargs[], int array_of_psize[], int order,
35                           MPI_Datatype oldtype, MPI_Datatype *newtype)
36
37 int MPI_Type_create_hindexed(int count, int array_of_blocklengths[],
38                             MPI_Aint array_of_displacements[], MPI_Datatype oldtype,
39                             MPI_Datatype *newtype)
40
41 int MPI_Type_create_hvector(int count, int blocklength, MPI_Aint stride,
42                             MPI_Datatype oldtype, MPI_Datatype *newtype)
43
44 int MPI_Type_create_indexed_block(int count, int blocklength,
45                                  int array_of_displacements[], MPI_Datatype oldtype,
46                                  MPI_Datatype *newtype)
47
48 int MPI_Type_create_resized(MPI_Datatype oldtype, MPI_Aint lb, MPI_Aint
49                             extent, MPI_Datatype *newtype)
50
51 int MPI_Type_create_struct(int count, int array_of_blocklengths[],
52                           MPI_Aint array_of_displacements[],
53                           MPI_Datatype array_of_types[], MPI_Datatype *newtype)
54
55 int MPI_Type_create_subarray(int ndims, int array_of_sizes[],
56                             int array_of_subsizes[], int array_of_starts[], int order,
57                             MPI_Datatype oldtype, MPI_Datatype *newtype)

```

```

        MPI_Datatype oldtype, MPI_Datatype *newtype)
1
MPI_Datatype MPI_Type_f2c(MPI_Fint datatype)
2
3
int MPI_Type_get_extent(MPI_Datatype datatype, MPI_Aint *lb,
4
5
6
7
8
MPI_Aint *extent)
9
int MPI_Type_get_true_extent(MPI_Datatype datatype, MPI_Aint *true_lb,
10
11
12
MPI_Aint *true_extent)
13
int MPI_Unpack_external(char *datarep, void *inbuf, MPI_Aint insize,
14
15
16
MPI_Aint *position, void *outbuf, int outcount,
17
18
19
MPI_Datatype datatype)
20
21
MPI_Fint MPI_Win_c2f(MPI_Win win)
22
23
int MPI_Win_create_errhandler(MPI_Win_errhandler_fn *function, MPI_Errhandler
24
25
26
*errhandler)
27
28
MPI_Win MPI_Win_f2c(MPI_Fint win)
29
30
int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
31
32
int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

A.6.2 Process Creation and Management

```

int MPI_Close_port(char *port_name)
24
25
int MPI_Comm_accept(char *port_name, MPI_Info info, int root, MPI_Comm comm,
26
27
28
MPI_Comm *newcomm)
29
30
int MPI_Comm_connect(char *port_name, MPI_Info info, int root,
31
32
33
MPI_Comm comm, MPI_Comm *newcomm)
34
35
int MPI_Comm_disconnect(MPI_Comm *comm)
36
37
int MPI_Comm_get_parent(MPI_Comm *parent)
38
39
int MPI_Comm_join(int fd, MPI_Comm *intercomm)
40
41
42
int MPI_Comm_spawn(char *command, char *argv[], int maxprocs, MPI_Info info,
43
44
45
int root, MPI_Comm comm, MPI_Comm *intercomm,
46
47
48
int array_of_errcodes[])
49
50
int MPI_Comm_spawn_multiple(int count, char *array_of_commands[],
51
52
53
char **array_of_argv[], int array_of_maxprocs[],
54
55
56
MPI_Info array_of_info[], int root, MPI_Comm comm,
57
58
59
MPI_Comm *intercomm, int array_of_errcodes[])
60
61
int MPI_Lookup_name(char *service_name, MPI_Info info, char *port_name)
62
63
int MPI_Open_port(MPI_Info info, char *port_name)
64
65
int MPI_Publish_name(char *service_name, MPI_Info info, char *port_name)
66
67
int MPI_Unpublish_name(char *service_name, MPI_Info info, char *port_name)
68
69

```

A.6.3 One-Sided Communications

```

1  int MPI_Accumulate(void *origin_addr, int origin_count,
2                      MPI_Datatype origin_datatype, int target_rank,
3                      MPI_Aint target_disp, int target_count,
4                      MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)
5
6
7  int MPI_Get(void *origin_addr, int origin_count, MPI_Datatype
8              origin_datatype, int target_rank, MPI_Aint target_disp, int
9              target_count, MPI_Datatype target_datatype, MPI_Win win)
10
11 int MPI_Put(void *origin_addr, int origin_count, MPI_Datatype
12             origin_datatype, int target_rank, MPI_Aint target_disp, int
13             target_count, MPI_Datatype target_datatype, MPI_Win win)
14
15 int MPI_Win_complete(MPI_Win win)
16
17 int MPI_Win_create(void *base, MPI_Aint size, int disp_unit, MPI_Info info,
18                   MPI_Comm comm, MPI_Win *win)
19
20 int MPI_Win_fence(int assert, MPI_Win win)
21
22 int MPI_Win_free(MPI_Win *win)
23
24 int MPI_Win_get_group(MPI_Win win, MPI_Group *group)
25
26 int MPI_Win_lock(int lock_type, int rank, int assert, MPI_Win win)
27
28 int MPI_Win_post(MPI_Group group, int assert, MPI_Win win)
29
30 int MPI_Win_start(MPI_Group group, int assert, MPI_Win win)
31
32 int MPI_Win_test(MPI_Win win, int *flag)
33
34 int MPI_Win_unlock(int rank, MPI_Win win)
35
36 int MPI_Win_wait(MPI_Win win)

```

A.6.4 Extended Collective Operations

```

34 int MPI_Alltoallw(void *sendbuf, int sendcounts[], int sdispls[],
35                  MPI_Datatype sendtypes[], void *recvbuf, int recvcounts[],
36                  int rdispls[], MPI_Datatype recvtypes[], MPI_Comm comm)
37
38 int MPI_Exscan(void *sendbuf, void *recvbuf, int count,
39               MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
40

```

A.6.5 External Interfaces

```

43 int MPI_Add_error_class(int *errorclass)
44
45 int MPI_Add_error_code(int errorclass, int *errorcode)
46
47 int MPI_Add_error_string(int errorcode, char *string)
48
49 int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)

```

```

int MPI_Comm_create_keyval(MPI_Comm_copy_attr_function *comm_copy_attr_fn, 1
                           MPI_Comm_delete_attr_function *comm_delete_attr_fn, 2
                           int *comm_keyval, void *extra_state) 3
4
int MPI_Comm_delete_attr(MPI_Comm comm, int comm_keyval) 5
6
int MPI_Comm_free_keyval(int *comm_keyval) 7
8
int MPI_Comm_get_attr(MPI_Comm comm, int comm_keyval, void *attribute_val, 9
                      int *flag) 10
11
int MPI_Comm_get_name(MPI_Comm comm, char *comm_name, int *resultlen) 12
13
int MPI_Comm_set_attr(MPI_Comm comm, int comm_keyval, void *attribute_val) 14
15
int MPI_Comm_set_name(MPI_Comm comm, char *comm_name) 16
17
int MPI_File_call_errhandler(MPI_File fh, int errorcode) 18
19
int MPI_Grequest_complete(MPI_Request request) 20
21
int MPI_Grequest_start(MPI_Grequest_query_function *query_fn, 22
                      MPI_Grequest_free_function *free_fn, 23
                      MPI_Grequest_cancel_function *cancel_fn, void *extra_state, 24
                      MPI_Request *request) 25
26
int MPI_Init_thread(int *argc, char *((*argv)[]), int required, 27
                   int *provided) 28
29
int MPI_Is_thread_main(int *flag) 30
31
int MPI_Query_thread(int *provided) 32
33
int MPI_Status_set_cancelled(MPI_Status *status, int flag) 34
35
int MPI_Status_set_elements(MPI_Status *status, MPI_Datatype datatype, 36
                           int count) 37
38
int MPI_Type_create_keyval(MPI_Type_copy_attr_function *type_copy_attr_fn, 39
                           MPI_Type_delete_attr_function *type_delete_attr_fn, 40
                           int *type_keyval, void *extra_state) 41
42
int MPI_Type_delete_attr(MPI_Datatype type, int type_keyval) 43
44
int MPI_Type_dup(MPI_Datatype type, MPI_Datatype *newtype) 45
46
int MPI_Type_free_keyval(int *type_keyval) 47
48
int MPI_Type_get_attr(MPI_Datatype type, int type_keyval, void
                      *attribute_val, int *flag)
49
int MPI_Type_get_contents(MPI_Datatype datatype, int max_integers,
                          int max_addresses, int max_datatypes, int array_of_integers[],
                          MPI_Aint array_of_addresses[],
                          MPI_Datatype array_of_datatypes[])
50
int MPI_Type_get_envelope(MPI_Datatype datatype, int *num_integers,
                          int *num_addresses, int *num_datatypes)
51

```

```

1         int *num_addresses, int *num_datatypes, int *combiner)
2
3 int MPI_Type_get_name(MPI_Datatype type, char *type_name, int *resultlen)
4
5 int MPI_Type_set_attr(MPI_Datatype type, int type_keyval,
6         void *attribute_val)
7
8 int MPI_Type_set_name(MPI_Datatype type, char *type_name)
9
10 int MPI_Win_create_keyval(MPI_Win_copy_attr_function *win_copy_attr_fn,
11         MPI_Win_delete_attr_function *win_delete_attr_fn,
12         int *win_keyval, void *extra_state)
13
14 int MPI_Win_delete_attr(MPI_Win win, int win_keyval)
15
16 int MPI_Win_free_keyval(int *win_keyval)
17
18 int MPI_Win_get_attr(MPI_Win win, int win_keyval, void *attribute_val,
19         int *flag)
20
21 int MPI_Win_get_name(MPI_Win win, char *win_name, int *resultlen)
22
23 int MPI_Win_set_attr(MPI_Win win, int win_keyval, void *attribute_val)
24
25 int MPI_Win_set_name(MPI_Win win, char *win_name)

```

A.6.6 I/O

```

26 int MPI_File_close(MPI_File *fh)
27
28 int MPI_File_delete(char *filename, MPI_Info info)
29
30 int MPI_File_get_amode(MPI_File fh, int *amode)
31
32 int MPI_File_get_atomicity(MPI_File fh, int *flag)
33
34 int MPI_File_get_byte_offset(MPI_File fh, MPI_Offset offset,
35         MPI_Offset *disp)
36
37 int MPI_File_get_group(MPI_File fh, MPI_Group *group)
38
39 int MPI_File_get_info(MPI_File fh, MPI_Info *info_used)
40
41 int MPI_File_get_position(MPI_File fh, MPI_Offset *offset)
42
43 int MPI_File_get_position_shared(MPI_File fh, MPI_Offset *offset)
44
45 int MPI_File_get_size(MPI_File fh, MPI_Offset *size)
46
47 int MPI_File_get_type_extent(MPI_File fh, MPI_Datatype datatype,
48         MPI_Aint *extent)
49
50 int MPI_File_get_view(MPI_File fh, MPI_Offset *disp, MPI_Datatype *etype,
51         MPI_Datatype *filetype, char *datarep)

```



```

int MPI_File_iread(MPI_File fh, void *buf, int count, MPI_Datatype datatype, 1
                    MPI_Request *request) 2
int MPI_File_iread_at(MPI_File fh, MPI_Offset offset, void *buf, int count, 3
                    MPI_Datatype datatype, MPI_Request *request) 4
int MPI_File_iread_shared(MPI_File fh, void *buf, int count, 5
                    MPI_Datatype datatype, MPI_Request *request) 6
int MPI_File_iwrite(MPI_File fh, void *buf, int count, 7
                    MPI_Datatype datatype, MPI_Request *request) 8
int MPI_File_iwrite_at(MPI_File fh, MPI_Offset offset, void *buf, int count, 9
                    MPI_Datatype datatype, MPI_Request *request) 10
int MPI_File_iwrite_shared(MPI_File fh, void *buf, int count, 11
                    MPI_Datatype datatype, MPI_Request *request) 12
int MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info, 13
                    MPI_File *fh) 14
int MPI_File_preallocate(MPI_File fh, MPI_Offset size) 15
int MPI_File_read(MPI_File fh, void *buf, int count, MPI_Datatype datatype, 16
                    MPI_Status *status) 17
int MPI_File_read_all(MPI_File fh, void *buf, int count, 18
                    MPI_Datatype datatype, MPI_Status *status) 19
int MPI_File_read_all_begin(MPI_File fh, void *buf, int count, 20
                    MPI_Datatype datatype) 21
int MPI_File_read_all_end(MPI_File fh, void *buf, MPI_Status *status) 22
int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count, 23
                    MPI_Datatype datatype, MPI_Status *status) 24
int MPI_File_read_at_all(MPI_File fh, MPI_Offset offset, void *buf, 25
                    int count, MPI_Datatype datatype, MPI_Status *status) 26
int MPI_File_read_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf, 27
                    int count, MPI_Datatype datatype) 28
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 29
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 30
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 31
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 32
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 33
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 34
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 35
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 36
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 37
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 38
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 39
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 40
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 41
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 42
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 43
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 44
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 45
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 46
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 47
int MPI_File_read_at_all_end(MPI_File fh, void *buf, MPI_Status *status) 48

```

```

1  int MPI_File_seek_shared(MPI_File fh, MPI_Offset offset, int whence)
2
3  int MPI_File_set_atomicity(MPI_File fh, int flag)
4
5  int MPI_File_set_info(MPI_File fh, MPI_Info info)
6
7  int MPI_File_set_size(MPI_File fh, MPI_Offset size)
8
9  int MPI_File_set_view(MPI_File fh, MPI_Offset disp, MPI_Datatype etype,
10                        MPI_Datatype filetype, char *datarep, MPI_Info info)
11
12 int MPI_File_sync(MPI_File fh)
13
14 int MPI_File_write(MPI_File fh, void *buf, int count, MPI_Datatype datatype,
15                   MPI_Status *status)
16
17 int MPI_File_write_all(MPI_File fh, void *buf, int count,
18                       MPI_Datatype datatype, MPI_Status *status)
19
20 int MPI_File_write_all_begin(MPI_File fh, void *buf, int count,
21                             MPI_Datatype datatype)
22
23 int MPI_File_write_all_end(MPI_File fh, void *buf, MPI_Status *status)
24
25 int MPI_File_write_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
26                      MPI_Datatype datatype, MPI_Status *status)
27
28 int MPI_File_write_at_all(MPI_File fh, MPI_Offset offset, void *buf,
29                          int count, MPI_Datatype datatype, MPI_Status *status)
30
31 int MPI_File_write_at_all_begin(MPI_File fh, MPI_Offset offset, void *buf,
32                                int count, MPI_Datatype datatype)
33
34 int MPI_File_write_at_all_end(MPI_File fh, void *buf, MPI_Status *status)
35
36 int MPI_File_write_ordered(MPI_File fh, void *buf, int count,
37                           MPI_Datatype datatype, MPI_Status *status)
38
39 int MPI_File_write_ordered_begin(MPI_File fh, void *buf, int count,
40                                 MPI_Datatype datatype)
41
42 int MPI_File_write_ordered_end(MPI_File fh, void *buf, MPI_Status *status)
43
44 int MPI_File_write_shared(MPI_File fh, void *buf, int count,
45                          MPI_Datatype datatype, MPI_Status *status)
46
47 int MPI_Register_datarep(char *datarep,
48                         MPI_Datarep_conversion_function *read_conversion_fn,
49                         MPI_Datarep_conversion_function *write_conversion_fn,
50                         MPI_Datarep_extent_function *dtype_file_extent_fn,
51                         void *extra_state)

```

A.6.7 Language Bindings

```

1  int MPI_Type_create_f90_complex(int p, int r, MPI_Datatype *newtype)

```

```

int MPI_Type_create_f90_integer(int r, MPI_Datatype *newtype)
int MPI_Type_create_f90_real(int p, int r, MPI_Datatype *newtype)
int MPI_Type_match_size(int typeclass, int size, MPI_Datatype *type)

```

A.6.8 User Defined Functions

```

typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm, int comm_keyval,
    void *extra_state, void *attribute_val_in,
    void *attribute_val_out, int *flag);
typedef int MPI_Comm_delete_attr_function(MPI_Comm comm, int comm_keyval,
    void *attribute_val, void *extra_state);
typedef void MPI_Comm_errhandler_fn(MPI_Comm *, int *, ...);
typedef int MPI_Datarep_conversion_function(void *userbuf,
    MPI_Datatype datatype, int count, void *filebuf,
    MPI_Offset position, void *extra_state);
typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
    MPI_Aint *file_extent, void *extra_state);
typedef void MPI_File_errhandler_fn(MPI_File *, int *, ...);
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);
typedef int MPI_Grequest_free_function(void *extra_state);
typedef int MPI_Grequest_query_function(void *extra_state,
    MPI_Status *status);
typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
    int type_keyval, void *extra_state, void *attribute_val_in,
    void *attribute_val_out, int *flag);
typedef int MPI_Type_delete_attr_function(MPI_Datatype type, int type_keyval,
    void *attribute_val, void *extra_state);
typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
    void *extra_state, void *attribute_val_in,
    void *attribute_val_out, int *flag);
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
    void *attribute_val, void *extra_state);
typedef void MPI_Win_errhandler_fn(MPI_Win *, int *, ...);

```

A.7 MPI-2 Fortran Bindings

A.7.1 Miscellany

```
MPI_ALLOC_MEM(SIZE, INFO, BASEPTR, IERROR)
```

```

1      INTEGER INFO, IERROR
2      INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR
3
4      MPI_COMM_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)
5      EXTERNAL FUNCTION
6      INTEGER ERRHANDLER, IERROR
7
8      MPI_COMM_GET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
9      INTEGER COMM, ERRHANDLER, IERROR
10
11     MPI_COMM_SET_ERRHANDLER(COMM, ERRHANDLER, IERROR)
12     INTEGER COMM, ERRHANDLER, IERROR
13
14     MPI_FILE_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)
15     EXTERNAL FUNCTION
16     INTEGER ERRHANDLER, IERROR
17
18     MPI_FILE_GET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
19     INTEGER FILE, ERRHANDLER, IERROR
20
21     MPI_FILE_SET_ERRHANDLER(FILE, ERRHANDLER, IERROR)
22     INTEGER FILE, ERRHANDLER, IERROR
23
24     MPI_FINALIZED(FLAG, IERROR)
25     LOGICAL FLAG
26     INTEGER IERROR
27
28     MPI_FREE_MEM(BASE, IERROR)
29     <type> BASE(*)
30     INTEGER IERROR
31
32     MPI_GET_ADDRESS(LOCATION, ADDRESS, IERROR)
33     <type> LOCATION(*)
34     INTEGER IERROR
35     INTEGER(KIND=MPI_ADDRESS_KIND) ADDRESS
36
37     MPI_INFO_CREATE(INFO, IERROR)
38     INTEGER INFO, IERROR
39
40     MPI_INFO_DELETE(INFO, KEY, IERROR)
41     INTEGER INFO, IERROR
42     CHARACTER*(*) KEY
43
44     MPI_INFO_DUP(INFO, NEWINFO, IERROR)
45     INTEGER INFO, NEWINFO, IERROR
46
47     MPI_INFO_FREE(INFO, IERROR)
48     INTEGER INFO, IERROR
49
50     MPI_INFO_GET(INFO, KEY, VALUELEN, VALUE, FLAG, IERROR)
51     INTEGER INFO, VALUELEN, IERROR
52     CHARACTER*(*) KEY, VALUE
53     LOGICAL FLAG
54
55     MPI_INFO_GET_NKEYS(INFO, NKEYS, IERROR)

```

```

    INTEGER INFO, NKEYS, IERROR
1
MPI_INFO_GET_NTHKEY(INFO, N, KEY, IERROR)
2
    INTEGER INFO, N, IERROR
3
    CHARACTER*(*) KEY
4
5
MPI_INFO_GET_VALUELEN(INFO, KEY, VALUELEN, FLAG, IERROR)
6
    INTEGER INFO, VALUELEN, IERROR
7
    LOGICAL FLAG
8
    CHARACTER*(*) KEY
9
10
MPI_INFO_SET(INFO, KEY, VALUE, IERROR)
11
    INTEGER INFO, IERROR
12
    CHARACTER*(*) KEY, VALUE
13
MPI_PACK_EXTERNAL(DATAREP, INBUF, INCOUNT, DATATYPE, OUTBUF, OUTSIZE,
14
    POSITION, IERROR)
15
    INTEGER INCOUNT, DATATYPE, IERROR
16
    INTEGER(KIND=MPI_ADDRESS_KIND) OUTSIZE, POSITION
17
    CHARACTER*(*) DATAREP
18
    <type> INBUF(*), OUTBUF(*)
19
20
MPI_PACK_EXTERNAL_SIZE(DATAREP, INCOUNT, DATATYPE, SIZE, IERROR)
21
    INTEGER INCOUNT, DATATYPE, IERROR
22
    INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
23
    CHARACTER*(*) DATAREP
24
MPI_REQUEST_GET_STATUS( REQUEST, FLAG, STATUS, IERROR)
25
    INTEGER REQUEST, STATUS(MPI_STATUS_SIZE), IERROR
26
    LOGICAL FLAG
27
28
MPI_TYPE_CREATE_DARRAY(SIZE, RANK, NDIMS, ARRAY_OF_GSIZES, ARRAY_OF_DISTRIBS,
29
    ARRAY_OF_DARGS, ARRAY_OF_PSIZEs, ORDER, OLDTYPE, NEWTYPE,
30
    IERROR)
31
    INTEGER SIZE, RANK, NDIMS, ARRAY_OF_GSIZES(*), ARRAY_OF_DISTRIBS(*),
32
    ARRAY_OF_DARGS(*), ARRAY_OF_PSIZEs(*), ORDER, OLDTYPE, NEWTYPE, IERROR
33
MPI_TYPE_CREATE_HINDEXED(COUNT, ARRAY_OF_BLOCKLENGTHS,
34
    ARRAY_OF_DISPLACEMENTS, OLDTYPE, NEWTYPE, IERROR)
35
    INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), OLDTYPE, NEWTYPE, IERROR
36
    INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
37
38
MPI_TYPE_CREATE_HVECTOR(COUNT, BLOCKLENGTH, STIDE, OLDTYPE, NEWTYPE, IERROR)
39
    INTEGER COUNT, BLOCKLENGTH, OLDTYPE, NEWTYPE, IERROR
40
    INTEGER(KIND=MPI_ADDRESS_KIND) STRIDE
41
MPI_TYPE_CREATE_INDEXED_BLOCK(COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS,
42
    OLDTYPE, NEWTYPE, IERROR)
43
    INTEGER COUNT, BLOCKLENGTH, ARRAY_OF_DISPLACEMENTS(*), OLDTYPE,
44
    NEWTYPE, IERROR
45
46
MPI_TYPE_CREATE_RESIZED(OLDTYPE, LB, EXTENT, NEWTYPE, IERROR)
47
    INTEGER OLDTYPE, NEWTYPE, IERROR
48

```

```

1      INTEGER(KIND=MPI_ADDRESS_KIND) LB, EXTENT
2
3      MPI_TYPE_CREATE_STRUCT(COUNT, ARRAY_OF_BLOCKLENGTHS, ARRAY_OF_DISPLACEMENTS,
4                             ARRAY_OF_TYPES, NEWTYPE, IERROR)
5      INTEGER COUNT, ARRAY_OF_BLOCKLENGTHS(*), ARRAY_OF_TYPES(*), NEWTYPE,
6      IERROR
7      INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_DISPLACEMENTS(*)
8
9      MPI_TYPE_CREATE_SUBARRAY(NDIMS, ARRAY_OF_SIZES, ARRAY_OF_SUBSIZES,
10                              ARRAY_OF_STARTS, ORDER, OLDTYPE, NEWTYPE, IERROR)
11      INTEGER NDIMS, ARRAY_OF_SIZES(*), ARRAY_OF_SUBSIZES(*),
12      ARRAY_OF_STARTS(*), ORDER, OLDTYPE, NEWTYPE, IERROR
13
14      MPI_TYPE_GET_EXTENT(DATATYPE, LB, EXTENT, IERROR)
15      INTEGER DATATYPE, IERROR
16      INTEGER(KIND = MPI_ADDRESS_KIND) LB, EXTENT
17
18      MPI_TYPE_GET_TRUE_EXTENT(DATATYPE, TRUE_LB, TRUE_EXTENT, IERROR)
19      INTEGER DATATYPE, IERROR
20      INTEGER(KIND = MPI_ADDRESS_KIND) TRUE_LB, TRUE_EXTENT
21
22      MPI_UNPACK_EXTERNAL(DATAREP, INBUF, INSIZE, POSITION, OUTBUF, OUTCOUNT,
23                          DATATYPE, IERROR)
24      INTEGER OUTCOUNT, DATATYPE, IERROR
25      INTEGER(KIND=MPI_ADDRESS_KIND) INSIZE, POSITION
26      CHARACTER*(*) DATAREP
27      <type> INBUF(*), OUTBUF(*)
28
29      MPI_WIN_CREATE_ERRHANDLER(FUNCTION, ERRHANDLER, IERROR)
30      EXTERNAL FUNCTION
31      INTEGER ERRHANDLER, IERROR
32
33      MPI_WIN_GET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
34      INTEGER WIN, ERRHANDLER, IERROR
35
36      MPI_WIN_SET_ERRHANDLER(WIN, ERRHANDLER, IERROR)
37      INTEGER WIN, ERRHANDLER, IERROR
38
39
40
41
42
43
44
45
46
47
48

```

A.7.2 Process Creation and Management

```

37      MPI_CLOSE_PORT(PORT_NAME, IERROR)
38      CHARACTER*(*) PORT_NAME
39      INTEGER IERROR
40
41      MPI_COMM_ACCEPT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
42      CHARACTER*(*) PORT_NAME
43      INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR
44
45      MPI_COMM_CONNECT(PORT_NAME, INFO, ROOT, COMM, NEWCOMM, IERROR)
46      CHARACTER*(*) PORT_NAME
47      INTEGER INFO, ROOT, COMM, NEWCOMM, IERROR
48
49      MPI_COMM_DISCONNECT(COMM, IERROR)

```

```

    INTEGER COMM, IERROR
1
MPI_COMM_GET_PARENT(PARENT, IERROR)
2
    INTEGER PARENT, IERROR
3
4
MPI_COMM_JOIN(FD, INTERCOMM, IERROR)
5
    INTEGER FD, INTERCOMM, IERROR
6
7
MPI_COMM_SPAWN(COMMAND, ARGV, MAXPROCS, INFO, ROOT, COMM, INTERCOMM,
8
    ARRAY_OF_ERRCODES, IERROR)
9
    CHARACTER*(*) COMMAND, ARGV(*)
10
    INTEGER INFO, MAXPROCS, ROOT, COMM, INTERCOMM, ARRAY_OF_ERRCODES(*),
11
    IERROR
12
MPI_COMM_SPAWN_MULTIPLE(COUNT, ARRAY_OF_COMMANDS, ARRAY_OF_ARGV,
13
    ARRAY_OF_MAXPROCS, ARRAY_OF_INFO, ROOT, COMM, INTERCOMM,
14
    ARRAY_OF_ERRCODES, IERROR)
15
    INTEGER COUNT, ARRAY_OF_INFO(*), ARRAY_OF_MAXPROCS(*), ROOT, COMM,
16
    INTERCOMM, ARRAY_OF_ERRCODES(*), IERROR
17
    CHARACTER*(*) ARRAY_OF_COMMANDS(*), ARRAY_OF_ARGV(COUNT, *)
18
19
MPI_LOOKUP_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
20
    CHARACTER*(*) SERVICE_NAME, PORT_NAME
21
    INTEGER INFO, IERROR
22
23
MPI_OPEN_PORT(INFO, PORT_NAME, IERROR)
24
    CHARACTER*(*) PORT_NAME
25
    INTEGER INFO, IERROR
26
27
MPI_PUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
28
    INTEGER INFO, IERROR
29
    CHARACTER*(*) SERVICE_NAME, PORT_NAME
30
31
MPI_UNPUBLISH_NAME(SERVICE_NAME, INFO, PORT_NAME, IERROR)
32
    INTEGER INFO, IERROR
33
    CHARACTER*(*) SERVICE_NAME, PORT_NAME
34

```

A.7.3 One-Sided Communications

```

MPI_ACCUMULATE(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK,
36
    TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR)
37
    <type> ORIGIN_ADDR(*)
38
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
39
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
40
    TARGET_DATATYPE, OP, WIN, IERROR
41
42
MPI_GET(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP,
43
    TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR)
44
    <type> ORIGIN_ADDR(*)
45
    INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
46
    INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
47
    TARGET_DATATYPE, WIN, IERROR
48

```

```

1  MPI_PUT(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP,
2          TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR)
3      <type> ORIGIN_ADDR(*)
4      INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP
5      INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT,
6      TARGET_DATATYPE, WIN, IERROR
7
8  MPI_WIN_COMPLETE(WIN, IERROR)
9      INTEGER WIN, IERROR
10
11 MPI_WIN_CREATE(BASE, SIZE, DISP_UNIT, INFO, COMM, WIN, IERROR)
12     <type> BASE(*)
13     INTEGER(KIND=MPI_ADDRESS_KIND) SIZE
14     INTEGER DISP_UNIT, INFO, COMM, WIN, IERROR
15
16 MPI_WIN_FENCE(ASSERT, WIN, IERROR)
17     INTEGER ASSERT, WIN, IERROR
18
19 MPI_WIN_FREE(WIN, IERROR)
20     INTEGER WIN, IERROR
21
22 MPI_WIN_GET_GROUP(WIN, GROUP, IERROR)
23     INTEGER WIN, GROUP, IERROR
24
25 MPI_WIN_LOCK(LOCK_TYPE, RANK, ASSERT, WIN, IERROR)
26     INTEGER LOCK_TYPE, RANK, ASSERT, WIN, IERROR
27
28 MPI_WIN_POST(GROUP, ASSERT, WIN, IERROR)
29     INTEGER GROUP, ASSERT, WIN, IERROR
30
31 MPI_WIN_START(GROUP, ASSERT, WIN, IERROR)
32     INTEGER GROUP, ASSERT, WIN, IERROR
33
34 MPI_WIN_TEST(WIN, FLAG, IERROR)
35     INTEGER WIN, IERROR
36     LOGICAL FLAG
37
38 MPI_WIN_UNLOCK(RANK, WIN, IERROR)
39     INTEGER RANK, WIN, IERROR
40
41 MPI_WIN_WAIT(WIN, IERROR)
42     INTEGER WIN, IERROR

```

A.7.4 Extended Collective Operations

```

41 MPI_ALLTOALLW(SENDBUF, SENDCOUNTS, SDISPLS, SENDTYPES, RECVBUF, RECVCOUNTS,
42               RDISPLS, RECVTYPES, COMM, IERROR)
43     <type> SENDBUF(*), RECVBUF(*)
44     INTEGER SENDCOUNTS(*), SDISPLS(*), SENDTYPES(*), RECVCOUNTS(*),
45     RDISPLS(*), RECVTYPES(*), COMM, IERROR
46
47 MPI_EXSCAN(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, COMM, IERROR)
48     <type> SENDBUF(*), RECVBUF(*)

```


INTEGER COUNT, DATATYPE, OP, COMM, IERROR	1
	2
A.7.5 External Interfaces	3
	4
MPI_ADD_ERROR_CLASS(ERRORCLASS, IERROR)	5
INTEGER ERRORCLASS, IERROR	6
	7
MPI_ADD_ERROR_CODE(ERRORCLASS, ERRORCODE, IERROR)	8
INTEGER ERRORCLASS, ERRORCODE, IERROR	9
	10
MPI_ADD_ERROR_STRING(ERRORCODE, STRING, IERROR)	11
INTEGER ERRORCODE, IERROR	12
CHARACTER*(*) STRING	13
	14
MPI_COMM_CALL_ERRHANDLER(COMM, ERRORCODE, IERROR)	15
INTEGER COMM, ERRORCODE, IERROR	16
	17
MPI_COMM_CREATE_KEYVAL(COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN, COMM_KEYVAL,	18
EXTRA_STATE, IERROR)	19
EXTERNAL COMM_COPY_ATTR_FN, COMM_DELETE_ATTR_FN	20
INTEGER COMM_KEYVAL, IERROR	21
INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE	22
	23
MPI_COMM_DELETE_ATTR(COMM, COMM_KEYVAL, IERROR)	24
INTEGER COMM, COMM_KEYVAL, IERROR	25
	26
MPI_COMM_FREE_KEYVAL(COMM_KEYVAL, IERROR)	27
INTEGER COMM_KEYVAL, IERROR	28
	29
MPI_COMM_GET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)	30
INTEGER COMM, COMM_KEYVAL, IERROR	31
INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL	32
LOGICAL FLAG	33
	34
MPI_COMM_GET_NAME(COMM, COMM_NAME, RESULTLEN, IERROR)	35
INTEGER COMM, RESULTLEN, IERROR	36
CHARACTER*(*) COMM_NAME	37
	38
MPI_COMM_SET_ATTR(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, IERROR)	39
INTEGER COMM, COMM_KEYVAL, IERROR	40
INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL	41
	42
MPI_COMM_SET_NAME(COMM, COMM_NAME, IERROR)	43
INTEGER COMM, IERROR	44
CHARACTER*(*) COMM_NAME	45
	46
MPI_FILE_CALL_ERRHANDLER(FH, ERRORCODE, IERROR)	47
INTEGER FH, ERRORCODE, IERROR	48
MPI_GREQUEST_COMPLETE(REQUEST, IERROR)	
INTEGER REQUEST, IERROR	
MPI_GREQUEST_START(QUERY_FN, FREE_FN, CANCEL_FN, EXTRA_STATE, REQUEST,	
IERROR)	

```

1      INTEGER REQUEST, IERROR
2      EXTERNAL QUERY_FN, FREE_FN, CANCEL_FN
3      INTEGER (KIND=MPI_ADDRESS_KIND) EXTRA_STATE
4
5      MPI_INIT_THREAD(REQUIRED, PROVIDED, IERROR)
6          INTEGER REQUIRED, PROVIDED, IERROR
7
8      MPI_IS_THREAD_MAIN(FLAG, IERROR)
9          LOGICAL FLAG
10         INTEGER IERROR
11
12      MPI_QUERY_THREAD(PROVIDED, IERROR)
13          INTEGER PROVIDED, IERROR
14
15      MPI_STATUS_SET_CANCELLED(STATUS, FLAG, IERROR)
16          INTEGER STATUS(MPI_STATUS_SIZE), IERROR
17          LOGICAL FLAG
18
19      MPI_STATUS_SET_ELEMENTS(STATUS, DATATYPE, COUNT, IERROR)
20          INTEGER STATUS(MPI_STATUS_SIZE), DATATYPE, COUNT, IERROR
21
22      MPI_TYPE_CREATE_KEYVAL(TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN, TYPE_KEYVAL,
23          EXTRA_STATE, IERROR)
24          EXTERNAL TYPE_COPY_ATTR_FN, TYPE_DELETE_ATTR_FN
25          INTEGER TYPE_KEYVAL, IERROR
26          INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
27
28      MPI_TYPE_DELETE_ATTR(TYPE, TYPE_KEYVAL, IERROR)
29          INTEGER TYPE, TYPE_KEYVAL, IERROR
30
31      MPI_TYPE_DUP(TYPE, NEWTYPE, IERROR)
32          INTEGER TYPE, NEWTYPE, IERROR
33
34      MPI_TYPE_FREE_KEYVAL(TYPE_KEYVAL, IERROR)
35          INTEGER TYPE_KEYVAL, IERROR
36
37      MPI_TYPE_GET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)
38          INTEGER TYPE, TYPE_KEYVAL, IERROR
39          INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL
40          LOGICAL FLAG
41
42      MPI_TYPE_GET_CONTENTS(DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
43          ARRAY_OF_INTEGERS, ARRAY_OF_ADDRESSES, ARRAY_OF_DATATYPES,
44          IERROR)
45          INTEGER DATATYPE, MAX_INTEGERS, MAX_ADDRESSES, MAX_DATATYPES,
46          ARRAY_OF_INTEGERS(*), ARRAY_OF_DATATYPES(*), IERROR
47          INTEGER(KIND=MPI_ADDRESS_KIND) ARRAY_OF_ADDRESSES(*)
48
49      MPI_TYPE_GET_ENVELOPE(DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES,
50          COMBINER, IERROR)
51          INTEGER DATATYPE, NUM_INTEGERS, NUM_ADDRESSES, NUM_DATATYPES, COMBINER,
52          IERROR
53
54      MPI_TYPE_GET_NAME(TYPE, TYPE_NAME, RESULTLEN, IERROR)

```

```

    INTEGER TYPE, RESULTLEN, IERROR                                1
    CHARACTER*(*) TYPE_NAME                                       2
MPI_TYPE_SET_ATTR(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, IERROR)      3
    INTEGER TYPE, TYPE_KEYVAL, IERROR                             4
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL                  5
                                                                    6
MPI_TYPE_SET_NAME(TYPE, TYPE_NAME, IERROR)                        7
    INTEGER TYPE, IERROR                                           8
    CHARACTER*(*) TYPE_NAME                                       9
                                                                    10
MPI_WIN_CALL_ERRHANDLER(WIN, ERRORCODE, IERROR)                  11
    INTEGER WIN, ERRORCODE, IERROR                                12
                                                                    13
MPI_WIN_CREATE_KEYVAL(WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN, WIN_KEYVAL,
    EXTRA_STATE, IERROR)                                         14
    EXTERNAL WIN_COPY_ATTR_FN, WIN_DELETE_ATTR_FN                15
    INTEGER WIN_KEYVAL, IERROR                                    16
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE                   17
                                                                    18
MPI_WIN_DELETE_ATTR(WIN, WIN_KEYVAL, IERROR)                      19
    INTEGER WIN, WIN_KEYVAL, IERROR                              20
                                                                    21
MPI_WIN_FREE_KEYVAL(WIN_KEYVAL, IERROR)                          22
    INTEGER WIN_KEYVAL, IERROR                                    23
                                                                    24
MPI_WIN_GET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, FLAG, IERROR)   25
    INTEGER WIN, WIN_KEYVAL, IERROR                              26
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL                 27
    LOGICAL FLAG                                                  28
                                                                    29
MPI_WIN_GET_NAME(WIN, WIN_NAME, RESULTLEN, IERROR)               30
    INTEGER WIN, RESULTLEN, IERROR                               31
    CHARACTER*(*) WIN_NAME                                       32
                                                                    33
MPI_WIN_SET_ATTR(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, IERROR)         34
    INTEGER WIN, WIN_KEYVAL, IERROR                              35
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL                 36
                                                                    37
MPI_WIN_SET_NAME(WIN, WIN_NAME, IERROR)                          38
    INTEGER WIN, IERROR                                           39
    CHARACTER*(*) WIN_NAME                                       40
                                                                    41
A.7.6 I/O                                                         42
                                                                    43
MPI_FILE_CLOSE(FH, IERROR)                                        44
    INTEGER FH, IERROR                                            45
                                                                    46
MPI_FILE_DELETE(FILENAME, INFO, IERROR)                          47
    CHARACTER*(*) FILENAME                                       48
    INTEGER INFO, IERROR                                          49
                                                                    50
MPI_FILE_GET_AMODE(FH, AMODE, IERROR)                            51
    INTEGER FH, AMODE, IERROR                                     52

```

```

1      INTEGER FH, AMODE, IERROR
2
3      MPI_FILE_GET_ATOMICITY(FH, FLAG, IERROR)
4          INTEGER FH, IERROR
5          LOGICAL FLAG
6
7      MPI_FILE_GET_BYTE_OFFSET(FH, OFFSET, DISP, IERROR)
8          INTEGER FH, IERROR
9          INTEGER(KIND=MPI_OFFSET_KIND) OFFSET, DISP
10
11     MPI_FILE_GET_GROUP(FH, GROUP, IERROR)
12         INTEGER FH, GROUP, IERROR
13
14     MPI_FILE_GET_INFO(FH, INFO_USED, IERROR)
15         INTEGER FH, INFO_USED, IERROR
16
17     MPI_FILE_GET_POSITION(FH, OFFSET, IERROR)
18         INTEGER FH, IERROR
19         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
20
21     MPI_FILE_GET_POSITION_SHARED(FH, OFFSET, IERROR)
22         INTEGER FH, IERROR
23         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
24
25     MPI_FILE_GET_SIZE(FH, SIZE, IERROR)
26         INTEGER FH, IERROR
27         INTEGER(KIND=MPI_OFFSET_KIND) SIZE
28
29     MPI_FILE_GET_TYPE_EXTENT(FH, DATATYPE, EXTENT, IERROR)
30         INTEGER FH, DATATYPE, IERROR
31         INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT
32
33     MPI_FILE_GET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, IERROR)
34         INTEGER FH, ETYPE, FILETYPE, IERROR
35         CHARACTER*(*) DATAREP, INTEGER(KIND=MPI_OFFSET_KIND) DISP
36
37     MPI_FILE_IREAD(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
38         <type> BUF(*)
39         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
40
41     MPI_FILE_IREAD_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)
42         <type> BUF(*)
43         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
44         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
45
46     MPI_FILE_IWRITE(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
47         <type> BUF(*)
48         INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
49
50     MPI_FILE_IWRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, REQUEST, IERROR)

```

```

    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_IWRITE_SHARED(FH, BUF, COUNT, DATATYPE, REQUEST, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, REQUEST, IERROR
MPI_FILE_OPEN(COMM, FILENAME, AMODE, INFO, FH, IERROR)
    CHARACTER*(*) FILENAME
    INTEGER COMM, AMODE, INFO, FH, IERROR
MPI_FILE_PREALLOCATE(FH, SIZE, IERROR)
    INTEGER FH, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) SIZE
MPI_FILE_READ(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
MPI_FILE_READ_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_READ_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_READ_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)

```

```

1      INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
2
3      MPI_FILE_READ_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
4          <type> BUF(*)
5      INTEGER FH, COUNT, DATATYPE, IERROR
6
7      MPI_FILE_READ_ORDERED_END(FH, BUF, STATUS, IERROR)
8          <type> BUF(*)
9      INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
10
11     MPI_FILE_READ_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
12         <type> BUF(*)
13     INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
14
15     MPI_FILE_SEEK(FH, OFFSET, WHENCE, IERROR)
16         INTEGER FH, WHENCE, IERROR
17         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
18
19     MPI_FILE_SEEK_SHARED(FH, OFFSET, WHENCE, IERROR)
20         INTEGER FH, WHENCE, IERROR
21         INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
22
23     MPI_FILE_SET_ATOMICITY(FH, FLAG, IERROR)
24         INTEGER FH, IERROR
25         LOGICAL FLAG
26
27     MPI_FILE_SET_INFO(FH, INFO, IERROR)
28         INTEGER FH, INFO, IERROR
29
30     MPI_FILE_SET_SIZE(FH, SIZE, IERROR)
31         INTEGER FH, IERROR
32         INTEGER(KIND=MPI_OFFSET_KIND) SIZE
33
34     MPI_FILE_SET_VIEW(FH, DISP, ETYPE, FILETYPE, DATAREP, INFO, IERROR)
35         INTEGER FH, ETYPE, FILETYPE, INFO, IERROR
36         CHARACTER*(*) DATAREP
37         INTEGER(KIND=MPI_OFFSET_KIND) DISP
38
39     MPI_FILE_SYNC(FH, IERROR)
40         INTEGER FH, IERROR
41
42     MPI_FILE_WRITE(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
43         <type> BUF(*)
44         INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
45
46     MPI_FILE_WRITE_ALL(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
47         <type> BUF(*)
48         INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
49
50     MPI_FILE_WRITE_ALL_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
51         <type> BUF(*)
52         INTEGER FH, COUNT, DATATYPE, IERROR
53
54     MPI_FILE_WRITE_ALL_END(FH, BUF, STATUS, IERROR)

```

```

    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_AT(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_WRITE_AT_ALL(FH, OFFSET, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_WRITE_AT_ALL_BEGIN(FH, OFFSET, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
    INTEGER(KIND=MPI_OFFSET_KIND) OFFSET
MPI_FILE_WRITE_AT_ALL_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_ORDERED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_ORDERED_BEGIN(FH, BUF, COUNT, DATATYPE, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, IERROR
MPI_FILE_WRITE_ORDERED_END(FH, BUF, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, STATUS(MPI_STATUS_SIZE), IERROR
MPI_FILE_WRITE_SHARED(FH, BUF, COUNT, DATATYPE, STATUS, IERROR)
    <type> BUF(*)
    INTEGER FH, COUNT, DATATYPE, STATUS(MPI_STATUS_SIZE), IERROR
MPI_REGISTER_DATAREP(DATAREP, READ_CONVERSION_FN, WRITE_CONVERSION_FN,
                     DTYPE_FILE_EXTENT_FN, EXTRA_STATE, IERROR)
    CHARACTER*(*) DATAREP
    EXTERNAL READ_CONVERSION_FN, WRITE_CONVERSION_FN, DTYPE_FILE_EXTENT_FN
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
    INTEGER IERROR

```

A.7.7 Language Bindings

```

MPI_SIZEOF(X, SIZE, IERROR)
    <type> X
    INTEGER SIZE, IERROR
MPI_TYPE_CREATE_F90_COMPLEX(P, R, NEWTYPE, IERROR)

```

```

1      INTEGER P, R, NEWTYPE, IERROR
2      MPI_TYPE_CREATE_F90_INTEGER(R, NEWTYPE, IERROR)
3      INTEGER R, NEWTYPE, IERROR
4
5      MPI_TYPE_CREATE_F90_REAL(P, R, NEWTYPE, IERROR)
6      INTEGER P, R, NEWTYPE, IERROR
7
8      MPI_TYPE_MATCH_SIZE(TYPECLASS, SIZE, TYPE, IERROR)
9      INTEGER TYPECLASS, SIZE, TYPE, IERROR

```

A.7.8 User Defined Subroutines

```

12     SUBROUTINE COMM_COPY_ATTR_FN(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
13         ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
14         INTEGER OLDCOMM, COMM_KEYVAL, IERROR
15         INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
16             ATTRIBUTE_VAL_OUT
17         LOGICAL FLAG
18
19     SUBROUTINE COMM_DELETE_ATTR_FN(COMM, COMM_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
20         IERROR)
21         INTEGER COMM, COMM_KEYVAL, IERROR
22         INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
23
24     SUBROUTINE COMM_ERRHANDLER_FN(COMM, ERROR_CODE, ...)
25         INTEGER COMM, ERROR_CODE
26
27     SUBROUTINE DATAREP_CONVERSION_FUNCTION(USERBUF, DATATYPE, COUNT, FILEBUF,
28         POSITION, EXTRA_STATE, IERROR)
29         <TYPE> USERBUF(*), FILEBUF(*)
30         INTEGER COUNT, DATATYPE, IERROR
31         INTEGER(KIND=MPI_OFFSET_KIND) POSITION
32         INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
33
34     SUBROUTINE DATAREP_EXTENT_FUNCTION(DATATYPE, EXTENT, EXTRA_STATE, IERROR)
35         INTEGER DATATYPE, IERROR
36         INTEGER(KIND=MPI_ADDRESS_KIND) EXTENT, EXTRA_STATE
37
38     SUBROUTINE FILE_ERRHANDLER_FN(FILE, ERROR_CODE, ...)
39         INTEGER FILE, ERROR_CODE
40
41     SUBROUTINE GREQUEST_CANCEL_FUNCTION(EXTRA_STATE, COMPLETE, IERROR)
42         INTEGER IERROR
43         INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
44         LOGICAL COMPLETE
45
46     SUBROUTINE GREQUEST_FREE_FUNCTION(EXTRA_STATE, IERROR)
47         INTEGER IERROR
48         INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
49
50     SUBROUTINE GREQUEST_QUERY_FUNCTION(EXTRA_STATE, STATUS, IERROR)
51         INTEGER STATUS(MPI_STATUS_SIZE), IERROR

```



```

    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE
SUBROUTINE TYPE_COPY_ATTR_FN(OLDTYPE, TYPE_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDTYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE,
        ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
SUBROUTINE TYPE_DELETE_ATTR_FN(TYPE, TYPE_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
    IERROR)
    INTEGER TYPE, TYPE_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
SUBROUTINE WIN_COPY_ATTR_FN(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
    INTEGER OLDWIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
SUBROUTINE WIN_DELETE_ATTR_FN(WIN, WIN_KEYVAL, ATTRIBUTE_VAL, EXTRA_STATE,
    IERROR)
    INTEGER WIN, WIN_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
SUBROUTINE WIN_ERRHANDLER_FN(WIN, ERROR_CODE, ...)
    INTEGER WIN, ERROR_CODE

```

A.8 MPI-2 C++ Bindings

A.8.1 Miscellany

```

void* MPI::Alloc_mem(MPI::Aint size, const MPI::Info& info)
static MPI::Errhandler
    MPI::Comm::Create_errhandler(MPI::Comm::Errhandler_fn*
        function)
MPI::Errhandler MPI::Comm::Get_errhandler() const
void MPI::Comm::Set_errhandler(const MPI::Errhandler& errhandler)
MPI::Datatype MPI::Datatype::Create_darray(int size, int rank, int ndims,
    const int array_of_gsizes[], const int array_of_distrib[],
    const int array_of_dargs[], const int array_of_psize[],
    int order) const
MPI::Datatype MPI::Datatype::Create_hindexed(int count,
    const int array_of_blocklengths[],
    const MPI::Aint array_of_displacements[]) const

```

```

1  MPI::Datatype MPI::Datatype::Create_hvector(int count, int blocklength,
2      MPI::Aint stride) const
3
4  MPI::Datatype MPI::Datatype::Create_indexed_block( int count,
5      int blocklength, const int array_of_displacements[]) const
6
7  static MPI::Datatype MPI::Datatype::Create_struct(int count,
8      const int array_of_blocklengths[], const MPI::Aint
9      array_of_displacements[], const MPI::Datatype array_of_types[])
10
11 MPI::Datatype MPI::Datatype::Create_subarray(int ndims,
12     const int array_of_sizes[], const int array_of_subsizes[],
13     const int array_of_starts[], int order) const
14
15 void MPI::Datatype::Get_extent(MPI::Aint& lb, MPI::Aint& extent) const
16
17 void MPI::Datatype::Get_true_extent(MPI::Aint& true_lb,
18     MPI::Aint& true_extent) const
19
20 void MPI::Datatype::Pack_external(const char* datarep, const void* inbuf,
21     int incout, void* outbuf, MPI::Aint outsize,
22     MPI::Aint& position) const
23
24 MPI::Aint MPI::Datatype::Pack_external_size(const char* datarep,
25     int incout) const
26
27 MPI::Datatype MPI::Datatype::Resized(const MPI::Aint lb,
28     const MPI::Aint extent) const
29
30 void MPI::Datatype::Unpack_external(const char* datarep, const void* inbuf,
31     MPI::Aint insize, MPI::Aint& position, void* outbuf,
32     int outcount) const
33
34 static MPI::Errhandler
35     MPI::File::Create_errhandler(MPI::File::Errhandler_fn*
36     function)
37
38 MPI::Errhandler MPI::File::Get_errhandler() const
39
40 void MPI::File::Set_errhandler(const MPI::Errhandler& errhandler)
41
42 void MPI::Free_mem(void *base)
43
44 MPI::Aint MPI::Get_address(void* location)
45
46 static MPI::Info MPI::Info::Create()
47
48 void MPI::Info::Delete(const char* key)
49
50 MPI::Info MPI::Info::Dup() const
51
52 void MPI::Info::Free()
53
54 bool MPI::Info::Get(const char* key, int valuelen, char* value) const
55
56 int MPI::Info::Get_nkeys() const
57
58

```

```

void MPI::Info::Get_nthkey(int n, char* key) const
bool MPI::Info::Get_valuelen(const char* key, int& valuelen) const
void MPI::Info::Set(const char* key, const char* value)
bool MPI::Is_finalized()
bool MPI::Request::Get_status() const
bool MPI::Request::Get_status(MPI::Status& status) const
static MPI::Errhandler MPI::Win::Create_errhandler(MPI::Win::Errhandler_fn*
    function)
MPI::Errhandler MPI::Win::Get_errhandler() const
void MPI::Win::Set_errhandler(const MPI::Errhandler& errhandler)

```

A.8.2 Process Creation and Management

```

void MPI::Close_port(const char* port_name)
void MPI::Comm::Disconnect()
static MPI::Intercomm MPI::Comm::Get_parent()
static MPI::Intercomm MPI::Comm::Join(const int fd)
MPI::Intercomm MPI::Intracomm::Accept(const char* port_name,
    const MPI::Info& info, int root) const
MPI::Intercomm MPI::Intracomm::Connect(const char* port_name,
    const MPI::Info& info, int root) const
MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
    const char* argv[], int maxprocs, const MPI::Info& info,
    int root) const
MPI::Intercomm MPI::Intracomm::Spawn(const char* command,
    const char* argv[], int maxprocs, const MPI::Info& info,
    int root, int array_of_errcodes[]) const
MPI::Intercomm MPI::Intracomm::Spawn_multiple(int count,
    const char* array_of_commands[], const char** array_of_argv[],
    const int array_of_maxprocs[], const MPI::Info array_of_info[],
    int root)
MPI::Intercomm MPI::Intracomm::Spawn_multiple(int count,
    const char* array_of_commands[], const char** array_of_argv[],
    const int array_of_maxprocs[], const MPI::Info array_of_info[],
    int root, int array_of_errcodes[])
void MPI::Lookup_name(const char* service_name, const MPI::Info& info,
    char* port_name)

```

```

1 void MPI::Open_port(const MPI::Info& info, char* port_name)
2 void MPI::Publish_name(const char* service_name, const MPI::Info& info,
3     const char* port_name)
4
5 void MPI::Unpublish_name(const char* service_name, const MPI::Info& info,
6     const char* port_name)
7
8
9

```

A.8.3 One-Sided Communications

```

10 void MPI::Win::Accumulate(const void* origin_addr, int origin_count, const
11     MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
12     target_disp, int target_count, const MPI::Datatype&
13     target_datatype, const MPI::Op& op) const
14
15 void MPI::Win::Complete() const
16
17 static MPI::Win MPI::Win::Create(const void* base, MPI::Aint size, int
18     disp_unit, const MPI::Info& info, const MPI::Intracomm& comm)
19
20 void MPI::Win::Fence(int assert) const
21
22 void MPI::Win::Free()
23
24 void MPI::Win::Get(const void *origin_addr, int origin_count, const
25     MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
26     target_disp, int target_count, const MPI::Datatype&
27     target_datatype) const
28
29 MPI::Group MPI::Win::Get_group() const
30
31 void MPI::Win::Lock(int lock_type, int rank, int assert) const
32
33 void MPI::Win::Post(const MPI::Group& group, int assert) const
34
35 void MPI::Win::Put(const void* origin_addr, int origin_count, const
36     MPI::Datatype& origin_datatype, int target_rank, MPI::Aint
37     target_disp, int target_count, const MPI::Datatype&
38     target_datatype) const
39
40 void MPI::Win::Start(const MPI::Group& group, int assert) const
41
42 bool MPI::Win::Test() const
43
44 void MPI::Win::Unlock(int rank) const
45
46 void MPI::Win::Wait() const
47
48

```

A.8.4 Extended Collective Operations

```

44 void MPI::Comm::Allgather(const void* sendbuf, int sendcount, const
45     MPI::Datatype& sendtype, void* recvbuf, int recvcount,
46     const MPI::Datatype& recvttype) const = 0
47
48

```

```

void MPI::Comm::Allgatherv(const void* sendbuf, int sendcount, const      1
    MPI::Datatype& sendtype, void* recvbuf,
    const int recvcnts[], const int displs[],
    const MPI::Datatype& recvtype) const = 0                               2
                                                                              3
                                                                              4
                                                                              5
void MPI::Comm::Allreduce(const void* sendbuf, void* recvbuf, int count,   6
    const MPI::Datatype& datatype, const MPI::Op& op) const = 0          7
                                                                              8
void MPI::Comm::Alltoall(const void* sendbuf, int sendcount, const        9
    MPI::Datatype& sendtype, void* recvbuf, int recvcnt,
    const MPI::Datatype& recvtype) const = 0                             10
                                                                              11
void MPI::Comm::Alltoallv(const void* sendbuf, const int sendcounts[],    12
    const int sdispls[], const MPI::Datatype& sendtype,
    void* recvbuf, const int recvcnts[], const int rdispls[],
    const MPI::Datatype& recvtype) const = 0                             13
                                                                              14
                                                                              15
void MPI::Comm::Alltoallw(const void* sendbuf, const int sendcounts[],    16
    const int sdispls[], const MPI::Datatype sendtypes[], void*
    recvbuf, const int recvcnts[], const int rdispls[], const
    MPI::Datatype recvtypes[]) const = 0                                17
                                                                              18
                                                                              19
void MPI::Comm::Barrier() const = 0                                       20
                                                                              21

void MPI::Comm::Bcast(void* buffer, int count,
    const MPI::Datatype& datatype, int root) const = 0                   22
                                                                              23
                                                                              24
void MPI::Comm::Gather(const void* sendbuf, int sendcount, const          25
    MPI::Datatype& sendtype, void* recvbuf, int recvcnt,
    const MPI::Datatype& recvtype, int root) const = 0                   26
                                                                              27

void MPI::Comm::Gatherv(const void* sendbuf, int sendcount, const         28
    MPI::Datatype& sendtype, void* recvbuf,
    const int recvcnts[], const int displs[],
    const MPI::Datatype& recvtype, int root) const = 0                   29
                                                                              30
                                                                              31
                                                                              32
void MPI::Comm::Reduce(const void* sendbuf, void* recvbuf, int count,     33
    const MPI::Datatype& datatype, const MPI::Op& op, int root)
    const = 0                                                             34
                                                                              35

void MPI::Comm::Reduce_scatter(const void* sendbuf, void* recvbuf,        36
    int recvcnts[], const MPI::Datatype& datatype,
    const MPI::Op& op) const = 0                                         37
                                                                              38
                                                                              39
void MPI::Comm::Scatter(const void* sendbuf, int sendcount, const         40
    MPI::Datatype& sendtype, void* recvbuf, int recvcnt,
    const MPI::Datatype& recvtype, int root) const = 0                   41
                                                                              42

void MPI::Comm::Scatterv(const void* sendbuf, const int sendcounts[],     43
    const int displs[], const MPI::Datatype& sendtype,
    void* recvbuf, int recvcnt, const MPI::Datatype& recvtype,
    int root) const = 0                                                  44
                                                                              45
                                                                              46
                                                                              47
MPI::Intercomm MPI::Intercomm::Create(const Group& group) const          48

```

```

1 MPI::Intercomm MPI::Intercomm::Split(int color, int key) const
2 MPI::Intracomm MPI::Intracomm::Create(const Group& group) const
3
4 void MPI::Intracomm::Exscan(const void* sendbuf, void* recvbuf, int count,
5     const MPI::Datatype& datatype, const MPI::Op& op) const
6
7 MPI::Intracomm MPI::Intracomm::Split(int color, int key) const
8

```

9 A.8.5 External Interfaces

```

10 int MPI::Add_error_class()
11
12 int MPI::Add_error_code(int errorclass)
13
14 void MPI::Add_error_string(int errorcode, const char* string)
15
16 void MPI::Comm::Call_errhandler(int errorcode) const
17
18 static int MPI::Comm::Create_keyval(MPI::Comm::Copy_attr_function*
19     comm_copy_attr_fn,
20     MPI::Comm::Delete_attr_function* comm_delete_attr_fn,
21     void* extra_state)
22
23 void MPI::Comm::Delete_attr(int comm_keyval)
24
25 static void MPI::Comm::Free_keyval(int& comm_keyval)
26
27 bool MPI::Comm::Get_attr(int comm_keyval, void* attribute_val) const
28
29 void MPI::Comm::Get_name(char* comm_name, int& resultlen) const
30
31 void MPI::Comm::Set_attr(int comm_keyval, const void* attribute_val) const
32
33 void MPI::Comm::Set_name(const char* comm_name)
34
35 static int MPI::Datatype::Create_keyval(MPI::Datatype::Copy_attr_function*
36     type_copy_attr_fn, MPI::Datatype::Delete_attr_function*
37     type_delete_attr_fn, void* extra_state)
38
39 void MPI::Datatype::Delete_attr(int type_keyval)
40
41 MPI::Datatype MPI::Datatype::Dup() const
42
43 static void MPI::Datatype::Free_keyval(int& type_keyval)
44
45 bool MPI::Datatype::Get_attr(int type_keyval, void* attribute_val) const
46
47 void MPI::Datatype::Get_contents(int max_integers, int max_addresses,
48     int max_datatypes, int array_of_integers[],
49     MPI::Aint array_of_addresses[],
50     MPI::Datatype array_of_datatypes[]) const
51
52 void MPI::Datatype::Get_envelope(int& num_integers, int& num_addresses,
53     int& num_datatypes, int& combiner) const
54
55 void MPI::Datatype::Get_name(char* type_name, int& resultlen) const
56

```

```

void MPI::Datatype::Set_attr(int type_keyval, const void* attribute_val)
void MPI::Datatype::Set_name(const char* type_name)
void MPI::File::Call_errhandler(int errorcode) const
void MPI::Grequest::Complete()
static MPI::Grequest
    MPI::Grequest::Start(const MPI::Grequest::Query_function
        query_fn, const MPI::Grequest::Free_function free_fn,
        const MPI::Grequest::Cancel_function cancel_fn,
        void *extra_state)
int MPI::Init_thread(int required)
int MPI::Init_thread(int& argc, char**& argv, int required)
bool MPI::Is_thread_main()
int MPI::Query_thread()
void MPI::Status::Set_cancelled(bool flag)
void MPI::Status::Set_elements(const MPI::Datatype& datatype, int count)
void MPI::Win::Call_errhandler(int errorcode) const
static int MPI::Win::Create_keyval(MPI::Win::Copy_attr_function*
    win_copy_attr_fn,
    MPI::Win::Delete_attr_function* win_delete_attr_fn,
    void* extra_state)
void MPI::Win::Delete_attr(int win_keyval)
static void MPI::Win::Free_keyval(int& win_keyval)
bool MPI::Win::Get_attr(const MPI::Win& win, int win_keyval,
    void* attribute_val) const
void MPI::Win::Get_name(char* win_name, int& resultlen) const
void MPI::Win::Set_attr(int win_keyval, const void* attribute_val)
void MPI::Win::Set_name(const char* win_name)

```

A.8.6 I/O

```

void MPI::File::Close()
static void MPI::File::Delete(const char* filename, const MPI::Info& info)
int MPI::File::Get_amode() const
bool MPI::File::Get_atomicity() const
MPI::Offset MPI::File::Get_byte_offset(const MPI::Offset disp) const

```

```

1  MPI::Group MPI::File::Get_group() const
2
3  MPI::Info MPI::File::Get_info() const
4
5  MPI::Offset MPI::File::Get_position() const
6
7  MPI::Offset MPI::File::Get_position_shared() const
8
9  MPI::Offset MPI::File::Get_size() const
10
11 MPI::Aint MPI::File::Get_type_extent(const MPI::Datatype& datatype) const
12
13 void MPI::File::Get_view(MPI::Offset& disp, MPI::Datatype& etype,
14                          MPI::Datatype& filetype, char* datarep) const
15
16 MPI::Request MPI::File::Iread(void* buf, int count,
17                               const MPI::Datatype& datatype)
18
19 MPI::Request MPI::File::Iread_at(MPI::Offset offset, void* buf, int count,
20                                  const MPI::Datatype& datatype)
21
22 MPI::Request MPI::File::Iread_shared(void* buf, int count,
23                                      const MPI::Datatype& datatype)
24
25 MPI::Request MPI::File::Iwrite(const void* buf, int count,
26                               const MPI::Datatype& datatype)
27
28 MPI::Request MPI::File::Iwrite_at(MPI::Offset offset, const void* buf,
29                                   int count, const MPI::Datatype& datatype)
30
31 MPI::Request MPI::File::Iwrite_shared(const void* buf, int count,
32                                       const MPI::Datatype& datatype)
33
34 static MPI::File MPI::File::Open(const MPI::Intracomm& comm,
35                                   const char* filename, int amode, const MPI::Info& info)
36
37 void MPI::File::Preallocate(MPI::Offset size)
38
39 void MPI::File::Read(void* buf, int count, const MPI::Datatype& datatype)
40
41 void MPI::File::Read(void* buf, int count, const MPI::Datatype& datatype,
42                     MPI::Status& status)
43
44 void MPI::File::Read_all(void* buf, int count,
45                          const MPI::Datatype& datatype)
46
47 void MPI::File::Read_all(void* buf, int count,
48                          const MPI::Datatype& datatype, MPI::Status& status)
49
50 void MPI::File::Read_all_begin(void* buf, int count,
51                                const MPI::Datatype& datatype)
52
53 void MPI::File::Read_all_end(void* buf)
54
55 void MPI::File::Read_all_end(void* buf, MPI::Status& status)
56
57 void MPI::File::Read_at(MPI::Offset offset, void* buf, int count,
58                         const MPI::Datatype& datatype)
59

```



```
1 void MPI::File::Write_all(const void* buf, int count,
2     const MPI::Datatype& datatype)
3
4 void MPI::File::Write_all(const void* buf, int count,
5     const MPI::Datatype& datatype, MPI::Status& status)
6
7 void MPI::File::Write_all_begin(const void* buf, int count,
8     const MPI::Datatype& datatype)
9
10 void MPI::File::Write_all_end(const void* buf, MPI::Status& status)
11
12 void MPI::File::Write_at(MPI::Offset offset, const void* buf, int count,
13     const MPI::Datatype& datatype)
14
15 void MPI::File::Write_at(MPI::Offset offset, const void* buf, int count,
16     const MPI::Datatype& datatype, MPI::Status& status)
17
18 void MPI::File::Write_at_all(MPI::Offset offset, const void* buf,
19     int count, const MPI::Datatype& datatype)
20
21 void MPI::File::Write_at_all(MPI::Offset offset, const void* buf,
22     int count, const MPI::Datatype& datatype, MPI::Status& status)
23
24 void MPI::File::Write_at_all_begin(MPI::Offset offset, const void* buf,
25     int count, const MPI::Datatype& datatype)
26
27 void MPI::File::Write_at_all_end(const void* buf)
28
29 void MPI::File::Write_at_all_end(const void* buf, MPI::Status& status)
30
31 void MPI::File::Write_ordered(const void* buf, int count,
32     const MPI::Datatype& datatype)
33
34 void MPI::File::Write_ordered(const void* buf, int count,
35     const MPI::Datatype& datatype, MPI::Status& status)
36
37 void MPI::File::Write_ordered_begin(const void* buf, int count,
38     const MPI::Datatype& datatype)
39
40 void MPI::File::Write_ordered_end(const void* buf)
41
42 void MPI::File::Write_ordered_end(const void* buf, MPI::Status& status)
43
44 void MPI::File::Write_shared(const void* buf, int count,
45     const MPI::Datatype& datatype)
46
47 void MPI::File::Write_shared(const void* buf, int count,
48     const MPI::Datatype& datatype, MPI::Status& status)
49
50 void MPI::Register_datarep(const char* datarep,
51     MPI::Datarep_conversion_function* read_conversion_fn,
52     MPI::Datarep_conversion_function* write_conversion_fn,
53     MPI::Datarep_extent_function* dtype_file_extent_fn,
54     void* extra_state)
```

A.8.7 Language Bindings

```

static MPI::Datatype MPI::Datatype::Create_f90_complex(int p, int r)
static MPI::Datatype MPI::Datatype::Create_f90_integer(int r)
static MPI::Datatype MPI::Datatype::Create_f90_real(int p, int r)
static MPI::Datatype MPI::Datatype::Match_size(int typeclass, int size)

```

A.8.8 User Defined Functions

```

typedef int MPI::Comm::Copy_attr_function(const MPI::Comm& oldcomm,
                                           int comm_keyval, void* extra_state, void* attribute_val_in,
                                           void* attribute_val_out, bool& flag);
typedef int MPI::Comm::Delete_attr_function(MPI::Comm& comm,
                                           int comm_keyval, void* attribute_val, void* extra_state);
typedef void MPI::Comm::Errhandler_fn(MPI::Comm &, int *, ...);
typedef MPI::Datatype MPI::Datatype::Datarep_conversion_function(void* userbuf,
                                                                MPI::Datatype& datatype, int count, void* filebuf,
                                                                MPI::Offset position, void* extra_state);
typedef MPI::Datatype MPI::Datatype::Datarep_extents_function(const MPI::Datatype& datatype,
                                                             MPI::Aint& file_extents, void* extra_state);
typedef int MPI::Datatype::Copy_attr_function(const MPI::Datatype& oldtype,
                                              int type_keyval, void* extra_state,
                                              const void* attribute_val_in, void* attribute_val_out,
                                              bool& flag);
typedef int MPI::Datatype::Delete_attr_function(MPI::Datatype& type,
                                              int type_keyval, void* attribute_val, void* extra_state);
typedef void MPI::File::Errhandler_fn(MPI::File &, int *, ...);
typedef int MPI::Grequest::Cancel_function(void* extra_state,
                                           bool complete);
typedef int MPI::Grequest::Free_function(void* extra_state);
typedef int MPI::Grequest::Query_function(void* extra_state,
                                          MPI::Status& status);
typedef int MPI::Win::Copy_attr_function(const MPI::Win& oldwin,
                                         int win_keyval, void* extra_state, void* attribute_val_in,
                                         void* attribute_val_out, bool& flag);
typedef int MPI::Win::Delete_attr_function(MPI::Win& win, int win_keyval,
                                         void* attribute_val, void* extra_state);
typedef void MPI::Win::Errhandler_fn(MPI::Win &, int *, ...);

```

Annex B

MPI-1 C++ Language Binding

B.1 C++ Classes

The following are the classes provided with the C++ MPI-1 language bindings:

```
namespace MPI {
    class Comm { ... };
    class Intracomm : public Comm { ... };
    class Graphcomm : public Intracomm { ... };
    class Cartcomm : public Intracomm { ... };
    class Intercomm : public Comm { ... };
    class Datatype { ... };
    class Errhandler { ... };
    class Exception { ... };
    class Group { ... };
    class Op { ... };
    class Request { ... };
    class Prerequest : public Request { ... };
    class Status { ... };
};
```

Note that several MPI-1 functions, constants, and typedefs have been deprecated and therefore do not have corresponding C++ bindings. All deprecated names have corresponding new names in MPI-2 (albeit probably with different semantics). See the table in Section 2.6.1 for a list of the deprecated names and their corresponding new names. The bindings for the new names are listed in Annex A.

B.2 Defined Constants

These are required constants, defined in the file `mpi.h`. For brevity, the types of the constants are defined below in the comments.

```
// return codes
// Type: const int (or unnamed enum)
MPI::SUCCESS
MPI::ERR_BUFFER
```

MPI::ERR_COUNT	1
MPI::ERR_TYPE	2
MPI::ERR_TAG	3
MPI::ERR_COMM	4
MPI::ERR_RANK	5
MPI::ERR_REQUEST	6
MPI::ERR_ROOT	7
MPI::ERR_GROUP	8
MPI::ERR_OP	9
MPI::ERR_TOPOLOGY	10
MPI::ERR_DIMS	11
MPI::ERR_ARG	12
MPI::ERR_UNKNOWN	13
MPI::ERR_TRUNCATE	14
MPI::ERR_OTHER	15
MPI::ERR_INTERN	16
MPI::ERR_PENDING	17
MPI::ERR_IN_STATUS	18
MPI::ERR_LASTCODE	19
	20
// assorted constants	21
// Type: const void *	22
MPI::BOTTOM	23
// Type: const int (or unnamed enum)	24
MPI::PROC_NULL	25
MPI::ANY_SOURCE	26
MPI::ANY_TAG	27
MPI::UNDEFINED	28
MPI::BSEND_OVERHEAD	29
MPI::KEYVAL_INVALID	30
	31
// Error-handling specifiers	32
// Type: MPI::Errhandler (see below)	33
MPI::ERRORS_ARE_FATAL	34
MPI::ERRORS_RETURN	35
MPI::ERRORS_THROW_EXCEPTIONS	36
	37
// Maximum sizes for strings	38
// Type: const int	39
MPI::MAX_PROCESSOR_NAME	40
MPI::MAX_ERROR_STRING	41
	42
// elementary datatypes (C / C++)	43
// Type: const MPI::Datatype	44
MPI::CHAR	45
MPI::SHORT	46
MPI::INT	47
MPI::LONG	48

```
1  MPI::SIGNED_CHAR
2  MPI::UNSIGNED_CHAR
3  MPI::UNSIGNED_SHORT
4  MPI::UNSIGNED
5  MPI::UNSIGNED_LONG
6  MPI::FLOAT
7  MPI::DOUBLE
8  MPI::LONG_DOUBLE
9  MPI::BYTE
10 MPI::PACKED
11
12 // elementary datatypes (Fortran)
13 // Type: const MPI::Datatype
14 MPI::INTEGER
15 MPI::REAL
16 MPI::DOUBLE_PRECISION
17 MPI::F_COMPLEX
18 MPI::F_DOUBLE_COMPLEX
19 MPI::LOGICAL
20 MPI::CHARACTER
21
22 // datatypes for reduction functions (C / C++)
23 // Type: const MPI::Datatype
24 MPI::FLOAT_INT
25 MPI::DOUBLE_INT
26 MPI::LONG_INT
27 MPI::TWOINT
28 MPI::SHORT_INT
29 MPI::LONG_DOUBLE_INT
30
31 // datatype for reduction functions (Fortran)
32 // Type const MPI::Datatype
33 MPI::TWOREAL
34 MPI::TWODOUBLE_PRECISION
35 MPI::TWOINTEGER
36
37 // optional datatypes (Fortran)
38 // Type: const MPI::Datatype
39 MPI::INTEGER1
40 MPI::INTEGER2
41 MPI::INTEGER4
42 MPI::REAL2
43 MPI::REAL4
44 MPI::REAL8
45
46 // optional datatypes (C / C++)
47 // Type: const MPI::Datatype
48
```

```

MPI::LONG_LONG 1
MPI::UNSIGNED_LONG_LONG 2
3
4
// special datatypes for construction derived datatypes 5
// Type: const MPI::Datatype 6
MPI::UB 7
MPI::LB 8
9
// C++ datatypes 10
// Type: const MPI::Datatype 11
MPI::BOOL 12
MPI::COMPLEX 13
MPI::DOUBLE_COMPLEX 14
MPI::LONG_DOUBLE_COMPLEX 15
16
// reserved communicators 17
// Type: MPI::Intracomm 18
MPI::COMM_WORLD 19
MPI::COMM_SELF 20
21
// results of communicator and group comparisons 22
// Type: const int (or unnamed enum) 23
MPI::IDENT 24
MPI::CONGRUENT 25
MPI::SIMILAR 26
MPI::UNEQUAL 27
28
// environmental inquiry keys 29
// Type: const int (or unnamed enum) 30
MPI::TAG_UB 31
MPI::IO 32
MPI::HOST 33
MPI::WTIME_IS_GLOBAL 34
35
// collective operations 36
// Type: const MPI::Op 37
MPI::MAX 38
MPI::MIN 39
MPI::SUM 40
MPI::PROD 41
MPI::MAXLOC 42
MPI::MINLOC 43
MPI::BAND 44
MPI::BOR 45
MPI::BXOR 46
MPI::LAND 47
MPI::LOR 48

```

```

1  MPI::LXOR
2
3  // Null handles
4  // Type: const MPI::Group
5  MPI::GROUP_NULL
6
7  // Type: See Section 10.1.7 regarding the MPI::Comm class hierarchy and
8  // the specific type of MPI::COMM_NULL.
9
10 MPI::COMM_NULL
11 // Type: const MPI::Datatype
12 MPI::DATATYPE_NULL
13 // Type: const MPI::Request
14 MPI::REQUEST_NULL
15 // Type: const MPI::Op
16 MPI::OP_NULL
17 // Type: MPI::Errhandler
18 MPI::ERRHANDLER_NULL
19
20 // Empty group
21 // Type: const MPI::Group
22 MPI::GROUP_EMPTY
23
24 // Topologies
25 // Type: const int (or unnamed enum)
26 MPI::GRAPH
27 MPI::CART
28
29 // Predefined functions
30 // Type: MPI::Copy_function
31 MPI::NULL_COPY_FN
32 MPI::DUP_FN
33 // Type: MPI::Delete_function
34 MPI::NULL_DELETE_FN

```

B.3 Typedefs

The following are defined C++ types, also included in the file `mpi.h`.

```

39 // Typedef
40 MPI::Aint

```

The rest of this annex uses the `namespace` notation because all the functions listed below are prototypes. The `namespace` notation is not used previously because the lists of constants and types above are not actual declarations.

```

46 // prototypes for user-defined functions
47 namespace MPI {
48     typedef void User_function(const void *invec, void* inoutvec, int len,

```



```

        const Datatype& datatype);
};

```

B.4 C++ Bindings for Point-to-Point Communication

Except where specifically noted, all non-static member functions in this annex are **virtual**. For brevity, the keyword **virtual** is omitted.

```

namespace MPI {

    void Comm::Send(const void* buf, int count, const Datatype& datatype,
                    int dest, int tag) const

    void Comm::Recv(void* buf, int count, const Datatype& datatype,
                    int source, int tag, Status& status) const

    void Comm::Recv(void* buf, int count, const Datatype& datatype,
                    int source, int tag) const

    int Status::Get_count(const Datatype& datatype) const

    void Comm::Bsend(const void* buf, int count, const Datatype& datatype,
                     int dest, int tag) const

    void Comm::Ssend(const void* buf, int count, const Datatype& datatype,
                     int dest, int tag) const

    void Comm::Rsend(const void* buf, int count, const Datatype& datatype,
                     int dest, int tag) const

    void Attach_buffer(void* buffer, int size)

    int Detach_buffer(void*& buffer)

    Request Comm::Isend(const void* buf, int count, const
                        Datatype& datatype, int dest, int tag) const

    Request Comm::Ibsend(const void* buf, int count, const
                         Datatype& datatype, int dest, int tag) const

    Request Comm::Issend(const void* buf, int count, const
                         Datatype& datatype, int dest, int tag) const

    Request Comm::Irsend(const void* buf, int count, const
                         Datatype& datatype, int dest, int tag) const

    Request Comm::Irecv(void* buf, int count, const Datatype& datatype,
                        int source, int tag) const

    void Request::Wait(Status& status)

    void Request::Wait()

    bool Request::Test(Status& status)

    bool Request::Test()

```

```

1  void Request::Free()
2
3  static int Request::Waitany(int count, Request array_of_requests[],
4                               Status& status)
5
6  static int Request::Waitany(int count, Request array_of_requests[])
7
8  static bool Request::Testany(int count, Request array_of_requests[],
9                               int& index, Status& status)
10
11 static bool Request::Testany(int count, Request array_of_requests[],
12                               int& index)
13
14 static void Request::Waitall(int count, Request array_of_requests[],
15                               Status array_of_statuses[])
16
17 static void Request::Waitall(int count, Request array_of_requests[])
18
19 static bool Request::Testall(int count, Request array_of_requests[],
20                               Status array_of_statuses[])
21
22 static bool Request::Testall(int count, Request array_of_requests[])
23
24 static int Request::Waitsome(int incount, Request array_of_requests[],
25                               int array_of_indices[], Status array_of_statuses[])
26
27 static int Request::Waitsome(int incount, Request array_of_requests[],
28                               int array_of_indices[])
29
30 bool Comm::Iprobe(int source, int tag, Status& status) const
31
32 bool Comm::Iprobe(int source, int tag) const
33
34 void Comm::Probe(int source, int tag, Status& status) const
35
36 void Comm::Probe(int source, int tag) const
37
38 void Request::Cancel() const
39
40 bool Status::Is_cancelled() const
41
42 Prequest Comm::Send_init(const void* buf, int count, const
43                           Datatype& datatype, int dest, int tag) const
44
45 Prequest Comm::Bsend_init(const void* buf, int count, const
46                           Datatype& datatype, int dest, int tag) const
47
48 Prequest Comm::Ssend_init(const void* buf, int count, const
49                           Datatype& datatype, int dest, int tag) const

```

```

Prequest Comm::Recv_init(void* buf, int count, const Datatype& datatype,
    int source, int tag) const
void Prequest::Start()
static void Prequest::Startall(int count, Prequest array_of_requests[])
void Comm::Sendrecv(const void *sendbuf, int sendcount, const
    Datatype& sendtype, int dest, int sendtag, void *recvbuf,
    int recvcount, const Datatype& recvtype, int source,
    int recvtag, Status& status) const
void Comm::Sendrecv(const void *sendbuf, int sendcount, const
    Datatype& sendtype, int dest, int sendtag, void *recvbuf,
    int recvcount, const Datatype& recvtype, int source,
    int recvtag) const
void Comm::Sendrecv_replace(void* buf, int count, const
    Datatype& datatype, int dest, int sendtag, int source,
    int recvtag, Status& status) const
void Comm::Sendrecv_replace(void* buf, int count, const
    Datatype& datatype, int dest, int sendtag, int source,
    int recvtag) const
Datatype Datatype::Create_contiguous(int count) const
Datatype Datatype::Create_vector(int count, int blocklength, int stride)
    const
Datatype Datatype::Create_indexed(int count,
    const int array_of_blocklengths[],
    const int array_of_displacements[]) const
int Datatype::Get_size() const
void Datatype::Commit()
void Datatype::Free()
int Status::Get_elements(const Datatype& datatype) const
void Datatype::Pack(const void* inbuf, int incount, void *outbuf,
    int outsize, int& position, const Comm &comm) const
void Datatype::Unpack(const void* inbuf, int insize, void *outbuf,
    int outcount, int& position, const Comm& comm) const
int Datatype::Pack_size(int incount, const Comm& comm) const
};

```

B.5 C++ Bindings for Collective Communication

```
namespace MPI {
```

```

1  void Intracomm::Barrier() const
2
3  void Intracomm::Bcast(void* buffer, int count, const Datatype& datatype,
4      int root) const
5
6  void Intracomm::Gather(const void* sendbuf, int sendcount, const
7      Datatype& sendtype, void* recvbuf, int recvcount,
8      const Datatype& recvtpe, int root) const
9
10 void Intracomm::Gatherv(const void* sendbuf, int sendcount, const
11     Datatype& sendtype, void* recvbuf, const int recvcounts[],
12     const int displs[], const Datatype& recvtpe, int root) const
13
14 void Intracomm::Scatter(const void* sendbuf, int sendcount, const
15     Datatype& sendtype, void* recvbuf, int recvcount,
16     const Datatype& recvtpe, int root) const
17
18 void Intracomm::Scatterv(const void* sendbuf, const int sendcounts[],
19     const int displs[], const Datatype& sendtype, void* recvbuf,
20     int recvcount, const Datatype& recvtpe, int root) const
21
22 void Intracomm::Allgather(const void* sendbuf, int sendcount, const
23     Datatype& sendtype, void* recvbuf, int recvcount,
24     const Datatype& recvtpe) const
25
26 void Intracomm::Allgatherv(const void* sendbuf, int sendcount, const
27     Datatype& sendtype, void* recvbuf, const int recvcounts[],
28     const int displs[], const Datatype& recvtpe) const
29
30 void Intracomm::Alltoall(const void* sendbuf, int sendcount, const
31     Datatype& sendtype, void* recvbuf, int recvcount,
32     const Datatype& recvtpe) const
33
34 void Intracomm::Alltoallv(const void* sendbuf, const int sendcounts[],
35     const int sdispls[], const Datatype& sendtype, void* recvbuf,
36     const int recvcounts[], const int rdispls[],
37     const Datatype& recvtpe) const
38
39 void Intracomm::Reduce(const void* sendbuf, void* recvbuf, int count,
40     const Datatype& datatype, const Op& op, int root) const
41
42 void Op::Init(User_function* function, bool commute)
43
44 void Op::Free()
45
46 void Intracomm::Allreduce(const void* sendbuf, void* recvbuf, int count,
47     const Datatype& datatype, const Op& op) const
48
49 void Intracomm::Reduce_scatter(const void* sendbuf, void* recvbuf,
50     int recvcounts[], const Datatype& datatype, const Op& op)
51     const
52
53 void Intracomm::Scan(const void* sendbuf, void* recvbuf, int count,
54     const Datatype& datatype, const Op& op) const

```

```
};
```

B.6 C++ Bindings for Groups, Contexts, and Communicators

For both syntactic and semantic reasons, the `Dup()` functions listed below are not virtual. Syntactically, they must each have a different return type. `Dup()` and `Clone` are discussed in Section 10.1.7, page 278.

```
namespace MPI {
    int Group::Get_size() const
    int Group::Get_rank() const
    static void Group::Translate_ranks (const Group& group1, int n,
                                       const int ranks1[], const Group& group2, int ranks2[])
    static int Group::Compare(const Group& group1, const Group& group2)
    Group Comm::Get_group() const
    static Group Group::Union(const Group& group1, const Group& group2)
    static Group Group::Intersect(const Group& group1, const Group& group2)
    static Group Group::Difference(const Group& group1, const Group& group2)
    Group Group::Incl(int n, const int ranks[]) const
    Group Group::Excl(int n, const int ranks[]) const
    Group Group::Range_incl(int n, const int ranges[][3]) const
    Group Group::Range_excl(int n, const int ranges[][3]) const
    void Group::Free()
    int Comm::Get_size() const
    int Comm::Get_rank() const
    static int Comm::Compare(const Comm& comm1, const Comm& comm2)
    Intracomm Intracomm::Dup() const
    Intercomm Intercomm::Dup() const
    Cartcomm Cartcomm::Dup() const
    Graphcomm Graphcomm::Dup() const
    Comm& Comm::Clone() const = 0
    Intracomm& Intracomm::Clone() const
    Intercomm& Intercomm::Clone() const
    Cartcomm& Cartcomm::Clone() const
```

```

1      Graphcomm& Graphcomm::Clone() const
2      Intracomm Intracomm::Create(const Group& group) const
3
4      Intracomm Intracomm::Split(int color, int key) const
5
6      void Comm::Free()
7
8      bool Comm::Is_inter() const
9
10     int Intercomm::Get_remote_size() const
11
12     Group Intercomm::Get_remote_group() const
13
14     Intercomm Intracomm::Create_intercomm(int local_leader, const
15         Comm& peer_comm, int remote_leader, int tag) const
16
17     Intracomm Intercomm::Merge(bool high) const
18
19 };

```

B.7 C++ Bindings for Process Topologies

```

20 namespace MPI {
21
22     Cartcomm Intracomm::Create_cart(int ndims, const int dims[],
23         const bool periods[], bool reorder) const
24
25     void Compute_dims(int nnodes, int ndims, int dims[])
26
27     Graphcomm Intracomm::Create_graph(int nnodes, const int index[],
28         const int edges[], bool reorder) const
29
30     int Comm::Get_topology() const
31
32     void Graphcomm::Get_dims(int nnodes[], int nedges[]) const
33
34     void Graphcomm::Get_topo(int maxindex, int maxedges, int index[],
35         int edges[]) const
36
37     int Cartcomm::Get_dim() const
38
39     void Cartcomm::Get_topo(int maxdims, int dims[], bool periods[],
40         int coords[]) const
41
42     int Cartcomm::Get_cart_rank(const int coords[]) const
43
44     void Cartcomm::Get_coords(int rank, int maxdims, int coords[]) const
45
46     int Graphcomm::Get_neighbors_count(int rank) const
47
48     void Graphcomm::Get_neighbors(int rank, int maxneighbors, int
49         neighbors[]) const
50
51     void Cartcomm::Shift(int direction, int disp, int& rank_source,
52         int& rank_dest) const
53
54 }

```

```

    Cartcomm Cartcomm::Sub(const bool remain_dims[]) const
    int Cartcomm::Map(int ndims, const int dims[], const bool periods[])
        const
    int Graphcomm::Map(int nnodes, const int index[], const int edges[])
        const
};

```

B.8 C++ Bindings for Environmental Inquiry

```

namespace MPI {
    void Get_processor_name(char* name, int& resultlen)
    void Errhandler::Free()
    void Get_error_string(int errorcode, char* name, int& resultlen)
    int Get_error_class(int errorcode)
    double Wtime()
    double Wtick()
    void Init(int& argc, char**& argv)
    void Init()
    void Finalize()
    bool Is_initialized()
    void Comm::Abort(int errorcode)
};

```

B.9 C++ Bindings for Profiling

```

namespace MPI {
    void Pcontrol(const int level, ...)
};

```

B.10 C++ Bindings for Status Access

```

namespace MPI {

```

```

1      int Status::Get_source() const
2      void Status::Set_source(int source)
3
4      int Status::Get_tag() const
5      void Status::Set_tag(int tag)
6
7      int Status::Get_error() const
8      void Status::Set_error(int error)
9
10
11 };
12

```

B.11 C++ Bindings for New 1.2 Functions

```

13 namespace MPI {
14
15     void Get_version(int& version, int& subversion);
16
17
18 };
19

```

B.12 C++ Bindings for Exceptions

```

20 namespace MPI {
21
22     Exception::Exception(int error_code);
23
24     int Exception::Get_error_code() const;
25
26     int Exception::Get_error_class() const;
27
28     const char* Exception::Get_error_string() const;
29
30
31 };
32

```

B.13 C++ Bindings on all MPI Classes

The C++ language requires all classes to have four special functions: a default constructor, a copy constructor, a destructor, and an assignment operator. The bindings for these functions are listed below; their semantics are discussed in Section 10.1.5. The two constructors are *not virtual*. The bindings prototype functions using the type `<CLASS>` rather than listing each function for every MPI class; the token `<CLASS>` can be replaced with valid MPI-2 class names, such as `Group`, `Datatype`, etc., except when noted. In addition, bindings are provided for comparison and inter-language operability from Sections 10.1.5 and 10.1.9.

B.13.1 Construction / Destruction

```

namespace MPI {
    <CLASS>::<CLASS>()
    <CLASS>::~~<CLASS>()

};

```

B.13.2 Copy / Assignment

```

namespace MPI {
    <CLASS>::<CLASS>(const <CLASS>& data)
    <CLASS>& <CLASS>::operator=(const <CLASS>& data)

};

```

B.13.3 Comparison

Since `Status` instances are not handles to underlying MPI objects, the `operator==()` and `operator!=()` functions are not defined on the `Status` class.

```

namespace MPI {
    bool <CLASS>::operator==(const <CLASS>& data) const
    bool <CLASS>::operator!=(const <CLASS>& data) const

};

```

B.13.4 Inter-language Operability

Since there are no C++ `MPI::STATUS_IGNORE` and `MPI::STATUSES_IGNORE` objects, the results of promoting the C or Fortran handles (`MPI_STATUS_IGNORE` and `MPI_STATUSES_IGNORE`) to C++ is undefined.

```

namespace MPI {
    <CLASS>& <CLASS>::operator=(const MPI_<CLASS>& data)
    <CLASS>::<CLASS>(const MPI_<CLASS>& data)
    <CLASS>::operator MPI_<CLASS>() const

};

```

B.13.5 Function Name Cross Reference

Since some of the C++ bindings have slightly different names than their C and Fortran counterparts, this section maps each language neutral MPI-1 name to its corresponding C++ binding.

For brevity, the “MPI:” prefix is assumed for all C++ class names.

Where MPI-1 names have been deprecated, the `<none>` keyword is used in the “Member name” column to indicate that this function is supported with a new name (see Annex A).

Where non-void values are listed in the “Return value” column, the given name is that of the corresponding parameter in the language neutral specification.

MPI Function	C++ class	Member name	Return value	
MPI_ABORT	Comm	Abort	void	1
MPI_ADDRESS		<none>		2
MPI_ALLGATHERV	Intracomm	Allgatherv	void	3
MPI_ALLGATHER	Intracomm	Allgather	void	4
MPI_ALLREDUCE	Intracomm	Allreduce	void	5
MPI_ALLTOALLV	Intracomm	Alltoallv	void	6
MPI_ALLTOALL	Intracomm	Alltoall	void	7
MPI_ATTR_DELETE		<none>		8
MPI_ATTR_GET		<none>		9
MPI_ATTR_PUT		<none>		10
MPI_BARRIER	Intracomm	Barrier	void	11
MPI_BCAST	Intracomm	Bcast	void	12
MPI_BSEND_INIT	Comm	Bsend_init	Prequest request	13
MPI_BSEND	Comm	Bsend	void	14
MPI_BUFFER_ATTACH		Attach_buffer	void	15
MPI_BUFFER_DETACH		Detach_buffer	void* buffer	16
MPI_CANCEL	Request	Cancel	void	17
MPI_CARTDIM_GET	Cartcomm	Get_dim	int ndims	18
MPI_CART_COORDS	Cartcomm	Get_coords	void	19
MPI_CART_CREATE	Intracomm	Create_cart	Cartcomm newcomm	20
MPI_CART_GET	Cartcomm	Get_topo	void	21
MPI_CART_MAP	Cartcomm	Map	int newrank	22
MPI_CART_RANK	Cartcomm	Get_rank	int rank	23
MPI_CART_SHIFT	Cartcomm	Shift	void	24
MPI_CART_SUB	Cartcomm	Sub	Cartcomm newcomm	25
MPI_COMM_COMPARE	Comm	static Compare	int result	26
MPI_COMM_CREATE	Intracomm	Create	Intracomm newcomm	27
MPI_COMM_DUP	Intracomm	Dup	Intracomm newcomm	28
	Cartcomm	Dup	Cartcomm newcomm	29
	Graphcomm	Dup	Graphcomm newcomm	30
	Intercomm	Dup	Intercomm newcomm	31
	Comm	Clone	Comm& newcomm	32
	Intracomm	Clone	Intracomm& newcomm	33
	Cartcomm	Clone	Cartcomm& newcomm	34
	Graphcomm	Clone	Graphcomm& newcomm	35
	Intercomm	Clone	Intercomm& newcomm	36
MPI_COMM_FREE	Comm	Free	void	37
MPI_COMM_GROUP	Comm	Get_group	Group group	38
MPI_COMM_RANK	Comm	Get_rank	int rank	39
MPI_COMM_REMOTE_GROUP	Intercomm	Get_remote_group	Group group	40
MPI_COMM_REMOTE_SIZE	Intercomm	Get_remote_size	int size	41
MPI_COMM_SIZE	Comm	Get_size	int size	42
MPI_COMM_SPLIT	Intracomm	Split	Intracomm newcomm	43
MPI_COMM_TEST_INTER	Comm	Is_inter	bool flag	44
MPI_DIMS_CREATE		Compute_dims	void	45

47

48

	MPI Function	C++ class	Member name	Return value
1	MPI_ERRHANDLER_CREATE		<none>	
2	MPI_ERRHANDLER_FREE	Errhandler	Free	void
3	MPI_ERRHANDLER_GET		<none>	
4	MPI_ERRHANDLER_SET		<none>	
5	MPI_ERROR_CLASS		Get_error_class	int errorclass
6	MPI_ERROR_STRING		Get_error_string	void
7	MPI_FINALIZE		Finalize	void
8	MPI_GATHERV	Intracomm	Gatherv	void
9	MPI_GATHER	Intracomm	Gather	void
10	MPI_GET_COUNT	Status	Get_count	int count
11	MPI_GET_ELEMENTS	Status	Get_elements	int count
12	MPI_GET_PROCESSOR_NAME		Get_processor_name	void
13	MPI_GRAPHDIMS_GET	Graphcomm	Get_dims	void
14	MPI_GRAPH_CREATE	Intracomm	Create_graph	Graphcomm newcomm
15	MPI_GRAPH_GET	Graphcomm	Get_topo	void
16	MPI_GRAPH_MAP	Graphcomm	Map	int newrank
17	MPI_GRAPH_NEIGHBORS_COUNT	Graphcomm	Get_neighbors_count	int nneighbors
18	MPI_GRAPH_NEIGHBORS	Graphcomm	Get_neighbors	void
19	MPI_GROUP_COMPARE	Group	static Compare	int result
20	MPI_GROUP_DIFFERENCE	Group	static Difference	Group newgroup
21	MPI_GROUP_EXCL	Group	Excl	Group newgroup
22	MPI_GROUP_FREE	Group	Free	void
23	MPI_GROUP_INCL	Group	Incl	Group newgroup
24	MPI_GROUP_INTERSECTION	Group	static Intersect	Group newgroup
25	MPI_GROUP_RANGE_EXCL	Group	Range_excl	Group newgroup
26	MPI_GROUP_RANGE_INCL	Group	Range_incl	Group newgroup
27	MPI_GROUP_RANK	Group	Get_rank	int rank
28	MPI_GROUP_SIZE	Group	Get_size	int size
29	MPI_GROUP_TRANSLATE_RANKS	Group	static Translate_ranks	void
30	MPI_GROUP_UNION	Group	static Union	Group newgroup
31	MPI_IBSEND	Comm	Ibsend	Request request
32	MPI_INITIALIZED		Is_initialized	bool flag
33	MPI_INIT		Init	void
34	MPI_INTERCOMM_CREATE	Intracomm	Create_intercomm	Intercomm newcomm
35	MPI_INTERCOMM_MERGE	Intercomm	Merge	Intracomm newcomm
36	MPI_IPROBE	Comm	Iprobe	bool flag
37	MPI_IRECV	Comm	Irecv	Request request
38	MPI_IRSEND	Comm	Irsend	Request request
39	MPI_ISEND	Comm	Isend	Request request
40	MPI_ISSEND	Comm	Issend	Request request
41	MPI_KEYVAL_CREATE		<none>	
42	MPI_KEYVAL_FREE		<none>	
43	MPI_OP_CREATE	Op	Init	void
44	MPI_OP_FREE	Op	Free	void
45	MPI_PACK_SIZE	Datatype	Pack_size	int size
46	MPI_PACK	Datatype	Pack	void

MPI Function	C++ class	Member name	Return value	
MPI_PCONTROL		Pcontrol	void	1
MPI_PROBE	Comm	Probe	void	2
MPI_RECV_INIT	Comm	Recv_init	Prequest request	3
MPI_RECV	Comm	Recv	void	4
MPI_REDUCE_SCATTER	Intracomm	Reduce_scatter	void	5
MPI_REDUCE	Intracomm	Reduce	void	6
MPI_REQUEST_FREE	Request	Free	void	7
MPI_RSEND_INIT	Comm	Rsend_init	Prequest request	8
MPI_RSEND	Comm	Rsend	void	9
MPI_SCAN	Intracomm	Scan	void	10
MPI_SCATTERV	Intracomm	Scatterv	void	11
MPI_SCATTER	Intracomm	Scatter	void	12
MPI_SENDRECV_REPLACE	Comm	Sendrecv_replace	void	13
MPI_SENDRECV	Comm	Sendrecv	void	14
MPI_SEND_INIT	Comm	Send_init	Prequest request	15
MPI_SEND	Comm	Send	void	16
MPI_SSEND_INIT	Comm	Ssend_init	Prequest request	17
MPI_SSEND	Comm	Ssend	void	18
MPI_STARTALL	Prequest	static Startall	void	19
MPI_START	Prequest	Start	void	20
MPI_TESTALL	Request	static Testall	bool flag	21
MPI_TESTANY	Request	static Testany	bool flag	22
MPI_TESTSOME	Request	static Testsome	int outcount	23
MPI_TEST_CANCELLED	Status	Is_cancelled	bool flag	24
MPI_TEST	Request	Test	bool flag	25
MPI_TOPO_TEST	Comm	Get_topo	int status	26
MPI_TYPE_COMMIT	Datatype	Commit	void	27
MPI_TYPE_CONTIGUOUS	Datatype	Create_contiguous	Datatype	28
MPI_TYPE_EXTENT		<none>		29
MPI_TYPE_FREE	Datatype	Free	void	30
MPI_TYPE_HINDEXED		<none>		31
MPI_TYPE_HVECTOR		<none>		32
MPI_TYPE_INDEXED	Datatype	Create_indexed	Datatype	33
MPI_TYPE_LB		<none>		34
MPI_TYPE_SIZE	Datatype	Get_size	int	35
MPI_TYPE_STRUCT		<none>		36
MPI_TYPE_UB		<none>		37
MPI_TYPE_VECTOR	Datatype	Create_vector	Datatype	38
MPI_UNPACK	Datatype	Unpack	void	39
MPI_WAITALL	Request	static Waitall	void	40
MPI_WAITANY	Request	static Waitany	int index	41
MPI_WAITSOME	Request	static Waitsome	int outcount	42
MPI_WAIT	Request	Wait	void	43
MPI_WTICK		Wtick	double wtick	44
MPI_WTIME		Wtime	double wtime	45

47

48

MPI Function Index

MPLACCUMULATE, 119
MPIADD_ERROR_CLASS, 181
MPIADD_ERROR_CODE, 182
MPIADD_ERROR_STRING, 182
MPIALLGATHER, 158
MPIALLGATHERV, 159
MPIALLOC_MEM, 48
MPIALLREDUCE, 162
MPIALLTOALL, 160
MPIALLTOALLV, 161
MPIALLTOALLW, 165
MPIBARRIER, 163
MPIBCAST, 153
MPLCLOSE_PORT, 96
MPICOMM_ACCEPT, 97
MPICOMM_C2F, 51
MPICOMM_CALL_ERRHANDLER, 183
MPICOMM_CLONE, 279
MPICOMM_CONNECT, 98
MPICOMM_COPY_ATTR_FUNCTION, 200
MPICOMM_CREATE, 146
MPICOMM_CREATE_ERRHANDLER, 62
MPICOMM_CREATE_KEYVAL, 199
MPICOMM_DELETE_ATTR, 202
MPICOMM_DELETE_ATTR_FUNCTION, 200
MPICOMM_DISCONNECT, 107
MPICOMM_DUP_FN, 199
MPICOMM_ERRHANDLER_FN, 62
MPICOMM_F2C, 51
MPICOMM_FREE_KEYVAL, 200
MPICOMM_GET_ATTR, 201
MPICOMM_GET_ERRHANDLER, 63
MPICOMM_GET_NAME, 178
MPICOMM_GET_PARENT, 88
MPICOMM_JOIN, 107
MPICOMM_NULL_COPY_FN, 199
MPICOMM_NULL_DELETE_FN, 200
MPICOMM_SET_ATTR, 201
MPICOMM_SET_ERRHANDLER, 62
MPICOMM_SET_NAME, 177
MPICOMM_SPAWN, 84
MPICOMM_SPAWN_MULTIPLE, 89
MPICOMM_SPLIT, 147
MPIDATAREP_CONVERSION_FUNCTION, 253
MPIDATAREP_EXTENT_FUNCTION, 252
MPIEXSCAN, 166
MPIFILE_C2F, 51
MPIFILE_CALL_ERRHANDLER, 184
MPIFILE_CLOSE, 213
MPIFILE_CREATE_ERRHANDLER, 64
MPIFILE_DELETE, 214
MPIFILE_ERRHANDLER_FN, 65
MPIFILE_F2C, 51
MPIFILE_GET_AMODE, 217
MPIFILE_GET_ATOMICITY, 258
MPIFILE_GET_BYTE_OFFSET, 235
MPIFILE_GET_ERRHANDLER, 65
MPIFILE_GET_GROUP, 216
MPIFILE_GET_INFO, 219
MPIFILE_GET_POSITION, 235
MPIFILE_GET_POSITION_SHARED, 240
MPIFILE_GET_SIZE, 216
MPIFILE_GET_TYPE_EXTENT, 249
MPIFILE_GET_VIEW, 223
MPIFILE_IREAD, 233
MPIFILE_IREAD_AT, 229
MPIFILE_IREAD_SHARED, 237
MPIFILE_IWRITE, 234
MPIFILE_IWRITE_AT, 229
MPIFILE_IWRITE_SHARED, 237
MPIFILE_OPEN, 211
MPIFILE_PREALLOCATE, 215
MPIFILE_READ, 230
MPIFILE_READ_ALL, 231
MPIFILE_READ_ALL_BEGIN, 243
MPIFILE_READ_ALL_END, 244
MPIFILE_READ_AT, 227

MPI_FILE_READ_AT_ALL, 227	MPI_INFO_GET, 45	1
MPI_FILE_READ_AT_ALL_BEGIN, 242	MPI_INFO_GET_NKEYS, 46	2
MPI_FILE_READ_AT_ALL_END, 242	MPI_INFO_GET_NTHKEY, 46	3
MPI_FILE_READ_ORDERED, 238	MPI_INFO_GET_VALUELEN, 46	4
MPI_FILE_READ_ORDERED_BEGIN, 245	MPI_INFO_SET, 44	5
MPI_FILE_READ_ORDERED_END, 245	MPI_INIT_THREAD, 196	6
MPI_FILE_READ_SHARED, 236	MPI_IS_THREAD_MAIN, 198	7
MPI_FILE_SEEK, 234	MPI_LOOKUP_NAME, 100	8
MPI_FILE_SEEK_SHARED, 239	MPI_OP_C2F, 51	9
MPI_FILE_SET_ATOMICITY, 257	MPI_OP_F2C, 51	10
MPI_FILE_SET_ERRHANDLER, 65	MPI_OPEN_PORT, 96	11
MPI_FILE_SET_INFO, 218	MPI_PACK_EXTERNAL, 78	12
MPI_FILE_SET_SIZE, 215	MPI_PACK_EXTERNAL_SIZE, 79	13
MPI_FILE_SET_VIEW, 221	MPI_PUBLISH_NAME, 99	14
MPI_FILE_SYNC, 258	MPI_PUT, 114	15
MPI_FILE_WRITE, 232	MPI_QUERY_THREAD, 197	16
MPI_FILE_WRITE_ALL, 232	MPI_REDUCE, 162	17
MPI_FILE_WRITE_ALL_BEGIN, 244	MPI_REDUCE_SCATTER, 163	18
MPI_FILE_WRITE_ALL_END, 244	MPI_REGISTER_DATAREP, 252	19
MPI_FILE_WRITE_AT, 228	MPI_REQUEST_C2F, 51	20
MPI_FILE_WRITE_AT_ALL, 228	MPI_REQUEST_F2C, 51	21
MPI_FILE_WRITE_AT_ALL_BEGIN, 242	MPI_REQUEST_GET_STATUS, 41	22
MPI_FILE_WRITE_AT_ALL_END, 243	MPI_SCAN, 164	23
MPI_FILE_WRITE_ORDERED, 239	MPI_SCATTER, 156	24
MPI_FILE_WRITE_ORDERED_BEGIN, 245	MPI_SCATTERV, 157	25
MPI_FILE_WRITE_ORDERED_END, 246	MPI_SIZEOF, 297	26
MPI_FILE_WRITE_SHARED, 236	MPI_STATUS_C2F, 54	27
MPI_FINALIZED, 43	MPI_STATUS_F2C, 54	28
MPI_FREE_MEM, 48	MPI_STATUS_SET_CANCELLED, 177	29
MPI_GATHER, 154	MPI_STATUS_SET_ELEMENTS, 176	30
MPI_GATHERV, 155	MPI_TYPE_C2F, 51	31
MPI_GET, 116	MPI_TYPE_COPY_ATTR_FUNCTION, 205	32
MPI_GET_ADDRESS, 67	MPI_TYPE_CREATE_DARRAY, 73	33
MPI_GET_VERSION, 21	MPI_TYPE_CREATE_F90_COMPLEX, 294	34
MPI_GREQUEST_CANCEL_FUNCTION,	MPI_TYPE_CREATE_F90_INTEGER, 295	35
172	MPI_TYPE_CREATE_F90_REAL, 294	36
MPI_GREQUEST_COMPLETE, 173	MPI_TYPE_CREATE_HINDEXED, 66	37
MPI_GREQUEST_FREE_FUNCTION, 171	MPI_TYPE_CREATE_HVECTOR, 66	38
MPI_GREQUEST_QUERY_FUNCTION, 171	MPI_TYPE_CREATE_INDEXED_BLOCK,	39
MPI_GREQUEST_START, 170	40	40
MPI_GROUP_C2F, 51	MPI_TYPE_CREATE_KEYVAL, 205	41
MPI_GROUP_F2C, 51	MPI_TYPE_CREATE_RESIZED, 69	42
MPI_INFO_C2F, 51	MPI_TYPE_CREATE_STRUCT, 67	43
MPI_INFO_CREATE, 44	MPI_TYPE_CREATE_SUBARRAY, 70	44
MPI_INFO_DELETE, 45	MPI_TYPE_DELETE_ATTR, 207	45
MPI_INFO_DUP, 47	MPI_TYPE_DELETE_ATTR_FUNCTION,	46
MPI_INFO_F2C, 51	205	47
MPI_INFO_FREE, 47	MPI_TYPE_DUP, 207	48

1 MPI_TYPE_DUP_FN, 205
2 MPI_TYPE_F2C, 51
3 MPI_TYPE_FREE_KEYVAL, 206
4 MPI_TYPE_GET_ATTR, 207
5 MPI_TYPE_GET_CONTENTS, 187
6 MPI_TYPE_GET_ENVELOPE, 185
7 MPI_TYPE_GET_EXTENT, 68
8 MPI_TYPE_GET_NAME, 180
9 MPI_TYPE_GET_TRUE_EXTENT, 69
10 MPI_TYPE_MATCH_SIZE, 298
11 MPI_TYPE_NULL_COPY_FN, 205
12 MPI_TYPE_NULL_DELETE_FN, 205
13 MPI_TYPE_SET_ATTR, 206
14 MPI_TYPE_SET_NAME, 179
15 MPI_UNPACK_EXTERNAL, 79
16 MPI_UNPUBLISH_NAME, 100
17 MPI_WIN_C2F, 51
18 MPI_WIN_CALL_ERRHANDLER, 183
19 MPI_WIN_COMPLETE, 127
20 MPI_WIN_COPY_ATTR_FUNCTION, 203
21 MPI_WIN_CREATE, 110
22 MPI_WIN_CREATE_ERRHANDLER, 63
23 MPI_WIN_CREATE_KEYVAL, 202
24 MPI_WIN_DELETE_ATTR, 204
25 MPI_WIN_DELETE_ATTR_FUNCTION,
26 203
27 MPI_WIN_DUP_FN, 202
28 MPI_WIN_ERRHANDLER_FN, 63
29 MPI_WIN_F2C, 51
30 MPI_WIN_FENCE, 126
31 MPI_WIN_FREE, 111
32 MPI_WIN_FREE_KEYVAL, 203
33 MPI_WIN_GET_ATTR, 204
34 MPI_WIN_GET_ERRHANDLER, 64
35 MPI_WIN_GET_GROUP, 112
36 MPI_WIN_GET_NAME, 180
37 MPI_WIN_LOCK, 130
38 MPI_WIN_NULL_COPY_FN, 202
39 MPI_WIN_NULL_DELETE_FN, 202
40 MPI_WIN_POST, 128
41 MPI_WIN_SET_ATTR, 204
42 MPI_WIN_SET_ERRHANDLER, 64
43 MPI_WIN_SET_NAME, 180
44 MPI_WIN_START, 127
45 MPI_WIN_TEST, 129
46 MPI_WIN_UNLOCK, 131
47 MPI_WIN_WAIT, 128
48